

Windows Internals

Alon Fliess

Chief Architect, CodeValue (<http://www.codevalue.net>)

alonf@codevalue.net

<http://blogs.microsoft.co.il/blogs/alon>



About Me

- **Alon Fliess:**

- Chief Software Architect & Co-Founder at CodeValue Ltd.
- More than 25 years of hands-on experience
- Microsoft Regional Director & Microsoft MVP
- Active member of several Patterns & Practices councils
- Renowned speaker at both international and domestic events



About CodeValue

- A leading software company
- ~150 employees: more than 130 technology experts
- Provides high quality software development solutions
 - Turn-Key projects
 - Software development and consultation
 - Tailor-made courses and training
- Fields of expertise include:
 - Desktop & LOB applications
 - Cloud Computing
 - Advanced Mobile & Web Technologies
 - User Experience (UX) & User Interface (UI)
 - Application Lifecycle Management (ALM) and DevOps
 - Embedded & IoT

About OzCode

- An innovative debugging extension for Visual-Studio
- Simplify & visualize complex statements
- Compare objects and collections
- Search and filter collections
- Focus on relevant data
- <http://oz-code.com>



Course Details

- Objectives
 - Understand Windows features and architecture
 - Uncover the internal algorithms used by Windows relevant to developers
 - Enhance the ability to design and implement optimized software for the Windows platform
- Target audience
 - Developers & IT Pros

Resources

- Books
 - Windows Internals / David Solomon & Mark Russinovich / 4th, 5th, 6th ,... editions
 - Windows Sysinternals Administrator's Reference / Mark Russinovich & Aaron Margosis
 - Windows via C/C++ / Jeffery Richter & Christof Nassare / 5th edition
 - Programming the Windows Driver Model / Walter Oney / 2nd edition
 - Developing Drivers with the Windows Driver Foundation / Orwick & Smith
- Web
 - www.sysinternals.com
 - www.osr.com
 - www.microsoft.com/whdc
 - <http://nirsoft.net/>
 - <http://msdn.microsoft.com>
 - [https://msdn.microsoft.com/en-us/library/windows/hardware/dn614614\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/windows/hardware/dn614614(v=vs.85).aspx)
 - <http://www.alex-ionescu.com/>

Course Contents

- 1: System Architecture
- 2: Kernel Mechanisms
- 3: Management Mechanisms
- 4: Processes, Threads & Jobs
- 5: Memory Management
- 6: Security
- The I/O Subsystem
- Networking
- 7: Introduction to Windows Universal Apps
- A: Introduction to COM
- B: Introduction to WinDbg

Module 1

SYSTEM ARCHITECTURE

Windows Internals

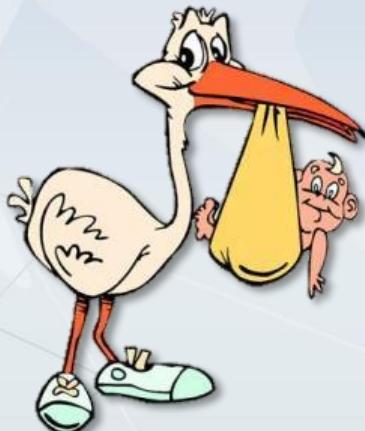
Agenda

- Windows NT History
- Tools
- Basic Concepts
- Windows NT Design Goals
- System Architecture
- Summary

Windows NT History

- Windows NT 3.1 (July 1993)
- Windows NT 3.5 (September 1994)
- Windows NT 3.51 (May 1995)
- Windows NT 4.0 (July 1996)
- Windows 2000 (December 1999)
- Windows XP (August 2001)
- Windows Server 2003 (March 2003)
- Windows Vista (January 2007)
- Windows Server 2008 (February 2008)
- Windows 7 & 2008 R2 (October 2009)
- Windows 8 & Windows Server 2012 (2012)
- Windows 8.1 & Server 2012 R2
- Windows 10 & Server 2016

Tools – Where are they Coming From?

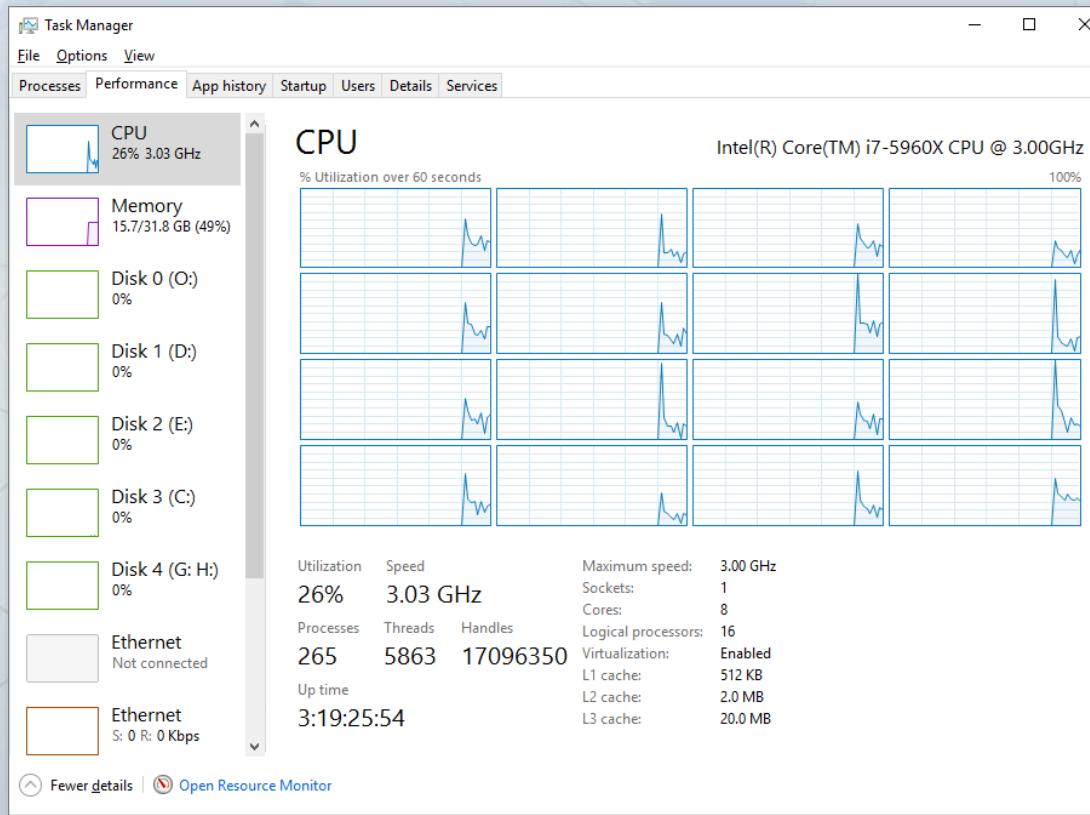


NirSoft

WINDOWS BASED TOOLS

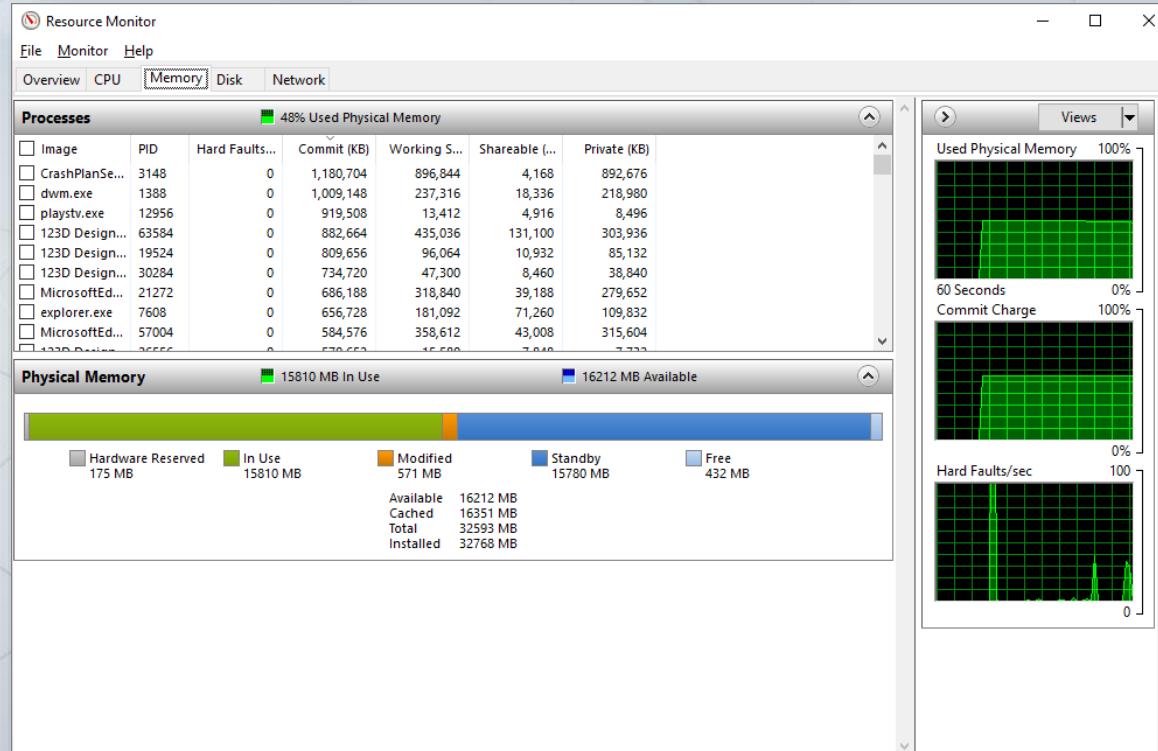
Windows Built-In Tools

Task Manager



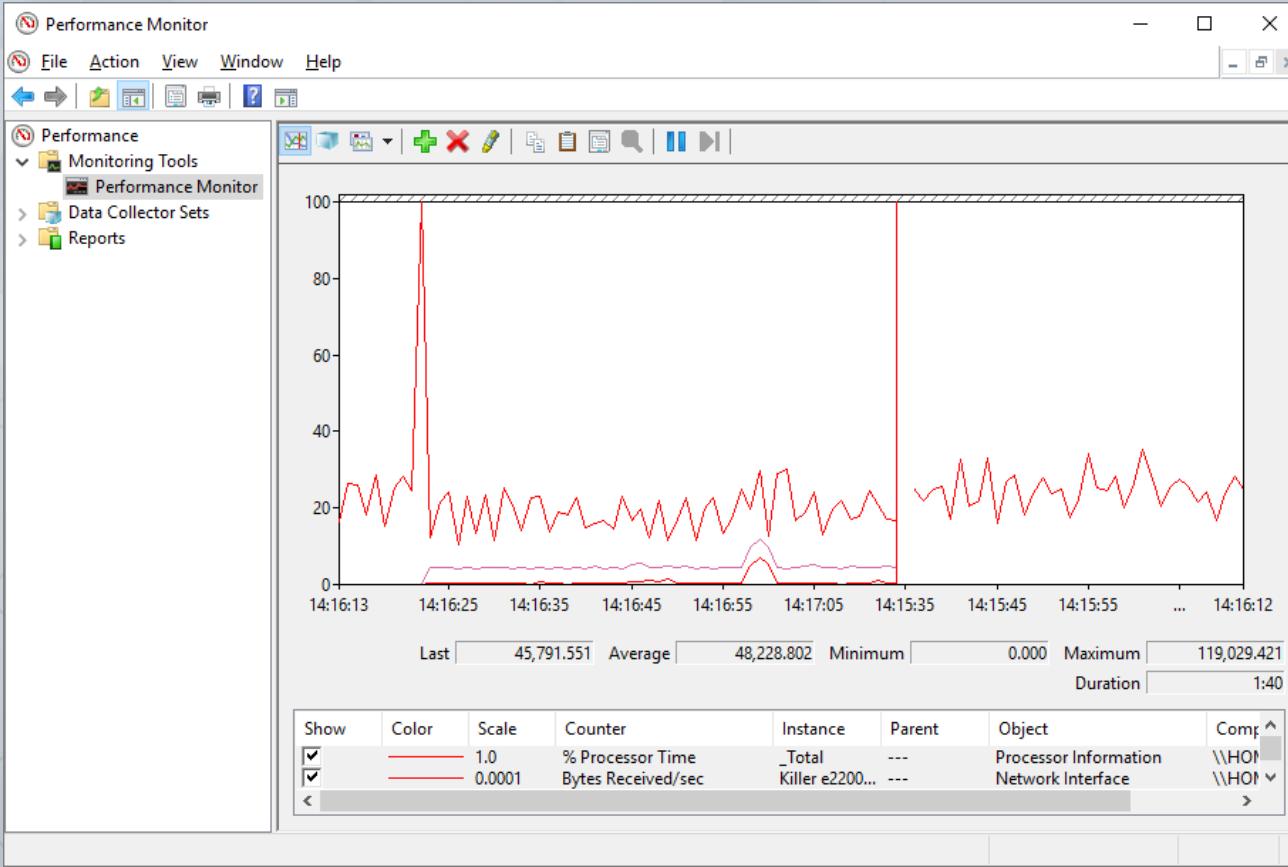
Windows Built-In Tools

Resource Monitor



Windows Built-In Tools

Performance Monitor



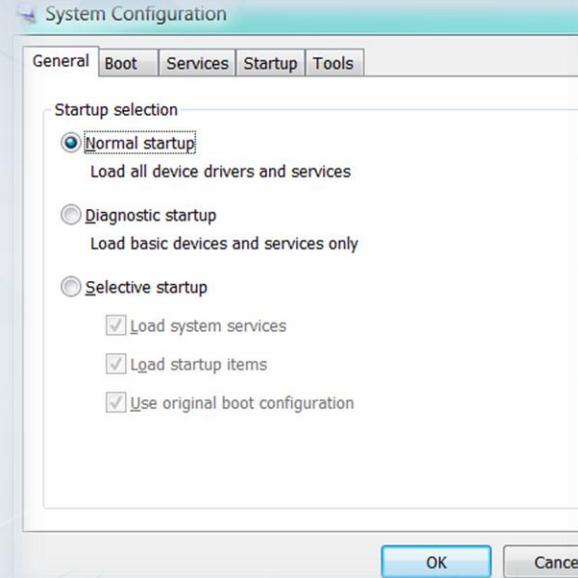
Windows Built-In Tools

Windows Troubleshooting Platform

The screenshot shows the Windows Troubleshooting Platform window. At the top, there's a navigation bar with icons for back, forward, search, and refresh, along with a search bar labeled "Search Tro...". Below the navigation bar, the title "Control Panel Home" is visible. On the left, a sidebar lists links: "View all", "View history", "Change settings", and "Get help from a friend". The main content area is titled "Troubleshoot computer problems" with the sub-instruction "Click on a task to automatically troubleshoot and fix common computer problems. To view more troubleshooters, click on a category or use the Search box.". It features four categories with icons: "Programs" (Run programs made for previous versions of Windows), "Hardware and Sound" (Configure a device | Use a printer | Troubleshoot audio recording | Troubleshoot audio playback), "Network and Internet" (Connect to the Internet | Access shared files and folders on other computers), and "System and Security" (Fix problems with Windows Update | Run maintenance tasks | Improve power usage). At the bottom left, under "See also", are links to "Security and Maintenance", "Help and Support", and "Recovery".

Windows Admin Tools

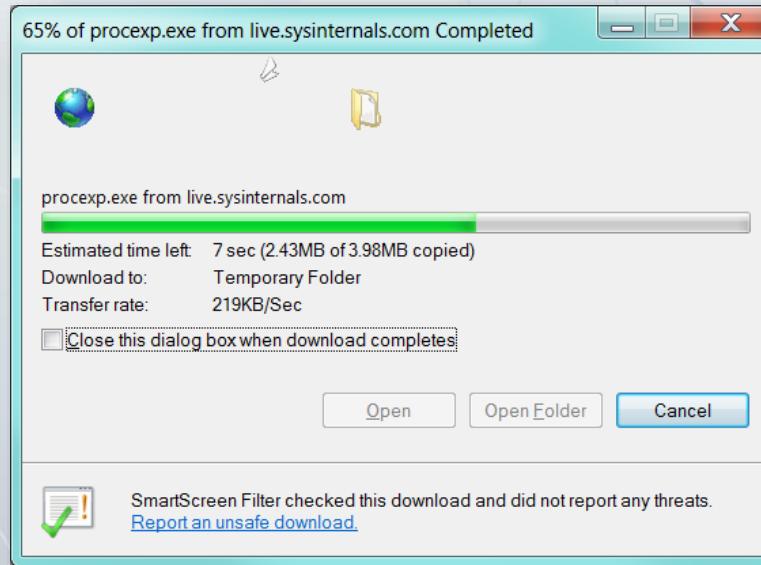
- Control Panel
- MMC
- System Configuration
- System Information
- RegEdit
- ...



SYSTEM INTERNALS TOOLS

<http://live.sysinternals.com>

<\\live.sysinternals.com>

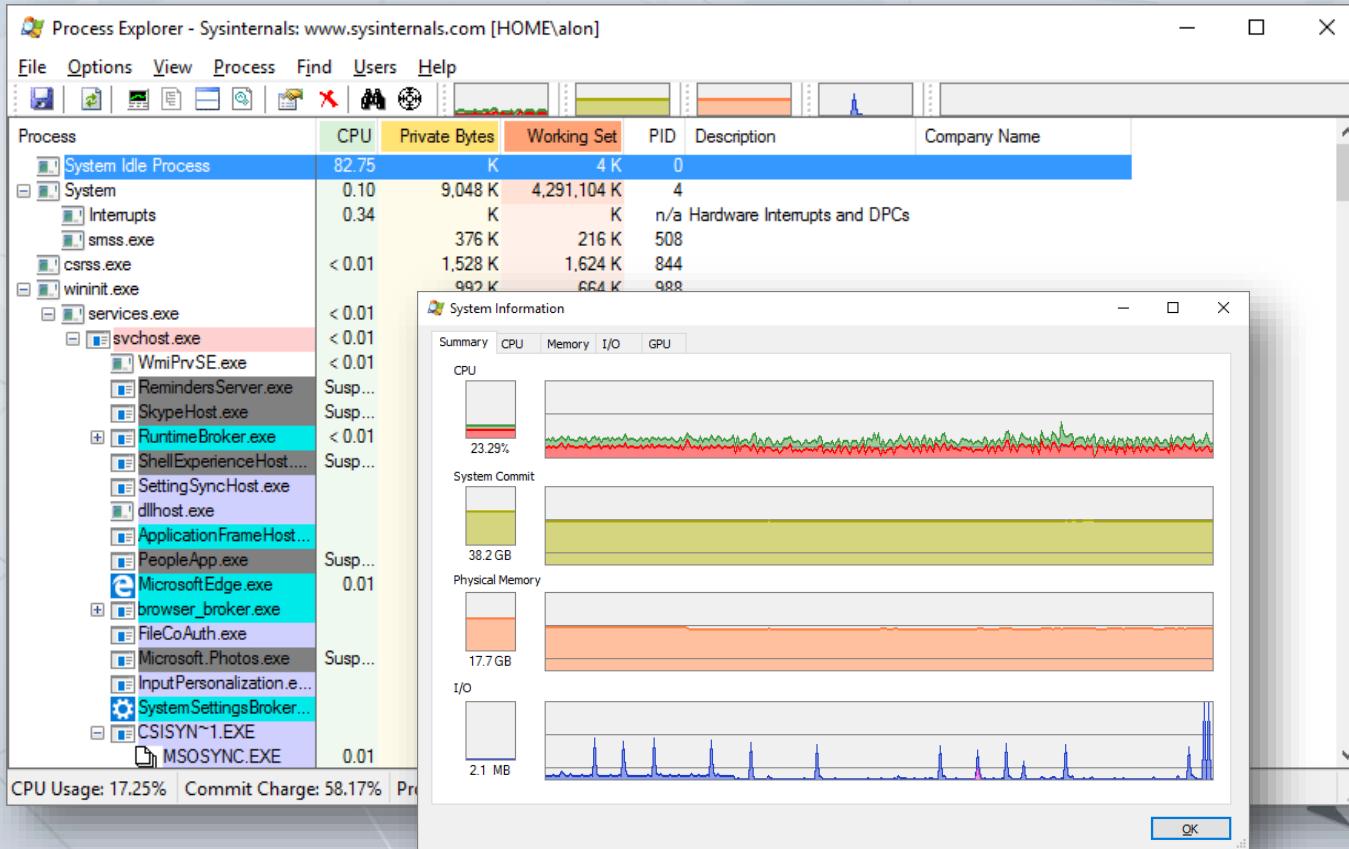


System Internals

- The Sysinternals web site was created in 1996 by Mark Russinovich and Bryce Cogswell
- Microsoft acquired Sysinternals in July, 2006
- Targets both IT pros and developers
- If you have a question about a tool or how to use it, please visit the [Sysinternals Forum](#)

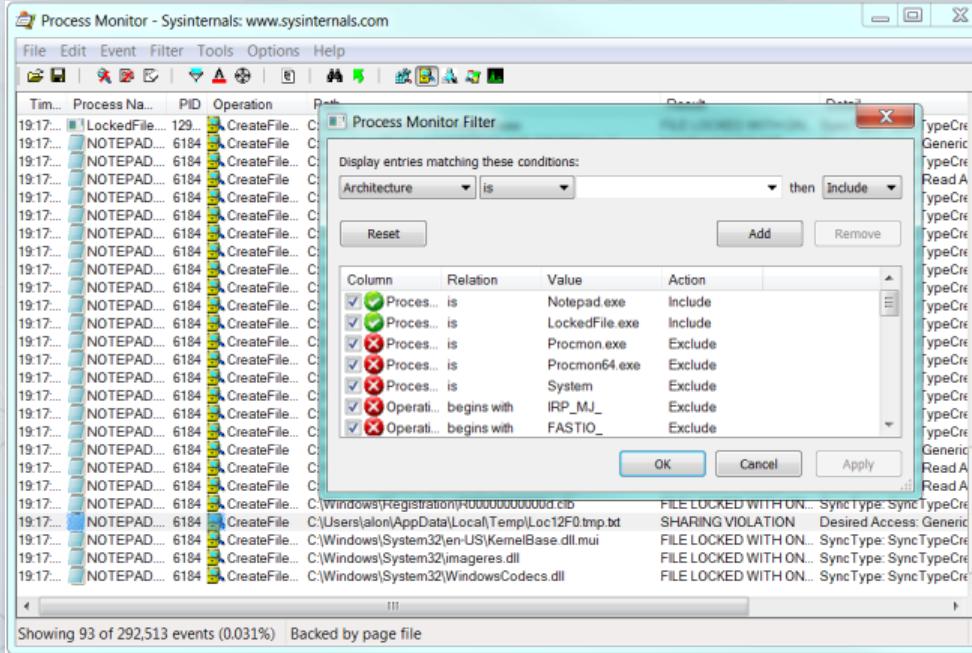
System Internals Tools

Process Explorer



System Internals Tools

Process Monitor



Additional Important SysInternals Tools

- [AccessEnum](#) - shows you who has what access to directories, files and Registry keys
- [Autoruns](#) – See who automatically loads a dll or auto starts an exe
- [DebugView](#) – Intercept debug output calls, very useful for debug sessions
- [PipeList](#) – Displays the named-pipes with their properties in the system
- [ProcDump](#) – Very useful process dump tool, dump process memory on various occasions
- [SysMon](#) - Monitor and log system activity to the Windows event log
- [PsTools](#) - CLI utilities for handling processes running on local or remote computers, rebooting computers, dumping event logs, and more.
- [RAMMap](#) - Physical memory usage analysis utility
- [TCPView](#) - Active socket command-line viewer
- [VMMMap](#) - A process virtual and physical memory analysis utility
- [WinObj](#) – The Object Manager namespace viewer

Windows SDK

DEBUGGING TOOLS FOR WINDOWS

Overview

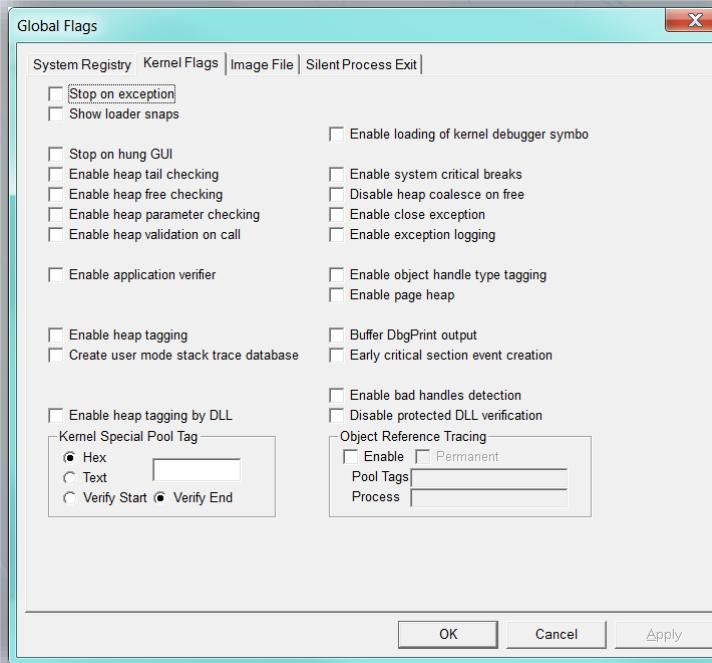
- Debugging Tools for Windows is a collection of debugging and diagnostic tools

It contains:

- Four Windows debuggers (WinDng, KD, CDB and NTSD)
- Logger and Log Viewer (Logger.exe, Logviewer.exe)
- ADPlus (Autodump+, Adplus.vbs) - Automatically create memory dump files and log files
- GFlags – Used to control System debug settings
- UMDH - User-Mode Dump Heap utility (Umdh.exe)
- Symbol & Source Server files
- Remote Debugging tools
- The debugger.chm help file
- In the next chapter we will deep dive into the debuggers

Global Flags

- GFlags (the Global Flags Editor), gflags.exe, enables and disables advanced debugging, diagnostic, and troubleshooting features



WINDOWS SDK AND WINDOWS DRIVER KIT TOOLS



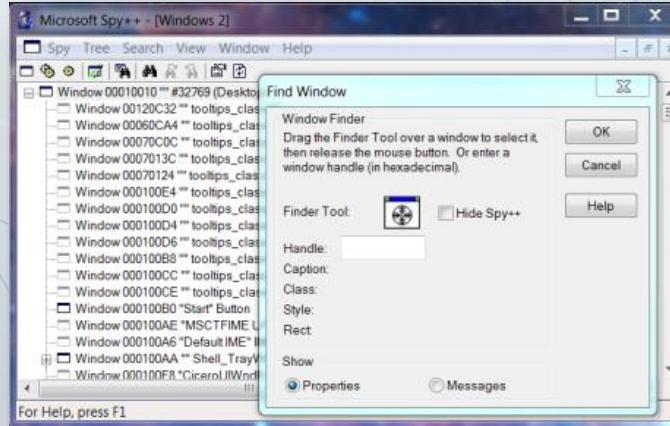
SDK/WDK Tools Overview

- The Windows SDK & WDK contain:
 - Tools
 - Code samples
 - Documentation
 - C/C++/.NET compilers, headers, and libraries
- Enable developers create applications & device drivers that run on the Windows platform

VISUAL STUDIO TOOLS

Visual Studio Tools Overview

- With Visual Studio comes a large set of tools for building, debugging (local & remote), trace and test applications
- Some of the tools are accessible from within Visual Studio
 - Others are command-line or external UI based tools
- Partial tool list:
 - Spy++
 - IntelliTrace
 - VS Profiler
 - C/C++ and .NET static Analysis
 - Error Lookup



Windows Internals

APPLICATION COMPATIBILITY TOOLKIT



ADK – Windows 10 & Deployment Kit

- Tools for:
 - Customize Windows images for large-scale deployment
 - Windows PE (Windows Preinstallation Environment)
 - Test the quality and performance of the system & applications
 - Contain the ACT – Application Compatibility Toolkit
 - Creating an inventory for your organization, including installed applications, computers, and devices
 - Collecting compatibility data
 - Creating mitigation packages to fix the compatibility issues

ADK

Windows Assessment and Deployment Kit - Windows 10

Select the features you want to install

Click a feature name for more information.

- Application Compatibility Toolkit (ACT)
- Deployment Tools
- Windows Preinstallation Environment (Windows PE)
- Imaging And Configuration Designer (ICD)
- User State Migration Tool (USMT)
- Volume Activation Management Tool (VAMT)
- Windows Performance Toolkit
- Windows Assessment Toolkit
- Windows Assessment Services - Client
- Microsoft SQL Server 2012 Express

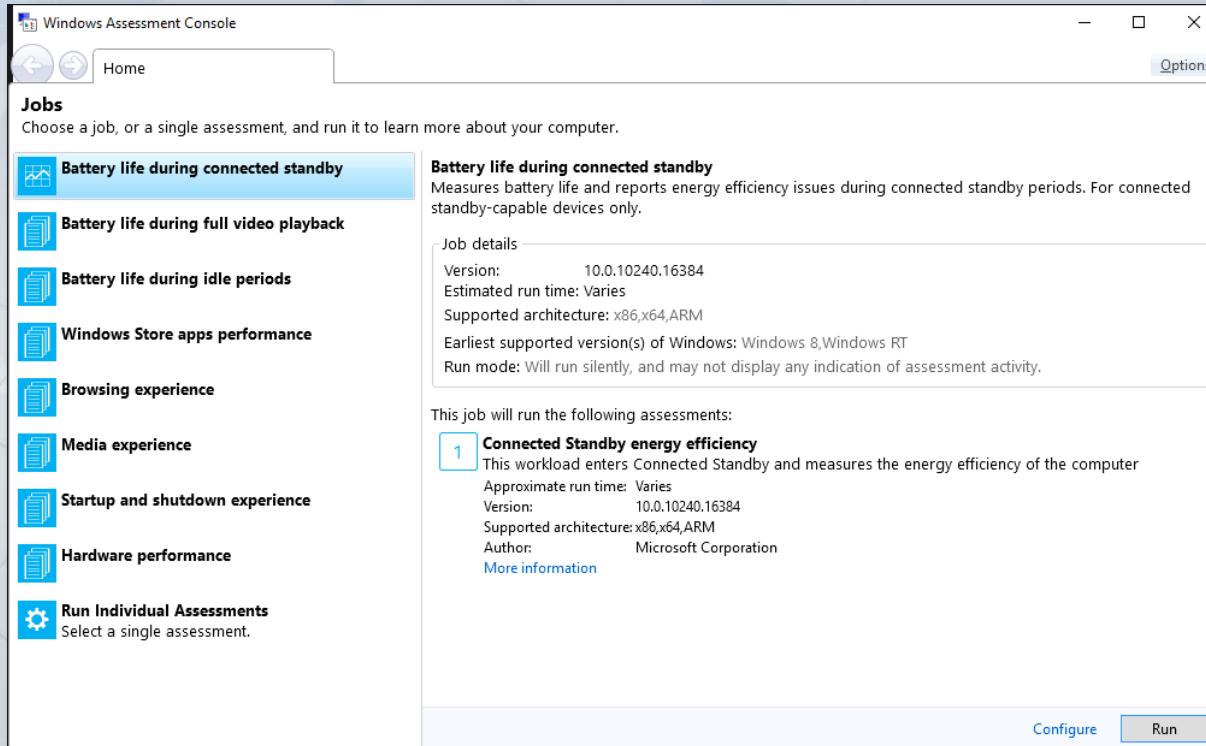
Application Compatibility Toolkit (ACT)
Size: 26.1 MB
Tools to evaluate and mitigate application compatibility issues before deploying a new version of Windows.

ACT requires access to a database. The database must be SQL Server 2005 (or Express Edition) or later. You can install SQL Server or use an existing installation.

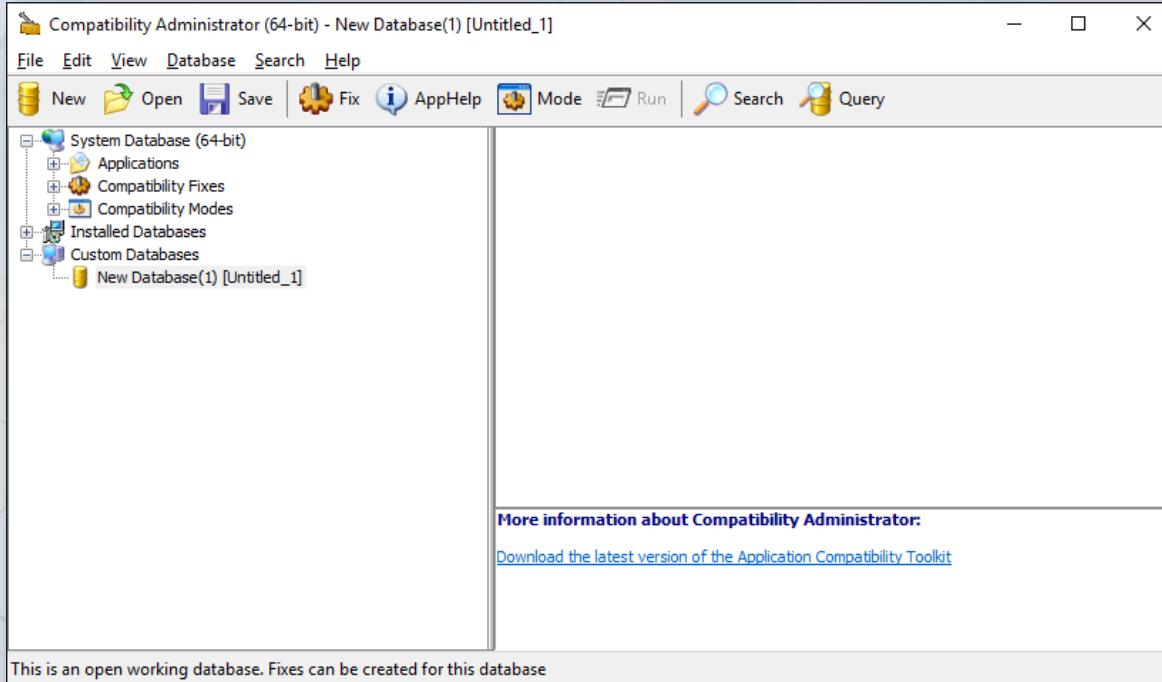
Estimated disk space required: 6.5 GB
Disk space available: 291.7 GB

Back  Install Cancel

Windows Assessment Console



Compatibility Administrator



Windows Internals

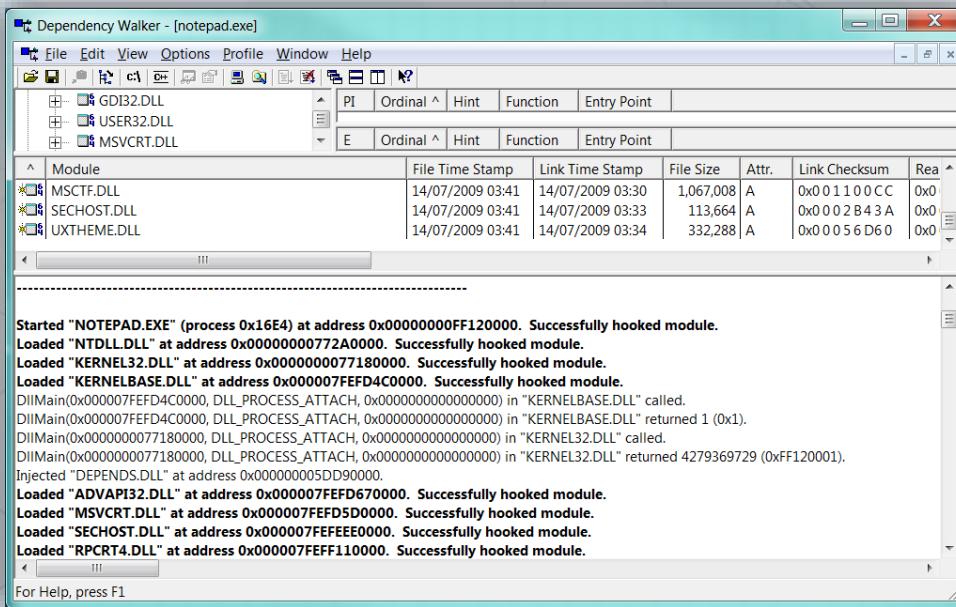
3RD PARTY TOOLS

3rd Party Tools Overview

- There are plenty of good 3rd party tools
 - Some of them cost money
 - Intel [VTune](#) profiler
 - SCI Tools – [Understand](#)
 - Parasoft [C/C++ Quality Solution](#)
 - [PVS-Studio](#)
 - .Net Tools such as JetBrains [Resharper](#) & CodeValue [OzCode](#)
 - Some of them are free
 - [NirSoft](#) tools
 - [Dependency Walker](#)
 - [MiniDumpWizard](#) & [MiniDumpView](#)

Dependency Walker

- This is an old utility that still works:



Tool Summary

- In this section we've been introduced to the Windows Developer Debugging Toolbox
- Tools come from many places, most of them are free
- Pick the set of tools that you are most comfortable with
 - Study it, know the pros and cons of each tool
- Understand the System mechanism that the tool analyzes

Basic Concepts

- Abstraction Hierarchy
- User Mode vs. Kernel Mode
- The Windows API
- Services and routines
- Processes, threads & Jobs
- Virtual Memory
- Objects and handles
- The Registry

The Computer

- No, it is not Computer Science 101
 - However there are some concepts that need dusting off
 - In order to fix the machine, we need to understand it!
- From the software point of view, the Computer is built from a collection of resources
 - The CPU is a resource responsible for executing code
 - The RAM is a resource that holds volatile information and lets the CPU execute code and consume data
 - I/O (Disk, Network) – is the (slow) resource that can provide or consume data

The CPU

- The CPU is the resource that can execute code
- Modern Computers usually have one or more CPUs that contain one or more cores
 - A core is a CPU inside the CPU chip
- The CPU can read instructions from memory (RAM) and execute them one-by-one
- The CPU fetches data from memory
- The CPU stores data into memory

The Random Access Memory

- Volatile memory stores information that came from I/O devices or from the CPU
- The CPU can work only against RAM (Physical Memory)
- The Memory is a resource, usually broken into pages
- I/O devices can work directly with the Memory
- I/O devices use interrupts to alert the CPU
- Code Instructions can also alert the CPU

Address Space, Bus & Bytes

- To fetch data from memory the CPU uses a bus
- The CPU puts the data address on the address bus
- The possible Address Space is derived from the number of lines in the address bus
 - To be more accurate, the size of the address register
- On 32 bit machines we have maximum 4GB ram
- On 64 bit machines we have maximum of 18 EB =
18,446,744,073,709,551,616
 - 1 EB = 10^{18} bytes = 1 billion gigabytes = 1 million terabytes
- Address is the location of the data
 - The data does not necessarily exist there!

Abstraction Layers

- People think and speak in abstractions
 - There are just too many tiny details
- When we say CPU, we actually mean the many components and busses and gates inside the CPU
- We say instruction, but a CPU instruction is built from many other tiny instructions that control the many components inside the CPU
- So we will continue to talk using abstractions and concepts, but...
- When a problem arises, sometime the abstraction breaks. We need to get to the real thing!

Assembly Language

- A collection of instructions that serve as the building blocks for a program for the CPU
- Most CPU Instructions are read/execute/write
 - Other instructions are: control, compute, etc.
- Each command can be executed in a certain CPU mode:
 - If the CPU reads a command that is forbidden to execute in the current CPU mode, an exception occurs
 - Exception is a software interrupt

The Operating System

"An operating system (OS) is software, consisting of programs and data, that runs on computers and manages computer hardware resources and provides common services for efficient execution of various application software."

-Wikipedia

User Mode vs. Kernel Mode

- Process access modes
- User mode
 - Allow access to non-operating system code & data only
 - No access to the hardware
 - Protects user applications from crashing the system
- Kernel mode
 - Privileged mode for use by the kernel and device drivers only
 - Allows access to all system resources
 - Can potentially crash the system

The Windows API

- Application Programming Interface for all Windows versions
- Documented in the Windows SDK (formerly Platform SDK)
- Each version implements a different subset of the API
- Now collectively called the Windows API
 - Previously referred to as the “Win32 API”
 - 64 bit windows introduced Win64
- Contains functions in the following areas
 - Base services, user interface services, component services, graphics and multimedia, messaging and collaboration, networking
- API style
 - Flat C functions
 - COM (Component Object Model)

The .NET Framework & Tools

- An higher abstraction layer
 - On top of Win32 API
- Object Oriented APIs
- Management:
 - Hosting, Lifetime, threading, memory, security, configuration, Metadata
- Support many programming language
- Call native code and APIs using .NET Interop

UWP & WinRT

- For Windows Store Application
- Windows 8/8.1 → WinRT → Windows Runtime
- Windows 10 → UWP → Universal Windows Platform
 - A superset of WinRT
- A set of unmanaged OO APIs
 - Hosting, Lifetime, threading, memory, security, configuration, Metadata
- Base on COM
 - Native & Managed language support

Processes

- Process
 - A set of resources used to execute a program
- A process consists of
 - A private virtual address space in which memory can be allocated and used
 - An executable program, referring to an image file on disk which contains the initial code and data to be executed
 - A table of handles to various objects, such as files, events, threads, and others
 - A security context, called an access token, used for security checks when accessing shared resources
 - One or more threads that execute code

Processes in Task Manager

Task Manager

File Options View

Processes Performance App history Startup Users Details Services

Name	Status	8% CPU	24% Memory	1% Disk	0% Network
Apps (12)					
Internet Explorer (12)	Running	0.2%	263.0 MB	0.1 MB/s	0 Mbp
Master MFC Application (32 bit)	Running	4.0%	21.0 MB	0 MB/s	0.8 Mbp
Microsoft Lync (2)	Running	1.8%	103.7 MB	0.1 MB/s	0.1 Mbp
Microsoft Outlook (2)	Running	0.4%	124.9 MB	0 MB/s	0 Mbp
Microsoft PowerPoint (3)	Running	0%	581.4 MB	0 MB/s	0 Mbp
Microsoft Silverlight Out-of-Bro...	Running	0%	13.4 MB	0 MB/s	0 Mbp
Skype (32 bit)	Running	0%	58.9 MB	0 MB/s	0 Mbp
Sticky Notes	Running	0%	2.9 MB	0 MB/s	0 Mbp
Task Manager	Running	0.1%	11.7 MB	0 MB/s	0 Mbp
Windows Explorer	Running	0%	12.0 MB	0 MB/s	0 Mbp
Windows Explorer (4)	Running	0%	60.8 MB	0 MB/s	0 Mbp
Windows Server Launchpad	Running	0%	46.6 MB	0 MB/s	0 Mbp
Background processes (62)					
μTorrent (32 bit)	Running	0.1%	27.7 MB	0.1 MB/s	0.1 Mbp
Adobe Acrobat Update Service ...	Running	0%	0.7 MB	0 MB/s	0 Mbp
Adobe® Flash® Player Utility	Running	0%	2.1 MB	0 MB/s	0 Mbp
AMD External Events Client Mo...	Running	0%	1.2 MB	0 MB/s	0 Mbp

[Fewer details](#) [End task](#)

Task Manager

File Options View

Processes Performance App history Startup Users Details Services

Name	PID	Status	User name	CPU	Memory (p...)	Description
ApacheMonitor.exe	7348	Running	alon	00	720 K	Apache HTTP Server Monitor
armsvc.exe	1632	Running	SYSTEM	00	672 K	Adobe Acrobat Update Service
aticebx.exe	1132	Running	SYSTEM	00	1,220 K	AMD External Events Client Module
aticesnx.exe	516	Running	SYSTEM	00	592 K	AMD External Events Service Module
audiogd.exe	6224	Running	LOCAL SE...	00	5,272 K	Windows Audio Device Graph Isolation
c2c_service.exe	3768	Running	SYSTEM	00	1,380 K	Skype C2C Service
CCC.exe	4648	Running	alon	00	4,480 K	Catalyst Control Center Host application
CLHNServiceForPowl...	8904	Running	SYSTEM	00	664 K	CLHNServiceForPowerDVD12 Module
ClientOperator.exe	6896	Running	alon	00	12,540 K	Windows Server Client File Backup Operator
CLMSMonitorService...	17564	Running	SYSTEM	00	624 K	CyberLink Media Server Monitor Service
CLMSMonitorServic...	1908	Running	SYSTEM	00	620 K	Media Server Monitor Service
CLMServerPDVD12...	1820	Running	SYSTEM	00	3,332 K	CyberLink Media Server Service
CLMServerPDVD13...	1952	Running	SYSTEM	00	1,580 K	Media Server Service
CSISYN-1.EXE	7864	Running	alon	00	9,048 K	Microsoft Office Document Cache Sync Client
crssrs.exe	596	Running	SYSTEM	00	1,524 K	Client Server Runtime Process
crssrs.exe	696	Running	SYSTEM	00	1,808 K	Client Server Runtime Process
dasHost.exe	884	Running	LOCAL SE...	00	6,756 K	Device Association Framework Provider Host
DivXUpdate.exe	1600	Running	alon	00	3,428 K	DivX Update
dwm.exe	684	Running	DWM-1	00	23,936 K	Desktop Window Manager
explorer.exe	1184	Running	alon	00	62,268 K	Windows Explorer
explorer.exe	14732	Running	alon	00	12,100 K	Windows Explorer
FABS.exe	8856	Running	SYSTEM	00	992 K	Verzeichnisüberwachung und Hilfsaufgaben f...
FlashUtil_ActiveX.exe	10948	Running	alon	00	2,112 K	Adobe® Flash® Player Utility
GROOVE.EXE	6708	Running	alon	00	36,824 K	Microsoft SkyDrive Pro
HD-LogRotatorServic...	1704	Running	SYSTEM	00	1,808 K	BlueStacks Log Rotator Service
HydraDM.exe	7476	Running	alon	00	916 K	HydraDM
HydraDM64.exe	7488	Running	alon	00	752 K	HydraDM64
iexplore.exe	10024	Running	alon	00	19,940 K	Internet Explorer
iexplore.exe	8716	Running	alon	00	124,132 K	Internet Explorer

[Fewer details](#) [End task](#)

Threads

- Thread
 - Entity that is scheduled by the kernel to executes code
- A thread contains
 - The state of CPU registers
 - Current access mode (user mode or kernel mode)
 - Two stacks, one in user mode and one in kernel mode
 - A private storage area, called Thread Local Storage (TLS)
 - Optional security token
 - Optional message queue and Windows the thread creates
 - A priority, used in thread scheduling
 - A state: running, ready, waiting

Jobs

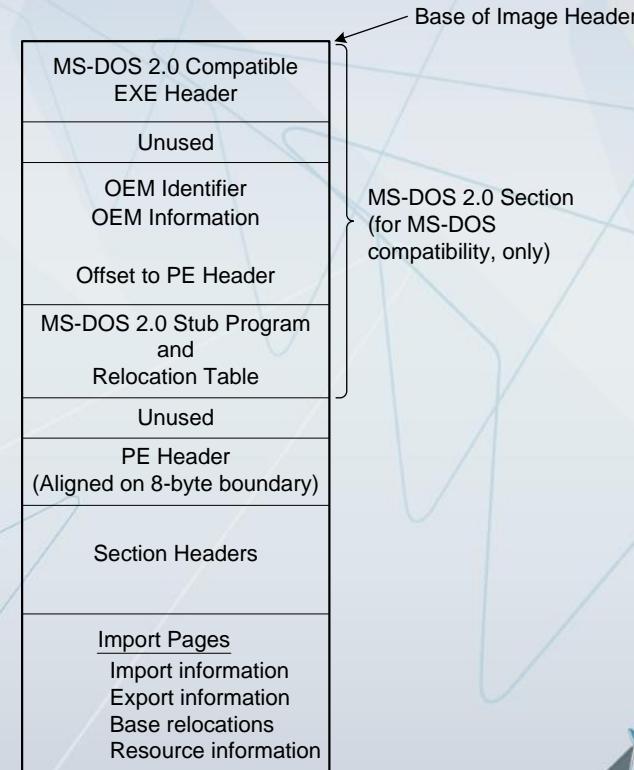
- A job object allows groups of processes to be managed as a unit
- Job Limits:
 - Memory, execution time, number of active processes, CPU affinity, priority, UI restriction, and more
- Job Notifications (process has exited)
- Resource Accounting (I/O counters)
- Nested Jobs (Windows 8/8.1/10)

The Windows Loader

- The Windows loader routines reside in ntdll.dll
 - In WinDbg type: `x ntdll!Ldr*`
- Whenever there is a need to map an exe or Dll file to the process, the loader does it
 - The public API to handle images from code are in the [ImageHlp.dll](#)
- In order to load a Dll or Executable into runnable program we need a protocol
 - A way to extract information from a file and create from it a live program

The Windows PE Format

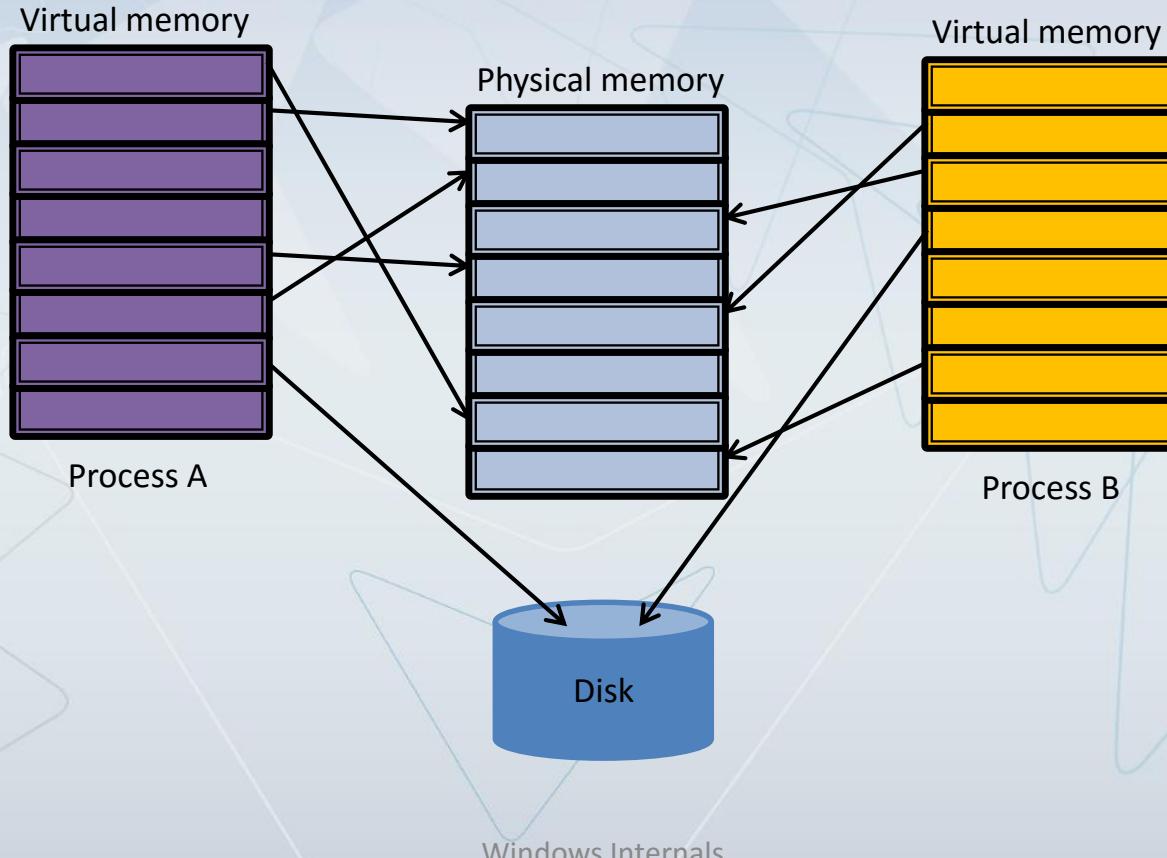
- The [Portable Executable](#) is the format for all DLL, EXE, SYS, OCX, and even PDB files
- The PE itself is a container for other format named [Common Object File Format](#) (originated in the Unix O/S)
- The data is saved in headers followed by sections
 - Headers contain metadata
 - Sections contain data
- The ImageHlp.dll provides APIs to handle and manipulate PE files



Virtual Memory

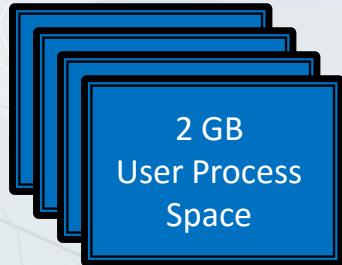
- Each process “sees” a flat linear memory
- Internally, virtual memory may be mapped to physical memory, but may also (temporarily) stored on disk
- Processes access memory regardless of where it actually exists
 - The memory manager handles mapping of virtual to physical pages
 - Processes cannot (and need not) know the actual physical address of a given address in virtual memory

Virtual Memory Mapping

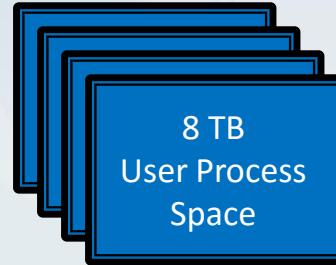


Virtual Memory Layout

x86 (32 bit)



x64 <=Windows 7 (44 bit)

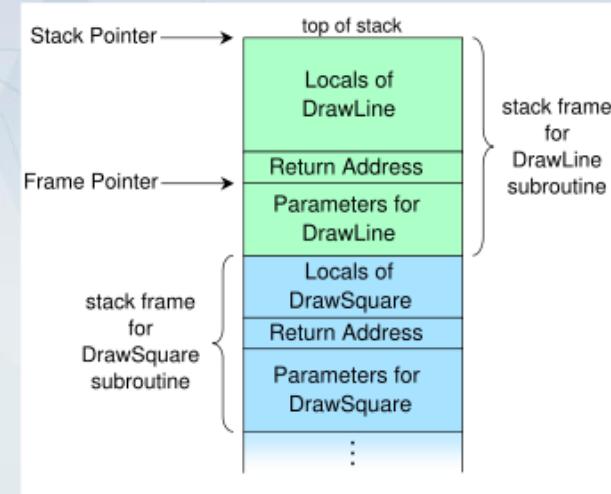


X64 >=Windows 8.1 (48 bit)



The (Call) Stack

- The stack is used to:
 - Store local variables
 - Handle arguments passing to a function
 - Keeping the return address
 - Housekeeping for exception handling (x86)
- Each thread has a stack with predefined size of 1 MB
 - This is the reserved size of the stack
 - The default size for the reserved and initially committed is specified in the PE header (linker: /STACK:reserve,commit)
 - CreateThread can allocate bigger stack size



The Call Stack – Deep Dive

- The stack expands and shrinks during the thread execution
 - The stack is built according to the function calling convention and the exception handling mechanism
 - There are 6 common calling conventions:
 - `__stdcall`
 - `__cdecl`
 - `this call`
 - `__fastcall`
 - `__naked`
 - X64 calling convention

Example of x86 __cdecl

```
Void Foo(int a, int b, int c)
{
    int x,y,z;
}
```



16(+ ebp)	third function argument (c)
12 + ebp)	second function argument (b)
8 + ebp)	first function argument (a)
4 + ebp)	old EIP (the function's "return address")
0(+ ebp)	old EBP (previous function's base pointer)
-4(+ ebp)	first local variable (x)
-8(+ ebp)	second local variable (y)
-12(+ ebp)	third local variable (z)

The x64 Calling Convention

Passing parameters

- All arguments are right justified in registers
- All stack parameters are 8 bytes aligned
- Any parameter that's not 1, 2, 4, or 8 bytes (including structs) is passed by reference
- structs and unions of 8, 16, 32, or 64-bits are passed as if they were integers of the same size
- The first 4 integer parameters are passed (in left to right order) in rcx, rdx, r8, r9
- Further integer parameters are passed on the stack by pushing them in right to left order (parameters to the left at lower addresses)
- The *this* pointer is passed in rcx

The x64 Calling Convention

Passing parameters

- All arguments are right justified in registers
- All stack parameters are 8 byte aligned
- Any parameter that's not 1, 2, 4, or 8 bytes (including structs) is passed by reference
- Structs and unions of 8, 16, 32, or 64-bits are passed as if they were integers of the same size
- The first 4 integer parameters are passed (in left to right order) in rcx, rdx, r8, r9
- Further integer parameters are passed on the stack by pushing them in right to left order (parameters to the left at lower addresses)
- The *this* pointer is passed in rcx

The x64 Calling Convention

Passing parameters

- The first floating point parameters are passed (left to right) in xmm0 to xmm3
- Further floating point parameters are passed on the stack, right to left order (parameters to the left at lower addresses)
- Return value:
 - On rax, if the return value is an integer or a pointer
 - On xmm0 if it is a floating point value
- Function must preserve: rbx, rbp, rdi, rsi, r12, r13, r14, r15, xmm6 - xmm15 and the x87 register stack
- Function may destroy: rax, rcx, rdx, r8, r9, r10, and r11 and xmm0 - xmm5

The x64 Calling Convention - Stack Handling

- The caller must reserve 32 bytes (4 64-bit values) on the stack
 - This space allows the rcx, rdx, r8, r9 registers to be easily copied to a well-known stack location if there is a need
- The caller is responsible for cleaning up the stack
 - Usually the compiler reserves enough stack space for the function that requires the most stack space and just adjust positioning within that stack space to fit all functions that it is calling
- A function can be a leaf or a frame function
 - Leaf functions don't need to support the stack unwinding process that is part of [exception handling](#) but it is limited (can't call other functions)
 - Frame functions must handle many tasks such as building a stack frame, and save some register values
- For more info look at: <http://software.intel.com/en-us/articles/introduction-to-x64-assembly/>

The Heap

- A heap is an abstraction for managing memory allocations
- In Windows we have several heaps:
 - The process global heap, default size is 1MB
 - you can change it using the linker flag `/Heap(reserve,commit)`
 - The process' custom heaps created by the `HeapCreate` API
 - The C++ Runtime heap
 - The kernel page and non-page pools

Objects and Handles

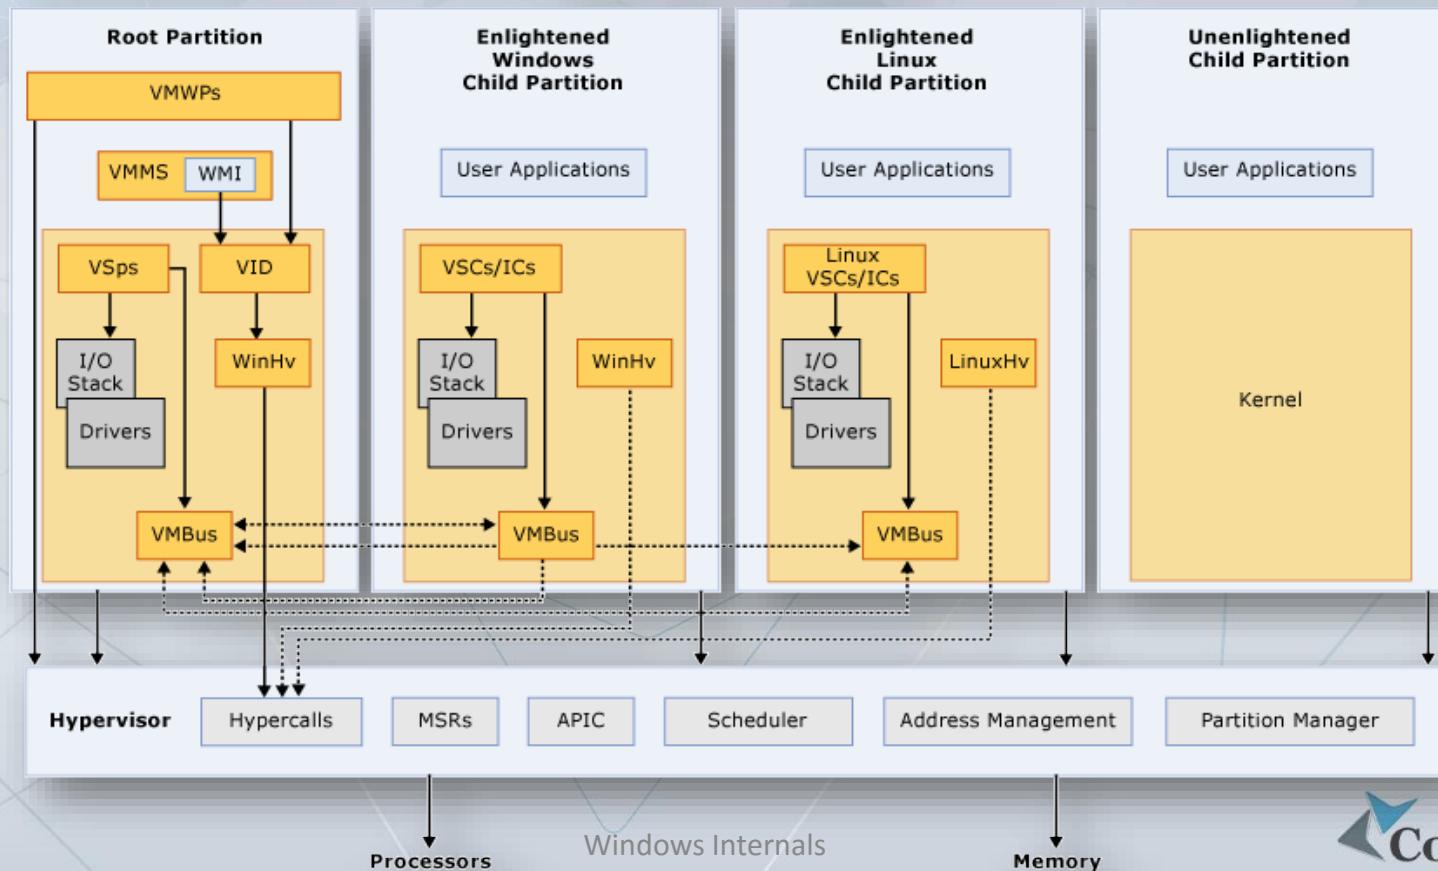
- Objects are runtime instances of static structures (object type)
 - Examples: process, mutex, event, desktop, file
- Reside in system memory space
- Kernel code can obtain direct pointer to an object
- User mode code can only obtain a handle to an object
 - Shields user code from directly accessing an object
- Objects are reference counted
- The Object Manager is the entity responsible for creating, obtaining and otherwise manipulating objects

The Registry

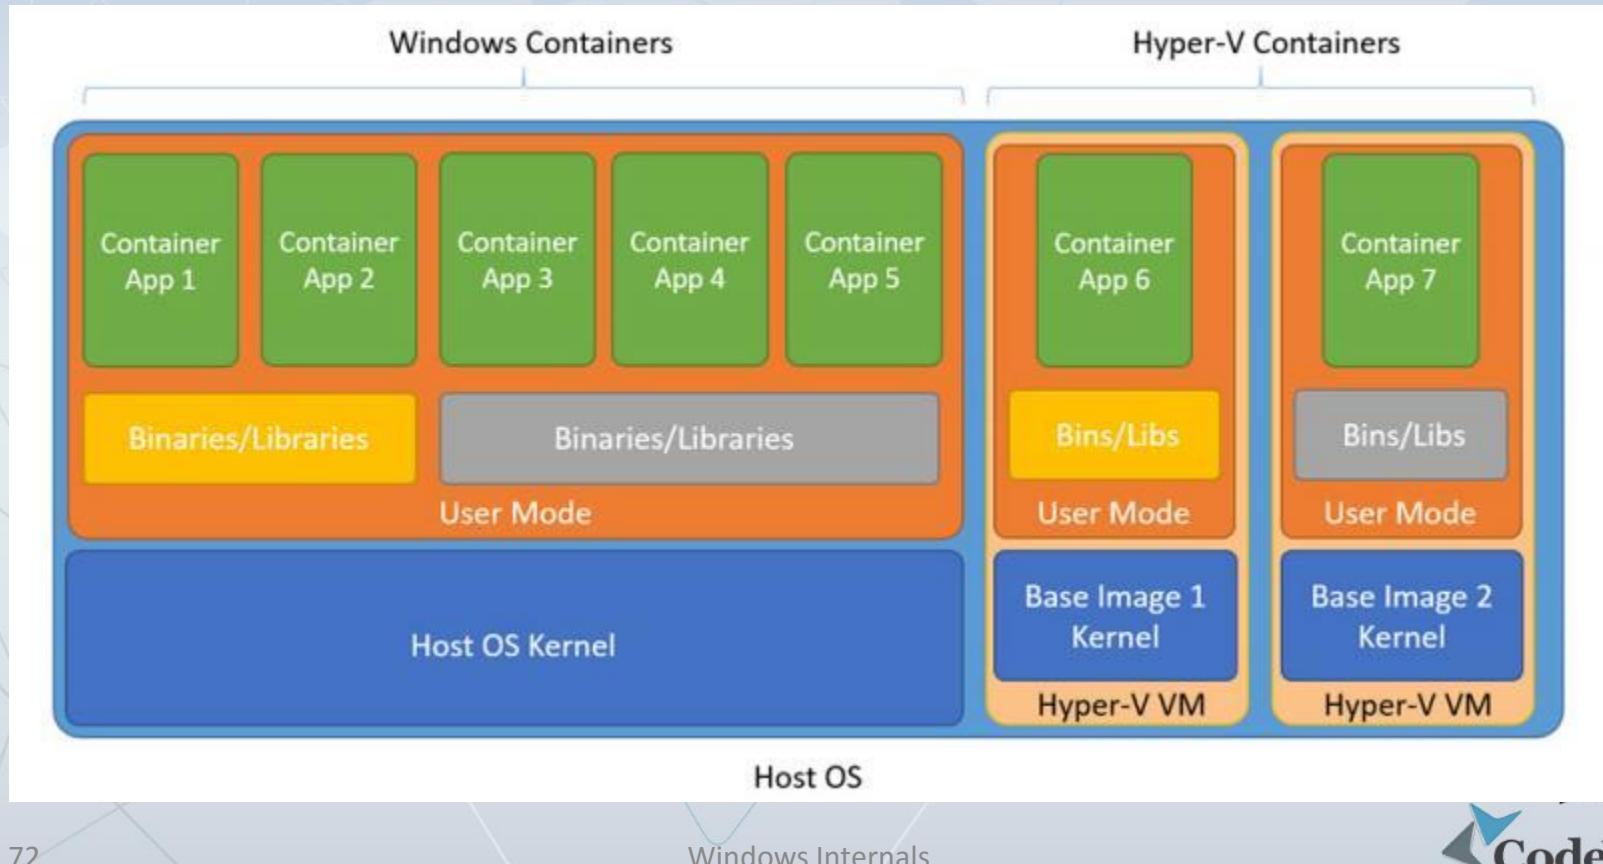
- Global hierarchical repository of data
- Machine wide data as well as user specific data
- Persistent as well as volatile data
- Also “contains” performance data
 - Not really in the registry
 - But the registry API is used to query this data
- Main hives
 - HKEY_LOCAL_MACHINE – machine wide settings
 - HKEY_CURRENT_USER – user specific settings

Virtualization (Hyper-V)

Hyper-V High Level Architecture



Containers (Windows 10/Server 2016)



STEVEN SPIELBERG PRESENTS

BACK TO THE FUTURE

PG

A ROBERT ZEMECKIS FILM



History

- Design started in 1989 by [David Cutler](#)
- First released as [Windows NT](#) 3.1 at 1993
- Many SKUs, Many flavors
 - From the smallest IoT device to the larger Surface Hub screen and the strong server

Windows NT Design Goals

- Separate address space per process
 - One process cannot (easily) corrupt another's memory
- Protected kernel
 - User mode applications cannot crash kernel
- Preemptive multitasking and multithreading
- Multiprocessing support
- Internationalization support using Unicode
- Security throughout the system
- Integrated networking

Windows NT Design Goals

- Powerful file system (NTFS)
 - Supports protection, compression and encryption
- Run most 16 bit Windows and DOS apps
 - On 32 bit systems (NTVDM is not supported in 64-bit versions of Windows)
- Run POSIX 1003.1 and OS/2 applications
- Portable across processors and platforms
- Be a great client as well as server platform

Windows Editions

- Windows XP Home
 - Designed as a replacement for the Windows 9x/ME family (“Consumer Windows”)
- Windows Professional (2000, XP), Vista, 7, 8, 8.1
 - Main desktop (client) OS
- Windows Server Standard, Advanced, Datacenter editions (Windows 2000, 2003, 2008, 2008 R2, 2012, 2012 R2)
 - Server platforms
- Other variants
 - XP starter, XP Home, Media center, 2003 & 2008 Server Web Edition, Vista / 7 Home / Premium, Ultimate, Business, Enterprise
- New Variants (Windows 10):
 - Xbox One, HoloLens, Surface Hub, Windows IoT

Windows Numeric Versions

- Windows NT 4 (NT 4.0)
- Windows 2000 (NT 5.0)
- Windows XP (NT 5.1)
- Windows Server 2003, 2003 R2, XP 64 (NT 5.2)
- Windows Vista, Server 2008 (NT 6.0)
- Windows 7, Server 2008 R2 (NT 6.1)
- Windows 8, Server 2012, Windows Phone 8 (NT 6.2)
- Windows 8.1, Server 2012 R2, Windows Phone 8.1 (NT 6.3)
- Windows 10, Server 2016(NT 10.0)
- These values are obtained using GetVersionEx

Professional vs. Server

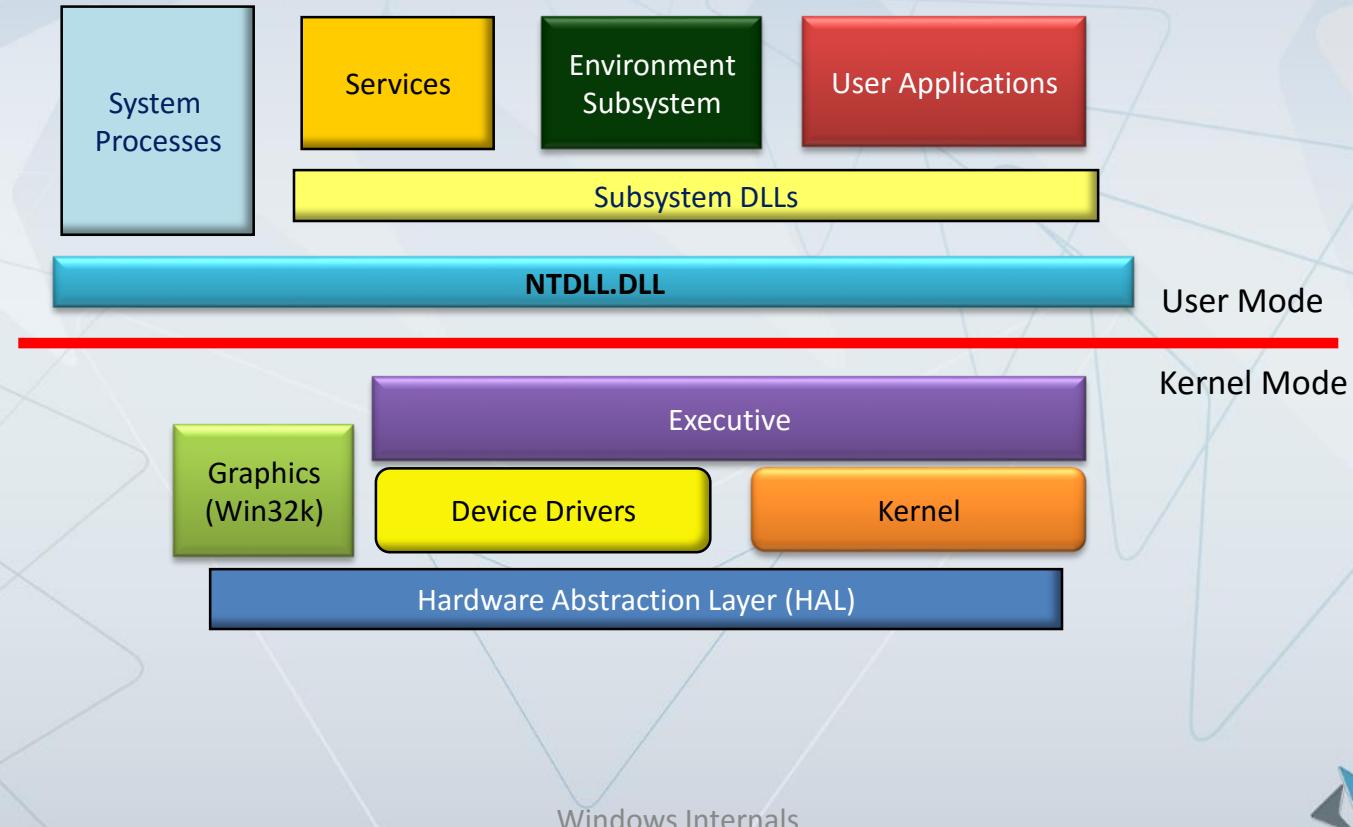
- Same core system files
- Differences
 - Number of processors supported
 - Maximum amount of RAM than can be used
 - Maximum of concurrent network connections supported for file and print sharing (10 on professional, 20 for Windows 7)
 - Some services only appear in Server versions
 - Other system policies and default settings (e.g. thread quantum)
- To query the OS type, use GetVersionEx (Win32) or RtlGetVersion(WDK)



Product Data Type

- The actual product type is stored in the registry
 - HKLM\SYSTEM\CCS\Control\ProductOptions
- Values
 - **ProductType**
 - **WinNT** – 2000 Professional, XP Home, XP professional, Vista, 8
 - **ServerNT** – Server (domain controller)
 - **LanmanNT** – Server (not a domain controller)
 - **ProductSuite**
 - Distinguishes the various server types, and between XP Home and Professional editions
- Memory limits for the various editions can be found at
[http://msdn.microsoft.com/en-us/library/aa366778\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/aa366778(v=vs.85).aspx)

General Architecture Overview



Kernel Mode Components

- Hardware abstraction Layer (HAL)
 - Isolates the kernel and device drivers from platform specific issues
- Kernel
 - Thread scheduling, interrupt & exception dispatching, multiprocessor support, synchronization primitives
- Device Drivers
 - Loadable kernel modules that handle I/O requests for hardware devices and buses
- Executive
 - Virtual memory manager, object manager, security, IPC, plug & play, power manager, configuration manager
- Win32K.SYS
 - The Windows subsystem kernel component
 - Handles user interface and graphics

User Mode Components

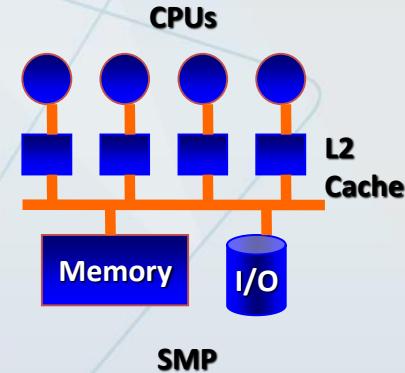
- User applications
 - Executables under one of 3 subsystems: Win32 (also Win 16, DOS), POSIX, OS/2
- System processes
 - Logon, Session manager, Service Control Manager
- Services
 - Normal Win32 processes that also interact with the SCM
- Subsystem process
 - A single process per subsystem (per session) handling subsystem specific issues
- Subsystem DLLs
 - Libraries implementing the API for a subsystem
 - Win32: kernel32.dll, user32.dll, gdi32.dll, advapi32.dll
- NTDLL.DLL
 - Implements the (undocumented) native API that subsystem DLLs use

Core System Files

- **Ntoskrnl.exe**
 - Executive and kernel
 - Original is NtOsKrnl.Exe (single CPU) or NtKrnlMp.Exe (multi CPU)
- **NtKrn1Pa.exe**
 - Executive and kernel (32 bit) with support for Physical Address Extension (PAE)
 - Original is NtKrn1Pa.Exe (single CPU) or NtKrPaMp.Exe (multi CPU)
- **Hal.dll**
 - Hardware Abstraction Layer
- **Win32k.sys**
 - Kernel component of the Win32/Win64 subsystem
- **NtD11.dll**
 - System support routines and Native API dispatcher to executive services
- **Kernel32.dll, user32.dll, gdi32.dll, advapi32.dll**
 - Core Windows subsystem DLLs
- **CSRSS.exe (Client Server Runtime SubSystem)**
 - The Win32/Win64 subsystem process

Symmetric Multiprocessing

- SMP
 - All CPUs are the same and share main memory and have equal access to peripheral devices (no master/slave)
- Basic architecture supports up to 32/64 CPUs
 - Windows 7 64 bit & 2008 R2 support up to 256 cores
 - Windows 8 / 2012 supports up to 640 cores
- Actual number of CPUs determined by licensing and product type
 - Prior NT 6.0: Number of licensed CPUs
HKLM\System\CCS\Control\Session Manager
 - LicensedProcessors (DWORD)
 - NT 6.0 and above ([Licensed Processors](#)):
 - HKLM\System\CCS\Control\ProductOptions\ProductPolicy



NUMA

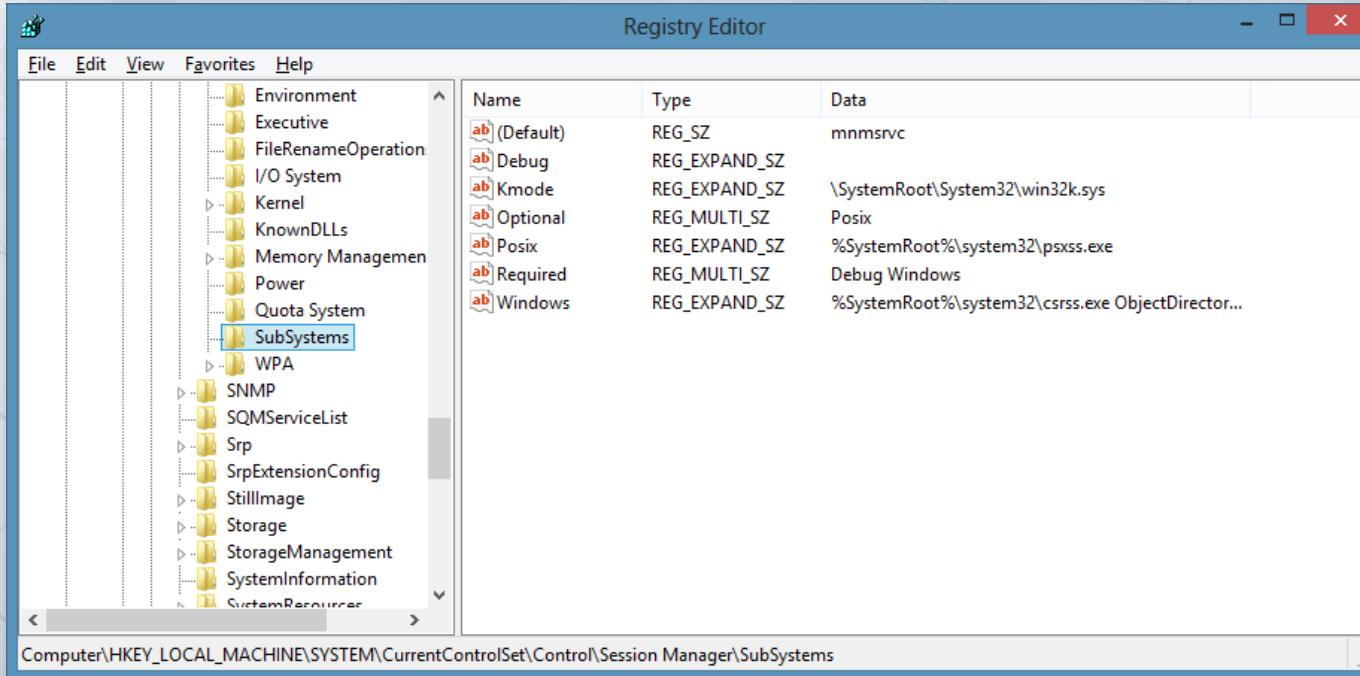
- NUMA (Non Uniform Memory Architecture) systems
 - Groups of physical processors (“nodes”) that have local memory
 - Connected to the larger system through a cache-coherent interconnect bus
 - Still an SMP system (e.g. any processor can access all of memory)
 - But node-local memory is faster
- Scheduling algorithms take this into account
 - Try to schedule threads on processors within the same node
 - Try to allocate memory from local memory for processes with threads on the node
- New Windows APIs to allow applications to optimize

Environment Subsystems

- A subsystem is a special view of the OS
 - Exposes services via subsystem DLLs
- Windows 2000 ships with Win32, OS/2 and POSIX 1003.1 (POSIX-1)
- Windows XP and later have Windows and a basic POSIX
 - An enhanced POSIX version is available with the “Services for UNIX” product
- Windows 10 has a new Ubuntu based capabilities
 - For providing Bash and other tools for developers
 - It does not implemented as a sub-system
- The Windows subsystem must always be running
 - Owner of keyboard, mouse and display
- Other subsystems configured to load on demand

Subsystems in the Registry

- Windows 8.1 snapshot



Subsystem DLLs

- Every image belongs to exactly one subsystem
 - Value stored in image PE header
 - Can view with the exetype.exe utility (originally from Win2K resource kit), can be found as a [GitHub](#) project
 - the dependency walker (depends.exe) or dump bin utility (dumpbin.exe)
 - Allows the Windows Loader to make correct decisions
- An image of a certain subsystem calls API functions exposed through the subsystem DLLs
 - E.g. kernel32.dll, user32.dll, etc. for the Windows subsystem
- Some images belong to no subsystem
 - “Native” images (Which API functions do they call?)

The Native API

- Implemented by NTDLL.DLL
 - Used by subsystem DLLs and “native” images
 - Undocumented interface
 - Lowest layer of user mode code
- Contains
 - Various support functions
 - Dispatcher to kernel services
 - Most of them accessible using Windows API “wrappers”

NTDLL.DLL Kernel Gate

- 32 bit dispatching code (XP & Pentium II and up)

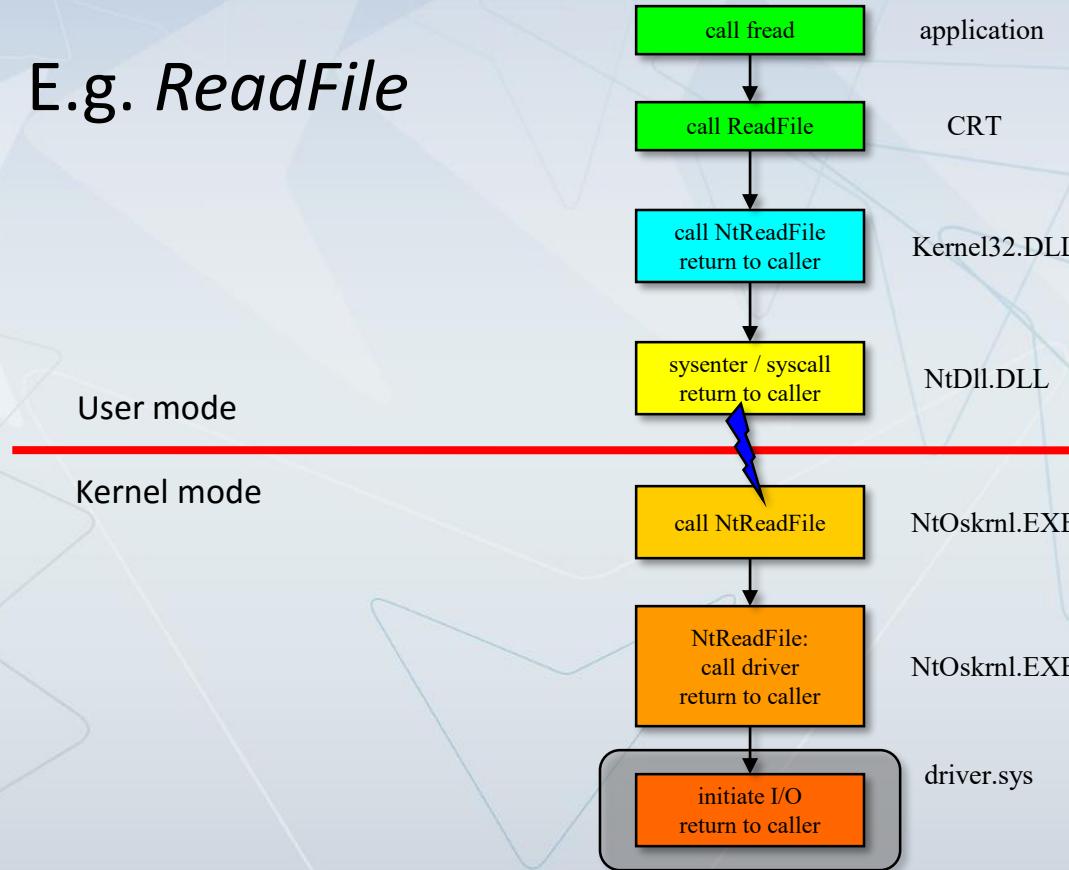
```
; ntdll!NtReadFile  
mov eax,0xb7  
mov edx, 0x7ffe0300  
call edx  
ret 0x24
```

```
; SharedUserData!SystemCallStub  
mov edx, esp  
syseenter  
ret
```

- 32 bit (XP Pentium I & Windows 2000)
 - Uses `int 0x2e` instead of `syseenter`
- 64 bit dispatching code is similar (using the `syscall` instruction)

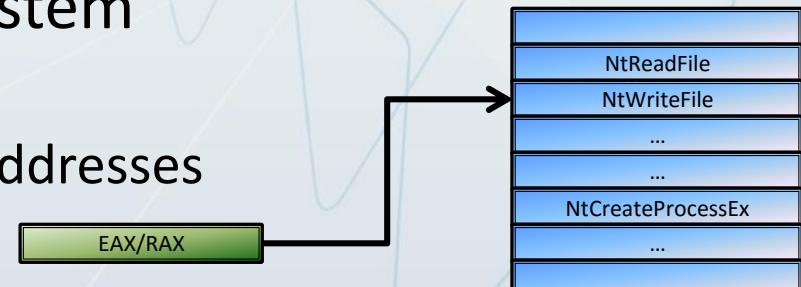
Function Call Flow

- E.g. *ReadFile*



System Service Table(s)

- 2 built-in, up to 4 supported
 - Bits 12-13 select the table
 - Lower 12 bits select the service
- Each thread has a pointer to its system service table
- EAX serves as an index to the system service required
 - 32 bit Windows holds the actual addresses
 - 64 bit Windows holds offsets
 - Lower 4 bits used as argument count
 - Must be masked off to get actual offset



Default System Tables

- [KeServiceDescriptorTable](#)
 - Contains only kernel related entries
- KeServiceDescriptorTableShadow
 - Contains both kernel related entries and USER and GDI entries
- When a thread is first created, its table pointer is set to KeServiceDescriptorTable
- The first time it makes any GDI or USER call, its table pointer is changed

Executive

- The upper layer of NtOsKrn1.Exe
- Functions
 - System service functions callable from user mode
 - Accessed through NtDll.dll
 - Most can be called from the Windows API
 - Callable by device drivers and documented in the WDK
 - Exported but undocumented for use by NtOsKrn1.Exe
 - Internal functions

Executive Components

- Configuration manager
 - Responsible for managing the system registry
- Object manager
 - Manages objects created by kernel or user code
- Process and thread manager
 - Creates and terminates processes and threads
 - The underlying support is provided by the lower-layer Kernel
- Security Reference Monitor (SRM)
 - Implements local security enforcement, including object protection and auditing
- I/O manager
 - Implements device independent I/O and dispatches calls to appropriate device drivers

Executive Components

- Plug & Play (PnP) manager
 - Handles hardware identification, enumeration, resource allocation and loading of appropriate drivers
- Power manager
 - Coordinates power events and generates I/O requests to drivers based on state changes
- Cache manager
 - Manages caching for file based I/O
- Memory manager
 - Implements virtual memory handling for use by other system components



Kernel

- The lower layer of the Kernel
- Implements all low level activity, such as interrupt dispatching, thread scheduling and processor synchronization
- Implements a set of kernel objects
 - Control objects
 - Controlling various OS functions, e.g. APC, DPC
 - Dispatcher objects
 - Have synchronization capabilities, e.g. mutex, event
 - Executive adds management over these primitive objects, e.g. security, handle management and reference counting

Hardware Abstraction Layer

- Isolates the kernel and device drivers from the underlying hardware platform
 - E.g. interrupt controller
- Allows easy access to hardware by device drivers
- Various HALs are on the Windows install media
 - The appropriate one is copied to the **System32** directory and is always named **Hal.Dll**

Device Drivers

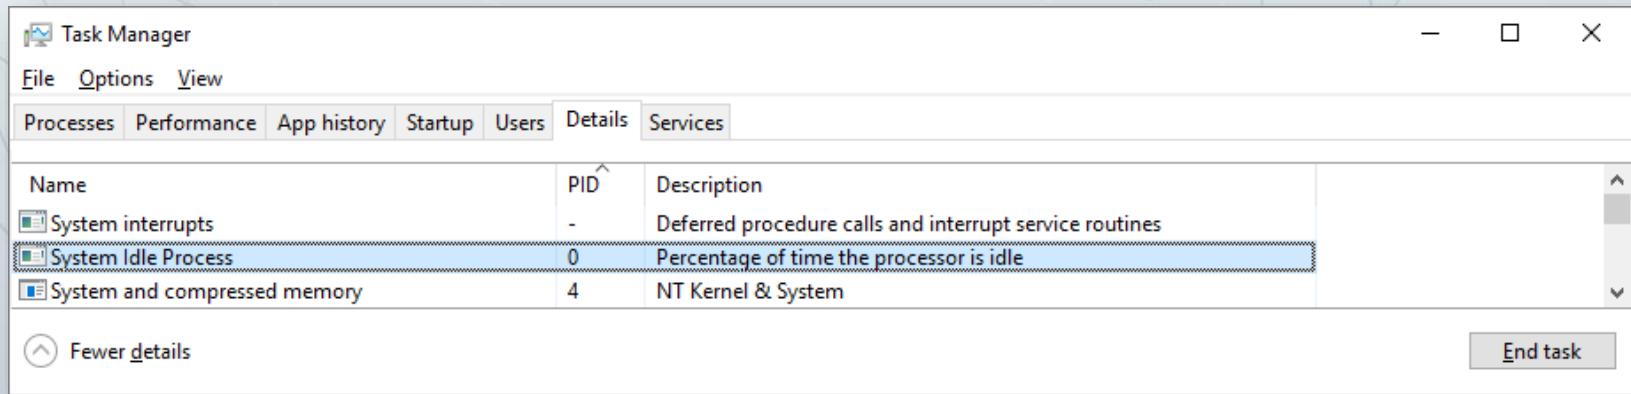
- Loadable kernel modules (usually with SYS extension)
- The only documented way to add code that runs in the kernel
- Types (partial list)
 - Hardware device drivers – manage a physical device
 - File system device drivers – translate file I/O to device specific requests
 - Protocol drivers – implement a specific network protocol (TCP/IP, IPX/SPX, etc.)
 - User programmed driver, simply for the ability to execute code in kernel mode

System Processes

- Idle process
- System process
- Session Manager (Smss.Exe)
- Windows subsystem (Csrss.Exe)
- Logon process (Winlogon.Exe)
- Service control manager (SCM) (Services.Exe)
- Local security authentication server (Lsass.Exe)

Idle Process

- Always has a PID of 0
- Not a real process (does not run any executable image)
- One thread per CPU (core)
- Accounts for idle time



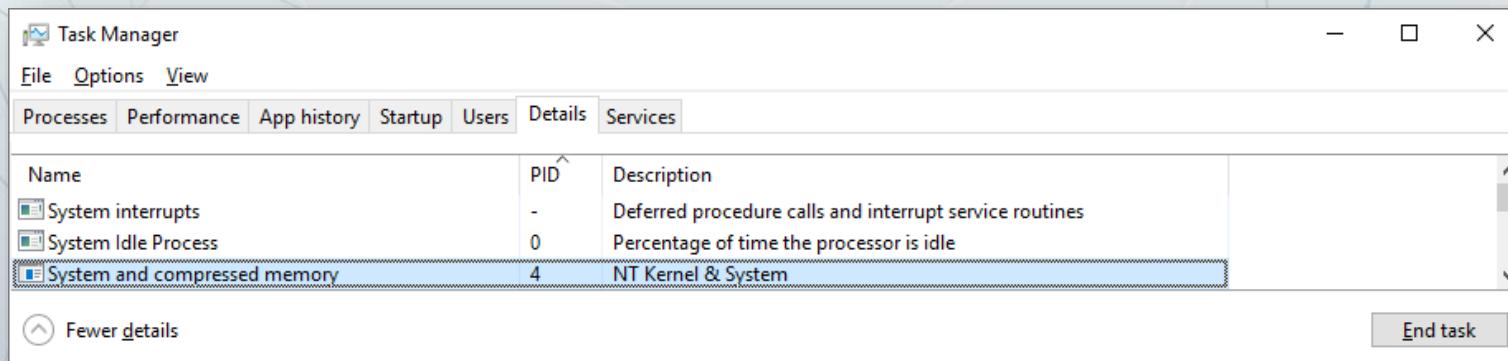
The screenshot shows the Windows Task Manager window. The title bar says "Task Manager". The menu bar includes "File", "Options", and "View". Below the menu is a tab bar with "Processes", "Performance", "App history", "Startup", "Users", "Details", and "Services". The "Details" tab is selected. A table follows, with columns "Name", "PID", and "Description". The table contains the following data:

Name	PID	Description
System interrupts	-	Deferred procedure calls and interrupt service routines
System Idle Process	0	Percentage of time the processor is idle
System and compressed memory	4	NT Kernel & System

At the bottom left is a link "Fewer details". At the bottom right is a button "End task".

System Process

- Has a fixed PID (8 on Win2K, 4 on XP and later)
- Home of system threads
 - Threads created by the kernel and device drivers
 - Execute code in system space only
 - Created using the [PsCreateSystemThread](#) kernel API (documented in the WDK)
 - Allocate memory from the system pools
 - Windows 10/Server 2016 – Memory size includes processes' compressed pages



Session Manager

- Running the image `\windows\system32\smss.exe`
- The first user mode process created by the system
- Tasks
 - Create additional page files
 - Creating system environment variables
 - Launches the subsystem processes (normally just `csrss.exe`)
 - Launches itself in other sessions
 - That instance loads `WINLOGON` and `CSRSS` in that session
 - Then terminates
- Finally
 - waits forever for `csrss.exe` instances to terminate
 - If any of them dies, crashes the system
 - Waits for subsystem creation requests
 - Waits for terminal services session creation requests

Winlogon

- Running the image `\windows\system32\winlogon.exe`
- Handles interactive logons and logoffs
- If terminated, logs off the user session
- Notified of a user request by the Secure Attention Sequence (SAS), typically `Ctrl+Alt+Del`
- Authenticates the user by presenting a username / password dialog
 - Can be replaced
- Sends captured username and password to LSASS
 - If successfully authenticated, initiates the user's session

LSASS

- Running the image `\windows\system32\Lsass.exe`
- Calls the appropriate authentication package
- Upon successful authentication, creates a token representing the user's security profile
- Returns information to Winlogon

Service Control Manager (SCM)

- Running the image `\windows\system32\services.exe`
- Responsible for starting, stopping and interacting with service processes
- Services
 - Similar to UNIX “daemon processes”
 - Normal Windows executables, that interact with the SCM
 - Can be started automatically when the system starts up without an interactive logon
 - Can run under “special” accounts
 - LocalSystem, NetworkService, LocalService

Windows Architecture Diagram

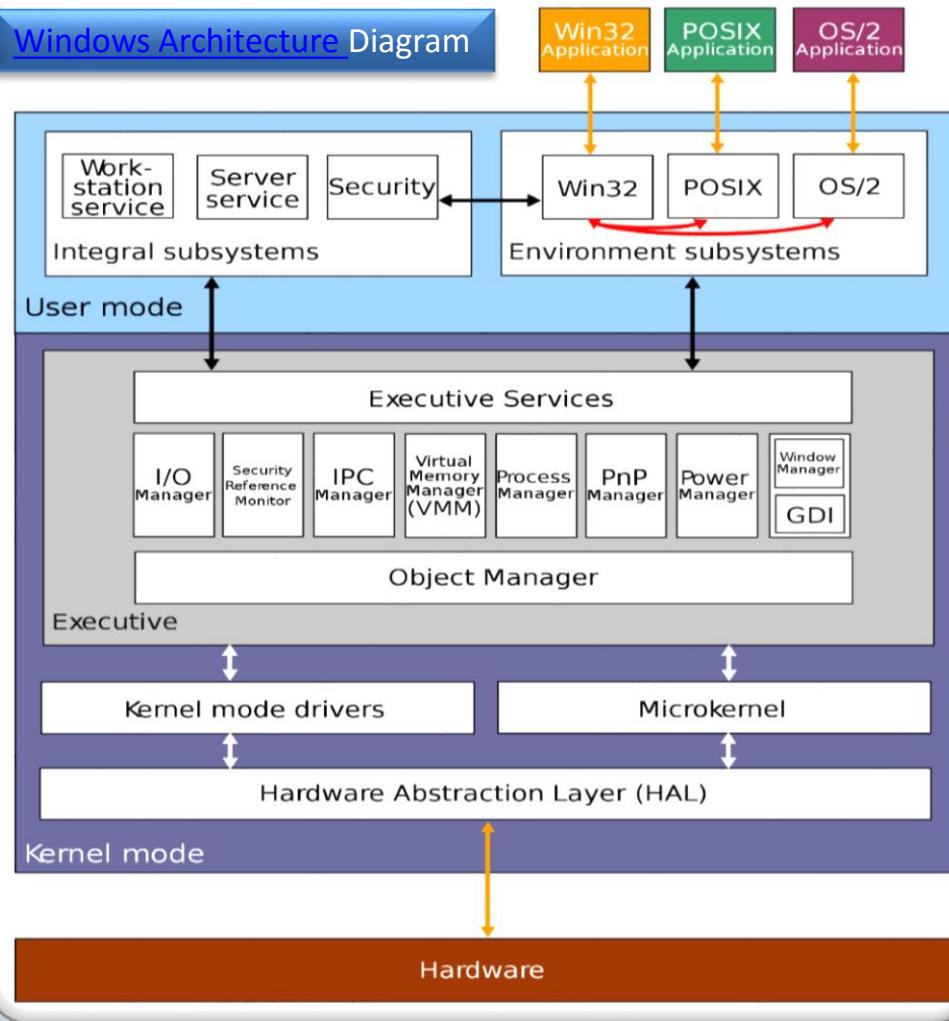


Image is taken
from [Wikipedia](#)

Summary

- User mode code has limited access to the system
- Kernel mode has full and total access to all system resources
- Applications run under a specific subsystem
- The primary subsystem is the Windows subsystem

Windows NT 6.x Versions

- Windows 6.0
 - Windows Vista
 - Windows Server 2008
 - Windows Home Server 2011
- Windows 6.1
 - Windows 7
 - Windows Server 2008 R2
 - Windows Thin PC
- Windows 6.2
 - Windows 8
- Windows 6.3
 - Windows 8.1
 - Windows Server 2012 R2

What's New in 6.0

- Lots of kernel improvements
 - In the memory manager, process scheduler and I/O scheduler
 - Many new security features such as [BitLocker](#), [ASLR](#), and an improved Windows Firewall
- Windows Aero desktop theme
- [DWM](#), [WDDM](#), DirectX 10, IE 7, [IIS 7](#)
- [UAC](#), [ReadyBoost](#), [SuperFetch](#), [IPv6](#)
- [Kernel Transaction](#), Transactional NTFS & Registry
- Restart Manager, I/O Priority, ETW
- Server Core, Self-healing NTFS
- And many [more...](#)

What's New in 6.1

- Many Kernel Changes
 - 256 Cores, Core & Socket Parking, Hot Locks Removal, Better Windows Memory Utilization, Faster Registry Access
- Power Consumption Improvements
 - Trigger Start Services, Timers, Power Consumption Profiling
- Client Side Enhancements
 - Taskbar, Multi-Touch, Ribbon, Shell Libraries, Federated Search, Direct2D & DirectWrite
- Tons of features
 - Compatibility, troubleshooting platform, Mount and boot from VHD, Sensor & Location

Power Consumption Profiling
Powercfg /energy

DEMO

What's New in 6.2

- Windows 8 and Windows Server 2012
- Better Tablet Support, ability to run on [ARM](#), Support 640 cores
- [WinRT](#) – A new Windows programming model
 - Object Oriented Class Library
 - Ref-Count lifetime management
 - .NET like Metadata Support
 - Natively supports C++, .NET and JavaScript
- [Windows Store App \(Metro\)](#) – A new UI approach for Windows applications
 - Chrome-less UI Tile based that supports touch and mouse
 - App development with XAML/C++|.NET or HTML5 + JavaScript
 - Handles application start, switch, suspend and shutdown
 - Support application inter-communication through contracts
 - App distribution & installation through Microsoft App Store

Windows Store Application

DEMO

What's New in 6.3

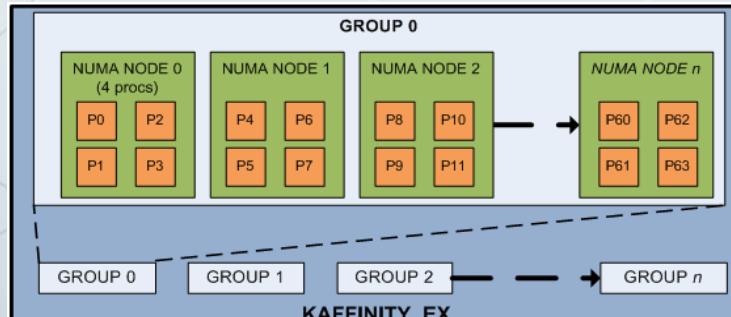
- Windows 8.1, Server 2012 R2
 - Internet Explorer 11 ([WebGL](#), [SPDY](#))
 - PowerShell 4.0
 - New commands for managing the Start screen, Windows Defender & Windows components
 - Start Screen Enhancement (Tile Sizes, All Apps, Search, Desktop Background)
 - Boot to desktop & Visible Desktop Start Button
 - New technology support
 - [NFC](#) printing, [Wi-Fi Direct](#) printing, [Miracast](#), [Mobile broadband tethering](#), Auto-triggered [VPN](#)
 - Many new (6000) XAML and WinJS APIs
 - [DirectX](#) 11.2
 - Better Windows Store Application Multitasking

NT 6.x Kernel Changes

- Windows 7 (x64) and Server 2008 R2 are based on same kernel (OS version number 6.1)
 - Windows 8/2012 (6.2), Windows 8.1/2012 R2 (6.3)
- Enhancements in many areas
 - Multi- and Many-Core Processing
 - Power Efficiency
 - Security
 - Management
 - Virtualization
 - Componentization and Layering
- And, lots of Windows 7 specific features
 - TaskBar, Direct Write/2D/3D, GPGPU, Touch, others

NUMA Processor Group Support

- Segmented specification – “groups” of CPUs
 - Windows 6.X [Supports](#) more than 64 Logical Processors
 - 256 on Windows 7/2008 R2 and 640 on Windows 8/2012
 - CPUs identified in software by Group#: CPU#
 - Allows backward compatibility with 64-bit affinity
 - New applications have full CPU range using new APIs
- Permits better locality during scheduling than a “flat” specification



Windows Internals

Remove Coarse-Grained Locks

- Locks serialize access to data structures
 - Prevents multiple threads from simultaneously modifying data
 - Inhibits scaling because threads must wait for their turn (contention)
- Examples:
 - Dispatcher Lock
 - Memory Manager PFN Lock
 - Cache Manager Virtual Address Control Block (VACB) Lock
 - Object Manager Type Lock
- Either replaced with lock-free algorithms or with finer-grained synchronization mechanisms

Memory Optimizations

- Desktop Window Manager (DWM) re-architecture reduces memory footprint per window by 50%
- Registry moved from mapped files to paged pool
 - Improves performance because views into registry file don't need to be mapped and unmapped
- Working set management improvements:
 - Working set is amount of RAM memory manager assigns to process or kernel memory type
 - Memory manager uses more information for better tuning
 - System cache, paged pool, and pageable system code each have own working set

Power Efficiency

- Keep idle and stay idle
 - Minimize running services and tasks
 - Avoid background processing
 - Let LPs and sockets stay idle so that they enter deep sleep (C states)
- Core Parking
 - Move load from one CPU to another
 - The CPU and the Socket go to sleep, conserving power
- Trigger-Started Services
 - To reduce the number of background running processes

Virtual Accounts

- Want better isolation than existing service accounts
 - Don't want to manage passwords
- Virtual accounts are like service accounts:
 - Process runs with virtual SID as principle
 - Can ACL objects to that SID
 - System-managed password
 - Show up as computer account when accessing network
- Services can specify a virtual account
 - Account name must be “NT SERVICE\<service>”
 - Service control manager verifies that service name matches account name
 - Service control manager creates a user profile for the account
- Also used by IIS app pool and SQL Server

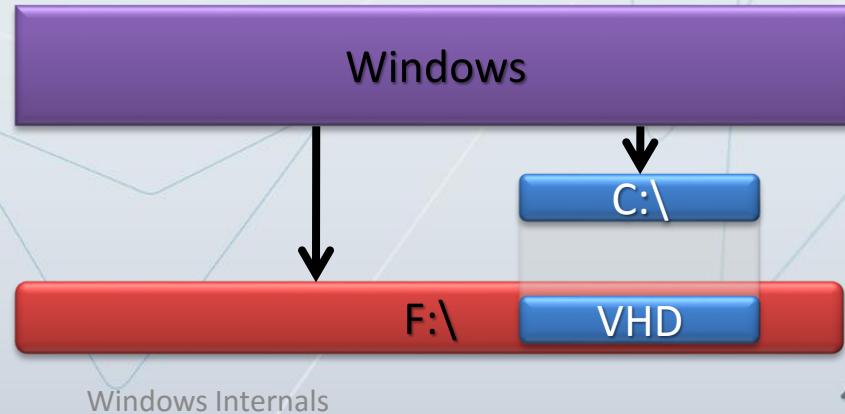
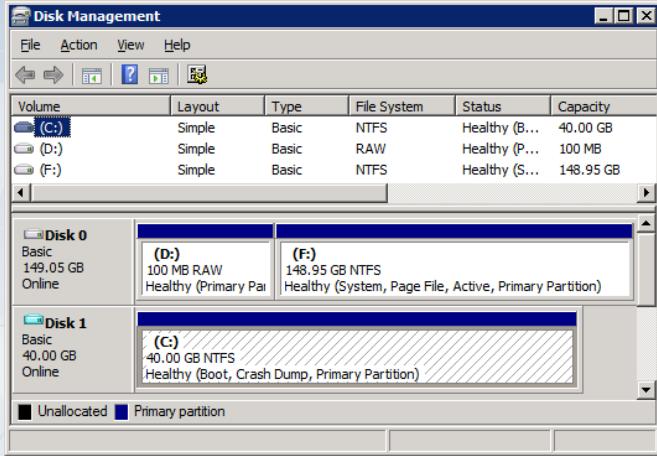
Native VHD Support

- Foundational support for booting from VHD and for Attach/Removal of VHDs
 - Orderly and surprise removal of volumes
 - Support for nested volumes
 - Servicing for mounted (offline) VHD volumes
- VHD operations
 - Create / Attach / Remove
- Tools and APIs:
 - Win32 APIs
 - VDS APIs (DCOM Remotable)
 - Hyper-V WMI for management operations
- Performance goal: within 10% of native

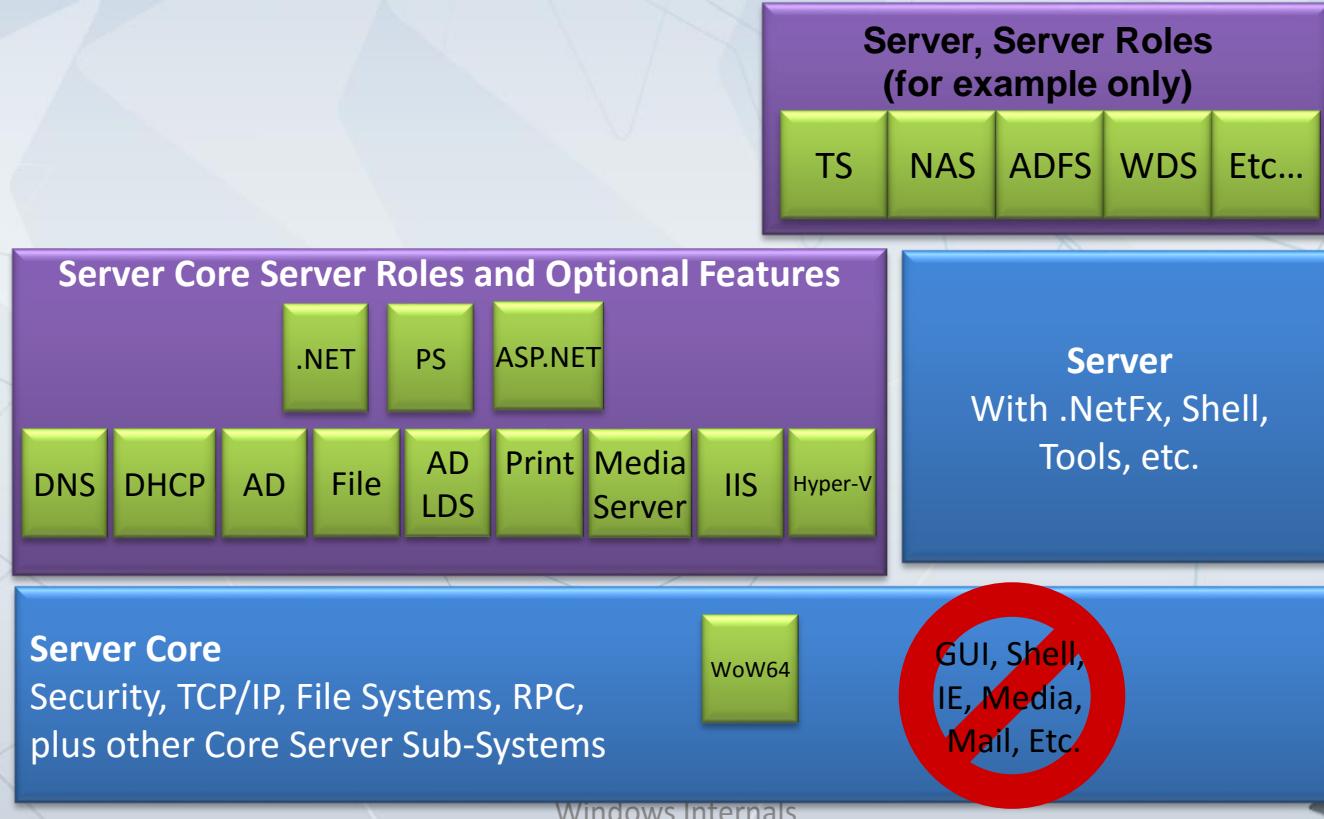
Create and Attach VHD

DEMO

VHD Boot in Windows



Server Core Architecture



What's New in Windows 10.X

Active pen

Accessory manager

Continuum

DirectLink

Print support on mobile

Custom Sensors

Media creation APIs

Passport and Windows Hello

Single, seamless tooling

DirectX12

Unified clipboard

Expanded API contracts

Action center APIs and triggers

App services

Backup / Restore

HLS, Dash, Closed Captioning

New tiles and actionable notifications

Volume purchase for apps

Bing services integration

Adaptive UI

Natural language support

2500+ new platform features

Video advertisements

Multi-adapter

Richer Analytics

Unified modern VOIP apps

Long running background tasks

Contacts APIs

AppLocker

Richer common controls

Xbox Live

Visual layer, Effects and Animation

Game DVR

Casting APIs and controls

Wifi developer tool connectivity

Calendar & Email APIs

Holographic

Drag/drop

Bluetooth beacons

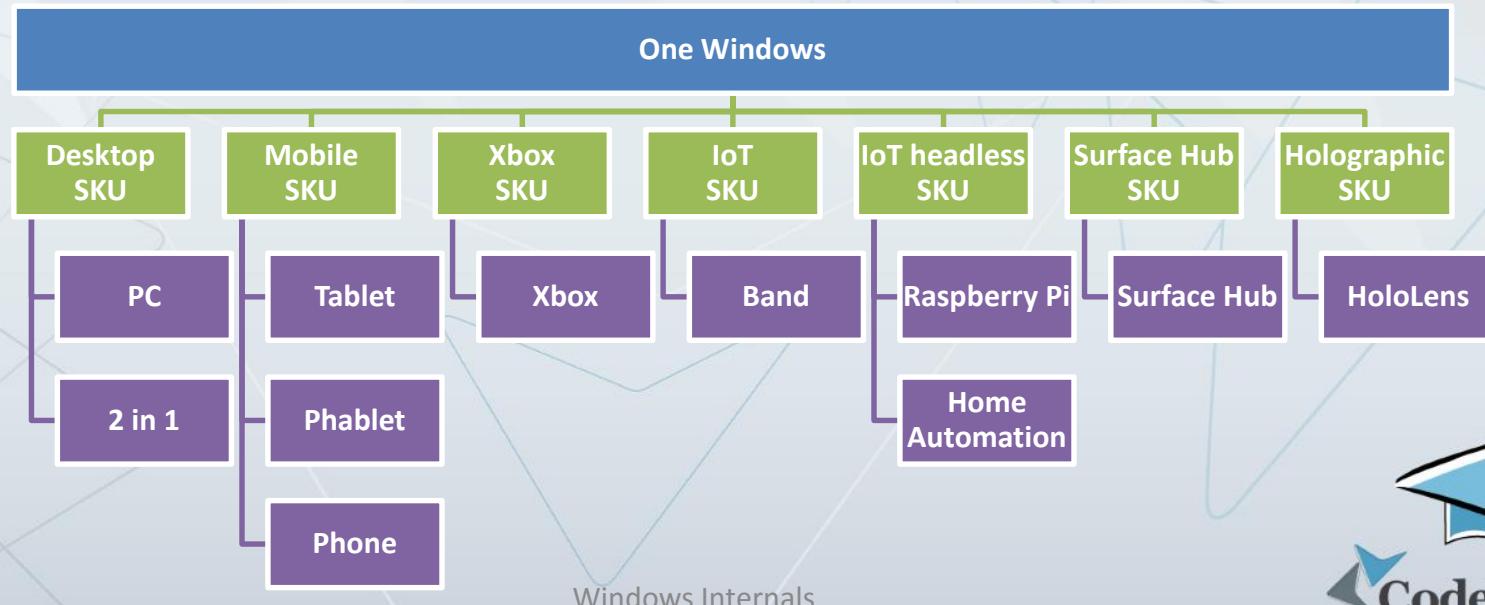
Cortana

Gesture, gaze, holograms

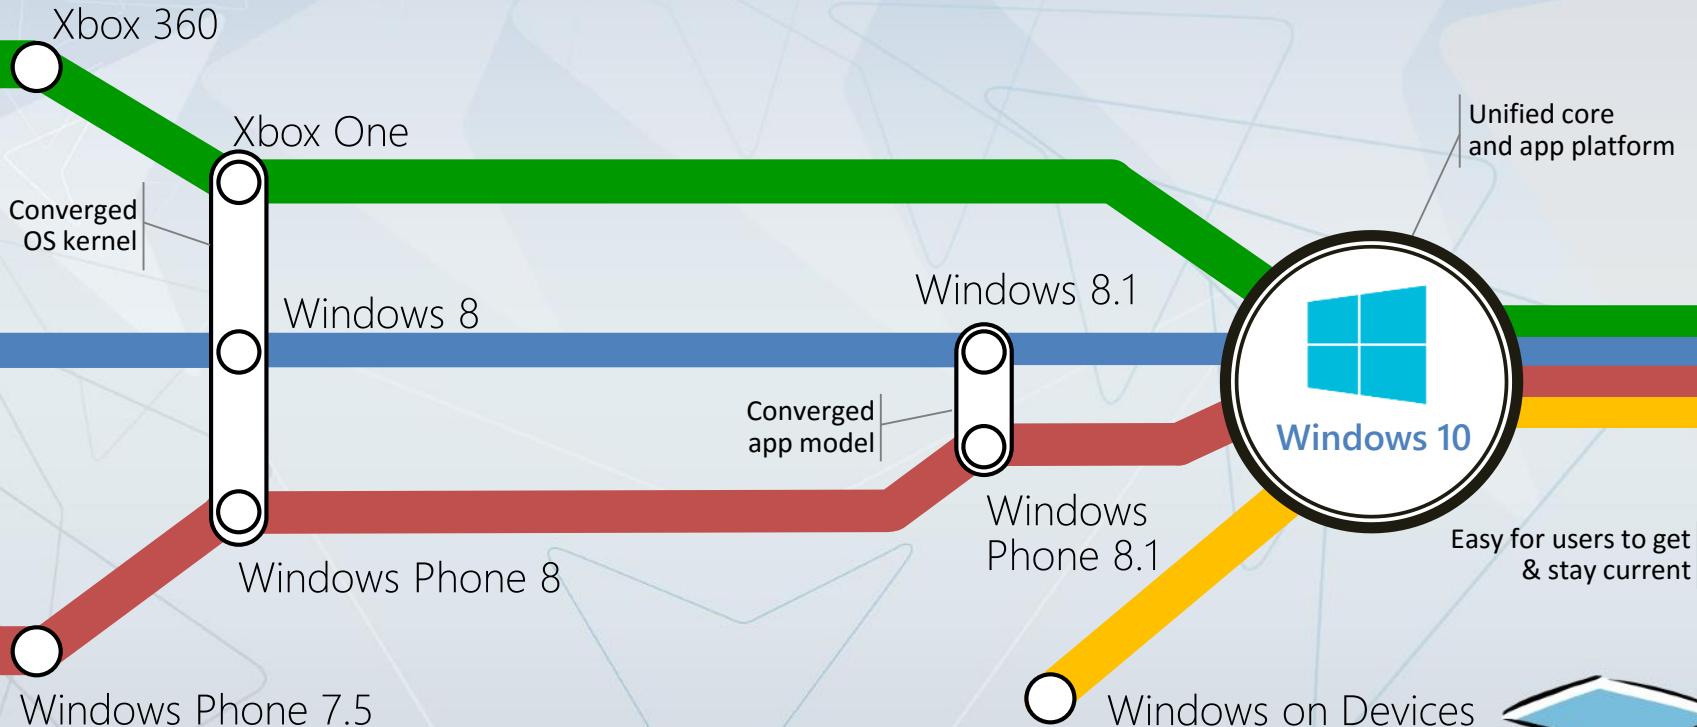
Activity detection APIs

Windows Core

- The refactored common core
 - One hardware platform, Universal hardware driver, Standard network and I/O



The convergence journey



Phone



Phablet



Small Tablet



Large Tablet



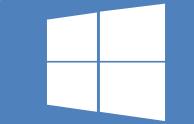
2-in-1s
(Tablet or Laptop)



Classic
Laptop



Desktops
& All-in-Ones



Windows 10

Surface Hub



Xbox



Holographic

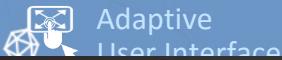


IoT

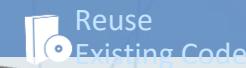




Multiple device families



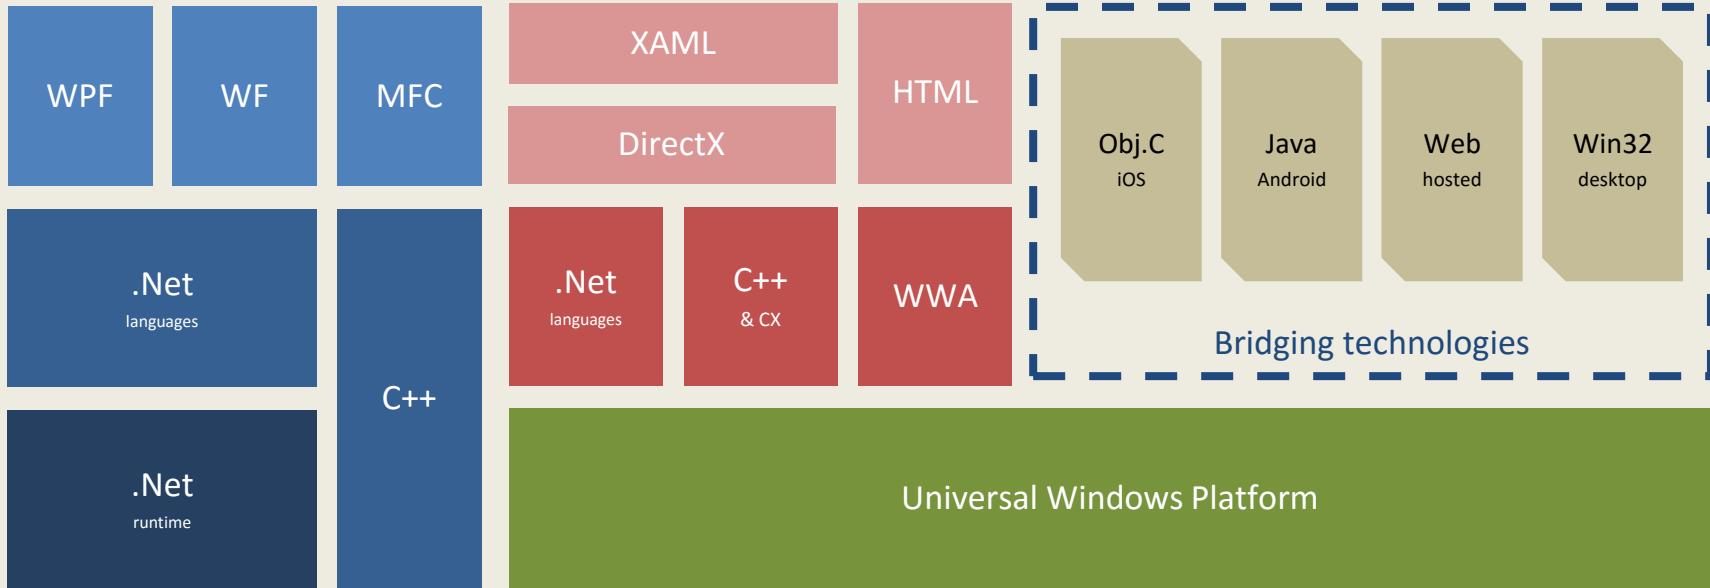
Adaptive
User Interface



Reuse
Existing Code



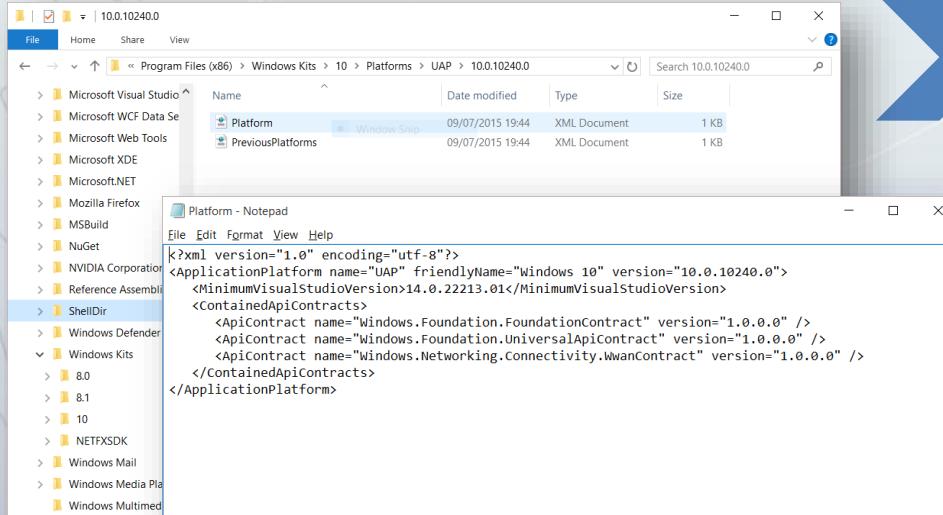
One Universal Windows Platform



Windows 10
operating system

Universal Windows Platform

- A single API surface
- A guaranteed API surface
- The same on all devices



Universal Windows Platform

Windows Core

Desktop
Device

Phone
Device

IoT
Device

Windows 10 for IoT Devices

Windows 10 IoT
for Industry
Devices

Windows 10 IoT
for Mobile
Devices

Windows 10 IoT
for Small
Devices

One Windows Platform

One Windows core
scalable across
devices

One universal app
and driver platform

One management
and deployment
approach

Secure

Enterprise-grade
security

Advanced
Customization and
lockdown
capabilities

Continued
innovation and
support

Connected

Local
connectivity

Cloud
connectivity

Device
services

Windows 10 IoT Core

- Use the same API, SDK, Visual Studio to develop Universal App
 - Extend UWP with device specific API
 - Run your desktop app on the PI
 - Extend it to control and to be controlled by hardware devices
 - You can Use [Arduino Wiring Pi](#) on Windows 10 IoT devices
 - C++ library that mimic the Arduino library

Nano Server (Server 2016)

- A remotely administered server operating system optimized for private clouds and datacenters
 - Similar to Windows Server in Server Core mode, but significantly smaller
 - Has no local logon capability
 - Only supports 64-bit applications, tools, and agents
 - Takes up far less disk space, sets up significantly faster, and requires far fewer updates and restarts
- Nano Server is ideal for a number of scenarios:
 - As a “compute” host for Hyper-V virtual machines
 - As a storage host for Scale-Out File Server
 - As a DNS server
 - As a web server running Internet Information Services (IIS)
 - As a host for applications that are developed using cloud application patterns

What's New in NT 10.X

- Security
 - VSM – Virtual Secure mode
 - Enables isolated User Mode (protect user mode from the kernel)
 - IO MMU – I/O Memory Management Unit
 - MMU for devices, as opposed to processors
 - Hardware based protection against DMA-access
 - Protects against buggy drivers and malicious code
 - CFG – Control Flow Guard
 - Windows Hello



What's New in NT 10.X

- Memory Compression
 - Can compress pages in memory and result in less Hard Page Faults
 - Does not prevent the use of pagefile but lowers the need to use it
- Lightweight Suspend & Deep Sleep
- Universal Drivers Development
- Cortana
- Continuum
- Developer Mode & Device Developer Portal
- Many internal mechanisms improvements such as Windows Loader, Kernel Timers & Scheduler



Module 2

KERNEL MECHANISMS

Agenda

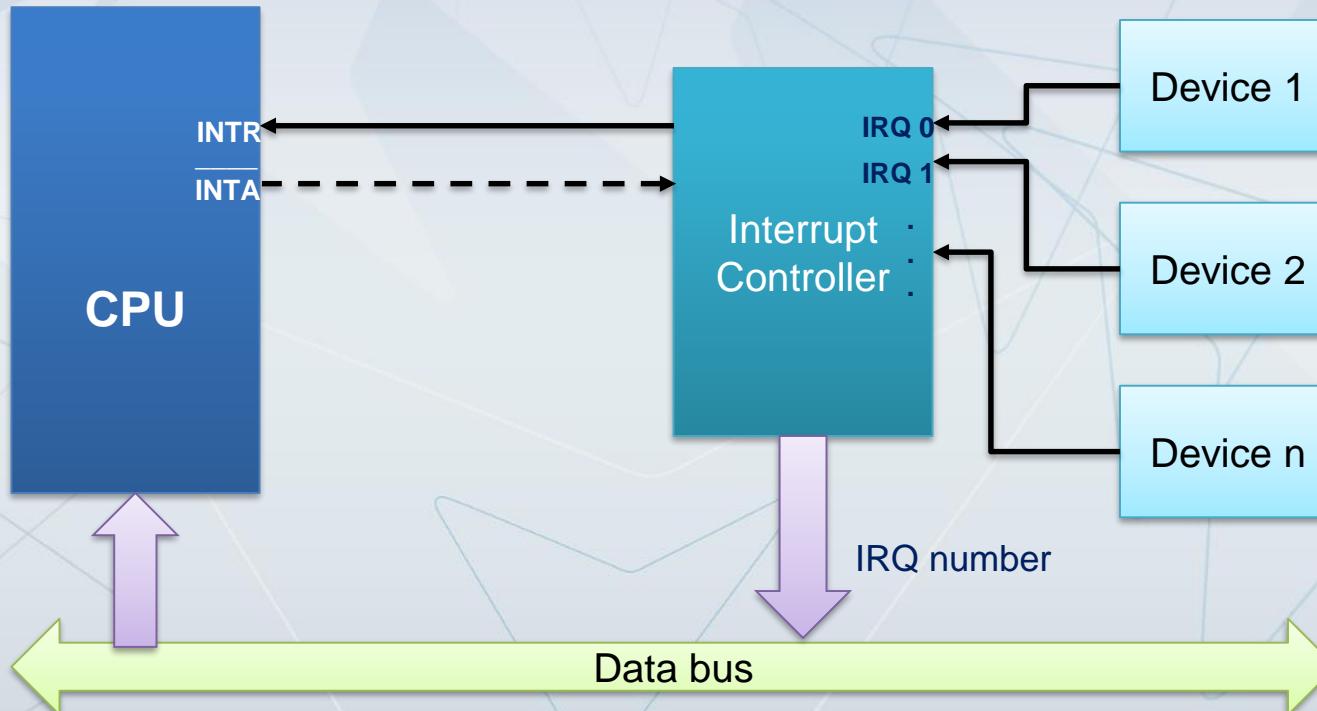
- Trap Dispatching
- Object Management
- Synchronization
- Kernel Event Tracing
- Wow64
- Kernel Transaction Manager
- Diagnostics
- Summary

Trap Dispatching

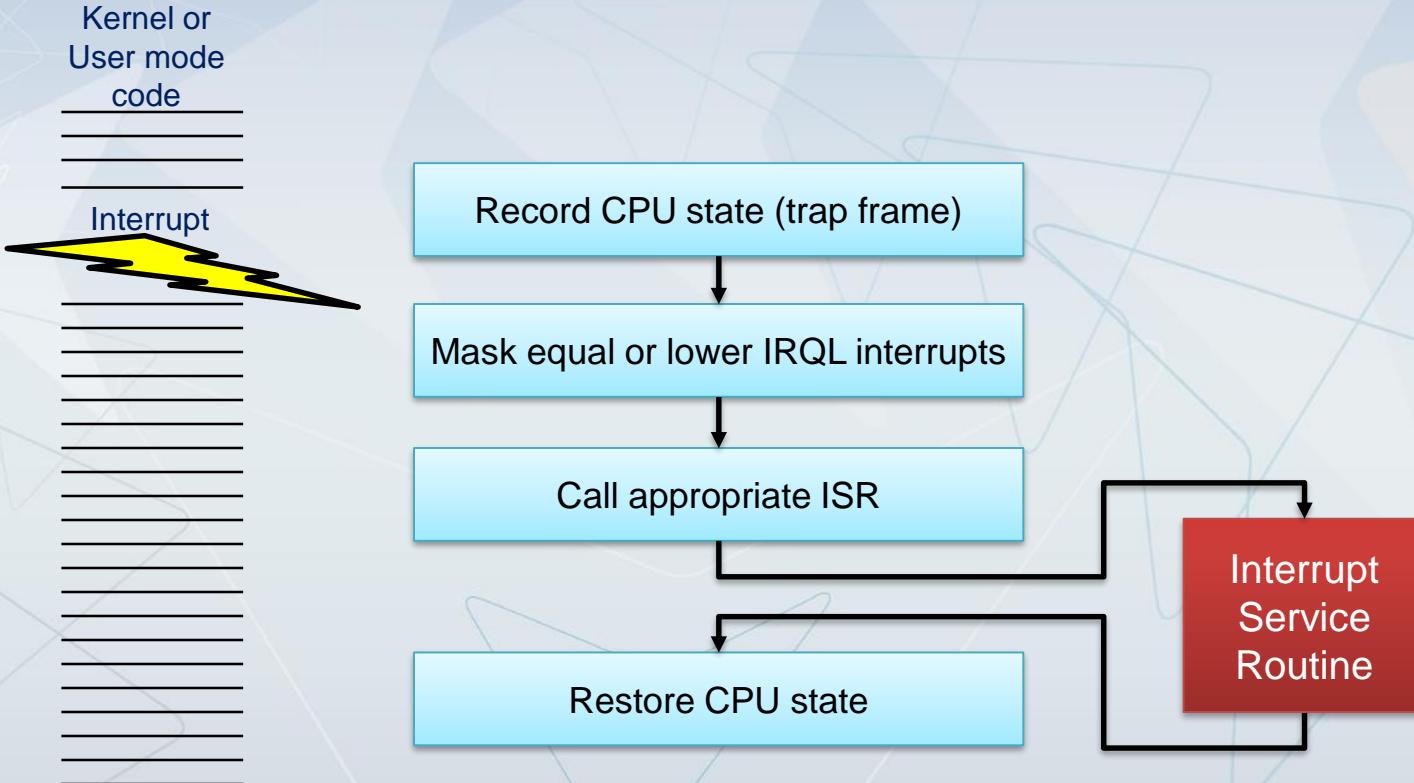
- Traps
 - Interrupts or exceptions
 - Divert code execution outside the normal flow
- Trap dispatching
 - Kernel mechanisms for capturing an executing thread when an interrupt or exception occurs and transferring control to a handling routine
- Interrupt
 - Asynchronous event, unrelated to the current executing code
- Exception
 - Synchronous call to certain instructions
 - Reproducible under the same conditions



Hardware Interrupts



Interrupt Dispatching



Hardware Interrupts Requests

- The device needing service signals the line connected to the interrupt controller (PIC or APIC)
- The interrupt controller signals the CPU on a single line
- The CPU requests the interrupt vector number from the interrupt controller
- The appropriate ISR is called based on its entry in the Interrupt Dispatch Table (IDT)

Interrupt Request Level (IRQL)

- Each interrupt has an associated Interrupt Request Level (IRQL)
 - Can be considered its priority
 - For hardware interrupts, mapped by the HAL
- Each processor's context includes its current IRQL
 - A CPU always runs the highest IRQL code
- Servicing an interrupt raises the processor IRQL to the level of the interrupt's IRQL
 - This masks all interrupts at that IRQL and lower
- Dismissing an interrupt restores the processor's IRQL to that prior to the interrupt
 - Allowing any previously masked interrupts to be serviced

IRQL Levels



IRQL Levels

- PASSIVE_LEVEL (0)
 - The “normal” IRQL level
 - User mode code always runs at this level
- APC_LEVEL (1)
 - Used for special kernel APCs (see I/O System module)
- DISPATCH_LEVEL or DPC_LEVEL (2)
 - The kernel scheduler runs at this level
 - If the CPU runs code at this (or higher) level, no context switching will occur on that CPU until IRQL drops below this level
 - Also no waiting on kernel objects (requires scheduler)
- Page faults can only be handled in $\text{IRQL} < \text{DPC_LEVEL}$
 - Code running at this or higher IRQL must always access non-paged memory

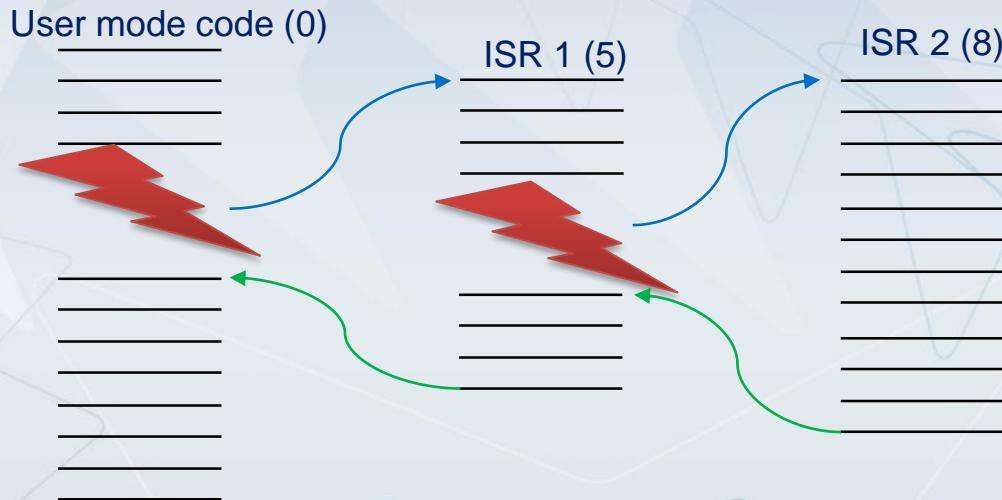
IRQL Levels

- Device IRQL (DIRQL) (3-11/26)
 - Reserved for hardware devices
 - The level that an ISR runs at
 - Always greater than DISPATCH_LEVEL (2)
- HIGH_LEVEL (x86=31, x64=15)
 - The highest level possible
 - If code runs at this level, nothing can interfere on that CPU
 - However, other CPUs are not affected
 - Use only when absolutely necessary!
- Other levels exist for kernel internal use

IRQL vs. Thread Priorities

- IRQL is an attribute of a CPU
- Thread priority is an attribute of a thread
- With $\text{IRQL} \geq \text{DISPATCH_LEVEL}(2)$
 - Thread priority has no meaning
 - The currently executing thread will execute forever, until IRQL drops below 2
- IRQLs can be changed (in kernel mode only) with [KeRaiseIrql](#) and [KeLowerIrql](#)

Interrupts and IRQLs



Mapping Interrupts to IRQLs

- IRQL is a software-only entity
- IRQL is not the same as IRQ
 - IRQ is an interrupt line connected to the interrupt controller
- The HAL is responsible for translating IRQs to vector numbers and IRQLs
- Vector is assigned in a round-robin fashion
- $\text{IRQL} = \text{Vector} / 16$

Software Interrupts

- The kernel uses software interrupts for various tasks
 - Initiating thread dispatching
 - Non-time critical interrupt processing
 - Handling timer expirations
 - Asynchronously execute a procedure in a context of a specific thread

Deferred Procedure Call (DPC)

- Used to defer processing from higher (device) IRQL to a lower (dispatch) level
- Implemented via DPC objects and software interrupts
- DPC object defines a procedure and arguments
- Executes specified procedure at Dispatch IRQL (2)
- Used heavily for driver "after interrupt" functions
 - ISR calls [KeInsertQueueDpc](#) or [ToRequestDpc](#) to queue a DPC

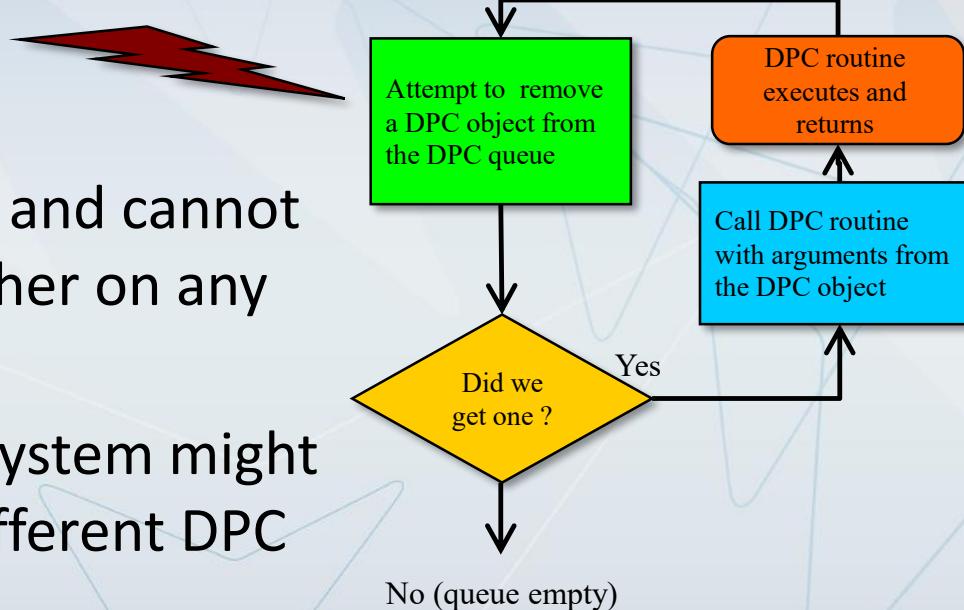
DPC Queue

- A list of "work requests"
 - One queue per CPU
 - But one CPU can process a DPC from another CPU's queue
 - Implicitly ordered by time (FIFO)
 - Can be somewhat manipulated with the Importance field
 - Each specifies procedure and arguments
- Processed after all higher-IRQL work (interrupts) completed

The DPC Processing Loop

- DPCs are queued and cannot interrupt each other on any one CPU
- Each CPU in the system might be executing a different DPC at the same time

Software interrupt
at DPC level



DPC Processing Notes

- Idle CPUs can help by processing DPCs from another CPU's DPC queue
- A DPC can be targeted at a different CPU than the executing one by calling
[KeSetTargetProcessorDPC/Ex API](#)
 - Still may be processed by another (idle) CPU
- DPC importance can be controlled by calling
[KeSetImportanceDpc](#)
 - Low, medium (default), high (goes to head of queue)

Exception Dispatching

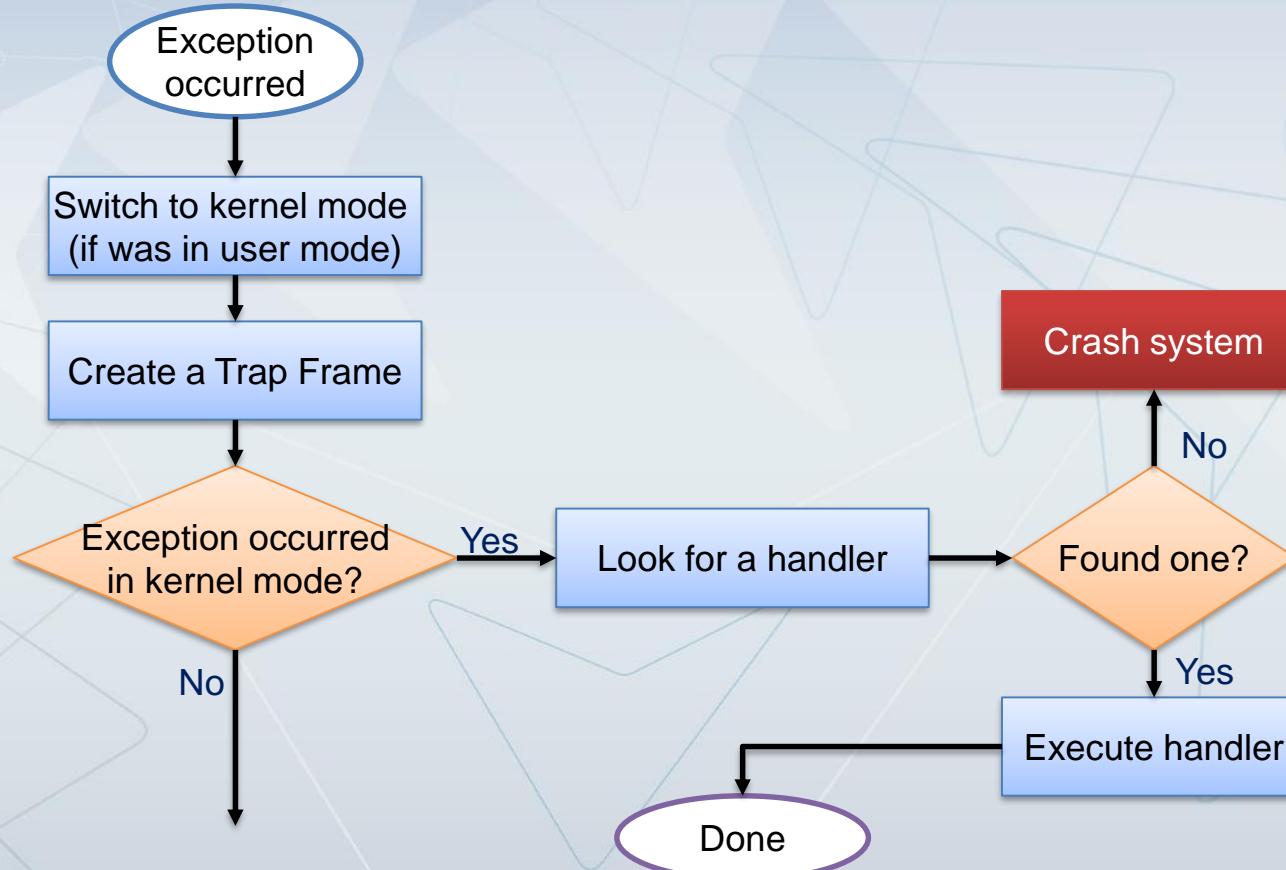
- Synchronous event resulting from certain code
- Examples
 - Divide by zero, access violation, stack overflow, invalid instruction
- Structured Exception Handling (SEH)
 - A mechanism used to handle, and possibly resolve, exceptions
- Exceptions are connected to entries in the IDT
 - An exception dispatcher in the kernel is responsible for “disposing” of the exception



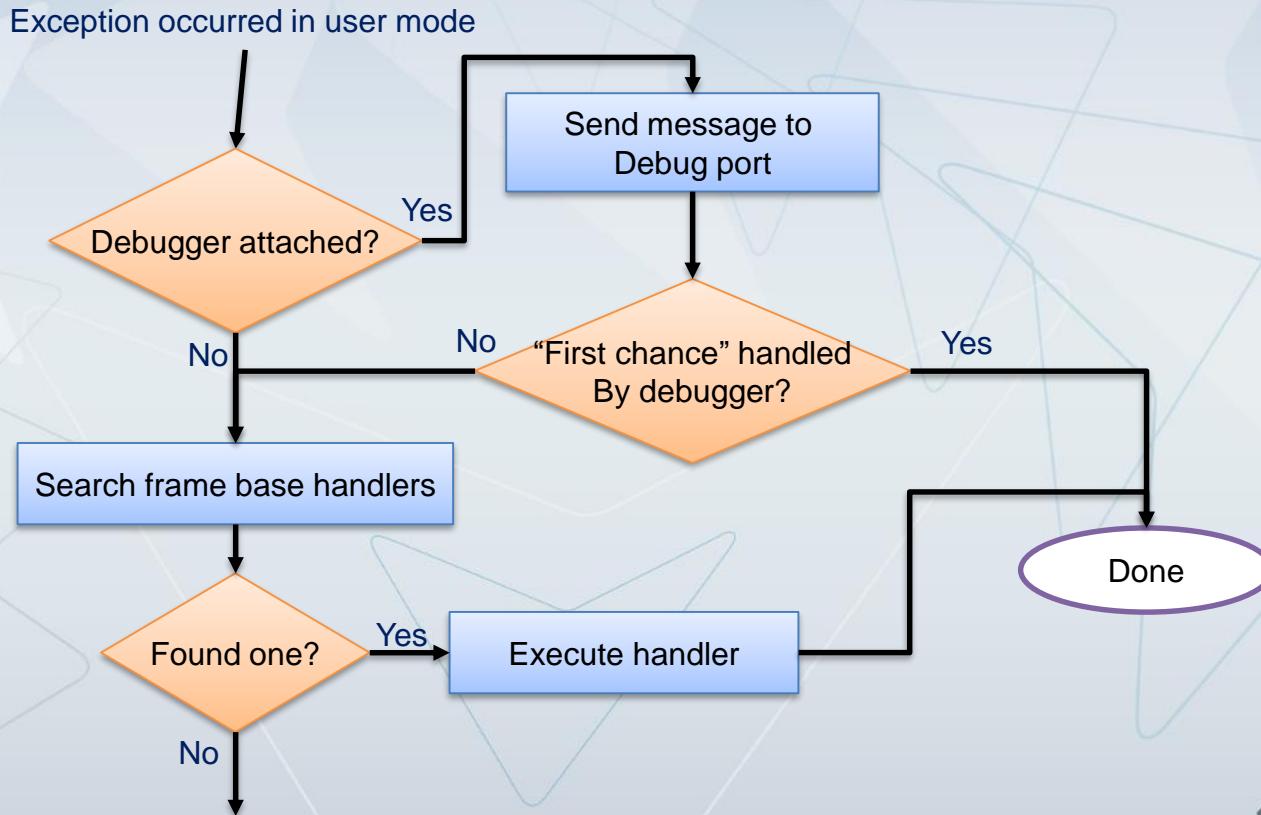
Exception Handling

- Some exceptions are handled transparently
 - E.g. a breakpoint is an exception, which is transferred to the appropriate debugger
- Some exceptions are filtered back to user mode for possible handling
 - E.g. accessing a user mode address that is not mapped
- Frame based exception handlers are searched
 - If the current frame has none, the previous frame is searched and so on
- Unhandled exceptions from kernel mode generate a “bug check” a.k.a. “Blue Screen of Death”

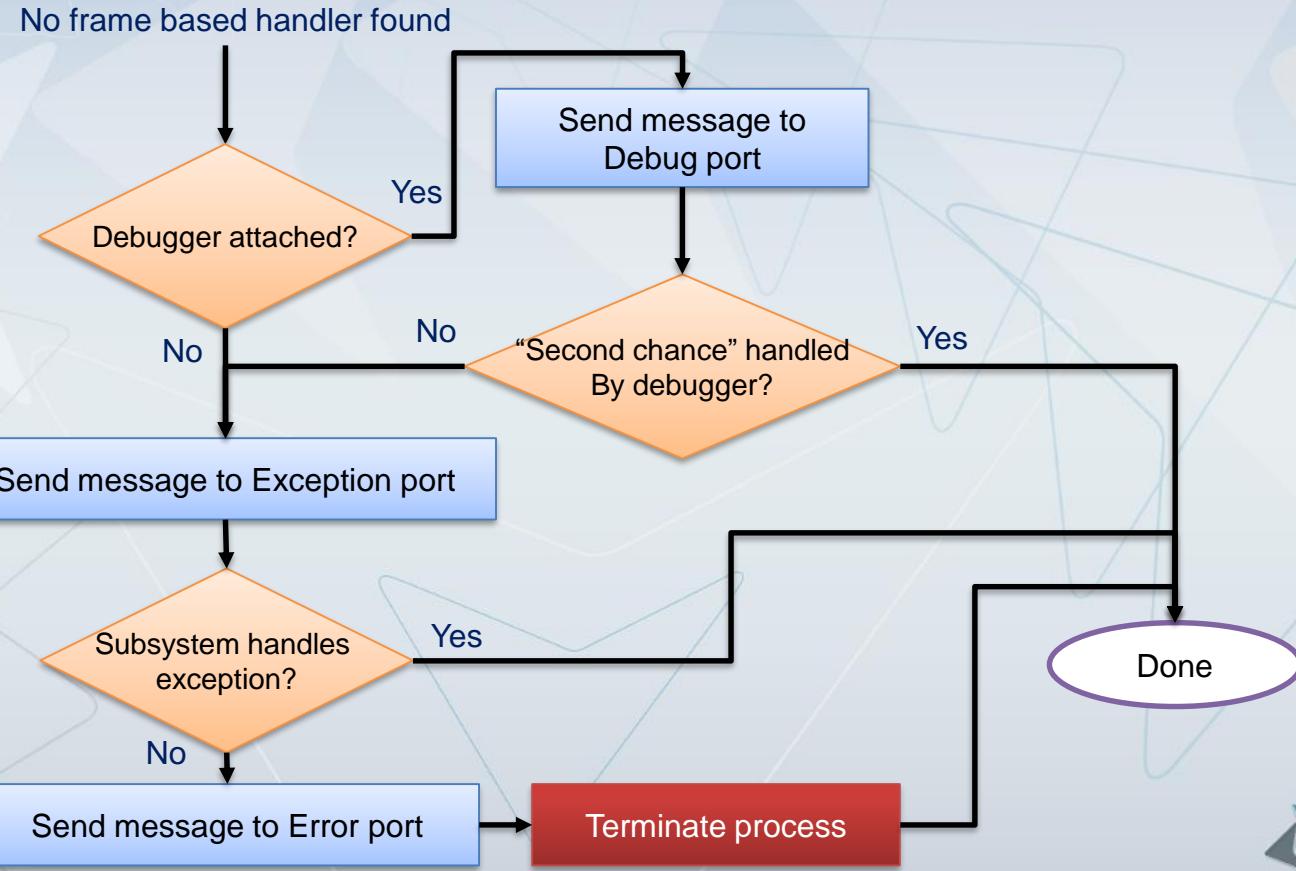
Resolving Exceptions



Resolving Exceptions



Resolving Exceptions



Structured Exception Handling (SEH)

- Exposed for developers by extended C keywords
 - `_try`
 - Wraps a block of code that may throw exceptions
 - `_except`
 - Possible block for handling exceptions in the preceding `_try` block
 - `_finally`
 - Execute code whether an exception occurred or not
 - `_leave`
 - Jumps to the `_finally` clause
 - Allowed blocks are `_try/_finally` and `_try/_except`
 - However, can be nested to any level
- Works in kernel mode and user mode
- Custom exceptions can be raised with [RaiseException](#) (Win32)

try / finally Example

```
CRITICAL_SECTION cs;

void f()
{
    EnterCriticalSection(&cs);
    try
    {
        // do work while critical section held
    }
    finally
    {
        LeaveCriticalSection(&cs);
    }
}
```

SEH and High Level Exceptions

- How does SEH relate to C++, .NET or Java exceptions?
- C++ (at least Microsoft's compiler), .NET and Java use SEH internally
 - With a custom exception code
 - The [EXCEPTION RECORD](#) holds the extra custom information that is understood by the particular runtime
- What about finally in C++?
 - No [**finally**](#) keyword
 - However, can be simulated with the idiom Resource Acquisition Is Initialization (RAII)

C++ RAII

```
CRITICAL_SECTION cs;

class AutoCS
{
    CRITICAL_SECTION* pcs;
public:
    AutoCS(CRITICAL_SECTION* c) : pcs(c)
    {
        EnterCriticalSection(pcs);
    }
    ~AutoCS()
    {
        LeaveCriticalSection(pcs);
    }
};
```

```
void f()
{
    AutoCS lock(&cs);

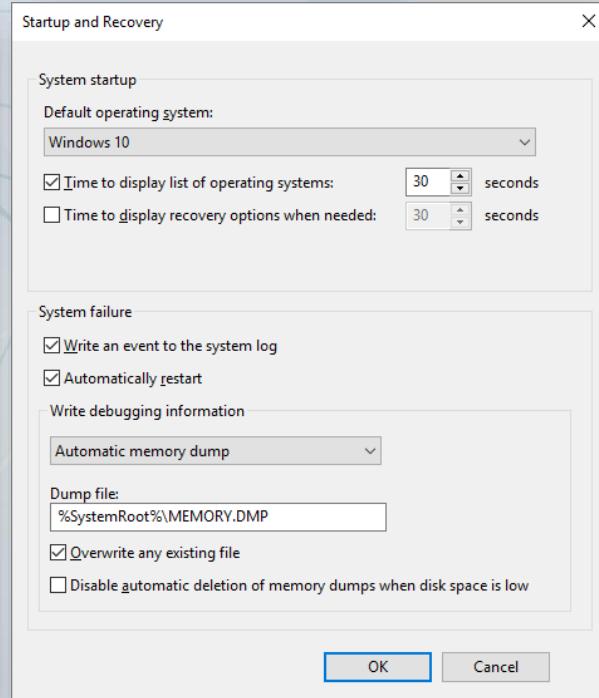
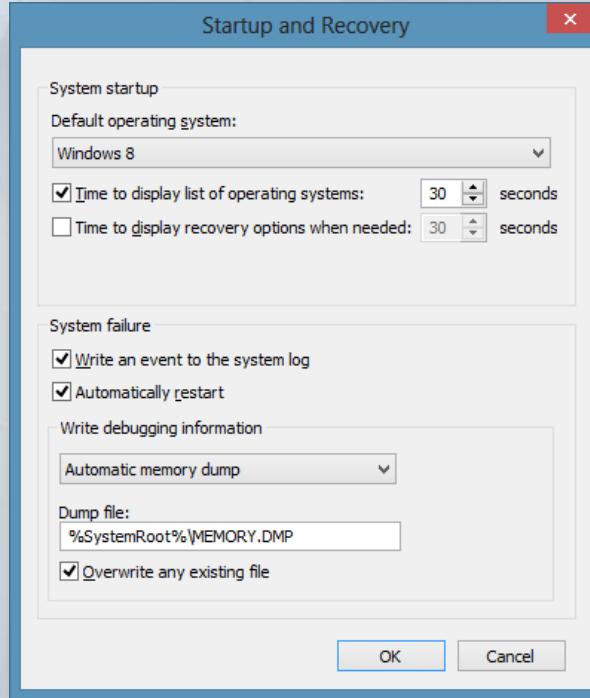
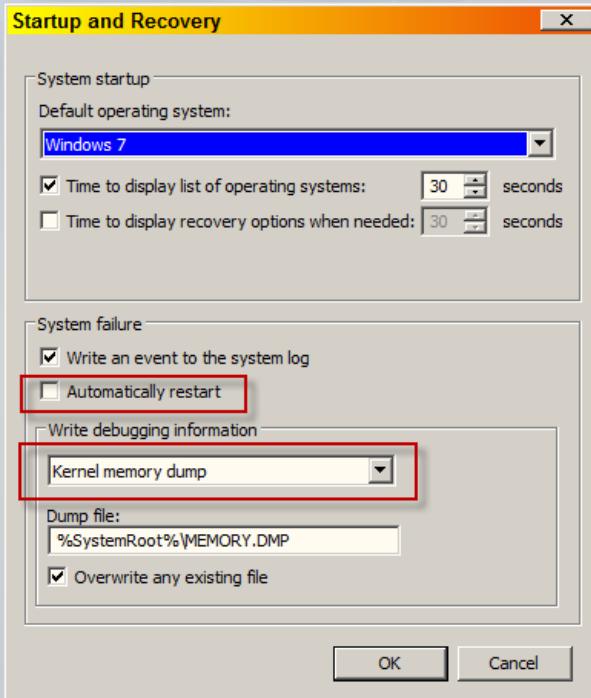
    // access shared resource...
}
```

System Crash

- “Blue Screen of Death”
- Occurs when code running in kernel mode raises an unhandled exception
- First steps in diagnosis
 - Make sure a crash dump file is written to disk
 - don’t automatically restart
- Set in the Startup Recovery dialog in System applet / Advanced Settings / Startup & Recovery

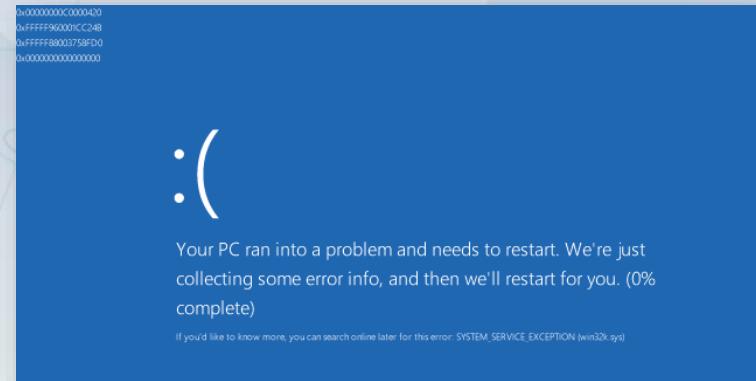


Startup And Recovery



“Blue Screen” Information

- A blue screen shows
 - The code of the crash (e.g. 0x0A)
 - The textual name of the code
 - IRQL NOT LESS OR EQUAL
 - Four numbers –
their meaning depends on the exact error code
 - Go to the debugging tools help and search for the stop code to learn more
 - A short stack trace leading to the crash
 - Typically ending in KeBugCheckEx call
 - For simple scenarios, culprit is on the stack



Object Manager

- Executive component responsible for managing executive objects
- Manages creating, deleting, protecting and tracking objects
- Maintains a hierarchical namespace of objects
 - Can be viewed partially with the [WinObj](http://www.sysinternals.com) utility (from www.sysinternals.com)
- User mode clients can obtain handles to objects
 - Cannot touch actual memory structure
- Kernel mode clients can do either

Object Manager Namespace

WinObj - Sysinternals: www.sysinternals.com

File View Help

Name	Type	Symlink
ArcName	Type	
BaseNamedObjects	Type	
Callback	Type	
Device	Type	
Driver	Type	
DxgkSharedResource	Type	
DxgkSharedSwapChainObject	Type	
DxgkSharedSyncObject	Type	
EtwConsumer	Type	
EtwRegistration	Type	
Event	Type	
File	Type	
FilterCommunicationPort	Type	
FilterConnectionPort	Type	
IoCompletion	Type	
IoCompletionReserve	Type	
IRTimer	Type	
Job	Type	
Key	Type	
KeyedEvent	Type	
Mutant	Type	
NetworkNamespace	Type	
Partition	Type	
PcwObject	Type	
PowerRequest	Type	
Process	Type	
Profile	Type	
PsSiloContextNonPaged	Type	
PsSiloContextPaged	Type	
RawInputManager	Type	

\ObjectTypes\VirtualKey

Object Kinds

- Executive objects
 - High level objects, usually encapsulating a kernel object
 - Add reference tracking, security, handle management
- Kernel objects
 - Primitive objects
 - Implement the actual functionality of the object
 - Provide synchronization mechanisms

Object Structure

Owned by Object Manager
Owned by kernel
Owned by Executive

Objects Name
Object Directory
Security Descriptor
Quota Charges
Open Handle Count
Open Handles List
Object Type
Reference Count

Kernel Object
Object Data

Windows Internals

Type Object

Type Name
Access Type
Synchronizable?
Pageable?
Object Methods

Object Types

- Symbolic link - Allows referencing objects indirectly
- Process - A virtual address space and control information necessary for execution of thread objects
- Thread - An executable entity within a process
- Job - A collection of processes managed as a group
- Section (File mapping) - A region of shared memory
- File - Represents an open file or another I/O device
- Token - A security profile of a thread or process
- Event - A flag object, used for synchronization or notification
- Semaphore - A counter providing a resource gate for maximum allowed threads accessing a shared resource



Object Types

- Mutex (mutant) - A synchronization object used to serialize access to a shared resource
- Timer - An object allowing notification after some time period elapsed
- IoCompletion (I/O completion port) - An object for notifications of I/O operation completion
- Key - An object representing registry data
- Window Station - A management object, containing desktops, a clipboard and an atom table
- Desktop - An object contained within a window station, contains a drawing surface, windows and hooks
- Transaction (many types) – Kernel Transaction Manager types
- And more... (Some undocumented yet)

Type Object Attributes

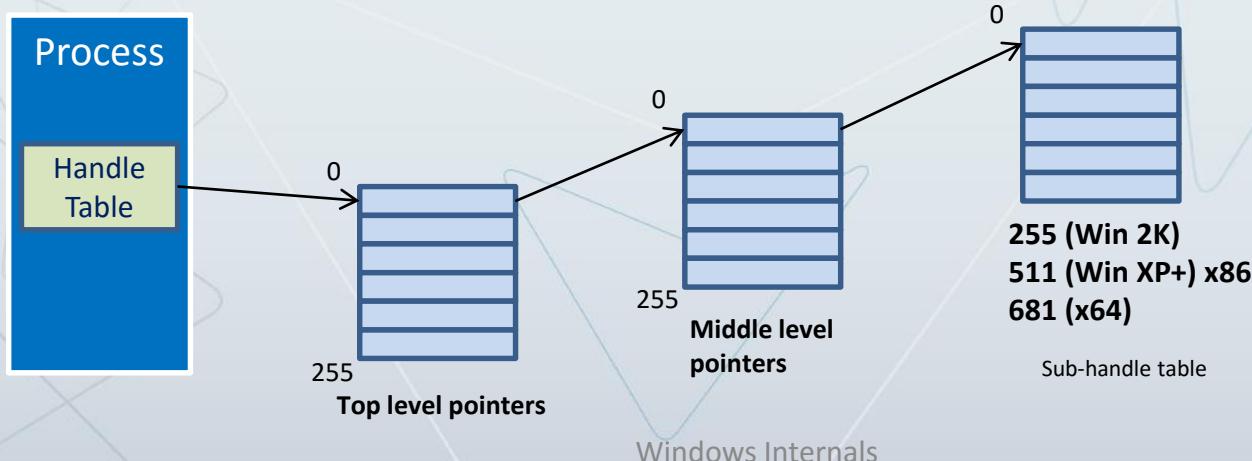
- Pool type
 - Indicates whether objects of this type need to be allocated from paged or non-paged pool
- Access type
 - The type of access a thread can request when opening an object
 - E.g. Read, Write, Synchronize
- Synchronization
 - Indicates whether a thread can wait on objects of this type
 - Job, Process, Thread, File, Timer, Event, Semaphore, Mutex, I/O Completion Port
- Methods
 - One or more routines that the Object Manager calls at certain points of the object's lifetime

Object Handles

- When a process creates or opens an object, it receives a handle to the object
 - Used as an opaque, indirect pointer to the underlying object
 - Allows sharing objects across processes
- Each process has a private handle table
 - Sharing is possible through
 - Process handle inheritance
 - Opening an object by name
 - Duplicating a handle ([DuplicateHandle](#) Windows API)
- Viewing process handles
 - Use Process Explorer (GUI), [handle.exe](#) (Console) (from [www.SysInternals.com](#))

Process Handle Table

- Pointed to by the [EPROCESS](#) object
 - A handle is an index to that table
 - Starting at 4 with increments of 4
- Implemented as a 3 level layout



Process Explorer Handle View

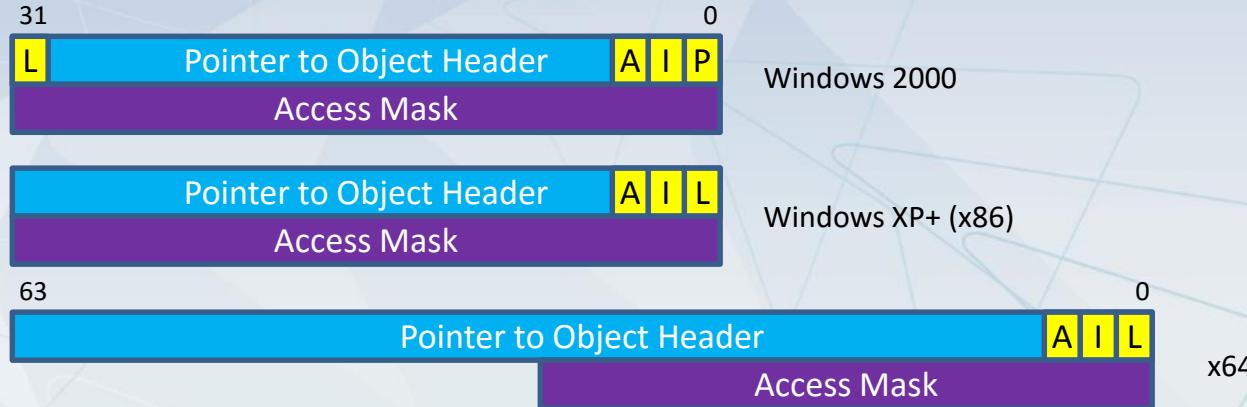
Process Explorer - Sysinternals: www.sysinternals.com [HOME\alon]

Process	PID	CPU	Private Bytes	Working Set	Description	User Name	Session
UcMapi.exe	8284	< 0.01	9,952 K	29,584 K	Microsoft Lync	HOME\alon	1
FlashUtil_ActiveX.exe	10948		3,108 K	8,792 K	Adobe® Flash® Player Utility	HOME\alon	1
explorer.exe	14732		19,316 K	45,384 K	Windows Explorer	HOME\alon	1
NAMECONTROLSER...	6604	< 0.01	8,420 K	17,896 K	Microsoft Office Contact Retr...	HOME\alon	1
Sums.exe	4204	Susp...	58,812 K	94,852 K	Sums	HOME\alon	1
WmiPrvSE.exe	14976		1,960 K	5,544 K	WMI Provider Host	NT AUTHORITY\...	0

Type	Name	Handle	Access	Object Address	S
Directory	\KnownDlls	0x4	0x00000003	0xFFFFF8A0012C6EB0	
Directory	\KnownDlls32	0x8	0x00000003	0xFFFFF8A00DBB9080	
File	C:\Windows	0xC	0x01000020	0xFFFFF8A802EE9F200	
Directory	\KnownDlls32	0x10	0x00000003	0xFFFFF8A00DBB9080	
File	C:\Program Files\Windows Apps\12149CodeValue.Sums_1.1.0.3_neutral_68yqwxhkd6p9c	0x14	0x01000020	0xFFFFF8A801E0339C0	
Key	HKLM\SYSTEM\ControlSet001\Control\Nls\Sorting\Versions	0x18	0x00020019	0xFFFFF8A042C697F0	
File	\Device\CNG	0x2C	0x0100001	0xFFFFF8A802D563B20	
Key	HKLM	0x34	0x00020019	0xFFFFF8A04A2ED270	
Key	HKCU	0x3C	0x00020019	0xFFFFF8A0427A59A0	
Key	HKLM\SOFTWARE\Wow6432Node\Microsoft\.NETFramework	0x40	0x00020019	0xFFFFF8A04B3749F0	
Window Station	\Sessions\1\Windows\WindowStations\WinSta0	0x50	0x00020327	0xFFFFF8A8012853900	
Desktop	\Default	0x54	0x000F0FF	0xFFFFF8A80157D8ED0	
Window Station	\Sessions\1\Windows\WindowStations\WinSta0	0x58	0x00020327	0xFFFFF8A8012853900	
Directory	\Sessions\1\AppContainerNamedObjects\S-1-15-2-1043270836-2805960744-1887870334-40...	0x58	0x0000000F	0xFFFFF8A046CCB870	
Key	HKCU\Software\Classes	0xA4	0x00020019	0xFFFFF8A04199A640	
Key	HKLM\SOFTWARE\Microsoft\Windows Runtime\ActivatableClassId	0xAC	0x00020019	0xFFFFF8A046738E90	
Key	HKLM\SOFTWARE\Microsoft\Windows Runtime\CLSID	0xB4	0x00020019	0xFFFFF8A0422A9B90	
Section	...\Cor_SxSPublic_IPCBlock	0xC8	0x000F0007	0xFFFFF8A042BD5670	
Thread	Sums.exe(4204): 12088	0x12C	0x001FFFFF	0xFFFFF8A802F670880	
Thread	Sums.exe(4204): 11672	0x130	0x001FFFFF	0xFFFFF8A802DEA6A00	
Event	\KernelObjects\LowMemoryCondition	0x14C	0x00100001	0xFFFFF8A801285B0F0	
Thread	Sums.exe(4204): 12120	0x1A0	0x001FFFFF	0xFFFFF8A802EEA3080	

CPU Usage: 51.45% | Commit Charge: 16.01% | Processes: 114 | Physical Usage: 25.56%

Handle Entry Layout



- Object header is always a multiple of 8 (x86) / 16 (x64) bytes
- Flags
 - P – protect from close
 - I – Is handle inherited in a child process?
 - A – auditing on close
 - L – synchronization lock
- Access Mask
 - Determines what actions can be performed using this handle

Sharing Objects

- A handle is private to its containing process
- Sometimes an object needs to be shared between processes
- Sharing is possible through
 - Process handle inheritance
 - Opening an object by name
 - Duplicating a handle

Handle Usage

- User mode processes retrieve a handle by calling an appropriate `Create*` function or an `Open*` function
 - A handle is returned upon success
- User mode processes close a handle using the [CloseHandle](#) API
 - Entry is deleted from the handle table
- Kernel mode code can obtain handles that reside in the kernel system space and is visible in any process context
- Alternatively, kernel mode code can obtain a direct pointer to the underlying object
 - Using [ObReferenceObjectByHandle](#)
 - Must release reference by calling [ObDereferenceObject](#)

Object Names

- Object can have names
 - Allows object lookup by name
 - Allows easy sharing between processes
- When an object is created, it can (optionally) have a name
 - Not allowed for all object types created from user mode
- Objects are stored in the object manager namespace
- Can be viewed with the [WinObj](#) tool

Object Namespace Directories

- **Global??** - Symbolic links
- **BaseNamedObjects** – Global Mutexes, semaphores, events, timers and sections that have a name (named objects)
- **Sessions/[0-...]/BaseNamedObjects** – Local named objects
- **Sessions/[0-...]/AppContainerNamedObject** – Local Universal Windows Platform named objects
- **Device** - Device objects
- **Driver** - Driver objects
- **FileSystem** - File system driver objects and file system recognizer device objects
- **KnownDLLs** - Section names and paths for known DLLs (stored also in the registry), loaded automatically at startup
- **ObjectTypes** - All known object types
- **Windows** - Windows subsystem ports and window stations

WinObj.Exe - BaseNamedObjects

WinObj - Sysinternals: www.sysinternals.com

File View Help

ArcName

BaseNamedObjects

- Restricted
- SensorLocationListeners

Callback

Device

Driver

FileSystem

GLOBAL??

KernelObjects

KnownDlls

KnownDlls32

NLS

ObjectTypes

RPC Control

Security

Sessions

- 0
 - AppContainerNamedObjects
 - >... DosDevices
- 1
 - AppContainerNamedObjects
 - BaseNamedObjects
 - Restricted
 - DosDevices
 - >... Windows
- 2
 - AppContainerNamedObjects
 - S-1-15-2-155514346-2573954481-755741238-1654018636-1.
 - S-1-15-2-1861897761-1695161497-2927542615-642690995-

Name	Type	Symlink
AppContainerNamedObjects	SymbolicLink	\Sessions\1\AppContain...
CicerServerStartedEvent	Event	
ComTaskPool:1080	Semaphore	
CTF.AsmListCache.FMPWin...	Section	
DBWinMutex	Mutant	
DWM_DX_FULLSCREEN_TR...	Event	
DwmComposedEvent_1	Event	
EventShutdownCSRSS	Event	
GdiPrimaryChanged	Event	
GfeInstallUninstall	Event	
GfeRegistryChange	Event	
Global	SymbolicLink	\BaseNamedObjects
HomePageEvent	Event	
HWNDInterface:10044	Section	
Local	SymbolicLink	\Sessions\1\BaseNamed...
MSCTF.Asm.MutexWinlogo...	Mutant	
MSCTF.CtfActivated.Winlo...	Event	
MSCTF.CtfDeactivated.Winl...	Event	
MSCTF.CtfMonitorInitialize...	Event	
MSCTF.CtfMonitorInstMute...	Mutant	
MSCTF.CtfServerMutexWinl...	Mutant	
Mutex_SRSPS_{3F6EBB2B-A...	Mutant	
Mutex_WMP_{7148E1CB-E4...	Mutant	
NVSV64.DLL	Mutant	
SaveSpanPersistence	Event	
ScNetDrvMsg	Event	
Session	SymbolicLink	\Sessions\BNOLINKS
SessionImmersiveColorMutex	Mutant	

Object Names and Sessions

- In a terminal session environment, each session should have its own objects
 - From Windows XP, each logon (switch user) has its own session
 - From Windows Vista, services run in session 0
 - From Windows 8, Windows Store app (AKA UWP on Windows 10) has their own sub folder:
- The object manager creates a Sessions directory with a session ID subdirectory under: Sessions/[0-...]/AppContainerNamedObject
 - Named objects are created under that directory for that session
 - Redirection is automatic
- Sessions can access the interactive session objects using the prefix “Global\” when using object names to retrieve a handle



User and GDI Objects

- The Object Manager is responsible for kernel objects only
- User and GDI objects are managed by Win32k.sys
- The API functions in user32.dll and gdi32.dll don't go through NtDll.dll
 - They invoke the sysenter/syscall instructions directly

User and GDI Objects

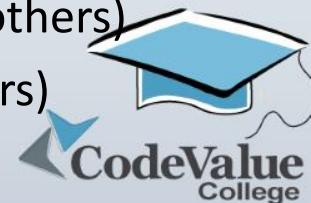
- User objects
 - Windows (HWND), menus (HMENU) and hooks (HHOOK)
- User object handles
 - No reference/handle counting
 - Private to a Window Station
- GDI objects
 - Device context (HDC), pen (HPEN), brush (HBRUSH), bitmap (HBITMAP) and others
- GDI object handles
 - No reference/handle counting
 - Private to a process
 - Cannot be shared across processes

Thread Synchronization

- Threads sometimes need to coordinate work
- Canonical example
 - Accessing a linked list concurrently from multiple threads
- Synchronization is based upon waiting for some condition to occur
- The kernel provides a set of synchronization (dispatcher) primitives on which threads can wait efficiently

Kernel Dispatcher Objects

- Maintain a state (signaled or non-signaled)
 - The meaning of “signaled” is dependent on object type
- Can be waited to change to the signaled state
 - Windows API: [WaitForSingleObject](#), [WaitForMultipleObjects](#) and their variants
 - Kernel mode: [KeWaitForSingleObject](#), [KeWaitForMultipleObjects](#)
- Dispatcher object types
 - Process, thread, event, mutex, semaphore, timer, file, I/O completion port
- Higher level wrappers exist
 - MFC: [CSyncObject](#) (abstract base of [CMutex](#), [CSemaphore](#) and others)
 - .NET: [WaitHandle](#) (abstract base of [Mutex](#), [Semaphore](#) and others)



“Signaled” Meaning

- Process - The process has terminated
- Thread - The thread has terminated
- Mutex - The mutex is free
- Event - The event flag is raised
- Semaphore - The semaphore count is below its maximum
- File, I/O completion port - I/O operation completed
- Timer - Interval time expires

Mutex

- Mutual exclusion
- Called Mutant in kernel terminology
- Allows a single thread to enter a critical region
- The thread that enters the critical region (its wait has succeeded) is the owner of the mutex
- Releasing the mutex allows one (single) thread to acquire it and enter the critical section
- Recursive acquisition is ok (increments a counter)
 - If the owning thread does not release the mutex before it terminates, the kernel releases it and the next wait succeeds with a special code (abandoned mutex)

Semaphore

- Maintains a counter (set at creation time)
- Allows x callers to “go through” a gate
- Does not maintain any ownership
- When a thread succeeds a wait, the semaphore counter decreases
 - When the counter reaches zero, subsequent waits do not succeed (state is non-signaled)
 - Releasing the semaphore increments its counter, releasing a thread that is waiting

Event

- Maintains a Boolean flag
- Event types
 - Manual reset (Notification)
 - Auto reset (Synchronization)
- When set (signaled) threads waiting for it succeed the wait
 - Manual reset event releases any number of threads
 - Auto reset event releases just one thread
 - And the event goes automatically to the non-signaled state
- Useful when no other object fits the bill



Critical Section

- User mode replacement for a mutex
- Can be used to synchronize threads within a single process
 - Operates on a structure of type [CRITICAL SECTION](#)
- Cheaper than a mutex when no contention exists
 - No transition to kernel mode in this case
- Uses [EnterCriticalSection](#) and [LeaveCriticalSection](#) API functions
 - No way to specify a timeout other than infinite or zero
 - Zero is accomplished with [TryEnterCriticalSection](#)
- .NET
 - A similar effect is achieved with the [lock C# keyword](#)
 - Calls the framework's [Monitor.Enter/Exit](#) in a [try/finally](#) block



Executive Resource

- Similar to a mutex
- Supports multiple-reader shared access
- Not exported to the Windows API
 - Win32 API has the [Slim Reader Writer Lock](#)
- Not a dispatcher object
 - Structure allocated from non-paged pool
 - Has a special API that can be used in kernel mode (documented in the WDK)
- Functions
 - [ExAcquireResourceExclusiveLite](#),
[ExAcquireResourceSharedLite](#), [ExReleaseResourceLite](#)

Push Locks

- Similar to executive resources
- Advantages
 - Acquisition for shared access is faster than executive resources
 - Storage for EX_PUSH_LOCK can be allocated from paged or non-paged pool
 - Structure is smaller than ERESOURCE
- Disadvantages
 - Cannot be acquired recursively
 - No function to check if lock is currently exclusively acquired
- API
 - [FltInitializePushLock](#), [FltDeletePushLock](#)
 - [FltAcquirePushLockExclusive](#), [FltAcquirePushLockShared](#)
 - [FltReleasePushLock](#)

Asynchronous Procedure Call (APC)

- An object representing a callback function to be called by a specific thread
- APC forces a specified thread to run a specified routine in that thread
 - Makes the thread ready to run if it was waiting
 - At the end of the APC routine, the thread might go back to waiting (if it was in that state earlier)
- APC takes precedence over "normal" thread code
 - Only within that thread (does not affect its priority)
- Kernel mode API is exported but not documented

APC Types

- User mode APCs
 - Used for I/O completion callback routines (e.g. [ReadFileEx](#), [WriteFileEx](#))
 - Only deliverable when thread goes into "alertable state" (e.g. with [SleepEx](#), [WaitForSingleObjectEx](#), [MsgWaitForMultipleObjectsEx](#))
 - Can be explicitly queued with [QueueUserApc](#) (Win32)
- Normal kernel APCs
 - Run in kernel mode, at IRQL PASSIVE_LEVEL (0)
 - Preempts user mode code (and user mode APCs)
- Special kernel APCs
 - Run in kernel mode, at IRQL APC_LEVEL (1)
 - Used for I/O completion report from "arbitrary thread context"

Critical Regions and Guarded Regions

- Critical region
 - User mode APCs and normal kernel mode APCs disabled
 - Enter with [KeEnterCriticalSection](#)
 - Leave with [KeLeaveCriticalSection](#)
- Guarded region
 - All APCs are disabled
 - Enter with [KeEnterGuardedRegion](#)
 - Leave with [KeLeaveGuardedRegion](#)
- Calls to Enter function and Leave functions must match



Fast Mutex

- Also called Executive Mutex
- Not exported to the Windows API
- Similar to a normal mutex
 - Better performance
- Differences
 - After acquisition IRQL is raised to APC_LEVEL
 - No recursive acquisition allowed (no ownership)
- API
 - [ExAcquireFastMutex](#), [ExAcquireFastMutexUnsafe](#)
 - Unsafe version requires normal kernel mode APCs to be disabled (IRQL is APC_LEVEL)
 - [ExReleaseFastMutex](#), [ExReleaseFastMutexUnsafe](#)

Guarded Mutexes

- Introduced in Windows Server 2003
- Similar to a fast mutex
 - Starting with Windows 8, implemented the same as fast mutex
 - Prior to Windows 8, a bit faster than a fast mutex
- Represented by the `KGUARDED_MUTEX` structure

Kernel Synchronization Objects

Object	Exposed to device drivers	Exposed to user mode	Disables normal kernel APCs	Disables special kernel APCs	Supports recursive acquisition	Supports shared and exclusive acquisition
Kernel Mutex	Yes	Yes	Yes	No	Yes	No
Kernel Semaphore	Yes	Yes	No	No	No	No
Fast Mutex	Yes	No	Yes	Yes	No	No
Guarded Mutex	Yes (2K3+)	No	Yes	Yes	No	No
Push lock	No	No	No	No	No	Yes
Executive Resource	Yes	No	Yes	No	Yes	Yes



High-IRQL Synchronization

- “Normal” kernel code runs at IRQL `PASSIVE_LEVEL`
- Interrupts may intervene at any time
 - Their IRQL is always $>$ `DISPATCH_LEVEL`
- Accessing global kernel data must be synchronized
- Solution: raise the CPU IRQL to at least that of the potentially accessing interrupt
 - Fine for one CPU

CPU Synchronization

- IRQL >= DISPATCH_LEVEL
- Thread scheduler cannot interrupt the current thread on the CPU
 - So another thread cannot be scheduled on that CPU
- Another CPU may try to access the shared resource
- Synchronization in this case involves using IRQL changes and spin locks

The Spin Lock

- Synchronization on MP systems uses IRQLs within each CPU and spin locks to coordinate among the CPUs
- A spin lock is just a data cell in memory
 - It is accessed with a test and modify operation, atomic across all processors
- Not exposed (and not needed) to user mode applications



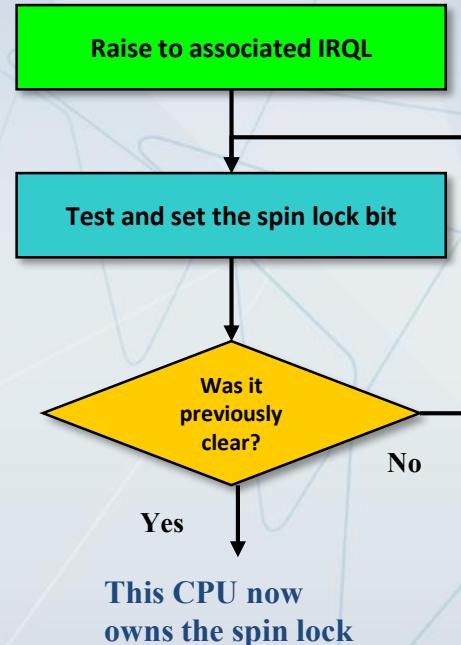
Windows Internals

Spin Lock Concepts

- Spin lock acquisition and release routines implement a one-owner-at-a-time algorithm
- Where do spin locks come from?
 - Some are defined in the system
 - Some are automatically associated with I/O devices and related objects
 - Device drivers can create and use additional spin locks
- A spin lock is either free or owned by a specific CPU
- A CPU should own a spin lock that protects shared data before manipulating it

Acquiring a Spin Lock

- IRQL is implicit in the choice of routine
 - [KeAcquireSpinLock](#) uses IRQL=DISPATCH_LEVEL
 - [KeAcquireSpinLockAtDpcLevel](#) does not change the IRQL
 - [KeSynchronizeExecution](#) and interrupt dispatcher use **Synclrql** found in interrupt object
 - [ExInterlockedXxx](#) routines use IRQL=HIGH_LEVEL
- Spin locks should not be requested if already owned
 - Causes a deadlock!



Spinlock “Busy Wait”

- Busy waiting causes repeated read attempts at the spinlock
- CPU is running continuously
- On Pentium 4 or newer
 - A new machine instruction (pause/yield) is inserted into such busy wait loops
 - Less power is consumed
 - Used as a hint to the CPU of a hyper threaded core, causing more CPU time to be allocated to the other core

Queued Spin Locks

- From Windows XP
- Efficient version of dispatch-level spin locks (IRQL DISPATCH_LEVEL)
- Ensures that the spin lock is acquired on a first-come first-serve CPU basis
- Check [KeAcquireInStackQueuedSpinLock](#) and [KeReleaseInStackQueuedSpinLock](#) in the WDK
- Useful when it's likely that more than one CPU will try to acquire the spin lock

Passive-Level Interrupt Service Routines

- Starting with Windows 8, a driver can use the [IoConnectInterruptEx](#) routine to register a passive-level Interrupt Service routine (ISR)
 - This can be used only for devices that communicate in low IRQL
 - Memory mapped I/O
- For example I²C
- This enables slow interrupt handling since the ISR execution time is very limited
- More Information [here](#)

System Worker Threads

- Threads created by Windows that reside under the System process
- Device drivers may create additional threads for their own use (by calling the function [PsCreateSystemThread](#))
- These threads always execute in kernel mode
- Work Item
 - A request by the system or a device driver to run a routine in IRQL 0
 - Carried out by a system worker thread
 - Requested by calling [ExQueueWorkItem](#) or [IoQueueWorkItem](#)

Worker Threads Creation

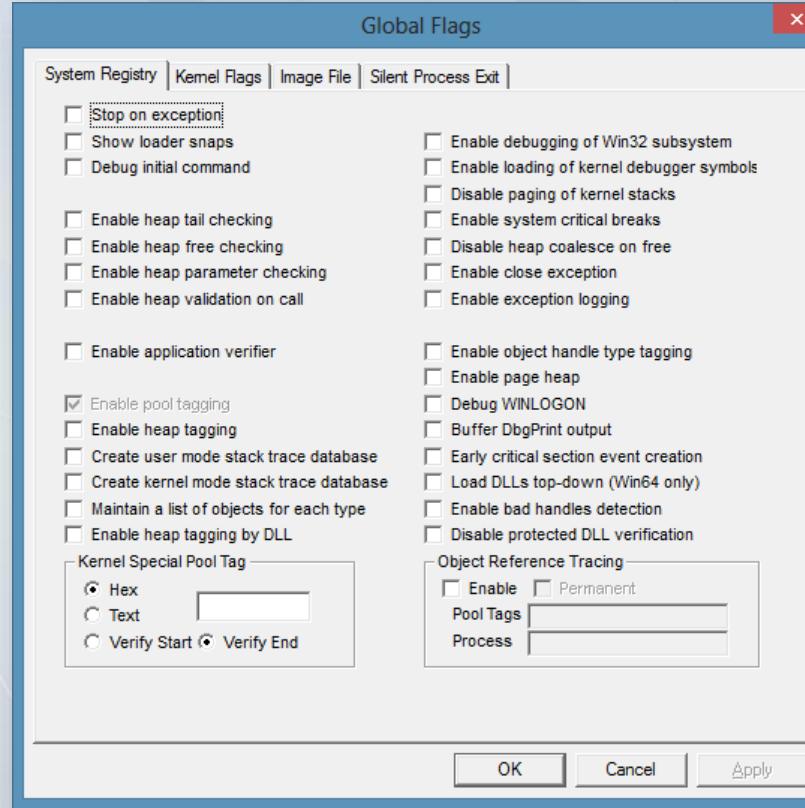
- Worker threads are created in three priority levels
 - Delayed worker threads (priority 12)
 - Process work items that are not considered time critical and can have their stack paged out while waiting for requests
 - Critical worker threads (priority 13)
 - Process time critical work items
 - On Server systems have their stack in memory at all times
 - Hypercritical worker thread (priority 15)
 - A single thread used to free terminated thread objects
- The number of critical and delayed worker threads is changing according to work item load
 - Initial number is based on OS version and the registry



Windows Global Flags

- Global flags maintained by the OS in a global variable named `NtGlobalFlags`, used for debugging, tracing and other information
 - Initialized from the registry at boot time
 - `HKEY_LOCAL_MACHINE\System\Control\Session Manager\GlobalFlags` value
- Each process can maintain its own global flags settings
- Flags can be changed with the [Gflags.exe](#) utility
(Windows SDK and Debugging Tools)

Global Flags



Kernel Event Tracing

- The kernel provides a mechanism to record data by kernel and user mode components, and exposed through [Event Tracing for Windows](#) (ETW)
- ETW entities
 - Provider
 - Produces event data
 - Defines a GUID for its event classes and registers them with ETW
 - Controller
 - Starts or stops logging sessions
 - Consumer
 - Selects trace sessions to receive data, either in real-time or from a log file

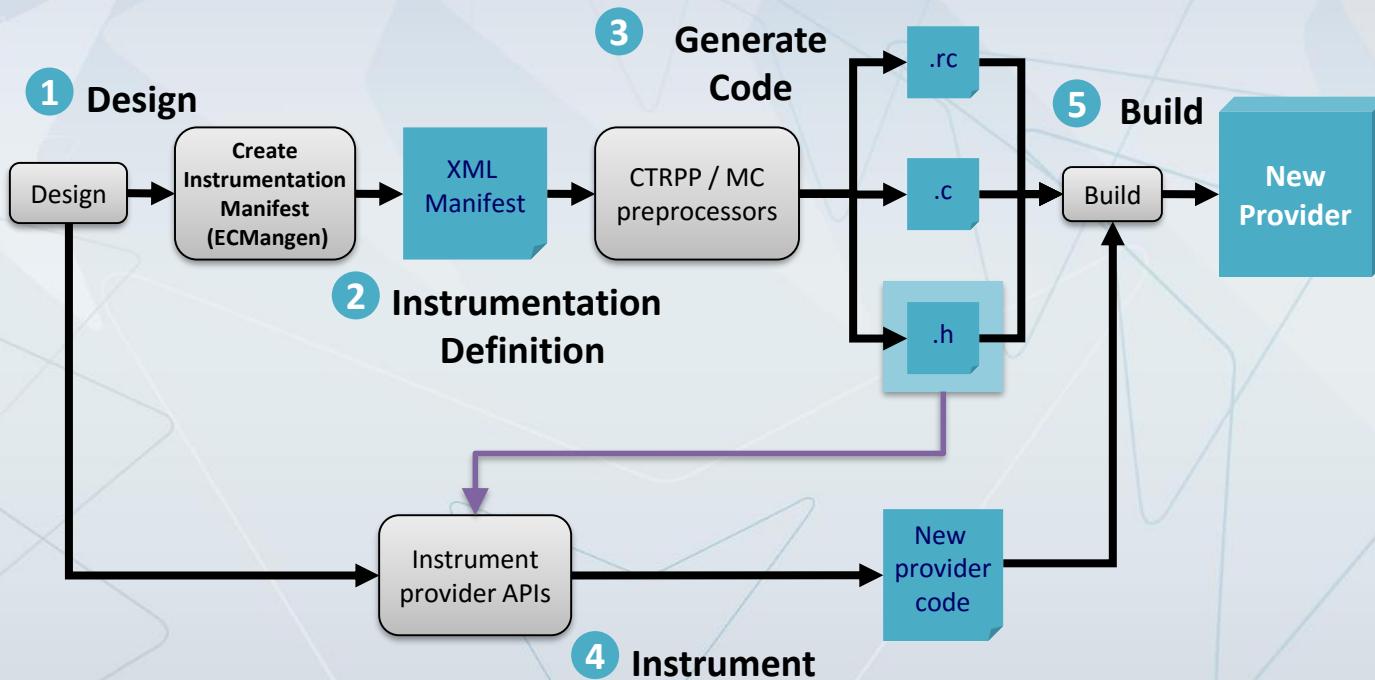
Example ETW Providers

- Disk I/O (disk class driver)
- File I/O (file system drivers)
- Hardware configuration (Plug & Play manager)
- Image Load/Unload (system image loader in the kernel)
- Page Faults (memory manager)
- Process create / delete (process manager)
- Registry activity (configuration manager)
- Context switches (kernel dispatcher)
- System calls (kernel dispatcher)
- Interrupts (kernel dispatcher)

Using ETW

- Can use the “Data Collector Sets” node in Performance Monitor
 - Can select provider(s)
 - Can start/stop a tracing session
- To generate a human readable output, use the **tracerpt.exe** command line tool
 - Can generate (e.g.) a CSV file, viewable in a text editor or MS Excel

Instrumentation Design Workflow



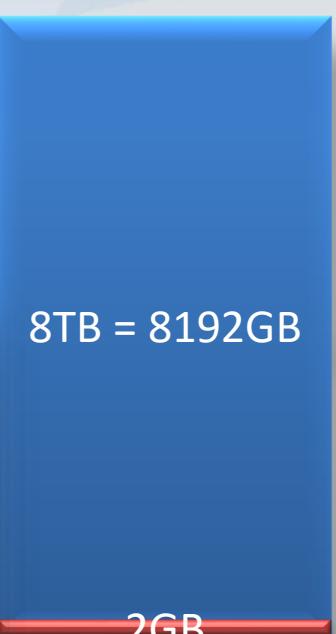
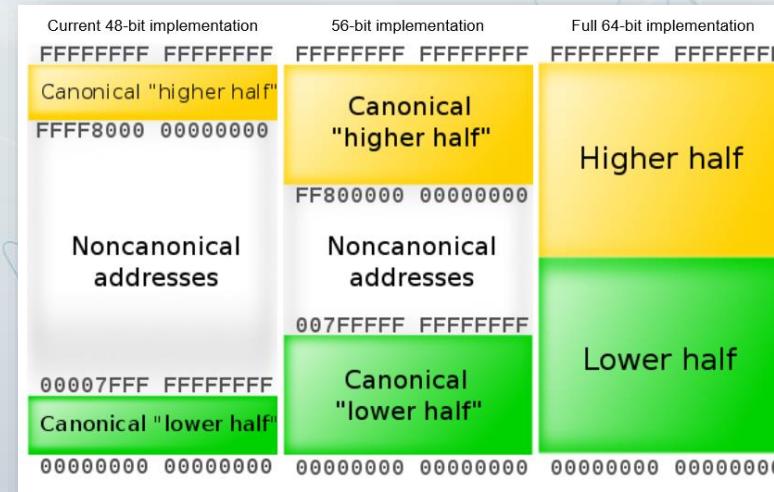
TraceLogging

- The new Windows 10 event tracing framework for user-mode applications and kernel-mode drivers
 - An abstraction on top of ETW
 - Tracing an event is as simple as an API call
 - Events are self-describing and do not require any additional binaries
 - All the metadata about the event is recorded in the event
 - Activities inside a single process can be easily expressed through start and stop events
 - TraceLogging is compatible with existing instrumentation support
 - The new ETW APIs continue to support the old providers
 - TraceLogging offers the same event tracing API for Windows 10, Xbox One, and Windows 10 Mobile
 - TraceLogging APIs are accessible from C, C++, .NET, and Windows Runtime

WINDOWS X64 & WOW 64

64?

- No more out of memory exceptions (address space)
 - 32 bit = 2GB + 2GB or 3GB + 1GB
 - 64 bit = 8TB + 8TB
- Actually x64 == 48bit today
- Canonical Form:
 - Bits 48-63 is signed extended of bit 47



x86-64 (A.K.A x64, x86_64 and amd64)

- The original specification was created by [AMD](#)
- Has been implemented by AMD, [Intel](#), [VIA](#), & others
- Fully backwards compatible with 16-bit and 32-bit x86 code
 - Existing x86 executables run with no compatibility or performance penalties
 - existing applications that are converted to x64 to take advantage of new features of the processor may achieve performance improvements

Windows 64

- First versions Server 2003/XP 64 (NT 5.2)
- Server 2008R2 is the first 64-bit only OS from MS
- No more VDM & Wow32
 - Virtualization (Hyper-V, VMWare, Oracle Virtual Box)
- Old (and new) x86 32bit applications run with no performance penalty
- 64bit development feels like 32bit development

Windows PE32+ File Format

- The Portable Executable is the format for all DLL, EXE, SYS, OCX, and even PDB files
- It has been changed to support x64 PE:
 - The file header Machine type has
IMAGE_FILE_MACHINE_[AMD64] | [IA64] | [ARM64]
- The optional header magic number determines whether an image is a PE32 or PE32+ executable.
- PE32+ images allow for a 64-bit address space while limiting the image size to 2 gigabytes

dumpbin /headers c:\windows\system32\notepad.exe
dumpbin /headers c:\windows\SysWow64\notepad.exe

DEMO

Wow64

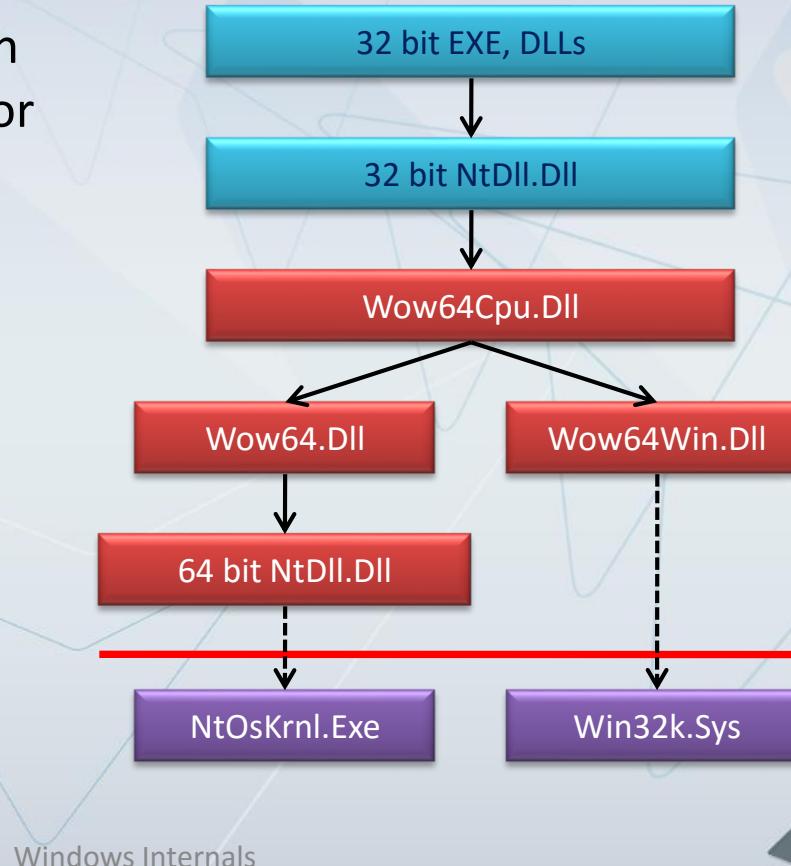
- Allows execution of Win32 binaries on 64-bit Windows
 - Wow64 intercepts system calls from the 32-bit application, Converts 32-bit data structure into 64-bit aligned structures and Issues the native 64-bit system call
- The [IsWow64Process](#) function can tell a 32-bit process if it is running under Wow64, in .NET:
 - [Environment.Is64BitOperatingSystem](#) && ![Environment.Is64BitProcess](#)
- Address space is 2GB or 4GB (if image is linked with the /LARGEADDRESSAWARE)
 - Use [Editbin](#) to edit the flag after compilation/linking
- Performance
 - On x64, 32 bit instructions executed by hardware
 - On IA64 (obsolete), instructions have to be emulated
- Device drivers must be native 64 bit

Wow64 Restrictions

- A 64 bit process cannot load a 32 bit DLL and vice versa
 - Except resource-only DLLs, which can be loaded cross-architecture
- Some APIs are not supported by Wow64 processes
 - E.g. [ReadFileScatter](#), [WriteFileGather](#), [AWE](#) functions

Wow64 Architecture

- Wow64.dll – core emulation infrastructure and thunks for Ntoskrnl.exe functions; exception dispatching
- Wow64win.dll – intercepts GUI system calls for Win32k.sys
- Wow64Cpu.dll – manages thread contexts, supports mode-switch instructions



Wow64 File System Redirection

- System directories names have not changed in 64 bit Windows (e.g. \Windows\System32 contains native 64 bit images)
- 32 bit applications must use their own directories
 - \Windows\System32 maps to \Windows\Syswow64
 - 32 bit apps installed in \Program Files (x86)
 - 64 bit apps installed in \Program Files
- Some directories are not redirected



Wow64 Registry Redirection

- Components trying to register as 32 bit and 64 bit will clash
- 32 bit components are redirected to the Wow64 registry node (`Wow6432Node`)
 - `HKLM\Software`
 - `HKEY_CLASSES_ROOT`
 - `HKCR\Software\Classes`
- New flags for Registry APIs allow access to the 64 bit or 32 bit nodes
 - `KEY_WOW64_64KEY` – open a 64 bit key
 - `KEY_WOW64_32KEY` – open a 32 bit key

Wow64 Registry Reflection

- Some keys used with COM (Component Object Model) are mirrored
 - HKLM\Software\Classes
 - HKLM\Software\Ole
 - HKLM\Software\Rpc
 - HKLM\Software\COM3
 - HKLM\Software\EventSystem
- HKLM\Software\Classes\CLSID is mirrored only for LocalServer32 keys only
 - i.e. Out of process servers (can be activated by 32 bit clients)
 - InProcServer32 keys are not mirrored

Avoid Potential Problems

- No NTVDM - 16-bit binaries are no longer supported
 - Except for some kinds of installers
- ALL drivers must be compiled and target 64-bit Windows
- 32-bit applications install to
 - "C:\Program Files (x86)"
- While 64 applications use "C:\Program Files"
 - Automatically handled by using
 - `SHGetFolderPath(CSIDL_PROGRAM_FILES)`
 - `Environment.GetFolderPath(Environment.SpecialFolder.ProgramFiles)`
 - Use `GetSystemWow64Directory` to get the SysWow64 (System32 of 32 bit applications)



IPC Wow64 and 64 bit processes

- It just works!!!
 - Sharing Kernel Objects (Handles)
 - Using HWND and Windows Messages
 - Using RPC, .NET Remoting, WCF, Sockets
 - D/COM (with x86/x64) proxy/stub
 - CreateProcess can create 32/64 bit process
 - Shared Memory (Content has to be pointer size independent)

Programming x64

- [Programming Guide for 64-bit Windows](#)
- 64bit development feels like 32bit development
- For native developer:
 - Pointer size has changed
 - Data alignment is even more important
 - Packing can waste a lot of memory
 - Pointer truncating is a real headache
 - Polymorphic Types and API helps targeting both 32 and 64 bit

X86/X64 Programming Model

Data model	char	short	int	long	long long	pointers	OSes
ILP32	8	16	32	32	64?	32	Windows, Linux, Solaris, OS X, *BSD, AIX, HP-UX, other UN*Xes
							Linux, Solaris, OS X, *BSD, AIX,

sizeof(size_t) != sizeof(int) and sizeof(size_t) != sizeof(long)

LLP64	8	16	32	32	64	64	Windows
-------	---	----	----	----	----	----	---------



Size Does Matter

- Pointers and variable that may hold pointer data (such as WPARAM & LPARAM) has changed to 64 bit
- Non-pointer types such int, long, DWORD, stay the same (32 bit)
- If you care about the exact size use:
 - DWORD32, DWORD64, INT32, INT64, LONG32, LONG64, UINT32, UINT64, ULONG32, ULONG64, POINTER_32, POINTER_64
- Polymorphic size types:
 - Void *, INT_PTR, UINT_PTR

Windows Header File are Type Polymorphic

- `LRESULT CALLBACK WndProc(
 HWND hWnd, UINT message, WPARAM wParam, LPARAM lParam)`
- From ...\\Include\\shared\\minwindef.h:
 - `/* Types use for passing & returning polymorphic values */`
 - `typedef UINT_PTR WPARAM;`
 - `typedef LONG_PTR LPARAM;`
 - `typedef LONG_PTR LRESULT;`
- From ...\\Include\\shared\\basetsd.h:
 - Next Page

Windows Header File are Type Polymorphic

- From ...\\Include\\shared\\basetsd.h:

```
#if defined(_WIN64)
    typedef __int64 INT_PTR, *PINT_PTR;
    typedef unsigned __int64 UINT_PTR, *PUINT_PTR;
    typedef __int64 LONG_PTR, *PLONG_PTR;
    typedef unsigned __int64 ULONG_PTR, *PULONG_PTR;
    #define __int3264 __int64
#else
    typedef _W64 int INT_PTR, *PINT_PTR;
    typedef _W64 unsigned int UINT_PTR, *PUINT_PTR;
    typedef _W64 long LONG_PTR, *PLONG_PTR;
    typedef _W64 unsigned long ULONG_PTR, *PULONG_PTR;
    #define __int3264 __int32
#endif
```

Predefined Macros

- The compiler defines these macros:

Macro	Meaning
_WIN64	A 64-bit platform.
_WIN32	A 32-bit platform. This value is also defined by the 64-bit compiler for backward compatibility.
_WIN16	A 16-bit platform
_M_IA64	Intel Itanium platform
_M_IX86	x86 platform
_M_X64	x64 platform

Data Alignment

- Allocate on aligned boundaries
 - Use `__declspec(align)` to align data to specific boundary
 - Example: `__declspec(align) int abc;`
- Use the C++ 11 `alignas` new keyword
- Can use dynamically with `aligned malloc()` family of runtime routines
- x64 handles alignment faults in hardware

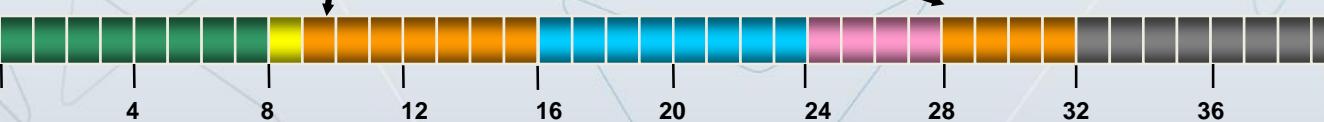


Be ware of Padding (Packing Difference)

32bit

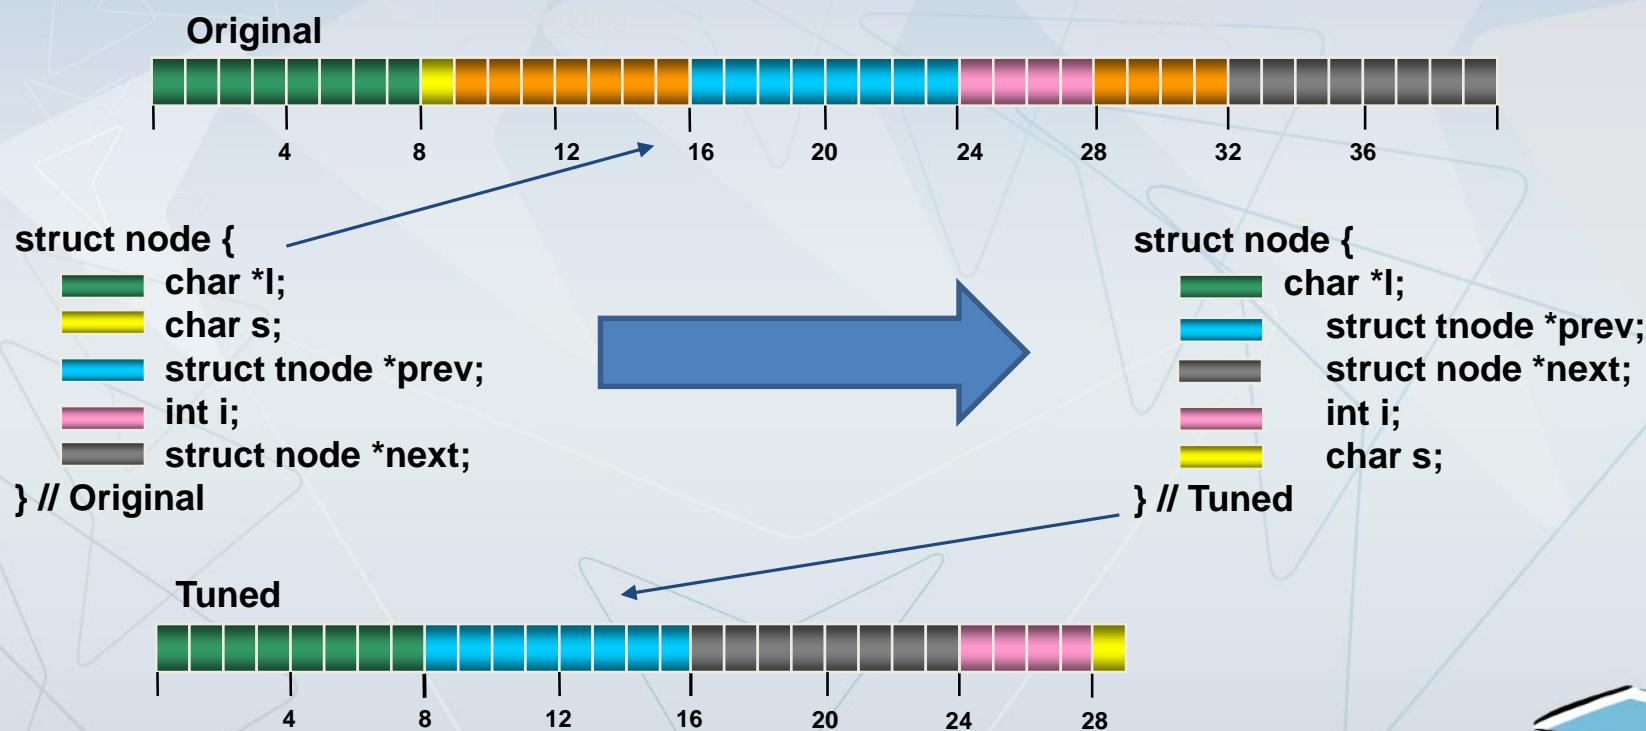


64bit



```
struct node {  
    char *l;  
    char s;  
    struct tnode *prev;  
    int i;  
    struct node *next;  
}
```

Smart Reordering Types Can Save a lot of Space



How to Correctly Cast a Pointer to int in a 64-bit Application

- The most general answer is – in no way
 - Putting a 64-bit pointer into a 32-bit variable causes cutting of high-order bits and therefore incorrect program behavior
- However, there is a specific case when you may store a pointer in 32-bit types
 - HANDLE, HWND, HBRUSH, Etc.
- Use the following helper function to safe convert:

Helper Functions

- `void * Handle64ToHandle(const void * POINTER_64 h)`
- `void * POINTER_64 HandleToHandle64(const void *h)`
- `long HandleToLong(const void *h)`
- `unsigned long HandleToUlong(const void *h)`
- `void * IntToPtr(const int i)`
- `void * LongToHandle(const long h)`
- `void * LongToPtr(const long l)`
- `void * Ptr64ToPtr(const void * POINTER_64 p)`
- `int PtrToInt(const void *p)`
- `long PtrToLong(const void *p)`
- `void * POINTER_64 PtrToPtr64(const void *p)`
- `short PtrToShort(const void *p)`
- `unsigned int PtrToUint(const void *p)`
- `unsigned long PtrToUlong(const void *p)`
- `unsigned short PtrToUshort(const void *p)`
- `void * UIntToPtr(const unsigned int ui)`
- `void * ULongToPtr(const unsigned long ul)`

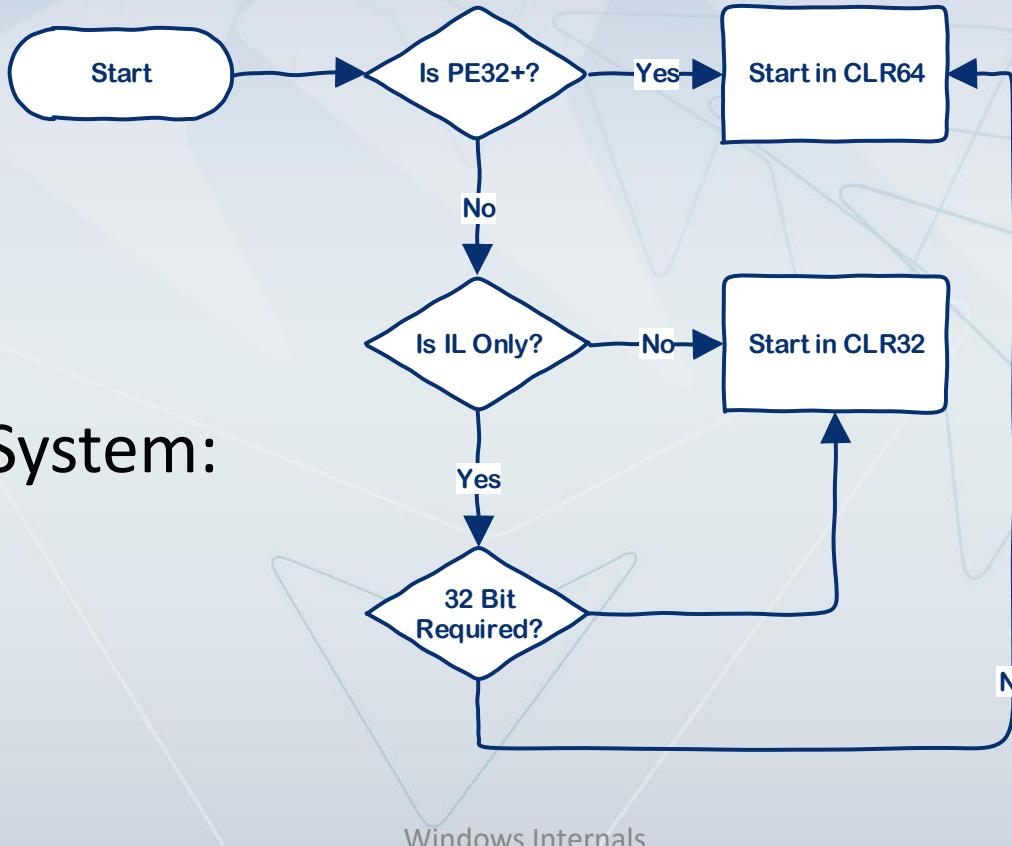
X64 Managed Applications

- You can choose to target Any-CPU, X86 or x64
- For assemblies (DLL) choose any-cpu unless you have a specific native code dependency or the assembly contains native code (C++/CLI)
- For executable, use specific hardware if you depend on native implementation (native DLL, P/Invoke, etc.)
 - Otherwise use Any-CPU
- The CorFlags.exe utility can show/change target CPU architecture

.NET Framework

- On x64 System there are two versions of each .NET framework (from framework 2.0)
 - A 32 bit framework (C:\Windows\Microsoft.NET\Framework)
 - V1.0, V1.1, V2.0, V3.0, V3.5 V4.0, V4.5, V4.5.2, V 4.5.1, V4.6, V4.6.1, V4.6.2
 - A 64 bit framework (C:\Windows\Microsoft.NET\Framework64)
 - V2.0, V3.0, V3.5 V4.0, V4.5, V4.5.2, V 4.5.1, V4.6, V4.6.1, V4.6.2
 - The GAC and NGEN also have different locations for Any-CPU, x86, x64

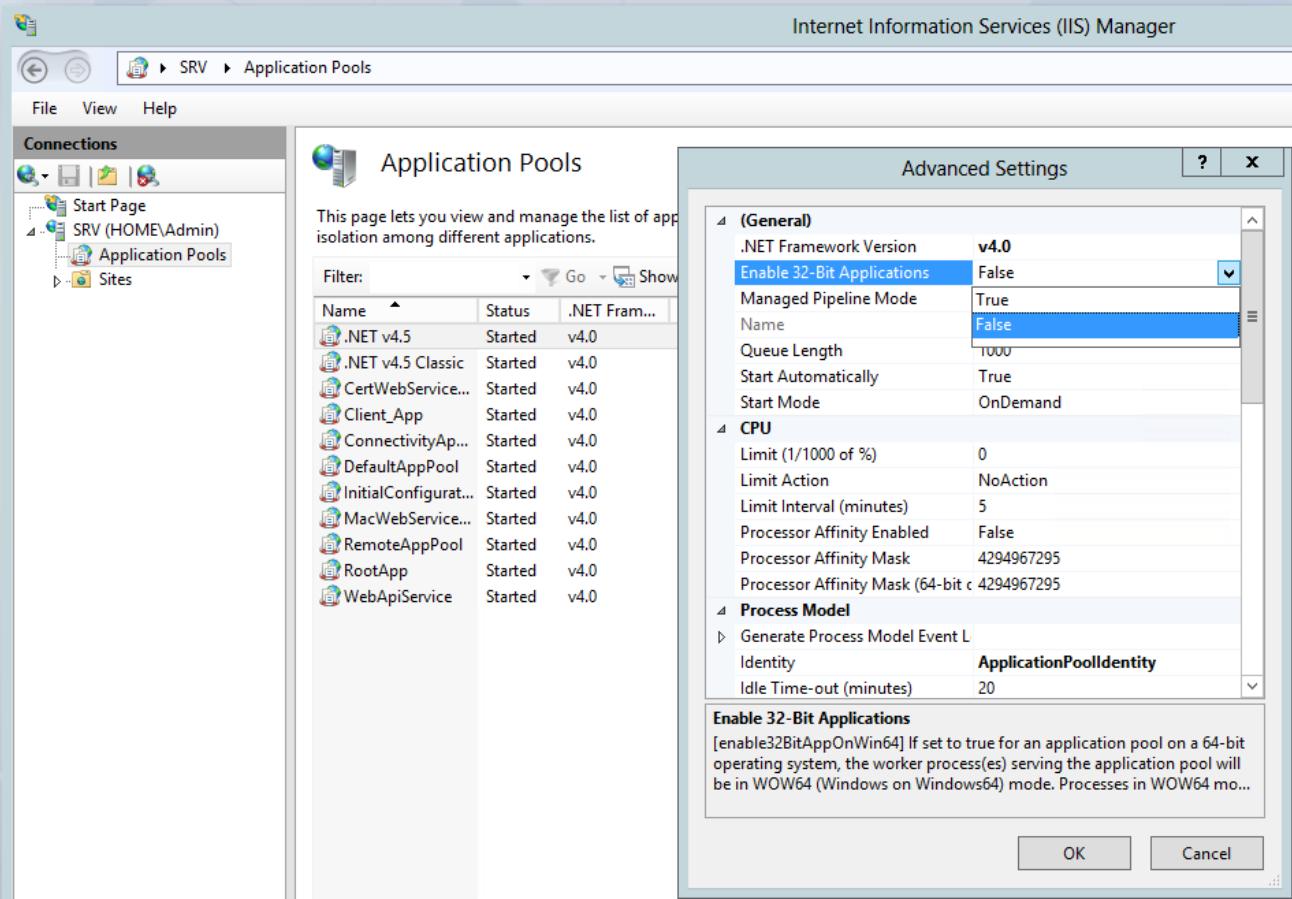
Which CLR Host My Application



- On x64 System:

IIS Settings

- .NET Sites and services are hosted in IIS under the application pool process
 - Choose the application pool bit-ness to match your needs



.NET X64 and Interop

- System.IntPtr & System.UIntPtr are polymorphic types
- Structures and Classes Layout has to be matched to the Native x86/x64 structure
 - Use polymorphic .NET types and/or the [StructLayout] attribute
- Be aware of x64 when you use Windows resources such as registry, shared memory, file-system

X64 Summary

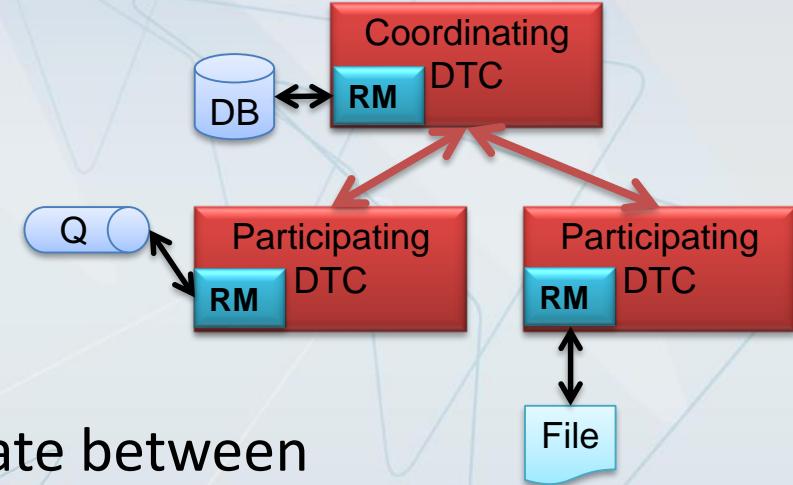
- We learned about Windows x64
 - The new hardware architecture
 - Windows 32 on Windows 64
 - The new PE32+ file format
 - Native development considerations
 - X64 .NET framework

Transaction

- Atomic transaction is a reliability principle
 - It is an all or nothing operation
 - the system will be in a good state after the execution of the transaction
- Transaction follows the ACID criteria:
 - Atomicity – all or nothing
 - Consistency – Application moves from one consistent state to another consistent state
 - Isolation – No transaction can “see” or “interfere” with another
 - Durability – Committed transaction will be committed!!!
- Transaction timeouts even solve deadlocks

Distributed Atomic Transaction

- A transaction with many parties, sometime across process and machine boundaries
 - Two phase commit protocol
- Transaction Roles
 - Transaction Manager (TM)
 - Microsoft Transaction Server,
Kernel Transaction Manager
 - Managers communicate between themselves
 - Resource Manager (RM)
 - The client of the TM, responsible to the state of the resource



Two Phase Commit

**Transaction
Manager**

Send Prepare to all RMs

Prepare

Force write "Commit"
record to TM log

Prepared

Lazy write "Committed"
record to TM log

Commit

**Resource
Manager**

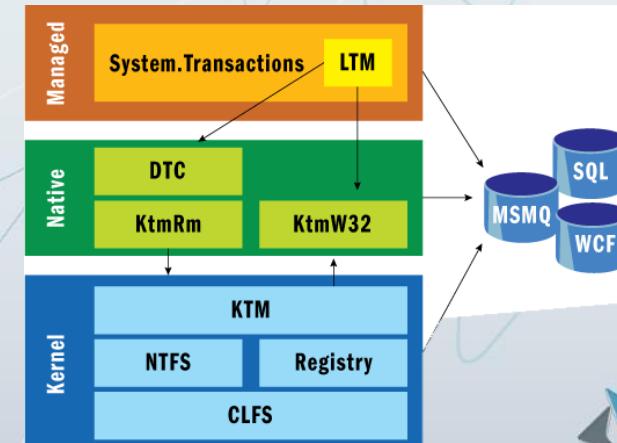
Prepare and Force write
"Prepared" record to RM log

Lazy write "Committed" record
to RM log and Release locks

Committed

Windows KTM

- A kernel based implementation of a distributed transaction system
 - Can participate in any MTS based transaction
 - Provide API to create your own Resource Manager
 - Two built in Kernel RMs:
 - Transactional NTFS
 - Transactional Registry
 - Play nice with the kernel handle concept



Creating Transaction

- CreateTransaction

- Returns handle to a kernel transaction object
- You can set the transaction timeout value
- You can ask for a non distributed transaction
- For distributed transaction:
 - you need to add the DTC's SID in the Security Descriptor that you can set using the **IpTransactionAttributes**
- CloseHandle closes the transaction handle
 - If you close the last handle without committing the transaction, the KTM will rollback the transaction

- OpenTransaction, GetTransactionId

- Opens transaction handle from transaction Id

Commit & Rollback

- To commit a transaction call [CommitTransaction](#)
- To Rollback a transaction call [RollbackTransaction](#)
 - Or close the handle without any commit
- To successfully commit a transaction, no party should call to Rollback and the transaction should be committed on time



Kernel Resource Managers

- KTM provides all you need to develop your own resource manager
- However in most cases you will use the KTM with one of the two existing managers
 - The Transactional NTFS
 - The Transactional Registry
- We will start with the registry, and then we will dive into transactional NTFS

Transactional Registry

- RegCreateKeyTransacted
 - Creates a registry key and associates it with a transaction
 - If the key already exists, the function opens it
- RegDeleteKeyTransacted
 - Deletes a sub-key and its values from the registry as a transacted operation
- RegOpenKeyTransacted
 - Opens the specified registry key and associates it with a transaction
- Having registry keys under a transaction means that all registry operations will be committed or rolled back

Transactional NTFS

- With transactional NTFS you can create a distributed transaction with a scope that includes a database, communication channels and the file system
 - Using TxNTFS you can update a UNC entry in a data base while creating this file – all in one atomic transaction
- There are two API types for TxNTFS
 - Old API that takes KTM handle
 - [ReadFile](#), [WriteFile](#), ...
 - New API with the *Transacted, that takes a transaction handle
 - [MoveFileTransacted](#)

TxNTFS Isolation

- TxNTFS is optimized for success
 - The Resource Manager makes the changes
 - Changes are hidden from non-transactional reader
 - At commit the Resource Manager make the changes available
 - Non transactional writer can't create a handle if the file is already open for transactional writer and vice versa
 - Transactional Reader can choose to see the dirty or committed state
 - [CreateFileTransacted\(...pusMiniVersion\)](#)

TXFS_MINIVERSION_*	Description
COMMITTED_VIEW	The view of the file as of its last commit
DIRTY_VIEW	The view of the file in the transaction (modify)
DEFAULT_VIEW	A transaction that is modifying the file gets the dirty view. A transaction that is not modifying the file gets the committed view.

Distributed Transaction

- Kernel transaction can take part with DTC based transaction
 - Combine TxNTFS & Transactional Registry with SQL Server, Oracle DB, MSMQ, IBM MQSeries
 - The famous scenario of using the file system to store large files and keeping its path in the DB becomes reliable
- Instead of creating kernel transaction handle
 - Ask the handle from the DTC
 - Query [IkernelTransaction](#) from the [ITransaction](#) interface
 - **HRESULT [GetHandle](#)(HANDLE * pHandle);**

IKernelTransaction (Native)

```
#include "stdafx.h"
#include <Windows.h>
#include <ktmw32.h>
#include <TxDtc.h>
#include <xolehlp.h>
#pragma comment (lib, "ktmw32.lib")
#pragma comment (lib, "xolehlp.lib")
int _tmain(int argc, _TCHAR* argv[])
{
    CoInitialize(NULL);
    ITransactionDispenser *pITransactionDispenser;
    ITransaction *pITransaction;
    IKernelTransaction *pKernelTransaction;
    HANDLE hTransactionHandle;
```

IKernelTransaction (Native)

```
HRESULT hr = DtcGetTransactionManagerEx( NULL, NULL,  
    IID_ITransactionDispenser, OLE_TM_FLAG_NONE,  
    NULL, (void**) &pITransactionDispenser);  
  
hr = pITransactionDispenser->BeginTransaction(  
    NULL, ISOLATIONLEVEL_READCOMMITTED,  
    ISOFLAG_RETAIN_BOTH, NULL, &pITransaction);  
  
hr = pITransaction->QueryInterface(IID_IKernelTransaction,  
    (void**) &pKernelTransaction);  
hr = pKernelTransaction->GetHandle(&hTransactionHandle);
```

IKernelTransaction (Native)

```
HANDLE hAppend = CreateFileTransacted(L"test.txt",  
    FILE_APPEND_DATA, FILE_SHARE_READ,  
    NULL, OPEN_ALWAYS, FILE_ATTRIBUTE_NORMAL, NULL,  
    hTransactionHandle, NULL, NULL);
```

```
DWORD nWritten;  
wchar_t text[] = L"Transaction Rocks!";  
WriteFile(hAppend, text, sizeof(text), &nWritten, nullptr);
```

IKernelTransaction (Native)

```
hr = pITransaction->Commit(FALSE, XACTTC_SYNC_PHASEONE, 0);
```

```
CloseHandle(hAppend);  
pKernelTransaction->Release();  
pITransaction->Release();  
pITransactionDispenser->Release();  
CoUninitialize();  
return 0;  
}
```

KTM & .NET

- .NET does not support KTM out of the box
 - However we can Interop to get [IKernelTransaction](#)
 - Using [System.Transaction interop mechanisms](#)
 - Once we have the [IKernelTransaction GetHandle](#) method, we can use a KTM transaction handle
- Interop to [CreateFileTransacted](#) can provide a transacted file handle
 - .NET stream supports creating a stream from a Win32 file handle
- Using **System.Transaction** + KTM enables strong scenarios such as distributed WCF based transaction with Databases and TxNTFS.

KTM Summary

Transaction == Retry Boundaries

- Windows NT 6.x provides new transaction capabilities such as KTM, TxRegisry, TxNTFS
- KTM can take part in DTC transaction
- Using System.Transaction + Interop, we can have a KTM transaction within .NET code

Enhanced I/O in NT 6.x

- NT 6.X provides new I/O capabilities
 - I/O Priority
 - I/O Bandwidth reservation
 - Better (asynchronous) I/O cancelation
 - A new ability to mount a Virtual Hard Drive
 - A new ability to boot from a VHD file

I/O Priority

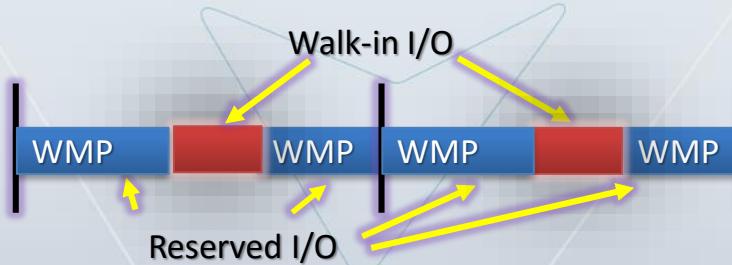
- In Windows I/O is bound to the thread that issues the API
 - Each thread has a list of pending I/O Request Packets ([IRP](#))
- Windows NT 6.x has added I/O priority level to the IRP, stored in the Flags field. Drivers should support this flag
 - There are 5 levels:
 - **Critical, High, Normal, Low, Very Low**
 - **High** not implemented
 - **Critical** is only for use by memory manager
- Windows execute at least one of **Low** or **Very Low** I/O request every second

I/O Priority APIs

- There are several methods to set the I/O priority:
 - Setting to thread or process to “Background Mode”
 - [SetThreadPriority](#), [SetPriorityClass](#)
 - **THREAD_MODE_BACKGROUND_BEGIN**
 - **THREAD_MODE_BACKGROUND_END**
 - **PROCESS_MODE_BACKGROUND_BEGIN**
 - **PROCESS_MODE_BACKGROUND_END**
 - Providing a hint for the file handle
 - [SetFileInformationByHandle](#)
 - With [FILE IO_PRIORITY_HINT INFO](#)
 - Using Multimedia Class Scheduling Service
 - Using Kernel call: [IoSetIoPriorityHint](#)

I/O Bandwidth Reservation

- Streaming applications can request I/O bandwidth guarantees
- Specified on individual files
- I/O system reports back to application
 - Optimal transfer size
 - Number of outstanding I/Os they should maintain



Reserving Bandwidth for Streaming

- An application that requires throughput of 200 bytes per second from the disk would make the following call:

```
SetFileBandwidthReservation( hFile, 1000, 200,  
    FALSE, &transferSize, &outstandingRequests );
```

- **TransferSize**: The minimum size of the I/O request
- **OutstandingRequests**: The number of outstanding TransferSize chunks. For maximum throughput
- To get the current setting
 - call the GetFileBandwidthReservation API

I/O Cancellation Support

- Until NT 6.x, opens could not be cancelled
 - Example: You browse to an off-line network share in a File Save dialog and hang for the duration of the network timeout
- In NT 6.x, opens and other synchronous I/O can be cancelled
 - [CancelSynchronousIo](#) cancels a pending synchronous I/O issued by another thread
 - [CancelIoEx](#) permits canceling I/Os from any thread
 - The old [CancelIo](#) could only cancel all I/Os issued by the calling thread, hence only for asynchronous operations
- Windows 7/2008 file open/save dialogs all implement cancellation

Summary

- Windows is an object based system managed by the Object Manager executive component
- Synchronization primitives exist for inter processor and inter thread scenarios
- The scheduler only works when IRQL is less than DISPATCH_LEVEL
- Transaction provides retry boundaries
- Windows 6.x solved and enhanced I/O operations

Module 3

MANAGEMENT & DIAGNOSTICS MECHANISMS

Windows Internals

275



Agenda

- The Registry
- Windows Services
- Windows Management Instrumentation
- Windows Error Report
- Summary

The Registry

- Hierarchical repository of system / user configuration data
 - Some stored in files, some built dynamically and stored in memory
- Access
 - REGEDIT.EXE
 - Originally written for Windows 95
 - Enhanced significantly in Windows 2000 and XP
 - Programmatically
 - APIs in Win32 and the Kernel
- Activity
 - Can watch with Process Monitor from SysInternals

Registry: The Hives

- HKEY_LOCAL_MACHINE (HKLM)
 - Contains machine specific configuration (not user related)
- HKEY_CURRENT_USER (HKCU)
 - Contains per-user information
- HKEY_CLASSES_ROOT (HKCR)
 - Contains file extension associations and COM (Component Object Model) registration
- HKEY_USERS
 - Contains sub-keys for each user ever logged on to the system
- HKEY_CURRENT_CONFIG
 - Contains current hardware configuration
- HKEY_PERFORMANCE_DATA
 - Contains performance counters data
 - Not really part of the registry

HKEY_CURRENT_USER

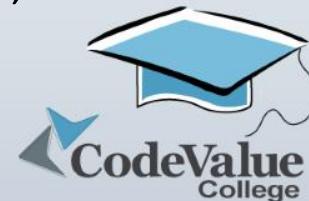
- Contains user specific information
- Maintained in a file named *NtUser.Dat* stored under \Documents and Settings\<user name>
- Sub-keys
 - AppEvents – sound/event associations
 - Console – command window settings
 - Control Panel – screen saver, desktop scheme, keyboard and mouse settings, etc.
 - Environment – environment variable definitions
 - Keyboard Layout
 - Network – network drives mappings
 - Printers – printers connections settings
 - Software – user specific software preferences

HKEY_CLASSES_ROOT

- Contains file extension association (used by the shell – Explorer.Exe) and COM servers registration (classes, interfaces, type libraries, applications, etc.)
- Actual data comes from two sources
 - HKLM\Software\Classes
 - Per user class registration in HKCU\Software\Classes stored in \Users\<username>\LocalSettings\Application Data\Microsoft\Windows\UsrClass.dat
 - The addition of per-user classes is new to Windows 2000
- Sub-keys
 - CLSID – COM classes registration
 - Interface – COM interfaces proxy\stub registration
 - TypeLib – Registered type libraries
 - AppID – COM server applications registration (not the same as COM+ applications)

HKEY_LOCAL_MACHINE

- Contains information relevant to the machine regardless of the logged on user
- Sub-keys
 - Hardware – contains device descriptions detected during the boot process
 - SAM – contains local users and group information
 - Actually a link to HKLM\Security\SAM
 - Security – system-wide security policy and user rights assignments
 - Software – System-wide configuration for system boot as well as third party software settings (directories, passwords, etc.)
 - System – system-wide configuration needed to boot the system, such as device drivers and Win32 services to load



Hive Store Path

- Non volatile registry data is stored in files
 - HKLM\System
 - System32\Config\System
 - HKLM\SAM
 - System32\Config\SAM
 - HKLM\Security
 - System32\Config\Security
 - HKLM\Software
 - System32\Config\Software

Registry Data Types

- Most useful types:

Data type	Description
REG_SZ	Null terminated, fixed length, Unicode string
REG_MULTI_SZ	Multiple null-terminated Unicode strings
REG_EXPAND_SZ	Variable length Unicode string that can hold embedded environment variables
REG_BINARY	Arbitrary length binary data
REG_DWORD	32 bit number
REG_DWORD_BIG_ENDIAN	32 bit number in big endian format (high byte first)
REG_LINK	Symbolic link to another key
REG_QWORD	64 bit number

Registry Access APIs



- User mode (Windows API)
 - [RegCreateKeyEx](#), [RegOpenKeyEx](#), [RegCloseKey](#)
 - [RegDeleteKey](#)
 - [RegSetValueEx](#), [RegQueryValueEx](#),
 - [RegEnumKeyEx](#), [RegEnumValue](#)
- Kernel mode
 - [ZwCreateKey](#), [ZwOpenKey](#), [ZwClose](#)
 - [ZwDeleteKey](#)
 - [ZwSetValueKey](#), [ZwQueryValueKey](#)
 - [ZwEnumerateKey](#), [ZwEnumerateValueKey](#)

Windows Services

- A Service is a program that provides some service without being tied to the logged on user
 - A service may run without any user logging in
- Service examples
 - IIS service (listening on HTTP port 80, etc.)
 - SQL server (listening for database requests via some mechanism)

Service Communication

- A service must register itself with the Service Control Manager (SCM)
- The SCM is responsible for starting, stopping and otherwise manipulating services
 - Communicate via a named pipe
- A Service Control Application (SCP) may start, stop and otherwise control a service by issuing commands to the SCM

Service Configuration

- A service application must be installed
 - By calling the [CreateService](#) API
- Inserts a new key into the registry under
HKLM\System\CCS\Services
 - Technically, the entries in that key correspond to services and device drivers
- Use the Services MMC plug-in to view service only information

Service (& Drivers) Registry Key

Registry Editor

File Edit View Favorites Help

Computer

HKEY_CLASSES_ROOT

HKEY_CURRENT_USER

HKEY_LOCAL_MACHINE

BCD00000000

HARDWARE

SAM

SECURITY

SOFTWARE

SYSTEM

ActivationBroker

ControlSet001

CurrentControlSet

Control

Enum

Hardware Profiles

NXPNFCcontroller

Policies

Services

.NET CLR Data

.NET CLR Networking

.NET CLR Networking 4.0.0.0

.NET Data Provider for Oracle

.NET Data Provider for SqlServer

.NET Memory Cache 4.0

.NET Framework

(31B3A958-3A7E-4D4D-850A-D36799A574BD)

(4E978C8-C0CB-4691-A3C9-76755A46B49C)

(7B32E89D-30AE-402B-9511-357113D6F5C1)

(E57BD837-74DB-48E6-BF11-B1D88D3E3A52)

(ED98B03E-CFDC-4106-BE2D-278CDC32B7CE)

1394ohci

3ware

ACPI

acpiasl

AcpiDev

acpix

acpipagr

AcpiPmi

acpitime

AdobeARMservice

ADQVMPPackage

ADP80XX

adsi

AFD

ahcache

AJRouter

ALG

...

Name	Type	Data
(Default)	REG_SZ	(value not set)
Description	REG_SZ	@%SystemRoot%\system32\AJRouter.dll,-1
DisplayName	REG_SZ	@%SystemRoot%\system32\AJRouter.dll,-2
ErrorControl	REG_DWORD	0x00000001 (1)
FailureActions	REG_BINARY	80 51 01 00 00 00 00 00 00 00 00 03 00 00 00 14 0...
ImagePath	REG_EXPAND_SZ	%SystemRoot%\system32\svchost.exe -k LocalSer...
ObjectName	REG_SZ	NT AUTHORITY\LocalService
ServiceSidType	REG_DWORD	0x00000001 (1)
Start	REG_DWORD	0x00000003 (3)
Type	REG_DWORD	0x00000020 (32)

Computer\HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Services\AJRouter

Services MMC Snap-in

Services

File Action View Help

Services (Local)

AllJoyn Router Service

[Start the service](#)

Description:
Routes AllJoyn messages for the local AllJoyn clients. If this service is stopped the AllJoyn clients that do not have their own bundled routers will be unable to run.

Name	Description	Status	Startup Type	Log On As
ActiveX Installer (AxInstSV)	Provides Us...	Manual	Local System	
Adobe Acrobat Update Serv...	Adobe Acro...	Running	Automatic	Local System
AllJoyn Router Service	Routes AllJo...	Manual (Trigger Start)	Local Service	
App Readiness	Gets apps re...	Manual	Local System	
Application Identity	Determines ...	Manual (Trigger Start)	Local Service	
Application Information	Facilitates t...	Running	Manual (Trigger Start)	Local System
Application Layer Gateway ...	Provides su...	Manual	Local Service	
Application Management	Processes in...	Manual	Local System	
AppX Deployment Service (...)	Provides inf...	Manual	Local System	
ASP.NET State Service	Provides su...	Manual	Network Service	
Auto Time Zone Updater	Automatica...	Manual (Trigger Start)	Local Service	
Background Intelligent Tran...	Transfers fil...	Manual	Local System	
Background Tasks Infrastru...	Windows in...	Running	Automatic	Local System
Base Filtering Engine	The Base Fil...	Running	Automatic	Local Service
BitLocker Drive Encryption ...	BDESVC hos...	Manual (Trigger Start)	Local System	
Block Level Backup Engine ...	The WBENG...	Manual	Local System	
Bluetooth Handsfree Service	Enables wir...	Manual (Trigger Start)	Local Service	
BranchCache	This service ...	Running	Manual	Network Service
bthserv	<Failed to R...	Running	Manual (Trigger Start)	Local Service
CDPUserSvc_2bd010		Running	Automatic (Delayed Start)	Local System
Certificate Propagation	Connec user...	Running	Manual	Local System

Extended Standard



Service Key Parameters

- Start
 - When to start the service / device driver?
- ErrorControl
 - What to do if starting fails?
- Type
 - Device driver or service?
- Group, Tag
 - Used to control load order within a start group
- ImagePath
 - The path of the service / device driver

Service Key Parameters

- **DependOnService**, **DependOnGroup**
 - Dependencies of the driver / service
- **ObjectName**
 - Account to run service in
- **DisplayName**
 - Name displayed in Services MMC
- **Description**
 - Service description displayed in Services MMC

Service Key Parameters

- **FailureActions**
 - Actions the SCM should take if service process crashes
- **FailureCommand**
 - Command line program to be executed, if appropriate failure action selected
- **Security**
 - Security descriptor of service / driver object

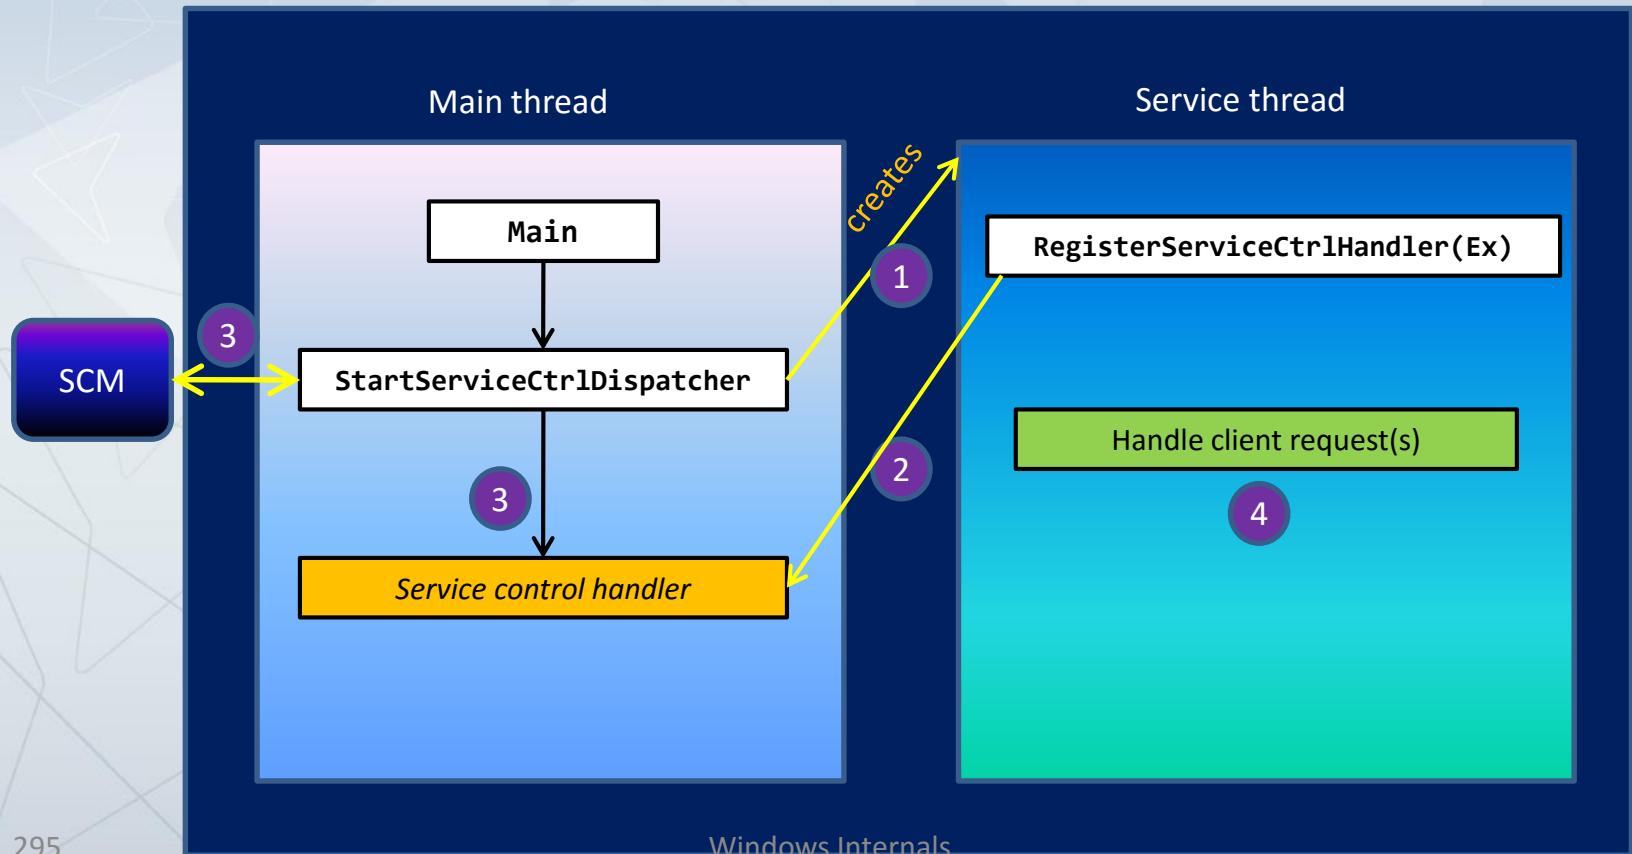
Service Key Parameters

Name	Value	Description
Start	SERVICE_BOOT_START (0)	Driver is loaded very early in the boot process
	SERVICE_SYSTEM_START (1)	Driver is loaded during kernel initialization
	SERVICE_AUTO_START (2)	Drivers and services started by the SCM after it loads (Services.exe)
	SERVICE_DEMAND_START (3)	Drivers and services that load on demand using the StartService API
	SERVICE_DISABLED (4)	Driver or service is not loaded
ErrorControl	SERVICE_ERROR_IGNORE (0)	Error is ignored
	SERVICE_ERROR_NORMAL (1)	A warning is displayed
	SERVICE_ERROR_SEVERE (2)	If last known good isn't used, the system boots with last known good configuration
	SERVICE_ERROR_CRITICAL (3)	If last known good not used, boot with last known good. If already in last known good – crash the system with a blue screen

Service Key Parameters

Name	Value	Description
Type	SERVICE_KERNEL_DRIVER (1)	Kernel device driver
	SERVICE_FILE_SYSTEM_DRIVER (2)	Kernel mode file system driver
	SERVICE_WIN32_OWN_PROCESS (16)	Service running in its own dedicated process
	SERVICE_WIN32_SHARE_PROCESS (32)	Service sharing a process with other services
	SERVICE_INTERACTIVE_PROCESS (256)	Flag indicating service is allowed to create a UI and interact with the logged on user
ImagePath	Path to .sys or .exe file	If empty, default for drivers is System32\Drivers folder and filename is registry key plus sys extension For services, looks using the PATH environment variable
ObjectName	Account to run service	May be a specific user, but usually a predefined account, e.g. LocalSystem, LocalService, NetworkService

Service Architecture



Controlling Services

- Service Control Programs, such as the Services MMC use the Windows API to control services
 - [OpenSCManager](#)
 - Opens a connection to the SCM
 - [OpenService](#), [CreateService](#)
 - Opens a connection to an existing service or installs a new service
 - [StartService](#)
 - Starts a service
 - [ControlServiceEx](#)
 - Sends other commands to the service (stop, pause, etc.)
 - [QueryServiceStatus](#), [DeleteService](#)
- Other tools
 - [sc.exe](#) - Command line tool used to control many aspects of services

Service Accounts

- LocalSystem
 - The most powerful on the local computer
 - Use with care
- NetworkService
 - Allows a service to authenticate outside the local computer
 - Has less privileges locally
- LocalService
 - Similar to NetworkService, but can only access network elements accepting anonymous access

Virtual Service Accounts

- Want better isolation than existing service accounts
 - Don't want to manage passwords
- Virtual accounts are like service accounts:
 - Process runs with virtual SID as principle
 - Can ACL objects to that SID
 - System-managed password
 - Show up as computer account when accessing network
- Services can specify a virtual account
 - Account name must be “NT SERVICE\<service>”
 - Service control manager verifies that service name matches account name
 - Service control manager creates a user profile for the account
- Also used by IIS app pool and SQL Server

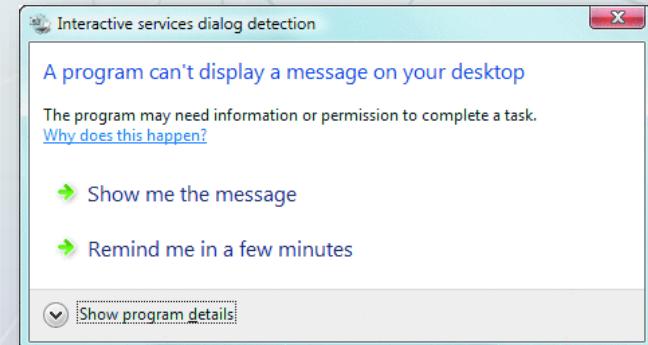


Interactive Services Before Vista

- Services are normally created under an invisible Window Station
 - Cannot interact with the user
- A flag can be used to create the service under `WinSta0` (the interactive Window Station)
 - May be dangerous for a `LocalSystem` account service because of its high privileges

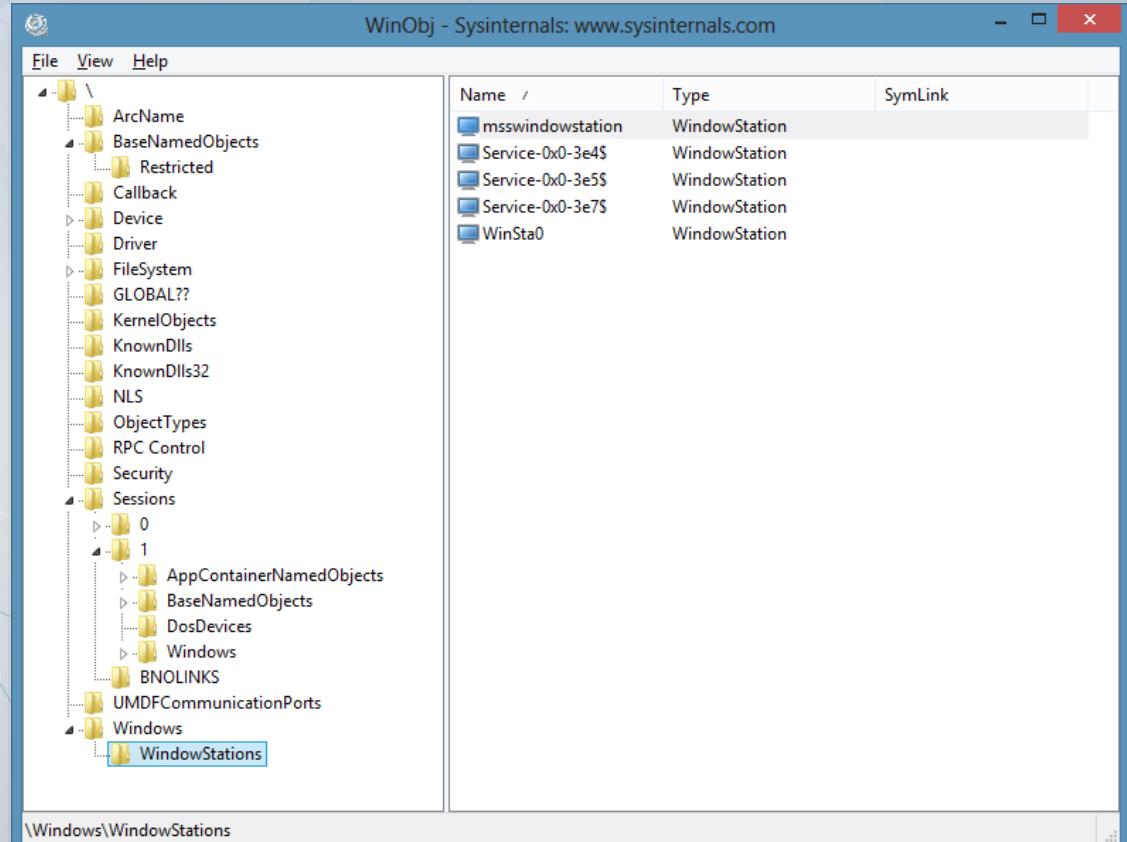
Interactive Services NT 6.0 and Above

- On Vista/2008 and up: If a Service wishes to interact with the user
 - The Interactive Service Detection service (`UI0Detect.exe`) launches an instance of itself in the user's desktop, displaying a message box
 - “Show me that message” switches the user's desktop to the Windows service desktop for the interaction



Window Station & Desktop

- A container of desktops, clipboard and atom table
 - A desktop is a container of windows, menus and hooks

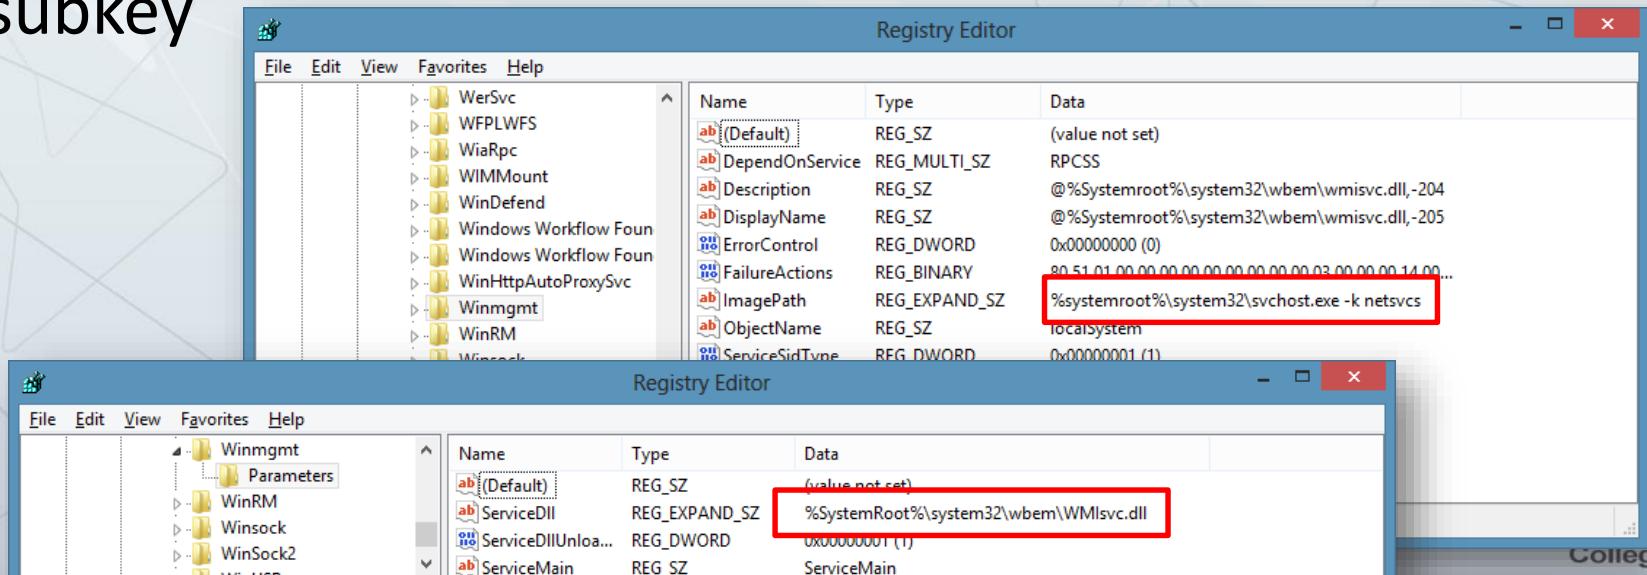


Shared Service Processes

- Some services run in their own process
- Some services are sharing a single process
 - Less system overhead of extra processes
 - If one service crashes, it brings down all other services in that process
 - All services running in a shared process run with the same account

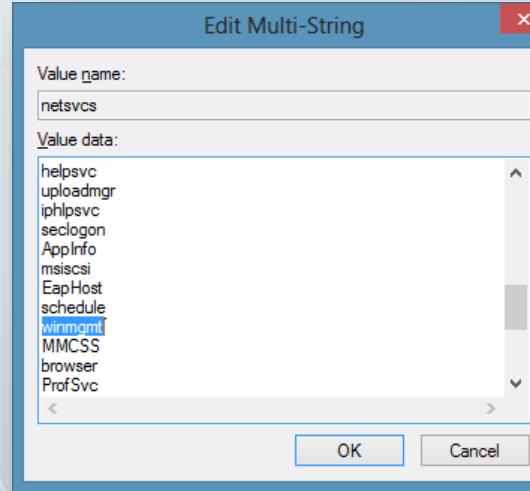
SvcHost.Exe

- Generic service host used by Windows
- The services themselves are DLLs
- Must have the DLL path under the *Parameters* subkey



SvcHost.Exe

- Determines which DLLs should share the same process by looking at `HKLM\Software\Microsoft\Windows NT\CurrentVersion\Svchost`
- Multi strings comprising of services under same service host



SvcHost.Exe

Registry Editor

Name	Type	Data
(Default)	REG_SZ	(value not set)
AxInstSvGroup	REG_MULTI_SZ	AxInstSV
DcomLaunch	REG_MULTI_SZ	Power LSM BrokerInfrastructure PlugPlay DeviceInstall DcomLaunch
defragsvc	REG_MULTI_SZ	defragsvc
ICService	REG_MULTI_SZ	vmicheartbeat vmicrdv
imgsvc	REG_MULTI_SZ	StiSvc
LocalService	REG_MULTI_SZ	nsi WdiServiceHost w32time EventSystem RemoteRegistry WinHttpAutoProxySvc Sstp TimeBroker SSDPSRV upnphost SCardSrv BthHFSrv QWAVE fdrespub WSService wcnc
LocalServiceAndNoImpersonation	REG_MULTI_SZ	DHCP eventlog AudioSrv wscsvc LmHosts ApplDSvc wcmsvc homegroupprovider WiFi DPS PLA BFE mpssc NcdAutoSetup WwanSvc
LocalServiceNetworkRestricted	REG_MULTI_SZ	PNRPSvc p2pimsvc p2psvc PnrrAutoReg
LocalServiceNoNetwork	REG_MULTI_SZ	WdiSystemHost WiaRpc Netman trkwks AudioEndpointBuilder WUDFSvc imron hidse
LocalServicePeerNet	REG_MULTI_SZ	AeLookupSvc CertPropSvc SCPolicySvc lanmanserver gpsvc IKEEXT FastUserSwitching CryptSvc nlasvc lanmanworkstation DHCP TermService NapAgent WinRM WECSV D KtmRm
LocalSystemNetworkRestricted	REG_MULTI_SZ	PolicyAgent
netsvcs	REG_MULTI_SZ	PeerDistSvc
NetworkService	REG_MULTI_SZ	PrintNotify
NetworkServiceAndNoImpersonation	REG_MULTI_SZ	RemoteRegistry
NetworkServiceNetworkRestricted	REG_MULTI_SZ	RpcEptMapper RpcSs
PeerDist	REG_MULTI_SZ	sdrsvc
print	REG_MULTI_SZ	swprv
regsvc	REG_MULTI_SZ	termsvcs
RPCSS	REG_MULTI_SZ	WbioSvcGroup
sdrsvc	REG_MULTI_SZ	wcssvc
swprv	REG_MULTI_SZ	WerSvcGroup
termsvcs	REG_MULTI_SZ	
WbioSvcGroup	REG_MULTI_SZ	
wcssvc	REG_MULTI_SZ	
WerSvcGroup	REG_MULTI_SZ	

Computer\HKEY_LOCAL_MACHINE\SOFTWARE\Microsoft\Windows NT\CurrentVersion\Svchost

Trigger Start Services

- Introduced in Windows 7
- Service can start with a certain “trigger”
- Cannot be configured using the Services MMC
 - Must call the [ChangeServiceConfig2](#) API function (or a comparable tool)
- Possible triggers
 - Computer joins a domain
 - Device arrival
 - Firewall port open
 - Group or user policy change
 - IP address availability
 - Network protocol (Windows 8 and later)
 - Custom, based on Event Tracing for Windows (ETW) event

Trigger Start API

```
SERVICE_TRIGGER trigger = {0};  
trigger.dwTriggerType =  
    SERVICE_TRIGGER_TYPE_IP_ADDRESS_AVAILABILITY;  
trigger.dwAction =  
    SERVICE_TRIGGER_ACTION_SERVICE_START;  
trigger.pTriggerSubtype =  
    (GUID*)&NETWORK_MANAGER_FIRST_IP_ADDRESS_ARRIVAL_GUID;  
SERVICE_TRIGGER_INFO info = {0};  
info.cTriggers = 1;  
info.pTriggers = &trigger;  
ChangeServiceConfig2 (hService,  
    SERVICE_CONFIG_TRIGGER_INFO, &info);
```

Scheduled Tasks

- Tasks have always been launched by triggers
- Windows NT 6.0 extends the variety of task **triggers** and **conditions**
- Task
 - **Actions** – what it does
 - **Triggers** – what triggers it
 - **Conditions** – what must hold

Task Triggers And Conditions

- Which user to use for launching the task?
- What triggers the task?
 - Schedule (calendar), delay, repeat, or auto-expire
 - At log on, start up, lock, or unlock
 - On an event log entry
- Start only if
 - Computer is idle, on AC power, or connected to a specific network connection
- Do what?
 - Run a program, send e-mail, show a message

Creating Tasks API

```
ITaskServicescheduler =  
    new TaskSchedulerClass();  
scheduler.Connect(null, null, null, null);  
  
ITaskFolderrootFolder =  
scheduler.GetFolder("\\");  
  
ITaskDefinition task = scheduler.NewTask(0);  
IExecAction action = (IExecAction)  
task.Actions.Create(  
    _TASK_ACTION_TYPE.TASK_ACTION_EXEC);
```

Task UI

Task Scheduler

File Action View Help

Task Scheduler (Local)

- Task Scheduler Library
 - BoilerControl
 - HomeAutomation
- Microsoft
- WPD

Name Status Triggers

Name	Status	Triggers
Close_All...	Ready	At 00:00 every day
Open_Eas...	Ready	At 07:05 every Sunday, Monday, Tuesday, Wednesday, Thursday, Friday of every week, startin
Open_Eas...	Ready	At 06:32 every Sunday, Monday, Tuesday, Wednesday, Thursday, Friday of every week, startin
Open_Nor...	Ready	
Open_Nor...	Ready	
Open_Shut...	Ready	
Open_Sou...	Ready	
Open_We...	Ready	
Turn_Off...	Ready	

Open_East_Shutter_Yarden_Room Properties (Local Computer)

General Triggers Actions Conditions Settings History

Action Details

Action	Details
Start a program	C:\Program Files\EIBExecuter\EIBExecuter.exe 7/11 1

New... Edit... Delete OK Cancel

When you create a task, you must specify the action that will occur when your task starts.

General Trigger

Name: Open_East_Shutter_Yarden_Room

Location: \\HOMESERVER\Yarden\Actions

Author: aaron.yarden@evaluecollege.com

Description:

Security options

When running

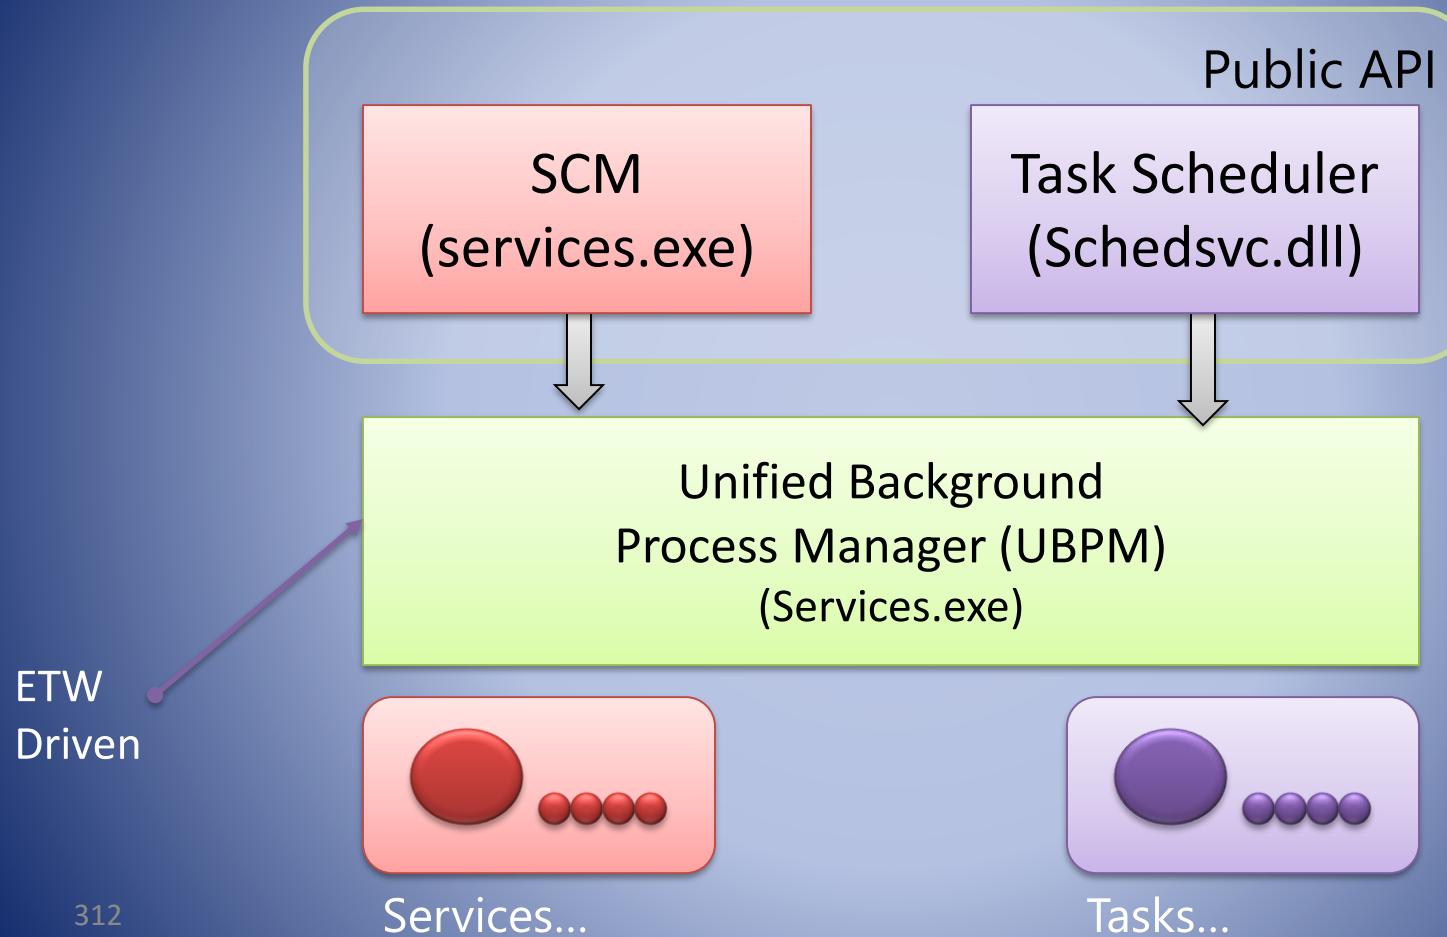
- alone
- Run only when user is logged on
- Run whether user is logged on or not
- Do not run if the computer is off

OK Cancel



eValue
College

Unified Background Processes Manager (UBPM)



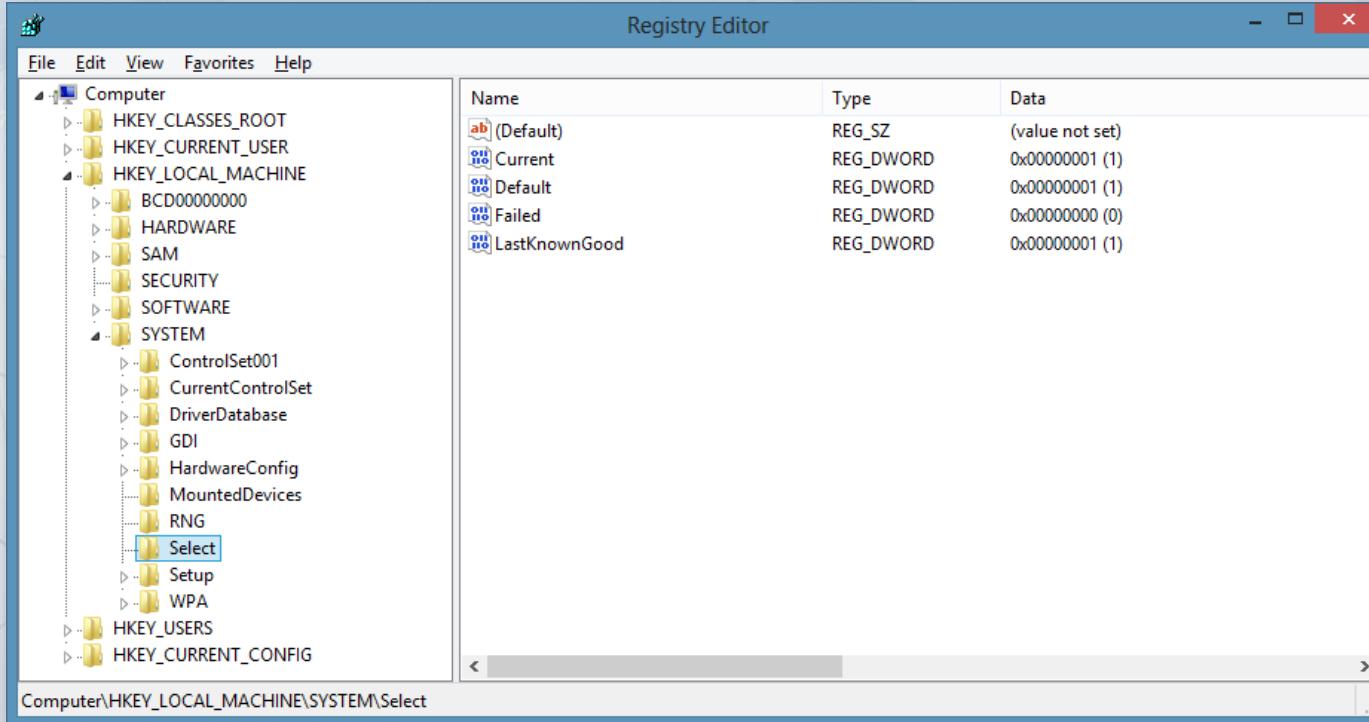
UBPM Summary

- Embrace service design goals: performance, security, power consumption
- Does your service have to run all the time?
- Retire “old” Windows XP-style services and consider: Automatic (Delayed Start) services, Trigger Start services, scheduled tasks
- Give your background activity the minimum security privileges
- Strive to minimize background work on idle

Control Sets

- A control set is a set of system configuration parameters
 - `HKLM\System\CurrentControlSet` is a symbolic link to one of the control sets
 - E.g. `ControlSet001`, `ControlSet002`
- One of the sets is the current
- Another set is the Last Known Good
- Values located at `HKLM\System\Select`

Control Sets in the Registry



Last Known Good

- The SCM is also responsible for saving the current control set as a successful boot
 - HKLM\System\CurrentControlSet
- A successful boot
 - Auto start services started successfully
 - Successful user logon
- Unsuccessful boot
 - An boot/system/auto-start driver has crashed the system
 - A service failed to start and has ErrorControl set to Severe or Critical

Last Known Good

- The meaning of a successful boot can be changed by a boot verification program
 - Pointed to by the registry key
HKLM\System\CurrentControlSet\Control\BootVerificationProgram
 - That program's installer must disable WinLogon's report of a successful boot by writing 0 to the key HKLM\SOFTWARE\Microsoft\Windows NT\CurrentVersion\Winlogon\ReportBootOk
 - Must call [NotifyBootConfigStatus](#) to indicate a successful or unsuccessful boot
 - An unsuccessful boot causes a restart

Windows Diagnostics & Management Overview

- What are my system and application doing?
- A short introduction of
 - Performance counters
 - Event Tracing for Windows (ETW)
 - WMI
 - Windows Error Report
- Integrating these diagnostics tools in your application can improve maintenance and reliability

Windows Performance Counters

- Provide information as to how well the operating system or an application, service, or driver is performing
- Numeric information grouped into **categories, counters and instances**
- Application can provide counter data, or consume this data
- Counter can be accessed from code, scripting & tools

Creating Counters

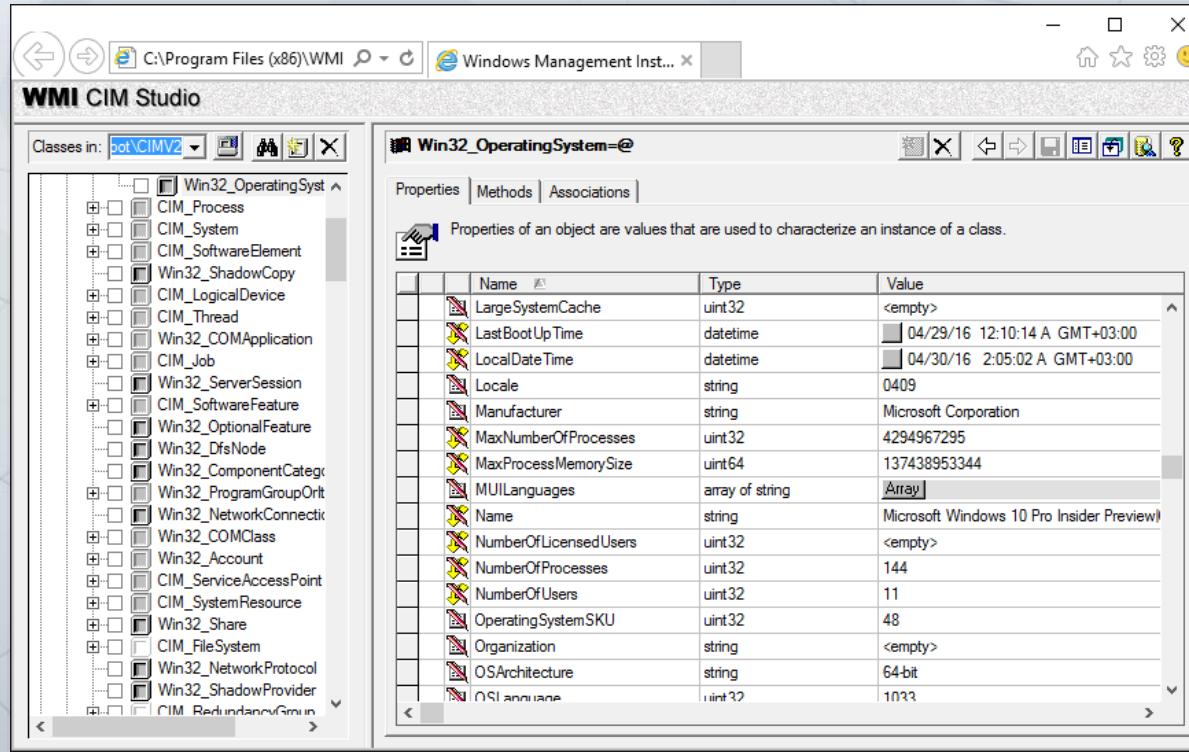
- Native:
 - Create a manifest
 - Use the CTRPP preprocessing tool
 - `PerfAutoInitialize`, `PerfCreateInstance`,
`PerfSet*Counter*Value`
- Managed:
 - **System.Diagnostics** namespace
 - `PerformanceCounterCategory` class
 - `CounterCreationDataCollection` class
 - `PerformanceCounter` class

Windows Management Instrumentation

- WMI is an implementation of the Web Based Enterprise Management (WBEM) standard
 - An extensible data management and collection facility with the required flexibility to connect to local and remote systems
- WMI entities are object based
 - Properties and methods
 - Events

WMI Namespace View

- WMI CIM Studio (downloadable as part of the [WMI Administrative Tools, fix it](#) for Windows 7 and above)



Windows Error Report

- WER collects mini dumps and other information and sends it to Microsoft
- WER is automatically activated when a process crashes
 - You can create and send a WER report even if your application didn't crash
- You can also disable a WER report from your application
([WerAddExcludedApplication](#))
- Open an account at [Winqual](#) to read these reports
- Demos: [C:\Program Files\Microsoft](#)
[SDKs\Windows\v7.0\Samples\winbase\windowserrorreporting](#)

Inside Software Tragic End-Of-Life

When The problem event occurs, Windows invokes WER

WER collects the data, builds a report, and prompts the user for consent and sends it to Microsoft (Watson Server)

If the Watson server requests additional data, WER collects the data and prompts the user for consent

If the application has registered for recovery and restart, WER executes the registered callback functions

If a response to the problem is available from Microsoft, the user is notified

Generating Custom Report

Call the [WerReportCreate](#) to create the report

Call the [WerReportSetParameter](#) to set the report parameters

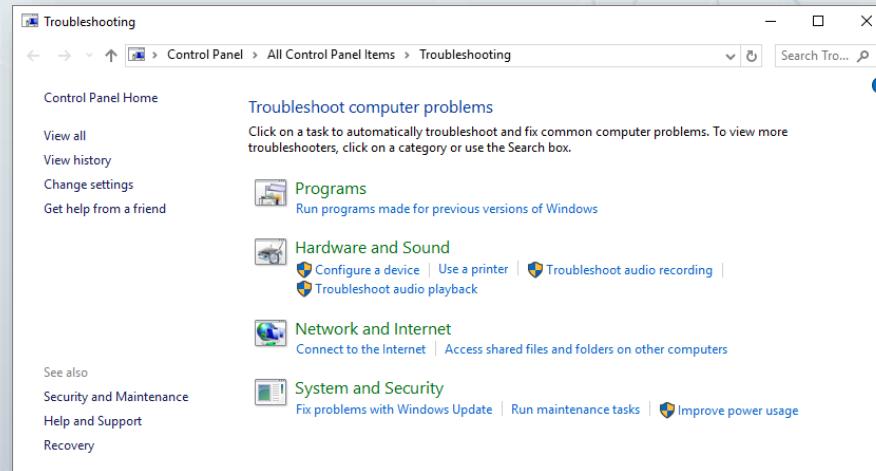
Call the [WerReportAddFile](#) to add files and/or the [WerReportAddDump](#) to add a minidump

Call the [WerReportSubmit](#) function to send the report

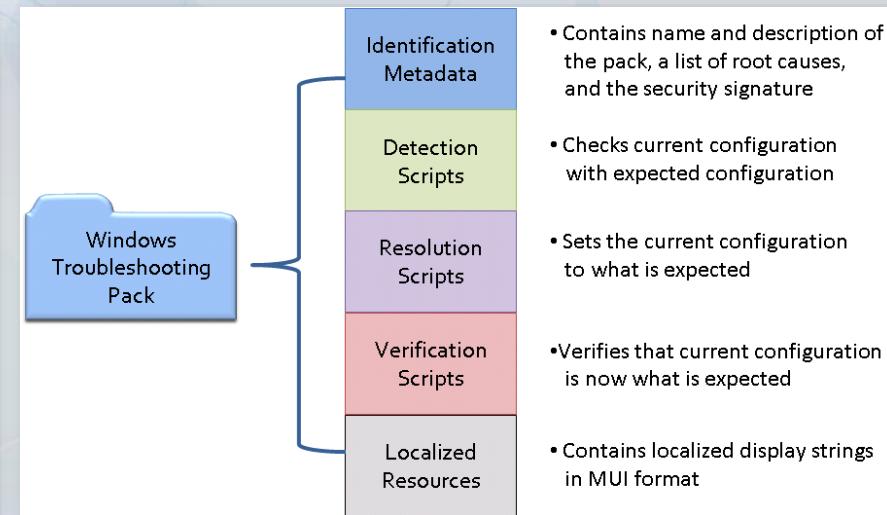
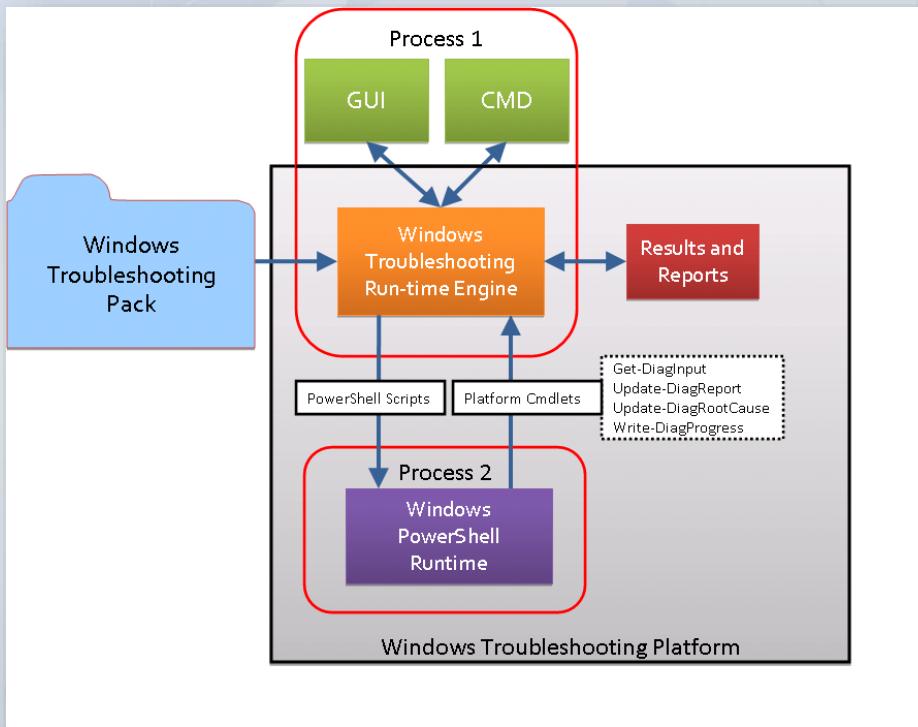
Call the [WerReportCloseHandle](#) to free resources

Windows Troubleshooting Platform

- WTP provides ISVs, OEMs, and administrators the ability to write troubleshooting packs
 - Used to discover and resolve issues found on the computer
- using WTP you can automate the process of fixing the most common detectable issues

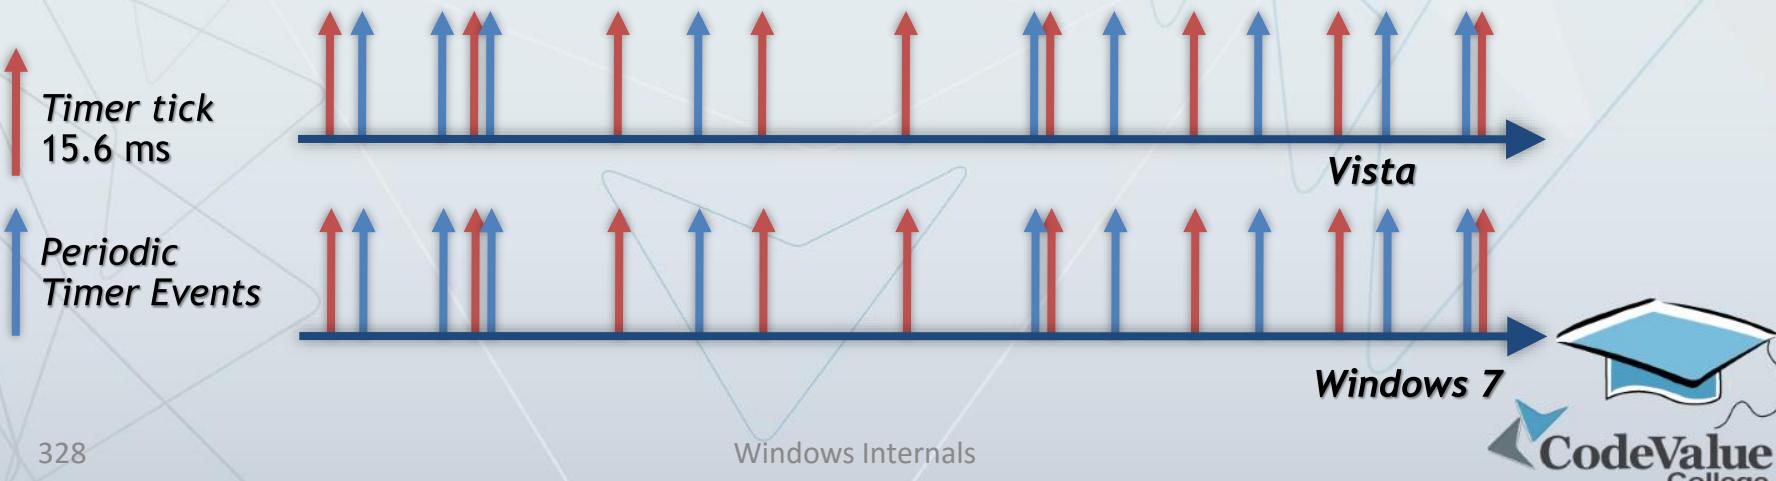


WTP Architecture



Timer Coalescing

- Staying idle requires minimizing timer interrupts
- Before, periodic timers had independent cycles even when period was the same
- New timer APIs permit timer coalescing
 - Application or driver specifies tolerable delay
- Timer system shifts timer firing to align periods on a coalescing interval:
 - 50ms, 100ms, 250ms, 1s



Coalesce wait able Timer (>= NT 6.1)

- BOOL SetWaitableTimerEx(

```
HANDLE hTimer,  
const LARGE_INTEGER *lpDueTime,  
LONG lPeriod,  
PTIMERAPCROUTINE pfnCompletionRoutine,  
LPVOID lpArgToCompletionRoutine,  
PREASON_CONTEXT WakeContext,  
ULONG TolerableDelay);
```

UI Coalescable Timer (>=NT 6.2)

- `UINT_PTR WINAPI SetCoalescableTimer(
 HWND hwnd,
 UINT_PTR nIDEvent,
 UINT uElapse,
 TIMERPROC lpTimerFunc,
 ULONG uToleranceDelay);`

Intelligent Timer Tick Distribution

- Before, primary timer interrupt on LP 0 propagated timer to all other LPs
 - Timer interrupt updates process and thread runtimes, checks for thread quantum end
 - Even if LP was idle, it had to service interrupt
- Now, timer system propagates timer only to processors that aren't idle (also called tick skipping)
 - Non-timer interrupts still wake LP

Wait Chain Traversal (>= NT 6.0)

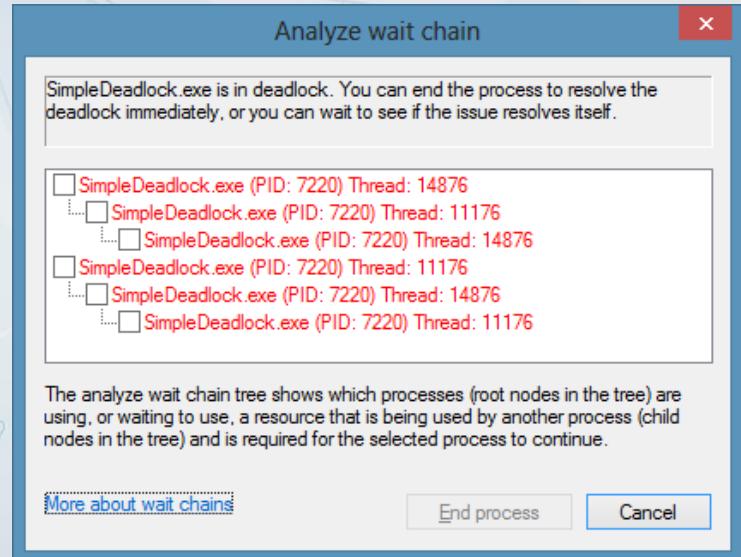
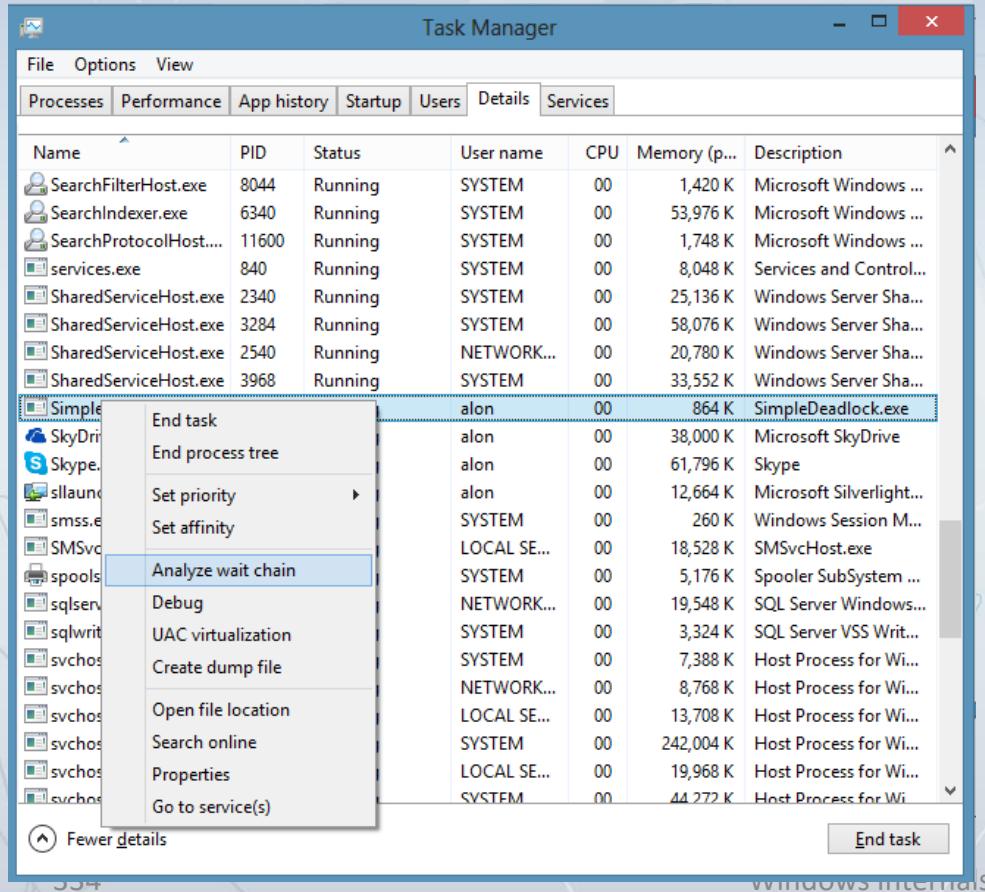
- WCT enables diagnosing of hangs and deadlocks
- A wait chain is an alternating sequence of threads and synchronization objects
 - A thread waits for an object using synchronization API
 - Once a thread took the object it owns it
 - If a thread waits for a lock that is owned by another thread, we can say that the first thread waits for the second thread
- WCT currently supports the following synchronization primitives:
 - Critical sections, Mutexes, GUI (*SendMessage*), Wait operations on processes and threads, ALPC and COM.



Deadlock Detection

```
BOOL WINAPI GetThreadWaitChain(  
    HWCT WctHandle,  
    DWORD_PTR Context,  
    DWORD Flags,  
    DWORD ThreadId,  
    LPDWORD NodeCount  
    PWAITCHAIN_NODE_INFO NodeInfoArray,  
    LPBOOL IsCycle  
);
```

Analyze Wait Chain



Summary

- The registry is a hierarchical database used to store machine wide and user based information
- Applications are encouraged to use other mechanisms if possible (e.g. XML files)
- Services are Windows processes that can (among other things) start without an interactive logon
- Windows implements most services as DLLs hosted in a common host (Svchost)
- We have also seen
 - performance counters
 - Windows Error Report
 - Windows Troubleshooting Platform
 - Background processing
 - Timer enhancement
 - Deadlock detection

Module 4

PROCESSES, THREADS & JOBS

Agenda

- Processes
- Access Tokens
- Threads
- Thread Scheduling
- Thread Pools
- Synchronization
- Jobs
- Summary

Process

- Management and containment object. Owns:
 - Private virtual address space (2GB/3GB on x86, 8TB / x64 – up to Windows 8, 128 TB Windows 8.1 and above)
 - Working set (physical memory owned by process)
 - Private handle table to kernel objects
 - Security token
- Has a priority class - Affects all threads running in that process
- Created from user mode by the Win32 APIs [CreateProcess](#), [CreateProcessAsUser](#), [CreateProcessWithLongonW](#)
- Terminated when
 - All threads within that process terminate
 - One of the threads in the process calls [ExitProcess](#) (Win32)
 - Killed with [TerminateProcess](#) (Win32)

Process Creation

- Flow of Process Creation
 - Open image file
 - Create kernel Executive Process object
 - Create initial thread
 - Notify CSRSS of new process and thread
 - Complete process and thread initialization
 - Load required DLLs and Initialize
 - DllMain function called with DLL_PROCESS_ATTACH reason
 - Start execution of main entry point (`_mainCRTStartup`
calls: `main`, `WinMain`)



Default DLL Search Paths

- Same logic as used by [LoadLibrary](#)
 - Directory of executable
 - The System directory ([GetSystemDirectory](#), e.g. c:\Windows\System32)
 - The 16 bit System directory (on 32 bit systems) (e.g. c:\Windows\System)
 - The Windows directory ([GetWindowsDirectory](#), e.g. c:\Windows)
 - The current directory of the process
 - Directories specified in the PATH environment variable
- If safe DLL search mode is not active (XP SP2+) or Windows version is XP SP1 or lower, the current directory is the second item in the above list
- A developer can alter the search path using [SetDefaultDllDirectories](#)
 - For Windows 7, Windows Server 2008 R2, Windows Vista, and Windows Server 2008, [KB2533623](#) must be installed
 - This provide a protection from most DLL injection related [security](#) problems

Creating Process Example



```
#include "stdafx.h"
#include <Windows.h>
#include <iostream>
using namespace std;
int main()
{
    STARTUPINFO si = { sizeof(si) };
    PROCESS_INFORMATION pi;

    wchar_t name[] = L"Notepad.exe";
    BOOL created = CreateProcess(nullptr, name,
        nullptr, 0, FALSE, 0, nullptr,
        nullptr, &si, &pi);
    if (!created)
    {
        DWORD dwErrorCode = GetLastError();
        wcerr <<
            L"Failed to start process (error number = "
            <<dwErrorCode<< ")" << endl;
        return dwErrorCode;
    } //else
}
```

```
//Continue here ...
wcout << L"Process started... pid = " <<
    pi.dwProcessId << L" Main thread id = " <<
    pi.dwThreadId << endl;
DWORD result = WaitForSingleObject(pi.hProcess,
    10000);

if (result == WAIT_TIMEOUT)
    wcout << L"Process still running..." << endl;
else
    wcout << L"Process exited." << endl;
CloseHandle(pi.hProcess);
CloseHandle(pi.hThread);
return 0;
}
```

Access Tokens

- Security context of a process or thread
- A process runs under an access token with the security settings of the specific user
 - [CreateProcess](#) (Win32) uses the logged on user
 - [CreateProcessAsUser](#) (Win32) uses whatever token is specified in the call (assuming a token for that user can be obtained)
- A thread by default runs under the access token of its process
 - Can assume another token temporarily by using *impersonation*

Thread

- Instance of a function executing code
 - Owns
 - Context (registers, etc.), 2 stacks (user mode and kernel mode)
 - Optionally, message queue and Windows
 - security token
 - Scheduling state
 - Priority (0-31), State (Ready, Wait, Running)
 - Current access mode (user or kernel)
 - Created using [CreateThread](#), [CreateRemoteThreadEx](#) (Win32)
 - Destroyed when
 - Thread function returns to win32 (ntdll.dll)
 - The thread calls [ExitThread](#) (Win32)
 - Terminated with [TerminateThread](#) (Win32)
 - The thread's process is exited or terminated

Thread Stacks

- Every user mode thread has two stacks
 - In kernel space (small, 12K (x86), 24K (x64))
 - Resides in physical memory (most of the time)
 - In user space (may be large)
 - By default 1MB is reserved, 64KB committed
 - A guard page is placed just below the last committed page, so that the stack can grow
 - Can change the initial size
 - Using linker settings as new defaults
 - On a thread by thread basis in the call to [CreateThread](#), [CreateRemoteThreadEx](#)



Creating Thread Example



```
DWORD WINAPI Thread(PVOID param)
{
    wcout << L"Thread started id = " << GetCurrentThreadId() << endl;
    PROCESSOR_NUMBER processorNumber;
    GetThreadIdealProcessorEx(GetCurrentThread(), &processorNumber);
    wcout << "Ideal processor is " << processorNumber.Number << endl;
    for (int i = 1; i <= 10; i++)
    {
        cout << i << endl; Sleep(rand() % 2000);
    }
    return 10;
}

void CreateThreadWithCreateThreadAPI()
{
    DWORD id, code;
    HANDLE hThread = CreateThread(nullptr, 0, Thread, nullptr, 0, &id);
    wcout << L"Thread " << GetCurrentThreadId() << L" free as a thread..." << endl;
    WaitForSingleObject(hThread, INFINITE);
    GetExitCodeThread(hThread, &code);
    wcout << "Thread finished. result = " << code << endl;
    CloseHandle(hThread);
}
```

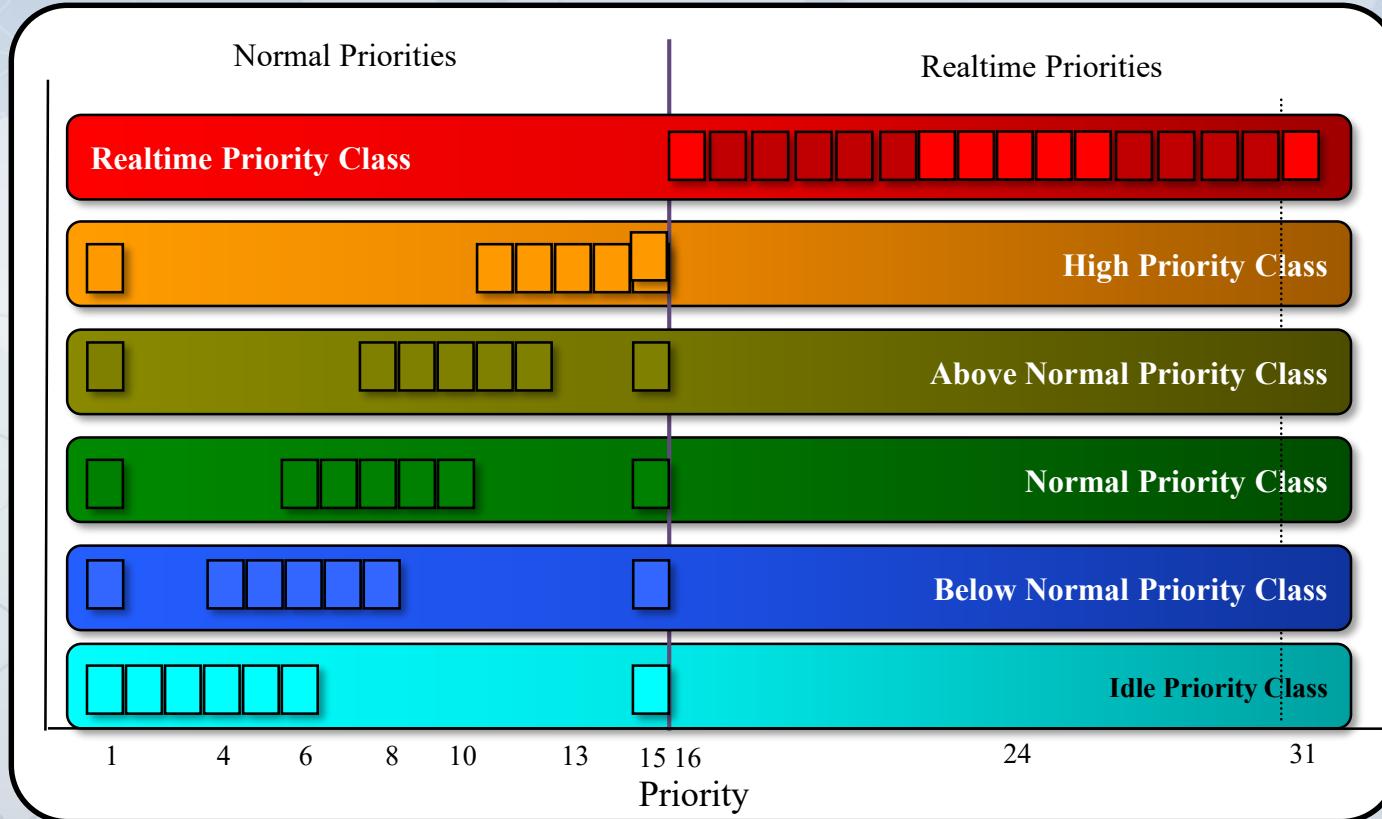


Creating Thread With Attribute



```
void CreateThreadWithCreateRemoteThreadAPI()
{
    DWORD id;
    LPPROC_THREAD_ATTRIBUTE_LIST pAttrlist;
    SIZE_T attrszie;
    //Get the required attribute list size
    InitializeProcThreadAttributeList(nullptr, 1, 0, &attrszie);
    pAttrlist = static_cast<LPPROC_THREAD_ATTRIBUTE_LIST>(_malloc(attrszie));
    memset(pAttrlist, 0, attrszie);
    InitializeProcThreadAttributeList(pAttrlist, 1, 0, &attrszie);
    PROCESSOR_NUMBER processorNumber; processorNumber.Group = 0;
    processorNumber.Number = 0;
    UpdateProcThreadAttribute(pAttrlist, 0, PROC_THREAD_ATTRIBUTE_IDEAL_PROCESSOR,
        &processorNumber, sizeof(processorNumber), nullptr, nullptr);
    //Create new thread with Ideal processor 0
    HANDLE hThread = CreateRemoteThreadEx(GetCurrentProcess(), nullptr, 0,
        Thread, nullptr, 0, pAttrlist, &id);
    ...
}
```

Thread Priorities (Win32 View)



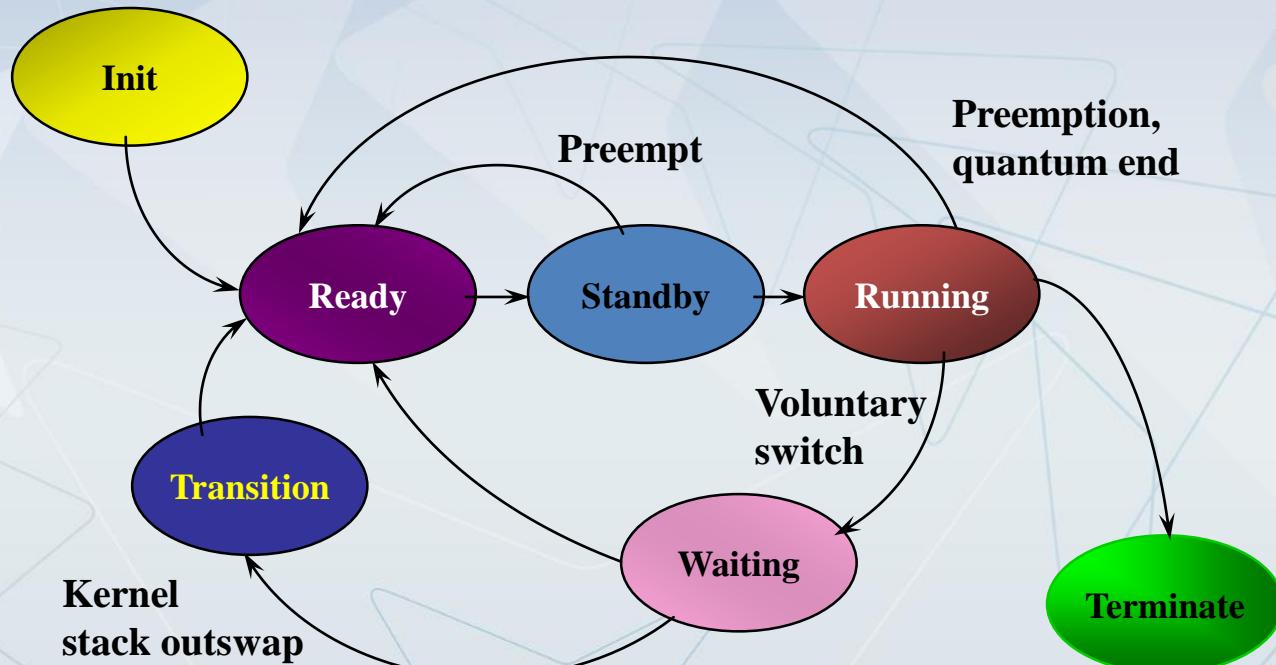
Thread Priorities

- Win32 view
 - Process has a priority class: Idle(4), Below Normal(6), Normal(8), Above Normal(10), High(13), Realtime(24)
 - Thread priority is an offset from that base priority (-2,-1,0,1,2 plus 2 special levels)
 - This is a Win32 constraint
 - More special levels for Real Time range
- Kernel view
 - Thread priority is an absolute value (0-31)
 - The scheduler does not care about processes

Thread Scheduling

- Priority based, preemptive, time-sliced
 - Highest priority thread runs first
 - If time slice (quantum) elapses, and there is another thread with the same priority in the Ready state – it runs
 - Otherwise, the same thread runs again
 - If thread A runs, and thread B (with a higher priority) receives something it waited upon (message, kernel object signaling, etc.), thread A is preempted and thread B becomes the Running thread
- Voluntary switch
 - A thread entering a wait state is dropped from the scheduler's Ready list
- Typical time slice
 - 20 msec (UP, Professional), 120 msec (UP, Servers)
 - 30 msec (MP, Professional), 180 msec (MP, Servers)
- On an MP system with n CPU cores, n concurrent threads may be running

Thread States



Thread States Details

- Init - Thread has just been created and is initializing
- Ready - Thread wants to run but no available CPU exists at the moment
- Standby - Represents a short state where context switch occurs
- Running - Thread is consuming a CPU and performing work
- Waiting - Thread is waiting for some event, kernel object, etc.
 - Does not consume any CPU time
- Transition - Thread becoming ready, but kernel stack needs to be read from page file back to RAM
- Terminate - Thread has terminated (for whatever reason)

The Scheduler

- Scheduling routines are called when scheduling events occur
 - Interval Timer interrupt checks for quantum end and timed wait completion
 - Hardware interrupt for I/O Completion
 - Changing state of waitable object other threads are waiting on
 - Entering a wait on one or more objects
 - Entering Sleep

The Quantum

- Each thread starts with a quantum number
 - 6 on Home/Professional, 36 on Servers
- Scheduler clock tick is typically
 - 10 msec (uniprocessor)
 - 15 msec (multiprocessor)
- Can determine with [`clockres.exe`](#) utility from SysInternals
- Quantum can be modified by using the registry or a Job (see later)

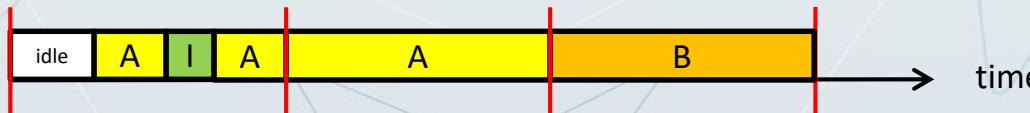
Quantum Accounting

- Pre NT 6.0
 - Thread quantum is decremented by 3 on clock tick
 - This means that the quantum is typically 30msec on Home/Professional and 180msec on Servers
 - Thread quantum is decremented by 1 when wait is complete
 - For threads with priority less than 14
 - Other threads get a replenished quantum
 - When Quantum reaches zero (or less)
 - Preempt thread and restore quantum value
 - Interrupt serviced during a thread's run
 - Implicitly counts against that thread's quantum



Quantum Accounting

- NT 6.0 and up
 - Quantum number itself is not used
 - Thread running time is used
 - Interrupt handling is not charged to the thread



Quantum Control

- Registry key: HKLM\SYSTEM\CCS\Control\PriorityControl
- Value: Win32PrioritySeparation

- Short vs. Long

- 1 = long, 2 = short
 - 0, 3 = default (long for Server, short for Client)

- Variable vs. Fixed

- 1=boost priority of foreground process, 2=don't boost
 - 0, 3=default (boost for Client, don't boost for Server)

- Foreground quantum boost

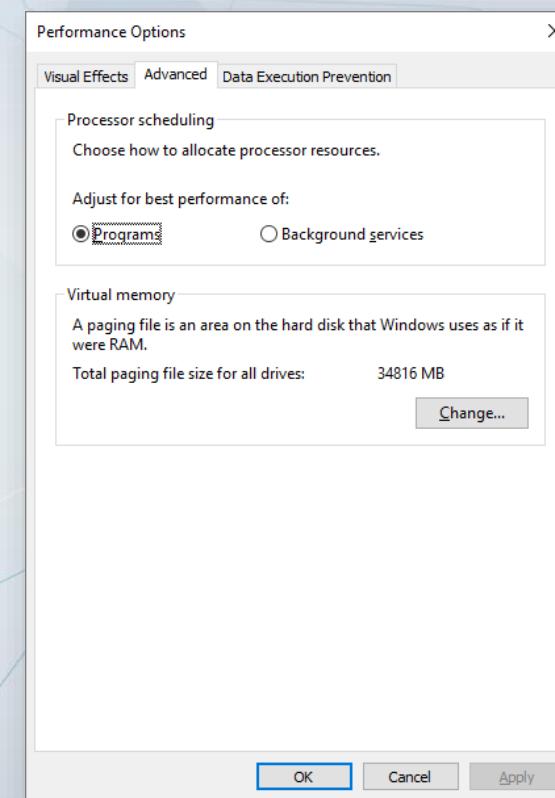
- Index into a table

4	2	0
Short vs. Long	Variable vs. Fixed	Foreground Priority Boost

	Short			Long		
Variable	6	12	18	12	24	36
Fixed	18	18	18	36	36	36

Performance Options

- System applet in Control Panel
- Programs
 - Short, variable quantum
 - Default for Professional
- Background services
 - Long, fixed quantum
 - Default for Servers



Quantum Boosting

- On a system configured for short, variable quantum
 - The foreground process gets triple quantum (6 clock ticks)
 - For any process with a priority class above Idle

Priority Boosts

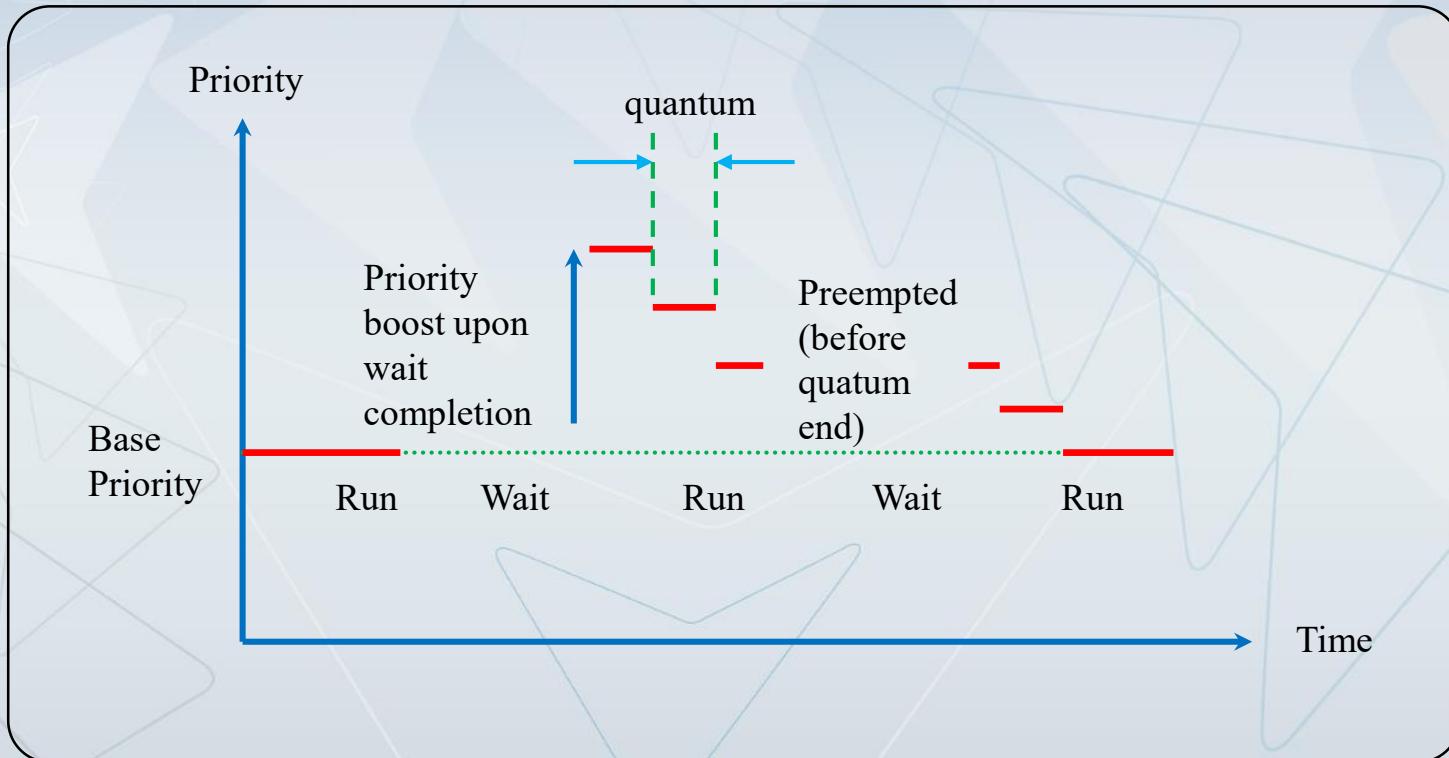
- Windows boosts the priority of threads in a number of scenarios
 - Completion of I/O operations
 - After waiting for Executive events or semaphores
 - During waiting for an executive resource
 - After threads in the foreground process complete a wait operation
 - When GUI threads wake up because of windowing activity
 - When a thread is starved

Completion of I/O Request or Wait

- Occurs when an I/O or wait completes
 - Can be specified by driver or Executive
 - [KeSetEvent](#) (Event, Increment...)
 - [IoCompleteRequest](#) (Irp, PriorityBoost)
- After a boost, thread runs for one quantum at that priority
 - Then drops one level, runs another quantum
 - Then drops another level, etc., until back to base priority
- Recommended boost values defined in <ntddk.h>

#define IO_SERIAL_INCREMENT	2
#define EVENT_INCREMENT	1
#define IO_KEYBOARD_INCREMENT	6

Thread Priority Boost and Decay



Foreground Process Wait Boost

- Foreground process
 - The process which contains the thread who is the owner (and creator) of the foreground window
- After a thread running in the foreground process completes a wait on a kernel object
- Receives a boost in the amount of the value set in the registry for foreground priority boost
 - +2 by default

GUI Thread Wakeup

- GUI threads receive a priority boost of 2 when they wake up due to a Window message arriving
- Provided by Win32k.sys
- Improves their chance of running sooner, giving a better responsiveness to the user

Priority Inversion Resolution

- Priority Inversion
 - High-priority thread waits on something locked by a lower priority thread which can't run because of a middle priority thread running
- Boosts thread to avoid priority inversion
 - Threads staying in ready state a long time (four seconds) get a big boost to priority 15
 - Get to run for 3 quantums at this special boost
 - Then priority drops to base
- Technically, starvation avoidance
- Implemented by the balance set manager
 - Scans at most 16 threads per pass
 - Boosts at most 10 threads per pass



Multi Processor Systems

- In a multi CPU / multicore / hyper-threaded system
 - Licensing accounts against physical processors (not logical)
- When choosing a processor for a thread
 - If a physical CPU's logical processors are idle, one of them would be selected
- NUMA systems should be handled with respect to the nodes

Multiprocessor Support

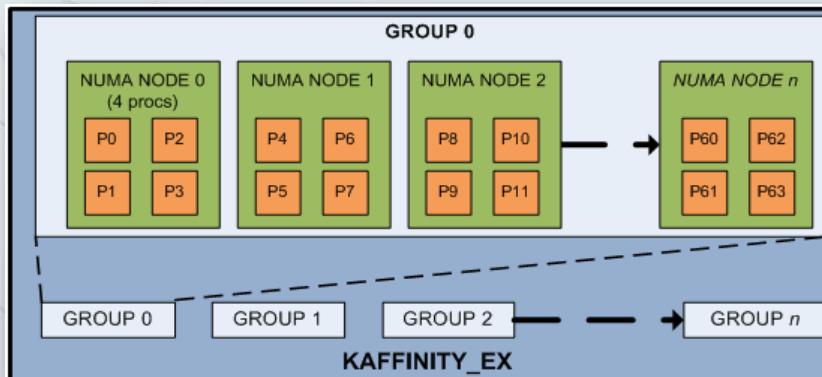
- Ideal Processor – Soft Affinity
 - Every thread has an ideal processor
 - Default value set in round-robin within each process
 - A random starting value
 - Can Override with [SetThreadIdealProcessor](#) or [SetThreadIdealProcessorEx](#)
 - On hyper-threaded systems, the next ideal processor selected is from the next physical CPU (not logical)

Hard Affinity

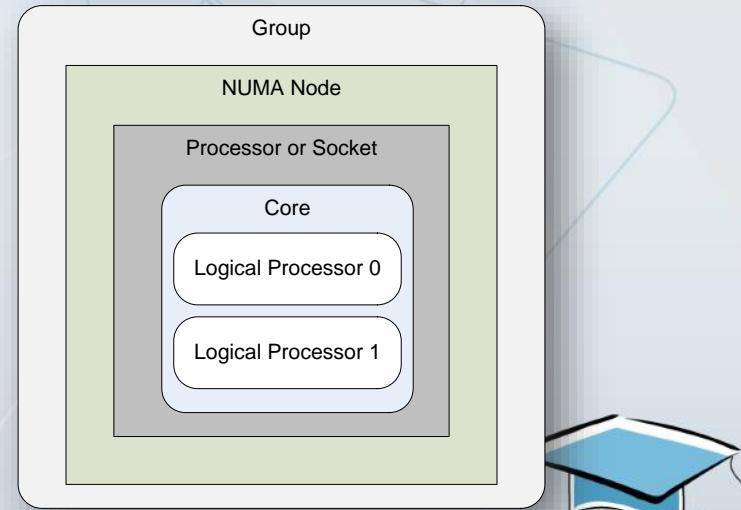
- Threads can run on any CPU unless hard affinity is set for that thread
 - [SetThreadAffinityMask](#)
 - The mask is a bit mask of allowed CPUs to run that thread
 - Default is process affinity mask, which defaults to all processors
 - Calling [SetProcessAffinityMask](#) changes priority mask for all threads running under that process
 - And future created thread in that process
- Using hard affinity may result in threads getting less CPU time

Processor Groups (>= NT 6.1)

- To support more than 64 logical cores, the scheduling mechanisms and APIs had to change
 - Instead of having one 64 bit flag to provide logical core id, we now have Processor Groups:



Windows Internals



Processor Group Characteristics

- Processor allocating to groups happens in boot time
- Each logical processor is assigned to a single group
- All the logical processors in a core, and all the cores in a physical processor, are assigned to the same group if possible
- Physical processors that are physically close to one another are assigned to the same group
- A process can have affinity for more than one group at a time
 - However, a thread can be assigned to only a single group at any time, and that group is always in the affinity of the thread's process
- An interrupt can only target processors of a single group
- In NUMA architectures, a group can contain processors from one or more nodes
 - All the processors in a node are assigned to the same group whenever possible



Group, Process, and Thread Affinity

- Windows initially assigns each process to a single group in a round-robin manner across the groups in the system
- The first thread of a process initially runs in the group to which Windows assigns the process
 - An application can override this default by using the [SetThreadGroupAffinity](#)
- Each newly created thread is by default assigned to the same group as the thread that created it
 - A new thread can be assigned to a group using the [CreateRemoteThreadEx](#) API
 - Pass the [PROC_THREAD_ATTRIBUTE_GROUP_AFFINITY](#) extended attribute together with a [GROUP_AFFINITY](#) structure
- Only the system process is assigned a multi-group affinity at startup time
- A single thread can never be assigned to more than one group at any time
 - However, a thread can change the group to which it is assigned

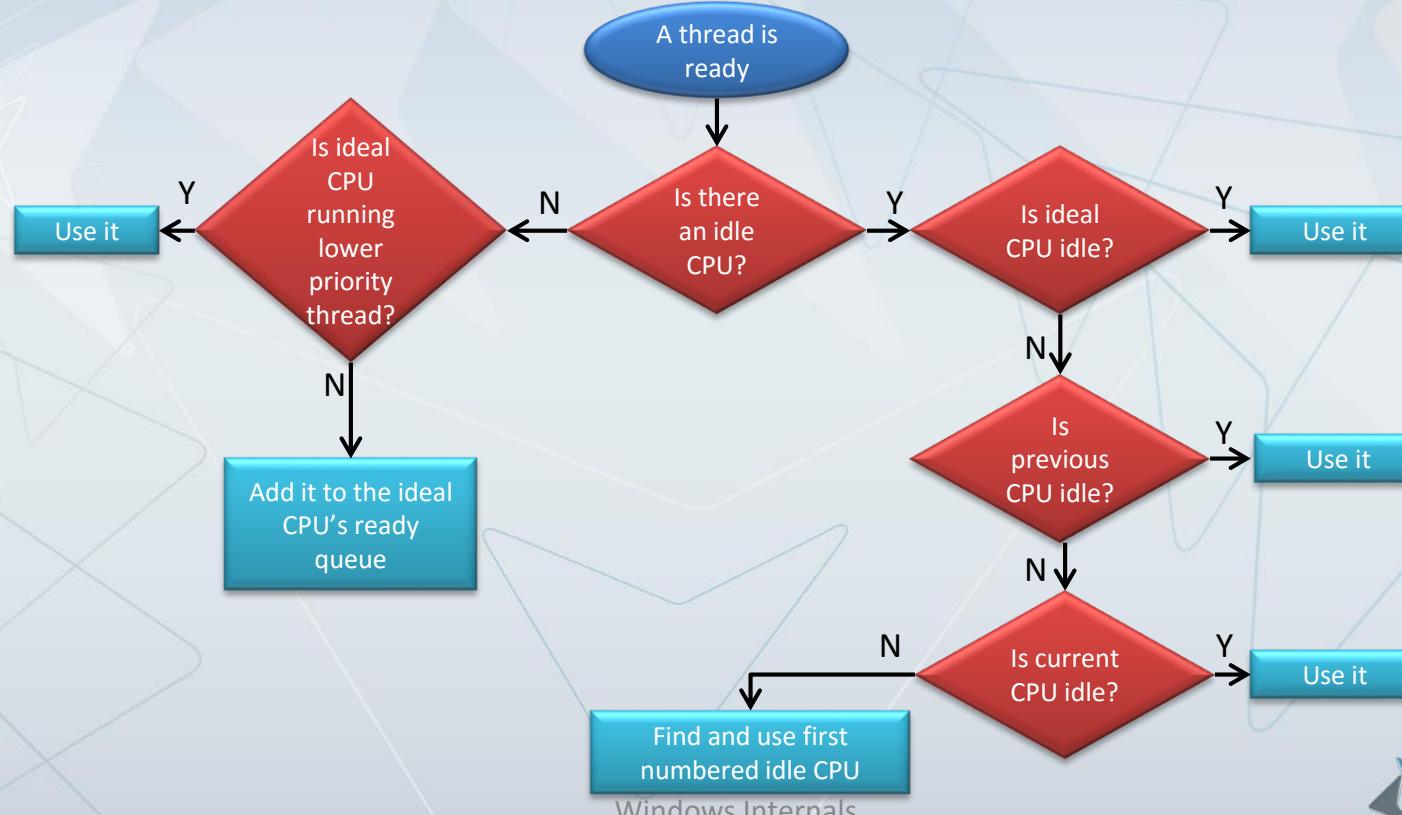
CPU Sets (NT 10.X)

- Provide granular “soft” affinity and information about CPU cores
 - On most systems each CPU Set ID will map directly to a single logical processor
 - Compatible with OS power management
 - Provides information about the energy consumption of a core ([EfficiencyClass](#))
 - Provide information about core parked status
 - Provide information about system reserved CPU Sets
 - Can't be assigned to the application threads
 - A thread affinitized to a given CPU Set will typically execute on one of the processors in its list of selected CPU Set IDs

CPU Sets APIs

- [GetProcessDefaultCpuSets](#)
- [GetSystemCpuSetInformation](#)
- [GetThreadSelectedCpuSets](#)
- [SetProcessDefaultCpuSets](#)
- [SetThreadSelectedCpuSets](#)
- [CPU SET INFORMATION TYPE](#)
- [SYSTEM CPU SET INFORMATION](#)

Scheduling on Multi-CPU System Per Processor Group (Simplified)



Waitable Objects

- Kernel supports various waitable objects, which can result in thread pre-emption
 - Process, Thread, Mutex (Mutant), Timer, File, Semaphore, Event (manual reset (notification) and auto reset (synchronization)), Job, Console, Change notification
- Waiting functions
 - Win32: [WaitForSingleObject\(Ex\)](#),
[WaitForMultipleObjects\(Ex\)](#),
[MsgWaitForMultipleObjects\(Ex\)](#), [SignalObjectAndWait](#)
 - Kernel: [KeWaitForSingleObject](#), [KeWaitForMultipleObjects](#)

Thread Pools

- Windows 2000 introduced a thread pool mechanism in user mode
 - One thread pool per process
 - Implemented entirely in user mode by ntdll.dll
 - Functions: [QueueUserWorkItem](#),
[RegisterWaitForSingleObject](#)
 - No way to prioritized work requests
- Windows Vista and up added support for private thread pools

Private Thread Pools

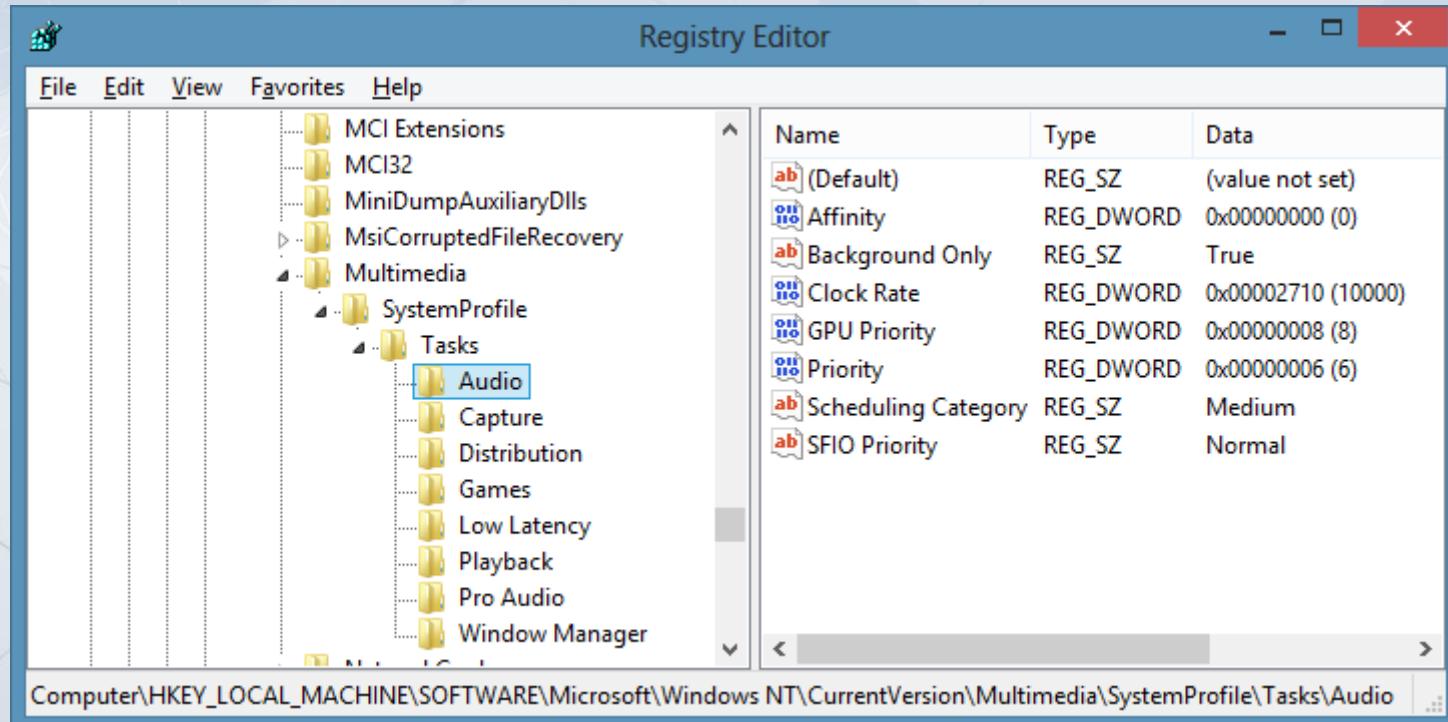
- Internally called Worker Factories
 - Kernel object: TpWorkerFactory
- More than one can be created per process
- Part of the implementation moved to kernel mode
 - Implemented behind the scenes using I/O completion ports
- Functions: [CreateThreadpool](#),
[CreateThreadpoolWork](#), [SubmitThreadpoolWork](#),
[SetThreadpoolWait](#), [many others](#)



MMCSS – Multimedia Class Scheduler Service (>= NT 6.0)

- Running as “Standard User” implies max thread priority – 15
- Standard Users want to play Audio, Video, Games and capture Audio and Video
 - They want great smooth experience
- MMCSS – A service that boost thread priorities and other characteristics
 - All the thread need to do is to announce that it is going to do Multimedia

Tasks & Registry



Using MMCSS

- HANDLE WINAPI [AvSetMmThreadCharacteristics](#)(
LPCTSTR TaskName, LPDWORD TaskIndex);
- HANDLE WINAPI [AvSetMmMaxThreadCharacteristics](#)(
LPCTSTR FirstTask, LPCTSTR SecondTask,
LPDWORD TaskIndex);
- BOOL WINAPI [AvRevertMmThreadCharacteristics](#)(
HANDLE AvrtHandle);

MMCSS Remarks

- The MMCSS service itself (mmcss.dll) runs as priority 27
- It boosts the multimedia threads into the range registered with the task as appears in the registry
 - This may be even to the real-time priority range
- MMCSS also deals with IO priority and networking throttling

MMCSS Remarks

- The MMCSS service itself (mmcss.dll) runs as priority 27
- It boosts the multimedia threads into the range registered with the task as appears in the registry
 - This may be even to the real-time priority range
- MMCSS also deals with IO priority and networking throttling

Thread Ordering (and timing) Service

- The mission:
 - Getting Real-time accuracy and task ordering under non-real-time O/S such as Windows
- The method:
 - Use MMCSS to get the resources
 - Use Thread Ordering Service to get timing and order
- The *thread ordering service* controls the execution of one or more client threads.
- It ensures that each client thread runs once during the specified period and in relative order



Using Thread Ordering Group

- The *parent thread* creates one or more thread ordering groups by calling the [AvRtCreateThreadOrderingGroup](#) API
 - Use this function to specify the period for the thread ordering group and a time-out interval
- Additional client threads call the [AvRtJoinThreadOrderingGroup](#) function to join an existing thread ordering group
 - Add a thread as a predecessor or successor to the parent thread in the execution order

Waiting to be executed

- Client threads should call [AvRtWaitOnThreadOrderingGroup](#) in their execution loop
 - the predecessor threads are executed one at a time in the order that they joined the group, then the parent thread and then all successor
 - Each time a thread return to the [AvRtWaitOnThreadOrderingGroup](#) the next thread starts
 - If all threads complete their execution before a period ends, all threads wait until the remainder of the period elapses before any are executed again

Remarks

- When the client need no longer run as part of the group, it calls the [AvRtLeaveThreadOrderingGroup](#)
- The parent thread calls the [AvRtDeleteThreadOrderingGroup](#) to delete the group
- The group is also destroyed if the parent thread does not complete its execution before the period plus the time-out interval elapse

Job Objects

- Kernel object introduced in Windows 2000
- Allows groups of one (or more) processes to be managed as a unit
- System enforces Job quotas and security
 - Total and process CPU time, working sets, CPU affinity and priority class, quantum length (for long, fixed quantums only)
 - Security limits
 - UI limits

Creating a Job Object

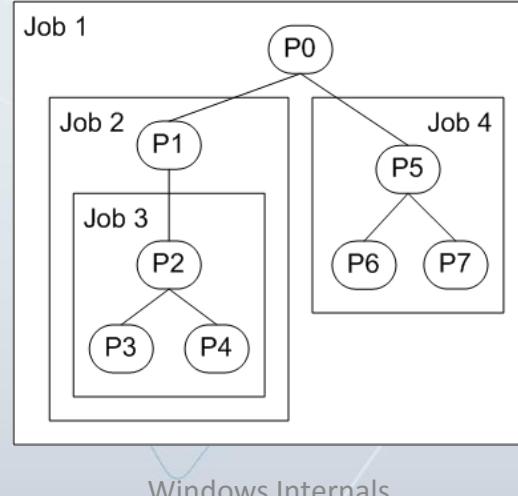
- Use the [CreateJobObject](#) function
 - When the job is created, no processes are associated with the it
- To associate a process with a job, use the [AssignProcessToJobObject](#) function
 - After a process is associated with a job, the association cannot be broken
 - All new sub-processes will be associated with the Job
 - [CreateProcess](#) can ask to break-away from the job, if allowed
 - (NT >= 6.2) A process can be associated with more than one job in a hierarchy of nested jobs

Using the Job object

- To determine if a process is running in a job: [IsProcessInJob](#)
- Kill all processes within the job: [TerminateJobObject](#)
- Setting limits: [SetInformationJobObject](#)
- Getting information: [QueryInformationJobObject](#)
 - the job's user time and kernel time
 - number of active processes
 - number of total processes created within the job
 - number of total processes terminated within the job
- Getting Notification: [Set Completion Port](#) using [SetInformationJobObject](#)

Nested Jobs (>= NT 6.2)

- An application can use nested jobs to manage subsets of processes
 - Nested jobs also enable an application that uses jobs to host other applications that also use jobs



Nested Jobs

- Nested jobs is created when a process that is already in a job assign a process to another Job
- The effective limit for child job can be more restrictive than the limit of its parent
 - It cannot be less restrictive
- For a nested job, event messages are sent to every I/O completion port associated with any job in the parent job chain of the job that triggered the message
- Parent Job termination terminates all nested jobs

Summary

- Process is a container of resources for running a program
- Thread is the actual entity scheduled by the kernel to run on a CPU
- Job allows managing process(es) as a single unit while applying constraints

Module 5

MEMORY MANAGEMENT

Agenda

- Memory Manager Features
- Virtual Memory vs. Physical Memory
- Virtual Memory Mapping
- Memory Mapped Files and Shared Memory
- Summary

Memory Manager Fundamentals

- Each process sees a virtual address space
 - 2 GB (32 bit), 8 TB (x64 <= NT 6.2), 128 TB (x64 >= 6.3)
- Memory Manager tasks
 - Mapping virtual addresses to physical addresses
 - Using page files to back up pages that cannot fit in physical memory
 - Provide memory management services to other system components

Managing Memory

- Memory is managed by chunks called Pages
- Page size is determined by CPU type
 - A compromise between fine and coarse
- Two page sizes are supported
- Allocations, de-allocations and other memory block attributes (e.g. protection) are always per page

Architecture	Small page size	Large page size
x86	4 KB	4 MB (2 MB on PAE systems)
x64	4 KB	2 MB
IA-64	8 KB	16 MB

Virtual Page States

- Each page in virtual memory may be in one of three states
- Free
 - Unallocated page
 - Any access causes an Access Violation exception
- Committed
 - Allocated page that can be accessed
 - Has a backup on disk
- Reserved
 - Unallocated page causing Access Violation on access
 - Address range will not be used for future allocations unless specifically requested
- Can view with [VMMMap](#) (SysInternals)

The Virtual Address Descriptor (VAD) Tree

- A range tree that describes memory ranges used by a process
 - The [!vad](#) WinDbg extension command and the [vmmap](#) SysInternals tool show this tree
- When process reserve or commit a range of pages, a new VAD entry is created
 - The real allocation will occur only on demand

Viewing Process Address Space

Process: notepad.exe		
PID: 42282		
Committed:		
 Private Bytes		
Working Set:		
32bit notepad 2GB		
Type		
Total	13	K
Image	50	K
Mapped File	19	K
Shareable	58	K
Heap	1	0 K
Managed Heap	2	8 K
Stack	2	8 K
Private Data	8	2 K
Page Table	4	8 K
Unusable	2	68 K
Free	1 961	728 K

Process: notepad.exe		
PID: 3744		
Committed:		
 Private Bytes		
Working Set:		
64bit notepad on W8 8TB		
Type		
Total	13	K
Image	4	12 K
Mapped File	1	72 K
Shareable	2	00 K
Heap	4	04 K
Managed Heap	2	12 K
Stack	1	30 K
Private Data	4	34 K
Page Table	1	78 K
Unusable	8 589	842 560 K
Free	8 589	842 560 K

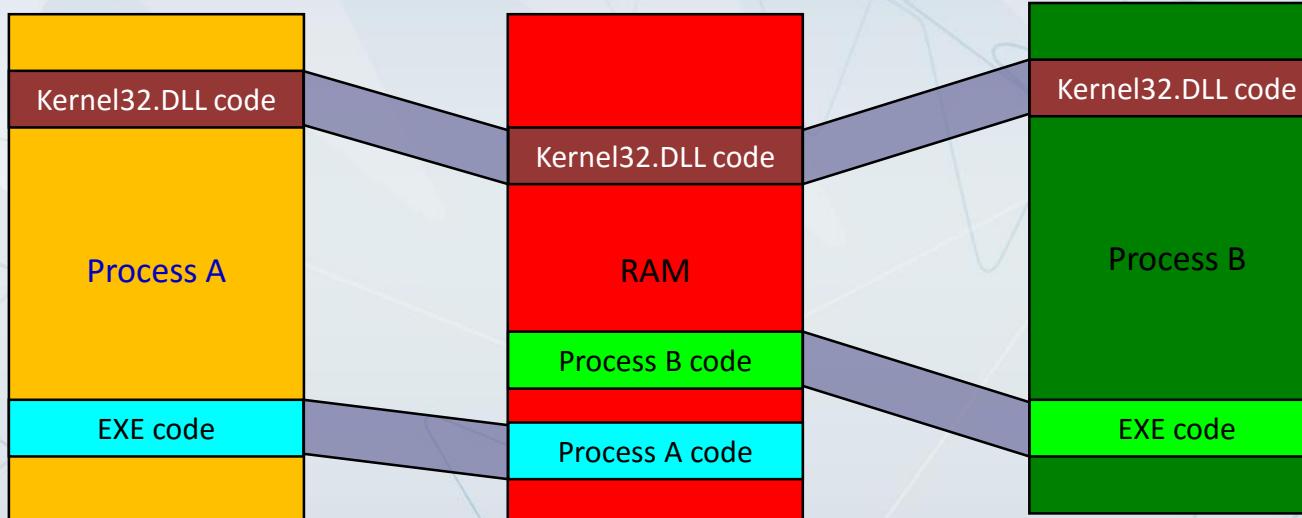
Process: notepad.exe		
PID: 42106		
Committed:		
 Private Bytes		
Working Set:		
64bit notepad on W8.1 128TB!		
Type		
Total	2	K
Image	60	148 K
Mapped File	19	912 K
Shareable	2	1 500 K
Heap	1	528 K
Managed Heap	2	560 K
Stack	1	416 K
Private Data	4	444 K
Page Table	1	3 272 K
Unusable	135	291 355 072 K
Free	135	291 355 072 K

Locking Memory

- Memory used by processes and kernel may be paged out temporarily
- The memory manager allows locking pages in RAM, preventing paging
- Processes can call [VirtualLock](#) to request a block of memory to stay in physical memory
 - The memory manager will comply if possible
 - Can increase maximum size with [SetProcessWorkingSetSize\(Ex\)](#)
- Device drivers can call [MmProbeAndLockPages](#)
 - No hard wired limits and will not be unlocked if memory pressure is high



Sharing Pages



Sharing Pages

- Code pages are shared between processes
 - 2 or more processes based on the same images
 - DLL code
 - However, DLLs must be loaded in same address
- Data pages (read/write) are shared at first
 - But with special protection called Copy-On-Write
 - If one process changes the data, an exception is caught by the Memory Manager, which creates a private copy of the accessed page for that process
 - Removing the Copy-On-Write protection

Sharing Data Without COW



- For use in DLLs only

```
#pragma data_seg("shared")
int count = 0;
__declspec(dllexport) double epsilon = .000001;
#pragma data_seg()

#pragma comment(linker, "/section:shared,RWS")
```

Protecting Memory

- Process memory is automatically protected against read/write from other processes
 - Unless the process shares pages, or
 - Another process has virtual memory read or write access handle and uses [ReadProcessMemory](#) or [WriteProcessMemory](#)
- Processors support hardware controlled memory protection
 - Can be set with the [VirtualAlloc/Ex](#), [VirtualProtect/Ex](#) functions



Protection Options

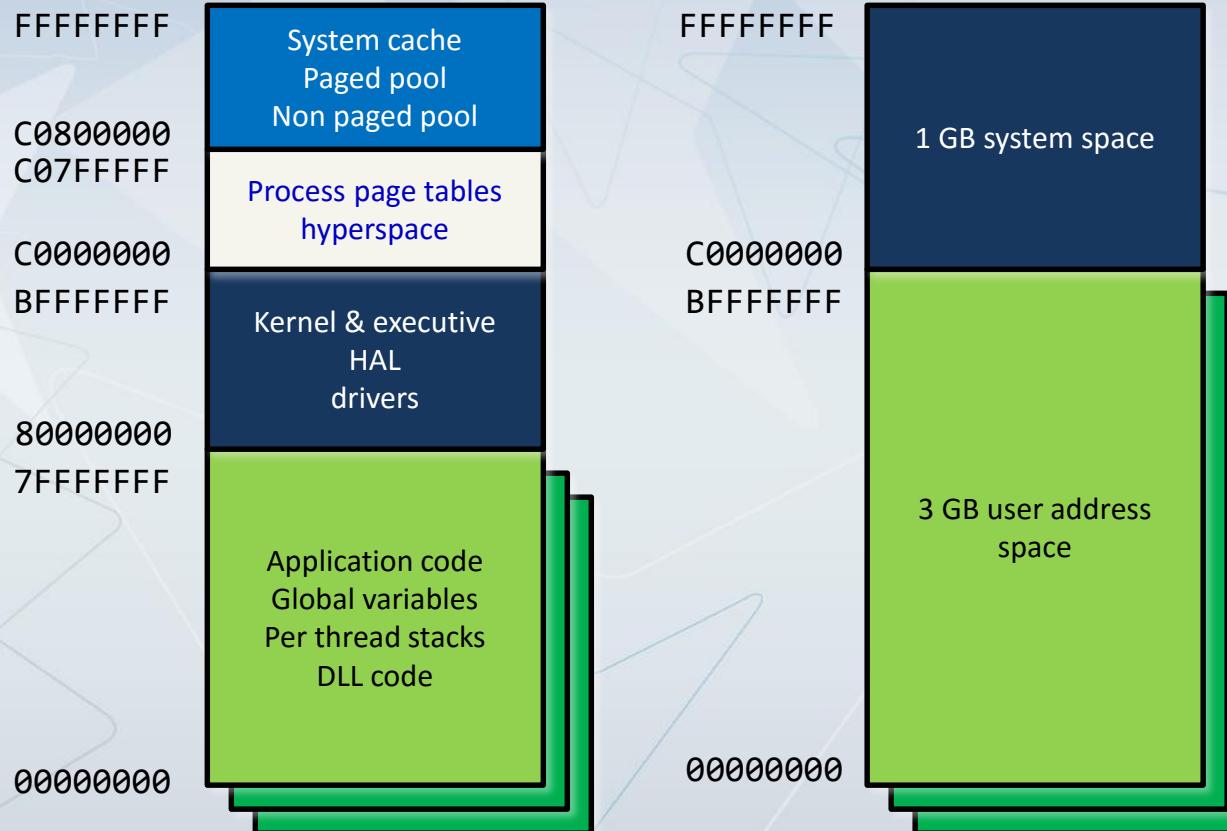
Attribute	Description
PAGE_NOACCESS	No access allowed
PAGE_READONLY	Only read is allowed. Other operations cause Access Violation exception
PAGE_READWRITE	Read and write allowed
PAGE_EXECUTE	Only code execution is allowed
PAGE_EXECUTE_READ	Execution and reading is allowed
PAGE_EXECUTE_READWRITE	All access is allowed
PAGE_WRITECOPY	Copy on write. Write access will cause a private copy to be set for the caller
PAGE_EXECUTE_WRITECOPY	Same as PAGE_WRITECOPY with execution support
PAGE_GUARD	Guard page. Access will cause EXCEPTION_GUARD_PAGE exception. Can be combined with other attributes. Used primarily for a thread's stack

- On processors that do not support execution protection, Execute is the same as Read

Virtual Address Space Layout

- Each process sees its own private address space
- System space is part of the entire address space that is visible (but not accessible by user mode code)
- The layout depends on the “bitness” of the OS and the specific process

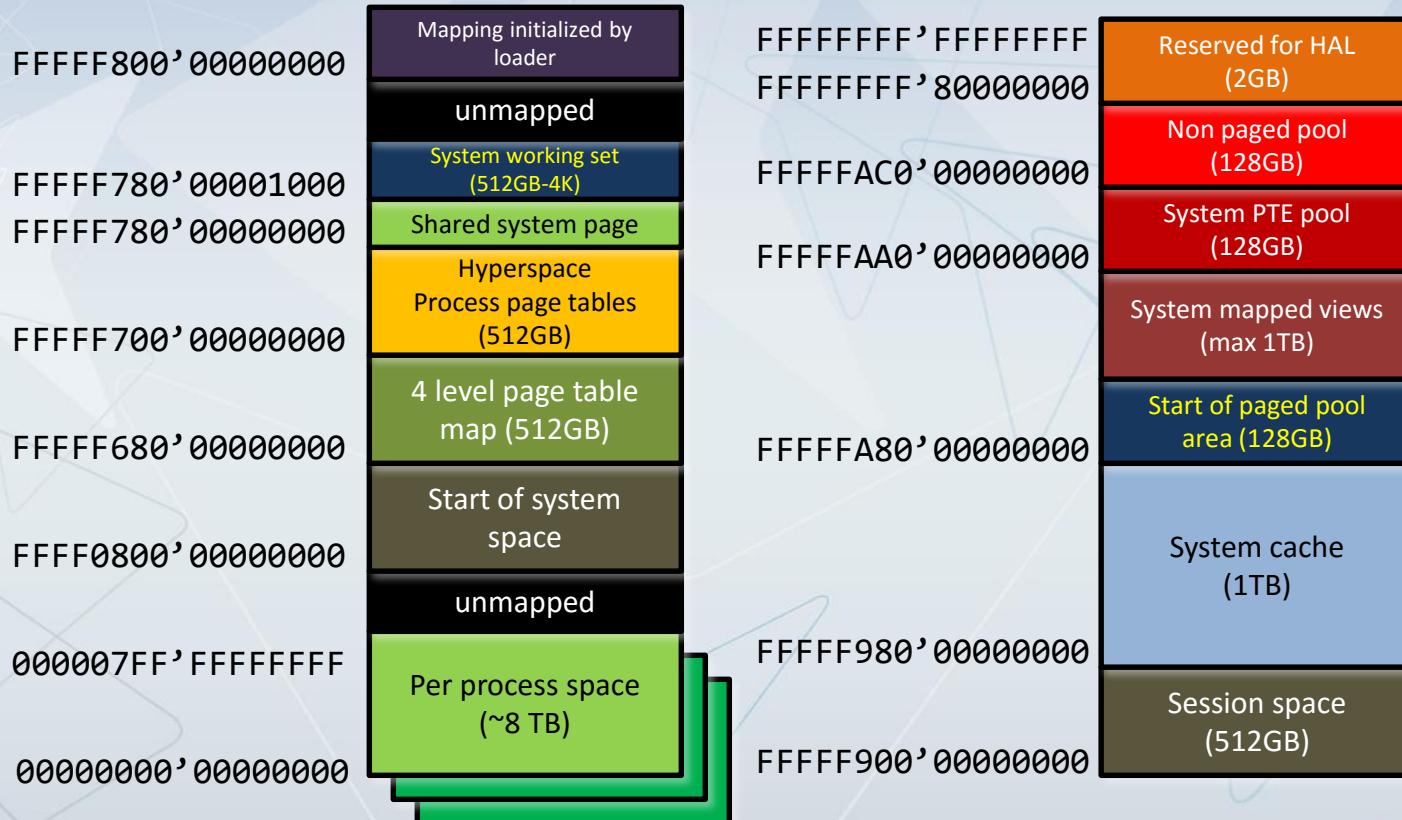
x86 Address Space Layout



3GB User Address Space (x86)

- Windows can be configured with the /3GB boot.ini switch (Windows 2003 and lower)
 - Or using the Boot Configuration Database (BCD) on Vista and up
- Each process gets a 3GB address space
- No returned pointer will be above 2GB unless the image is compiled with the LARGEADDRESSAWARE linker flag
 - To prevent breaking applications that depend on 31 bit addresses

x64 Address Space Layout (<= NT 6.2)



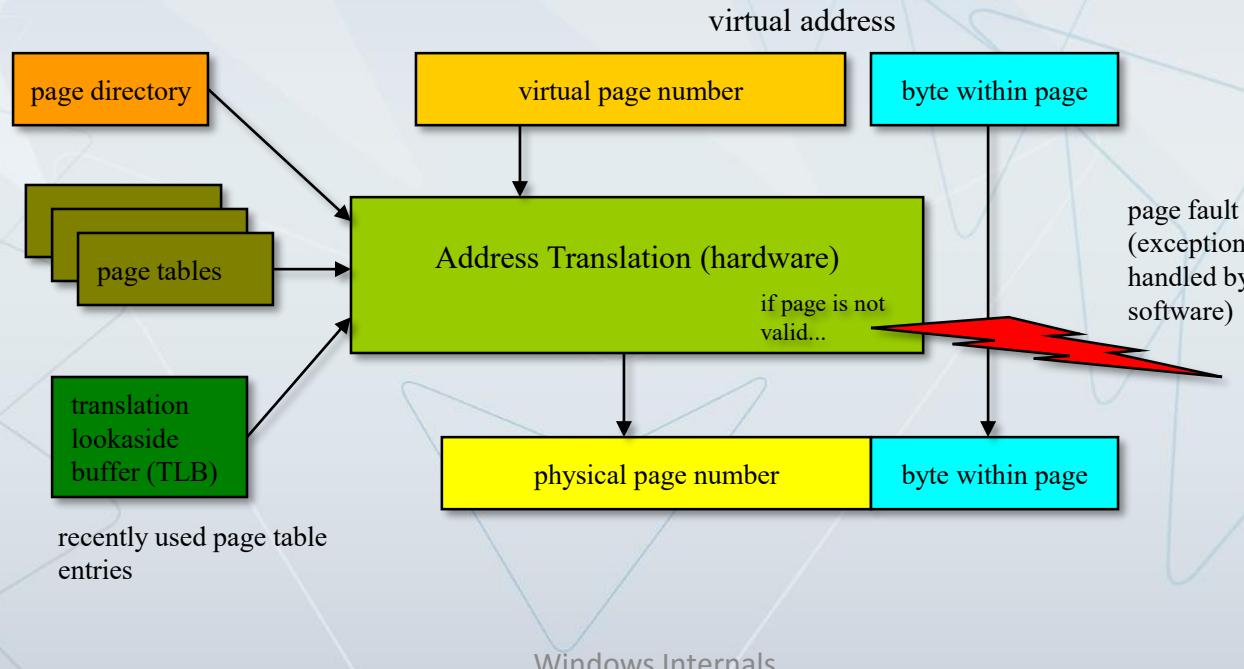
x64 Address Space Layout (>= NT 6.3)

Start	End	Size	Description
FFFF0000`00000000	FFFF07FF`FFFFFF	8TB	Memory Hole
FFFF0800`00000000	FFFFFAFF`FFFFFF	168TB	Unused Space
FFFFB000`00000000	FFFFBFFF`FFFFFF	16TB	System Cache
FFFC000`00000000	FFFCFFF`FFFFFF	16TB	Paged Pool
FFFD000`00000000	FFFDFFF`FFFFFF	16TB	System PTEs
FFFE000`00000000	FFFEFFF`FFFFFF	16TB	Nonpaged Pool
FFFFF000`00000000	FFFFF67F`FFFFFF	6.5TB	Unused Space
FFFFF680`00000000	FFFFF6FF`FFFFFF	512GB	PTE Space
FFFFF700`00000000	FFFFF77F`FFFFFF	512GB	HyperSpace
FFFFF780`00000000	FFFFF780`00000FFF	4K	Shared User Data
FFFFF780`00001000	FFFFF780`BFFFFFF	~3GB	System PTE WS

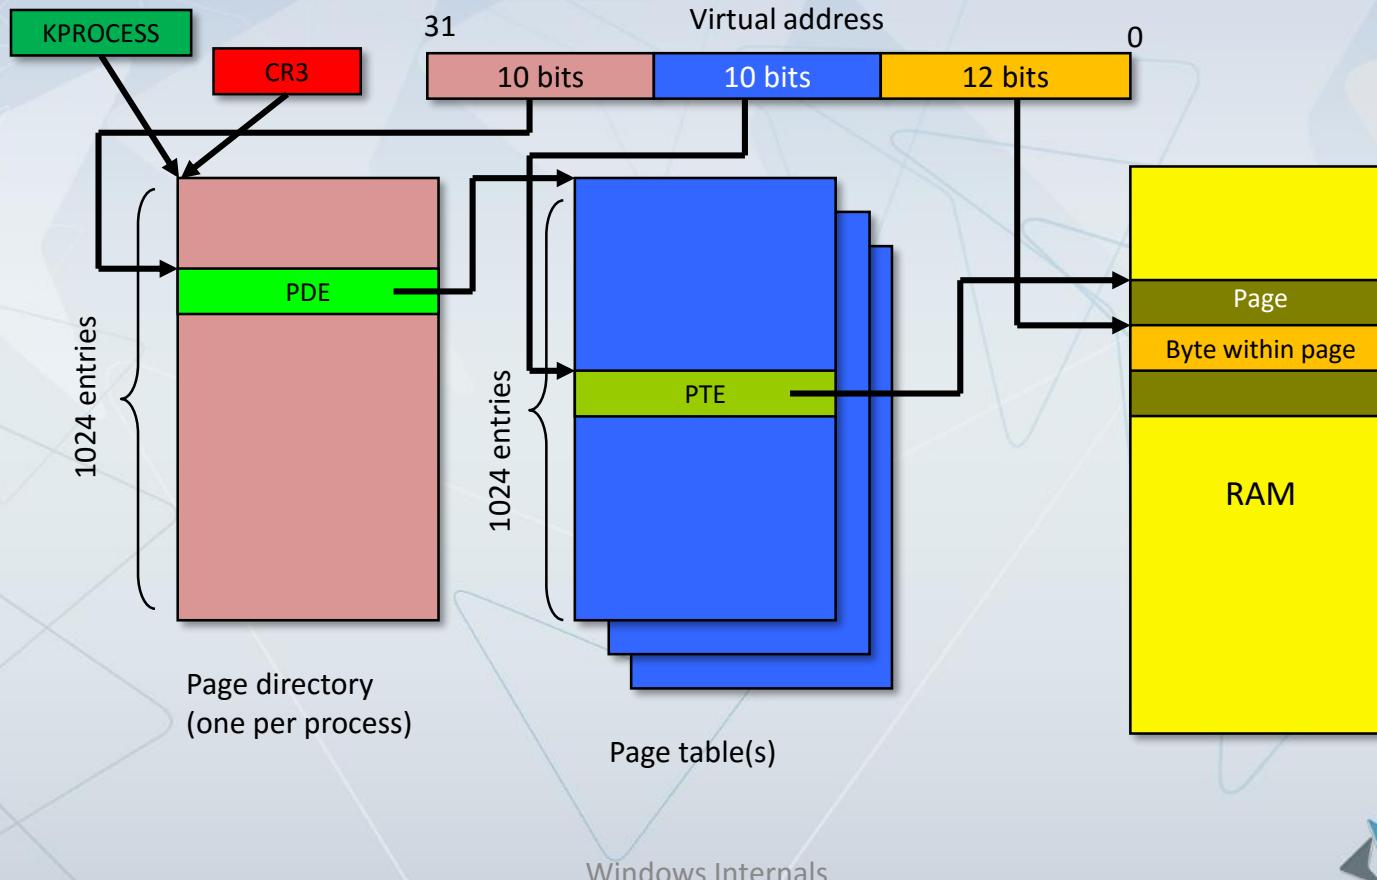
Start	End	Size	Description
FFFFF780`C0000000	FFFFF780`FFFFFF	1GB	WS Hash Table
FFFFF781`00000000	FFFFF791`3FFFFFF	65GB	Paged Pool WS
FFFFF791`40000000	FFFFF799`3FFFFFF	32GB	WS Hash Table
FFFFF799`40000000	FFFFF7A9`7FFFFFF	65GB	System Cache WS
FFFFF7A9`80000000	FFFFF7B1`7FFFFFF	32GB	WS Hash Table
FFFFF7B1`80000000	FFFFF7FF`FFFFFF	314GB	Unused Space
FFFFF800`00000000	FFFFF8FF`FFFFFF	1TB	System View PTEs
FFFFF900`00000000	FFFFF97F`FFFFFF	512GB	Session Space
FFFFF980`00000000	FFFFFA70`FFFFFF	1TB	Dynamic VA Space
FFFFFA80`00000000	FFFFFAFF`FFFFFF	512GB	PFN Database
FFFFFF`FFC00000	FFFFFF`FFFFFF	4MB	HAL Heap

Virtual Address Translation

- Hardware translates each virtual address to a physical address



x86 Virtual Address Translation



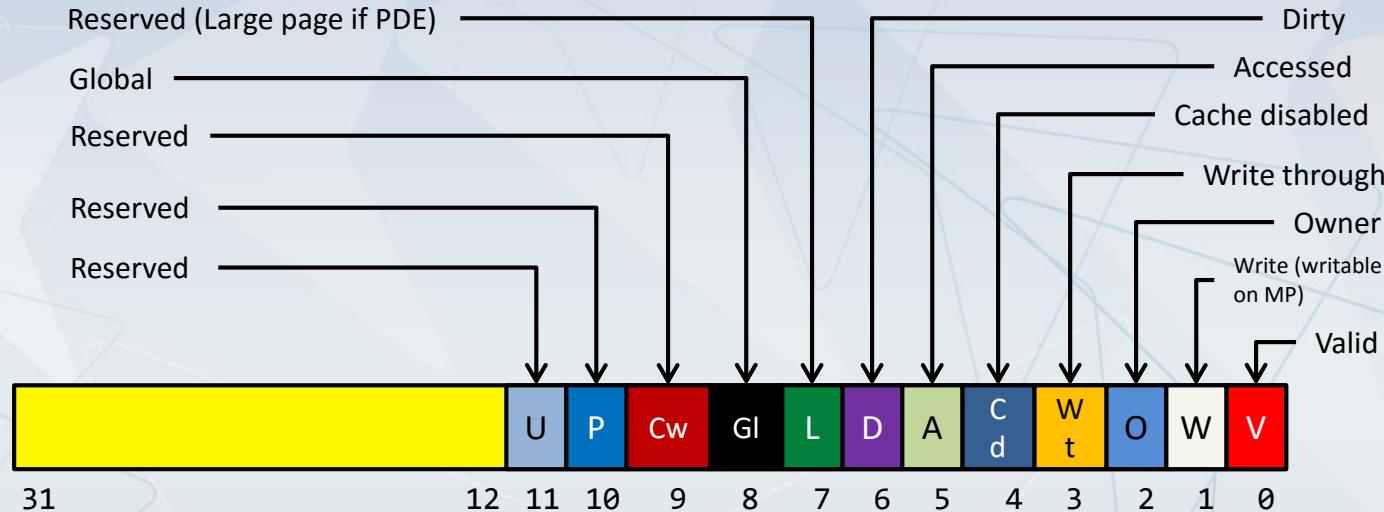
Page Directory

- One per process
 - Mapped to virtual address 0xC0300000 (0xC0600000 on PAE systems)
- Physical address of page directory stored in KPROCESS structure
- While a thread is executing, the CR3 register stores its address
 - When a thread context switch occurs between threads of different processes, **CR3** is reloaded from the appropriate **KPROCESS** instance

x86 PDE and PTE

- Each entry is 32 bits (64 bits on PAE)
- Upper 20 bits (25 bits on PAE) is the Page Frame Number (PFN)
- Bit 0 is the Valid bit
 - If set, indicates that the page is in RAM
 - Otherwise, accessing the page caused a page fault
- Other bits exist (see next slide)

Valid x86 PTE/PDE layout



- Dirty – page has been written to
- Large page – this maps a large page (2MB)
- Accessed – page has been read
- Owner – user mode or kernel mode accessible

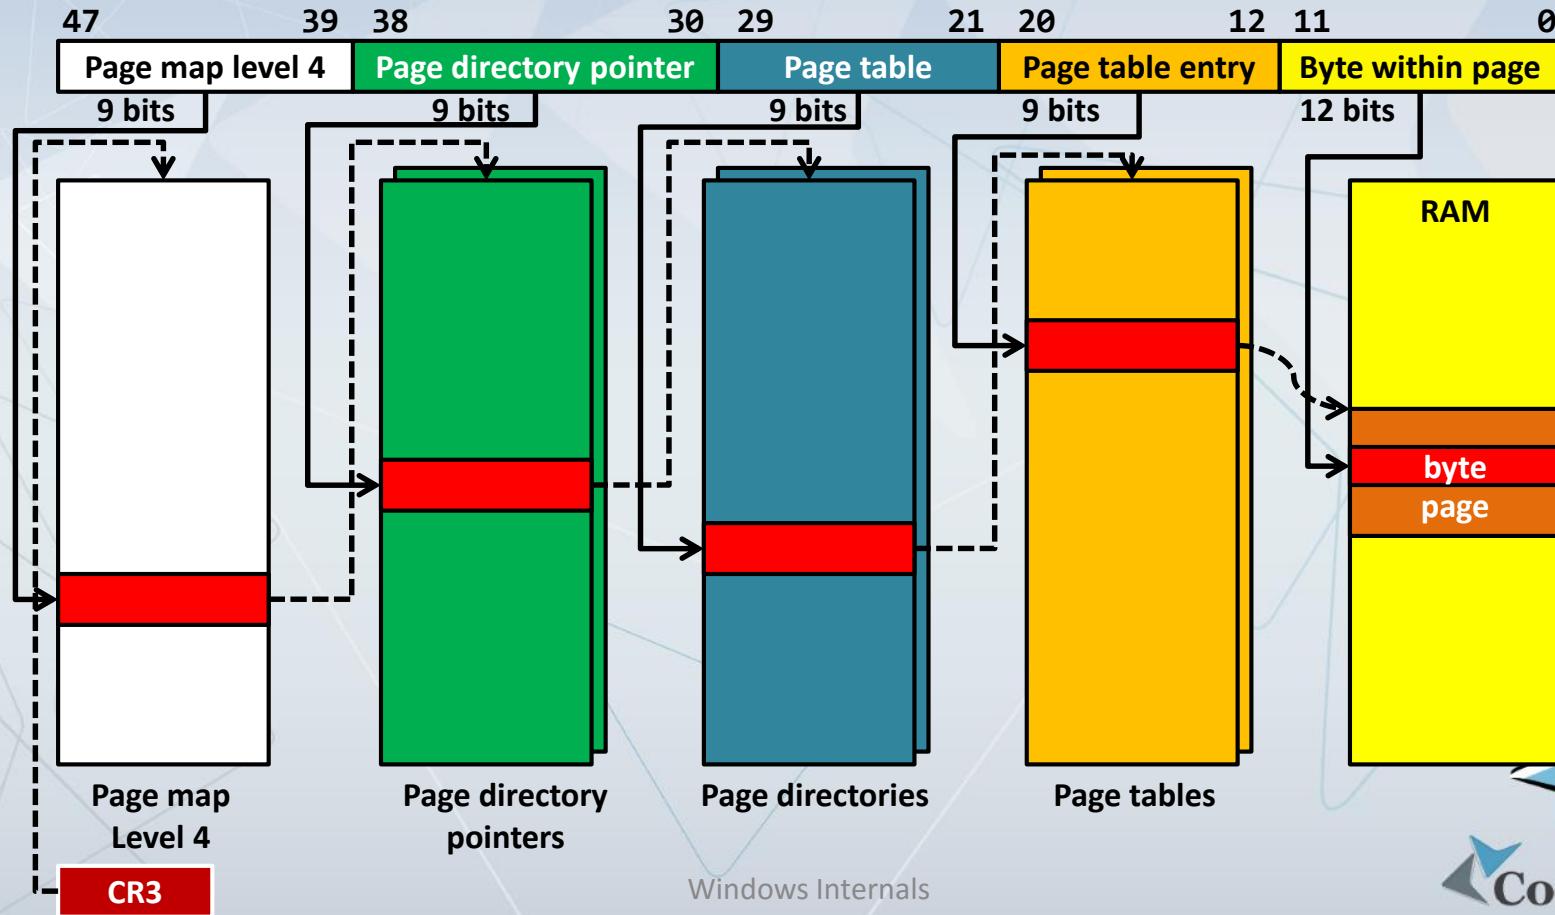
Physical Address Extensions (PAE)

- Intel Pentium Pro and up processors support a new PAE mode
- Current hardware allows access to up to 64GB of physical memory
- Virtual address translation contains an extra level of indirection
- Each PTE/PDE is 64 bits of which 25 are the PFN
 - Thus maximum is $2^{(25+12)} = 2^{37} = 128\text{GB}$

Address Windowing Extensions

- Temporary solution (on 32 bit) to providing access to large amounts of physical memory (>4GB)
- Requires the PAE kernel
- System cache uses the extra memory even if applications don't
- Applications can allocate the above physical memory and map it to a window in their process address space
 - Call [AllocateUserPhysicalPages](#)
 - Call [VirtualAlloc](#) with MEM_PHYSICAL and MEM_RESERVE flags
 - Map to a window with [MapUserPhysicalPages](#)

x64 Virtual Address Translation



Page Faults

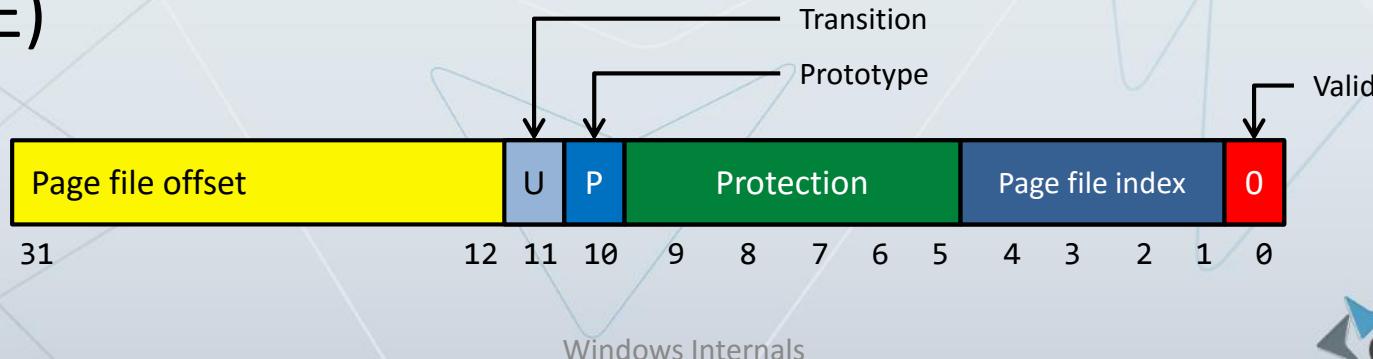
- Valid PTE/PDE results in the CPU accessing data in physical memory
- If PTE/PDE is invalid, the CPU throws a page fault exception
 - Expects the OS to handle it and instruct it to retry the memory access
- Windows has to get the data from disk, fix the required PTE/PDE and let the CPU try again
- Page fault types
 - Hard page fault – requires disk access
 - Soft page fault – does not require disk access
 - Example: a needed shared DLL is simply directed to the process by pointing PTE to it

Some Reasons For Faults

Reason for fault	Result
Accessing a page that is not in RAM but in a page file or mapped file	Allocate a physical page, read the data from disk and add page to the working set
Accessing a page that is in the modified or standby page list	Move the page to the working set
Accessing a page that is not committed	Access violation
Accessing a page from user mode that can only be accessed from kernel mode	Access violation
Writing to a page that is read only	Access violation
Accessing a demand zero page	Add a zeroed page to the working set
Writing to a guard page	Guard page violation (if part of a thread's stack, commit more memory and add to working set)
Writing to a copy-on-write page	Make a process private page copy and replace in working set
Executing code in page marked no-execute	Access violation

Invalid PTEs

- The CPU throws a page fault exception when the Valid bit (bit 0) in a PTE is clear
- Windows uses the other PTE bits to indicate where the required page can be found
- Example: a page that resides in a page file (x86 w/o PAE)



Page File(s)

- Backup storage for writeable, non-shareable committed memory
 - Up to 16 page files are supported
 - On different partitions
 - Initial size and maximum size can be set
 - Using the System applet in Control Panel
 - Named PageFile.Sys on disk (root partition)
 - Created contiguous on boot
 - Initial value should be maximum of normal usage
- Page file information in the registry

HKLM\SYSTEM\CurrentControlSet\Control\Session Manager\Memory Management\PagingFiles



Page File Sizes

- If selecting “system managed size” in System applet

RAM size	Initial page file size	Maximum page file size
< 1GB	$1.5 * \text{RAM}$	$3 * \text{RAM}$
$\geq 1\text{GB}$	$1 * \text{RAM}$	$3 * \text{RAM}$

- When page files space run low
 - “System running low on virtual memory”
 - First time: Before page file expansion
 - Second time: When committed bytes reaching commit limit
 - “System out of virtual memory”
 - Page files are full

Commit Charge

- Commit charge represents the memory than can be committed
 - In RAM and existing page file(s)
- Contributors to the commit charge
 - Private committed memory ([VirtualAlloc](#) with MEM_COMMIT flag)
 - No RAM or page file is used until memory is actually touched
 - Until then, considered demand zero pages
 - Page file backed memory mapped file allocated with [MapViewOfFile](#)
 - Copy on write mapped memory
 - Kernel non-paged and paged pools
 - Kernel mode stacks
 - Page tables and yet-to-be-created page tables
 - Allocations with Address Windowing Extensions (AWE) functions
- The commit limit is basically the amount of RAM plus maximum size of all page files

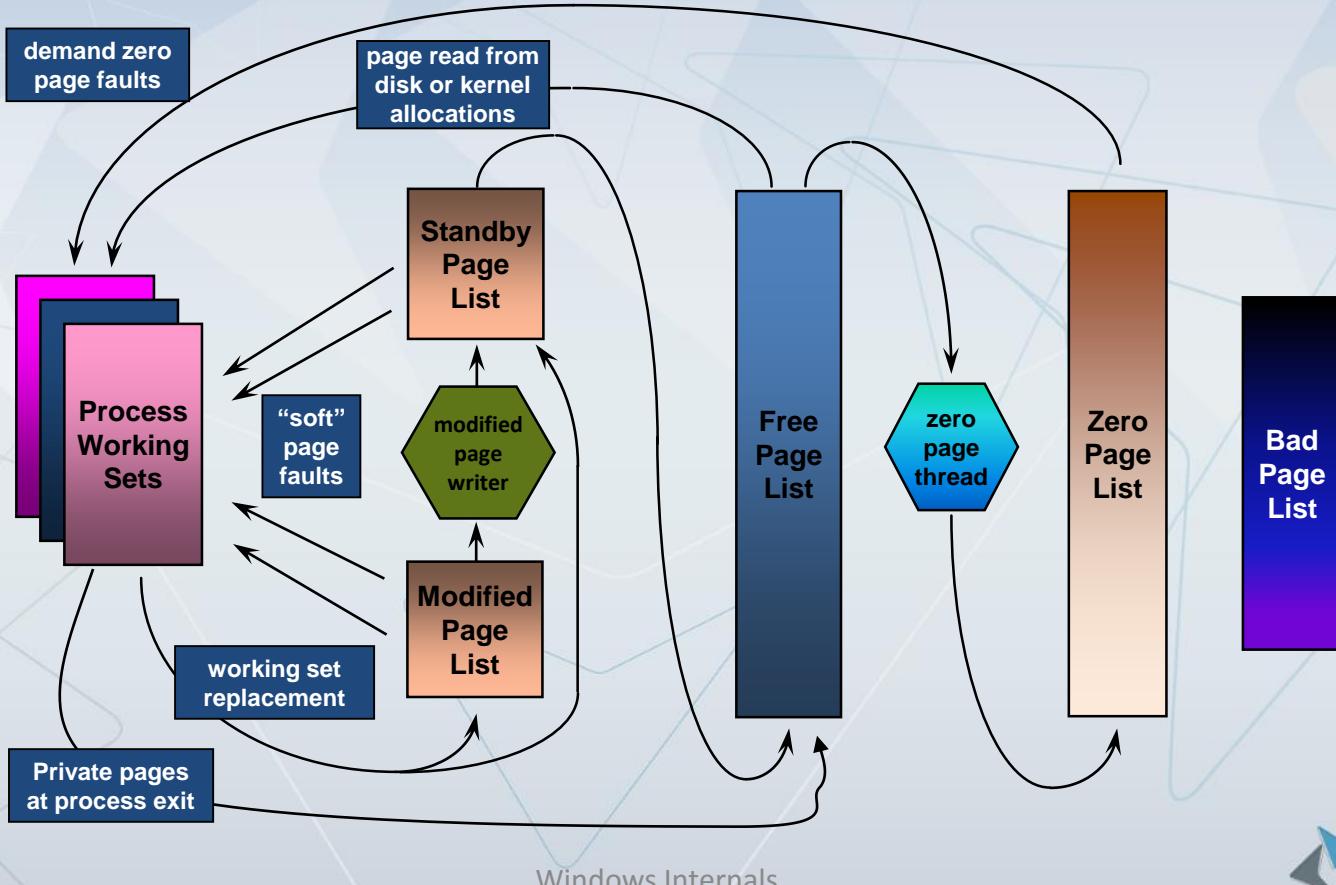
Working Sets

- Process Working Set
 - The subset of the process' committed memory that resides in physical memory
- System working set
 - The subset of system memory residing in physical memory
- Systems with terminal services
 - Some kernel memory is on a per session basis and as such working set as well
- Demand paging
 - When a page is needed from disk, more than one is read at a time to reduce I/O
 - 3 pages for code based images, 7 pages for others

Page Frame Number Database

- PFN database describes the state of all physical pages
- Valid PTEs point to entries in the PFN database, and a PFN entry points back to the PTE
- The structure layout of a PFN entry depends on the state of the page
- Kernel debugger: [!memusage](#), [!pfn](#)

Page Dynamics



Physical Page States

State	Description
Active (Valid)	The page is part of a working set (either process or kernel) or is not a part of any working set (kernel non-paged pool) with a valid PTE pointing to it
Transition	Temporary state when the page is in the middle of I/O. the PTE is encoded so that collided pages can be recognized and handled
Standby	The page belonged to a working set but was removed. It has not been modified since it was last written to disk. The PTE still refers to the page but is marked invalid and in transition
Modified	Similar to Standby, but the page was modified, so cannot be discarded until it's written back to disk
Modified no write	Similar to modified, but will not be written to disk. The cache manager does this at the request of file system drivers
Free	The page is free but contains garbage data in it
Zeroed	A free page that is zeroed out
Bad	Some hardware error, so page cannot be used

Standby and Modified Page Lists

- Used to
 - Avoid writing pages back to disk too soon & avoid releasing pages to the free list too soon
- The system can replenish the free page list by taking pages from the top of the standby page list
 - This breaks the association between the process and the physical page
- Pages move from the modified list to the standby list
 - Modified pages' contents are copied to the pages' backing stores (usually the paging file) by the modified page writer (see next slide)
 - The pages are then placed at the bottom of the standby page list
 - On Windows NT 10.X pages are compressed into the System process user space
 - Only in a case of low memory, compressed pages are written to the page file
- Pages can be faulted back into a process from the standby and modified page list
 - The SPL and MPL form a system-wide cache of “pages likely to be needed again”

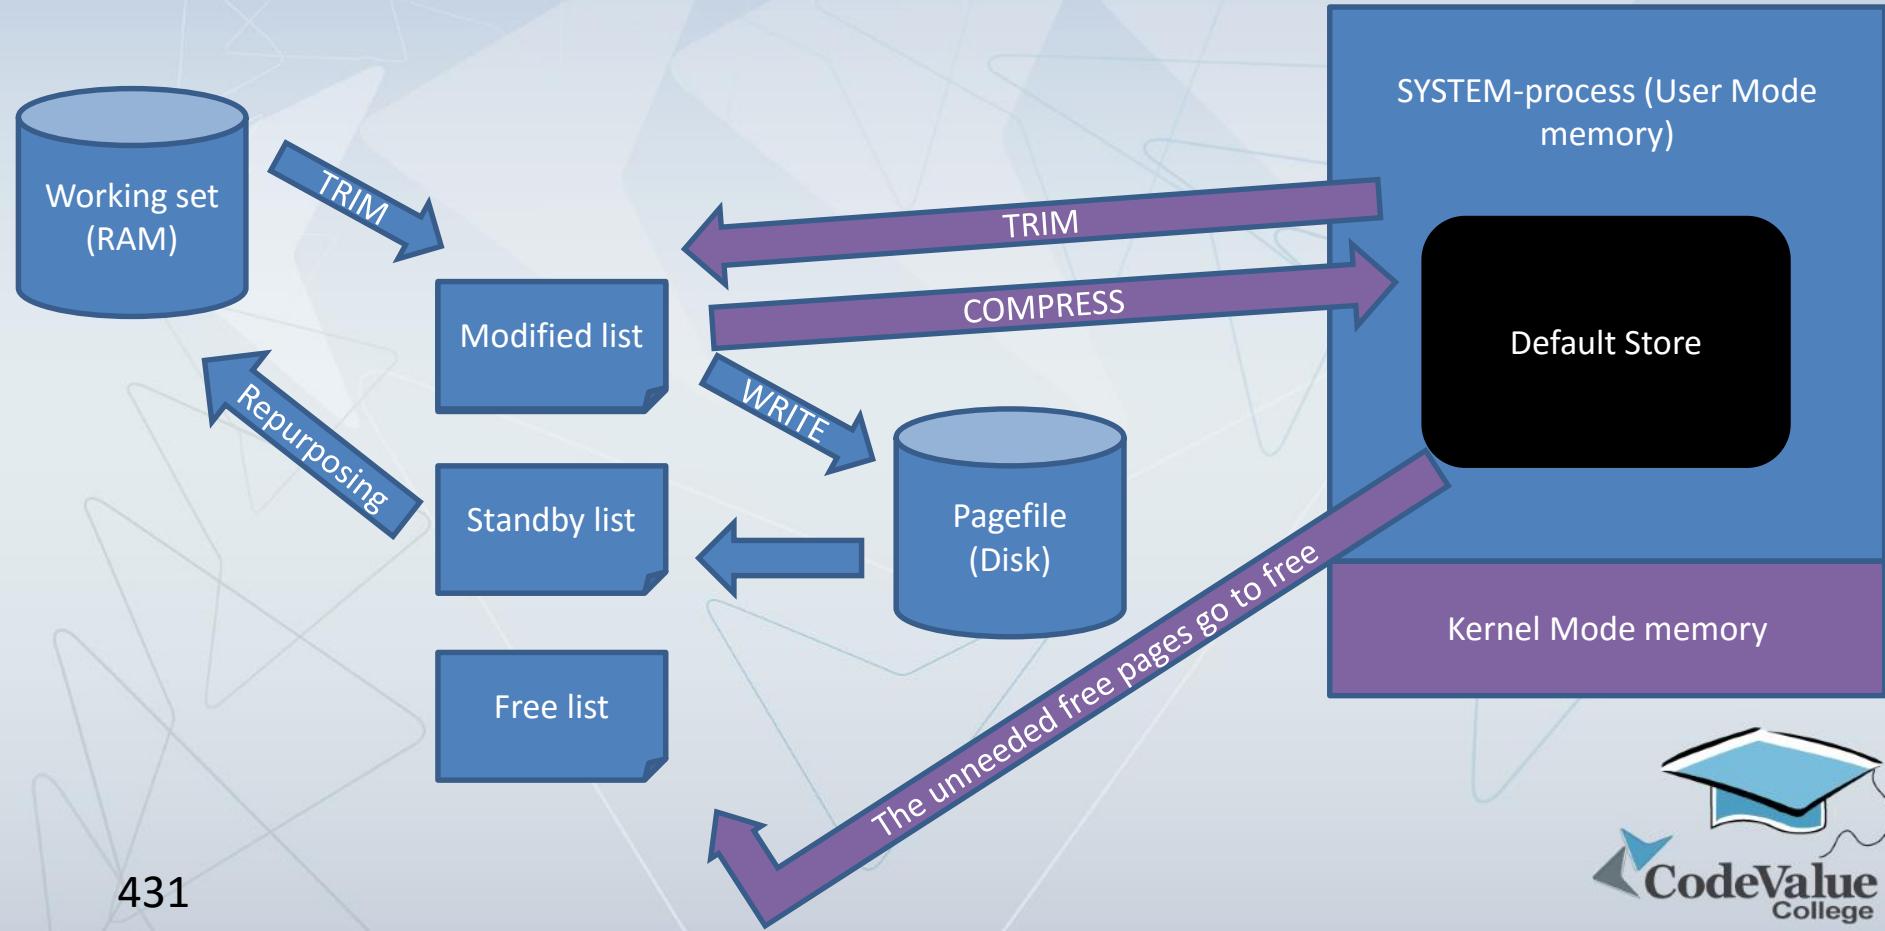
Modified Page Writer

- Moves pages from modified to standby list, and copies their contents to disk
 - i.e., this is what writes the paging file and updates mapped files (including the file system cache)
- Two system threads (priority 17)
 - One for mapped files, one for the paging file
- Triggered when
 - Memory is overcommitted (too few free pages)
 - Or modified page threshold is reached
 - Currently 800 pages
 - Does not flush entire modified page list

Zero Page List

- Large uninitialized data regions are mapped to demand zero pages
- On first reference to such a page, a page is allocated from the zero page list
 - No need to read zeroes from disk to provide the “data”
 - After modification, these pages are mapped to the paging file
- Zero page list is replenished by the “zero page thread”
 - The only thread running at priority 0
 - One per system (even on SMP)
 - Takes pages from the free page list, fills them with zeroes, and puts them on the zero page list while the CPU is otherwise idle
 - Usually is waiting on an event - which is set when, after resolving a fault, system notices that zero page list is too small

Memory Compression (NT 10.X)



Memory compression

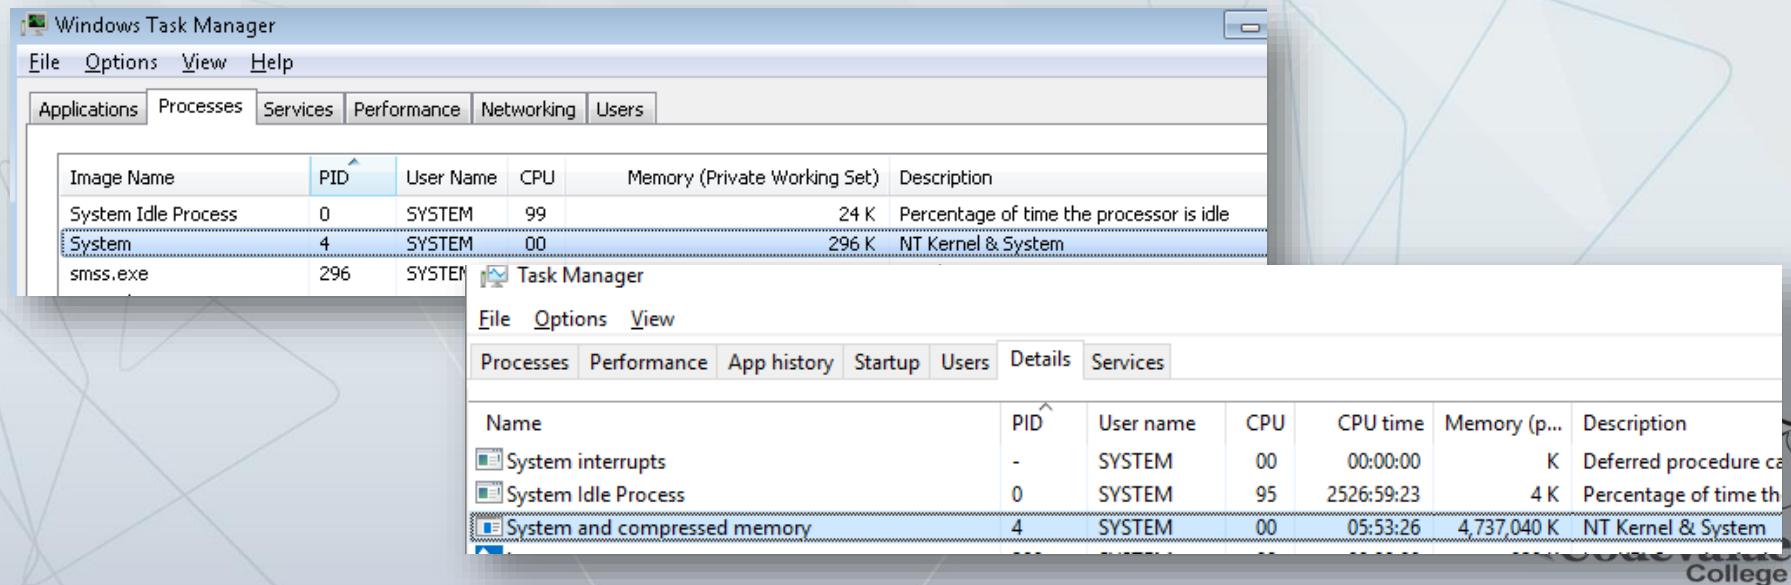
- Can compress pages in memory and result in less Hard Page Faults
 - Does not prevent the use of pagefile but lowers the need to use it
- History
 - Readyboost technology (Windows 7) used a compressed cache on a USB stick to improve performance to store data also written to the pagefile
 - Windows 8 evolved this by moving the cache into regular memory
 - But data was also still written out to the pagefile
 - The Windows 8.1 changed this by the cache switching to now being its own compressed pagefile in RAM that could be paged out to disk if required
 - Windows 10 changes this to now apply to all systems and all types of applications
 - Any data you now find in your disk-based pagefile is compressed after initially being compressed in the RAM based pagefile
 - It is only written to disk as a lack of resources required

Memory Compression – Why?

- Unused memory is unused resources
- Memory is 100000 times faster than a spindle and also 50 times faster than an SSD
- SSD's wear out
- After utilizing this most people don't need to page out to disk at all (statistically)

Memory Compression – How?

- Memory gets put into stores
 - Modern/Universal apps have their own store pre process
 - If app doesn't have its own then the SYSTEM-process' store is used
 - This might surprise some as the memory footprint of SYSTEM has been a few kilos for the past 25 years



The image shows two overlapping windows of the Windows Task Manager. The top window is titled 'Windows Task Manager' and displays the 'Processes' tab. It lists three processes: 'System Idle Process' (PID 0), 'System' (PID 4), and 'smss.exe' (PID 296). The 'System' process is highlighted. The bottom window is also titled 'Windows Task Manager' and displays the 'Details' tab. It lists three entries: 'System interrupts' (PID -), 'System Idle Process' (PID 0), and 'System and compressed memory' (PID 4). The 'System and compressed memory' entry is at the bottom of the list.

Image Name	PID	User Name	CPU	Memory (Private Working Set)	Description
System Idle Process	0	SYSTEM	99	24 K	Percentage of time the processor is idle
System	4	SYSTEM	00	296 K	NT Kernel & System
smss.exe	296	SYSTEM			

Name	PID	User name	CPU	CPU time	Memory (p...)	Description
System interrupts	-	SYSTEM	00	00:00:00	K	Deferred procedure call
System Idle Process	0	SYSTEM	95	2526:59:23	4 K	Percentage of time th
System and compressed memory	4	SYSTEM	00	05:53:26	4,737,040 K	NT Kernel & System

Memory Compression – Price?

- Compression/Decompression requires CPU so is it worth it?
 - Yes! It definitely is!!!
 - Decompression rates:
 - 3 Year old laptop = 300MB/s per core
 - Current laptop with 8 cores = 7GBps
 - Lowest end Windows phone = 250-300MB/s
- Compression and decompression are low priority tasks as long as the computer is not critically low on memory

Memory Compression

- Also organizes memory so that it's not fragmented
 - When pages enter a store that store is created a contiguous reservation in pagefile so that when that store is paged into pagefile they will not be fragmented
- Pages are compressed in page file as well
 - Less reads needed to get the same amount of data
 - Decompressed outside of the pagefile

Secure Enclave

en|clave

[ɛnkleɪv] 

NOUN

1. a portion of territory surrounded by a larger territory whose inhabitants are culturally or ethnically distinct.

"the besieged Muslim enclave of Srebrenica"

synonyms: area of land · area · region · enclave · country · state · [more]

- a place or group that is different in character from those surrounding it:

"the engineering department is traditionally a male enclave"

Powered by [Oxford Dictionaries](#) · © Oxford University Press · Translation by [Bing Translator](#)

Translations, word origin, and more definitions 

Enclave (NT 10 Update 2)

- Provide an APIs to use the new [Intel SGX Skylake](#) CPU capability – Software Guard Extensions
 - The system BIOS also needs to support the feature
- A method baked by the CPU to isolate a region of memory pages
- Only code executed inside the enclave can reach the data inside the enclaved region
 - First check: [IsEnclaveTypeSupported](#)
 - Then create: [CreateEnclave](#) (Kernel NtCreateEnclave)
 - Load data with: [LoadEnclaveData](#) (NtLoadEnclaveData)
 - Initialize it (begin code execution): [InitializeEnclave](#) (NtInitializeEnclave)
 - Destroy it: [VirtualFree](#) (NtVirtualFree)

Memory Manager Changes

- To support the feature In Windows 10 Update 2, the kernel adds a new Enclave Page List
- Internally, however, since the memory manager has actually run out of list identifiers, the kernel actually identifies these pages as being “bad” pages currently suffering an in-page error
 - This actually makes sense – the memory manager knows never to use bad pages, so calling enclave pages “bad” is another way to keep them at bay

System Memory Usage

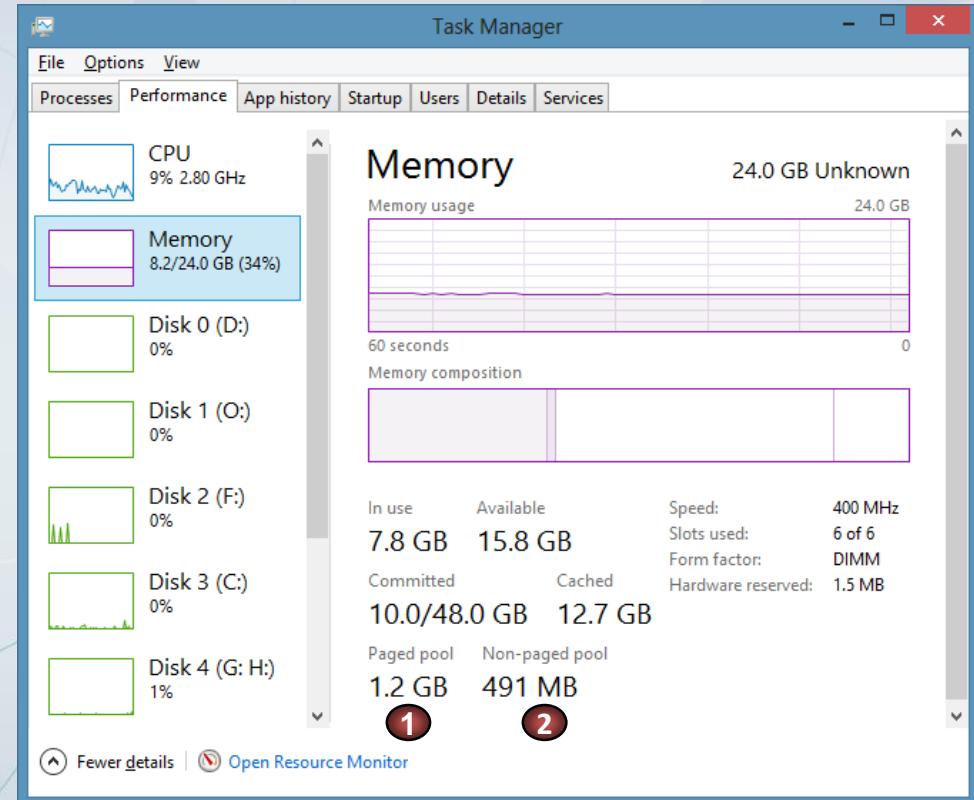
- Kernel and driver memory usage breaks down into:
 - Non-pageable code
 - Pageable code
 - File system cache
 - Non-paged pool
 - Paged pool

System Memory Pools

- The kernel provides two general memory pools for use by the kernel itself and device drivers
 - Non-paged pool
 - Memory always resides in RAM (never paged out)
 - Can be accessed at any IRQL
 - Paged pool
 - Memory can be swapped to disk
 - Should be accessed at $\text{IRQL} < \text{DPC_LEVEL (2)}$ only
- Pool sizes depend on the amount of RAM and the OS type (Professional vs. Servers)
 - Can be altered (up to some maxima) in registry
 - `HKLM\System\CurrentControlSet\Control\Session Manager\Executive`
- Task Manager displays current sizes (in RAM)

System Pools in Task manager

- 1 “Kernel Memory Paged” = physically resident size of paged pool
- 2 “Kernel Memory Nonpaged” = physical size of non-paged pool



System Memory Pools APIs

- [ExAllocatePool](#) (obsolete)
 - Allocate memory from the paged or non-paged pool
- [ExAllocatePoolWithTag](#)
 - Allocate memory and “tag” it with a 4-byte value
 - Used to track memory leaks
- [ExFreePoolWithTag](#)
 - deallocates a block of pool memory allocated with the specified tag
- Pool usage (and tags) can be viewed with [PoolMon.Exe](#)

C:\Program Files (x86)\Windows Kits\10\Tools\x64\poolmon.exe					
		Memory: 33376124K Avail: 15884660K PageFlts: 6382 InRam Krn! : 142416K P: 203773	Commit: 37505328K Limit: 69027708K Peak: 46959020K Pool N: 832660K P: 214		
		System pool information			
Tag	Type	Allocs	Frees	Diff	Bytes
dpG	Nonp	1 (0)	0 (0)	1	320 (0)
ger	Paged	439 (0)	437 (0)	2	8256 (0)
-UMD	Nonp	27 (0)	24 (0)	3	144 (0)
.UMD	Nonp	2 (0)	1 (0)	1	128 (0)
2UuQ	Nonp	67 (0)	63 (0)	4	16384 (0)
3DMN	Paged	1 (0)	0 (0)	1	144 (0)
ACPI	Nonp	3 (0)	1 (0)	2	96 (0)
ACPX	Paged	4 (0)	4 (0)	0	0 (0)
ACPX	Nonp	6 (0)	6 (0)	0	0 (0)
AEio	Nonp	36 (0)	36 (0)	0	0 (0)
AEt!i	Nonp	6 (0)	6 (0)	0	0 (0)
AEwi	Nonp	12 (0)	12 (0)	0	0 (0)
ALPC	Nonp	1503806 (2)	1498481 (0)	5325	3168816 (1184)
APL_	Paged	12476 (0)	4428 (0)	8048	1076144 (0)
APpt	Paged	9927306 (0)	9923921 (0)	3385	410400 (0)
APPt	Nonp	2455 (0)	960 (0)	1495	191360 (0)
ARFT	Paged	1273 (0)	1270 (0)	3	192 (0)
AVNw	Nonp	1 (0)	0 (0)	1	4480 (0)
AVTs	Nonp	1 (0)	0 (0)	1	352 (0)
AVNs	Nonp	1 (0)	0 (0)	1	48 (0)
AVpc	Nonp	1 (0)	0 (0)	1	4480 (0)
AVpd	Nonp	1 (0)	0 (0)	1	4480 (0)
AVvi	Nonp	2 (0)	0 (0)	2	1568 (0)
AaAu	Nonp	91 (0)	7 (0)	84	37184 (0)
AaCm	Paged	10 (0)	10 (0)	0	0 (0)
AaCm	Nonp	215 (0)	204 (0)	11	2880 (0)
AaFg	Nonp	6 (0)	2 (0)	4	400 (0)
AaMx	Nonp	9 (0)	2 (0)	7	1456 (0)
AaPx	Nonp	2 (0)	0 (0)	2	128 (0)
AaRT	Nonp	2 (0)	1 (0)	1	416 (0)



Memory Mapped Files

- Internally called *Section* objects
- Allows the creation of “views” into the file
 - Returns a memory pointer for data manipulation
- Implies shared memory capabilities
 - This is the usual case with mapping EXEs and DLLs
- Also can create “pure” shared memory
 - Backed up by the system paging file
 - When memory mapped file object destroyed, memory is “recycled”



Memory Mapped Files API

- Win32
 - [CreateFileMapping](#) – create a file mapping object based on a specific file (previously created with `CreateFile`) or based on system paging file
 - [OpenFileMapping](#) – open an existing MMF based on its name (NOT the filename)
 - [MapViewOfFile\(Ex\)](#) – create a “view” into the MMF
- Kernel
 - [ZwCreateSection](#) – creates a section object (if based on a file, call [ZwCreateFile](#) first)
 - [ZwMapViewOfSection](#) – map a “view” into system space

Large Pages

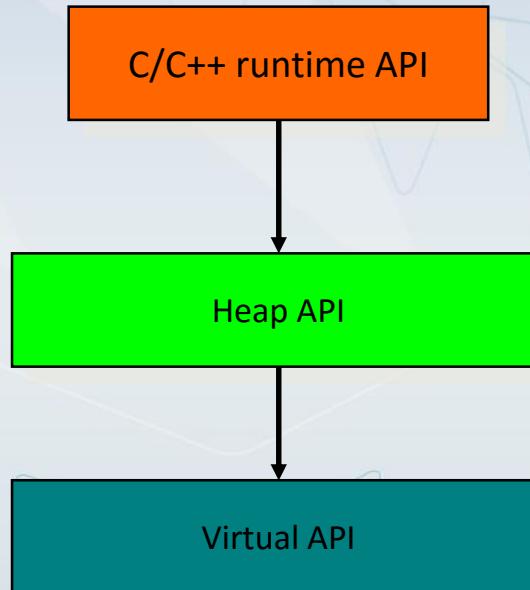
- Large pages allow mapping using a PDE only (no need for PTEs)
 - Advantage is better use of the translation look aside buffers
- Windows maps by default large pages for NtOsKrn1.Exe and Hal.Dll as well as core system data (initial non paged pool and PFN database)
- Potential disadvantage
 - Single protection to entire page
 - Thus, kernel is protected with read/write (can't protect with read only because data is there as well)

Applications Using Large Pages (>= NT 5.2)

- Applications can use large pages for their allocations by specifying the `MEM_LARGE_PAGE` in calls to the [VirtualAlloc](#) function
- Size and alignment must be multiple of large page size
 - Can be determined by calling the [GetLargePageMinimum](#) function
- Must have the [SeLockMemoryPrivilege](#) privilege

Memory APIs in User Mode

High Level



Low Level

Windows Internals

The Heap Manager

- Allocating in page granularity is sometimes too much
 - Need fine grained control
- The heap manager manages smaller allocations (8/16 bytes minimum)
- API exported from NtDll.Dll for user mode and NtOsKrn1.Exe from kernel mode
- The HeapXxx Windows API functions are a thin wrapper over the native NtDll.Dll functions

Heap Types

- One heap is always created with a process, called the Default Process Heap
 - Can be accessed using [GetProcessHeap](#)
- Additional heaps can be created using the [HeapCreate](#) function
- A heap can be fixed in size or growable
 - The default heap is growable
- Low Fragmentation Heap (LFH)

Windows Heaps



Main Heap API Functions

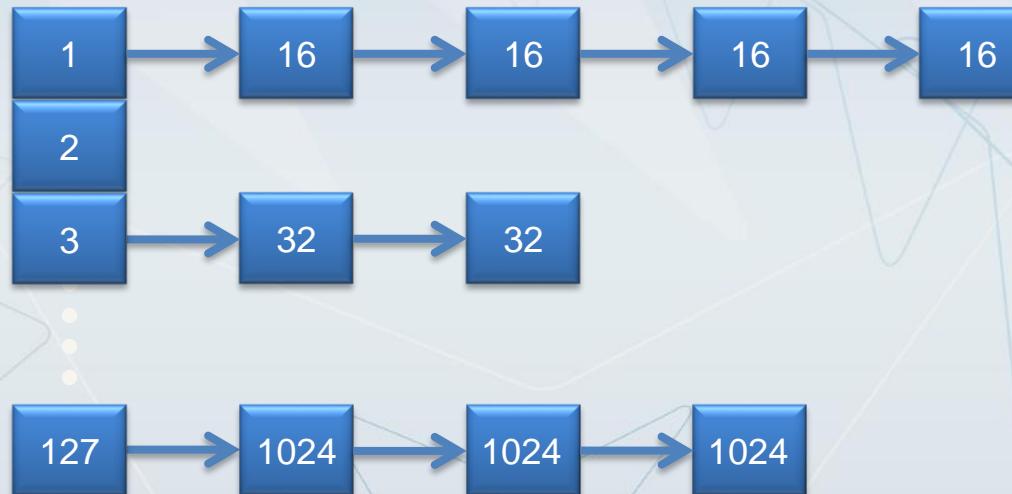
- HeapCreate
 - Creates a new heap with the specified committed and reserved sizes (can be expandable)
- HeapDestory
 - Destroy a heap
- HeapAlloc, HeapFree, HeapReAlloc
 - Allocate, free or re-allocate from a heap
- HeapWalk
 - Enumerate the contents of a heap
- GetProcessHeap, GetProcessHeaps
 - Return a handle to the default heap / an array of all process heaps

The Heap Data Structure

- The heap is built from two main components
 - Front-End Allocator
 - Fast, fixed size based memory allocator
 - Back-End Allocator
 - Complex, second chance memory allocator
- When a request is not satisfied by the Front-End allocator, it goes to the Back-End allocator
- Currently (\geq NT 6.x) there are two Front End allocators:
 - Look aside list allocator
 - Low fragmentation allocator
- Additional deep information:
 - [Understanding Low Fragmentation Heap](#)

The LAL Heap Data Structure

- In general the front-end allocator is built from an array of free look aside lists with fixed size blocks

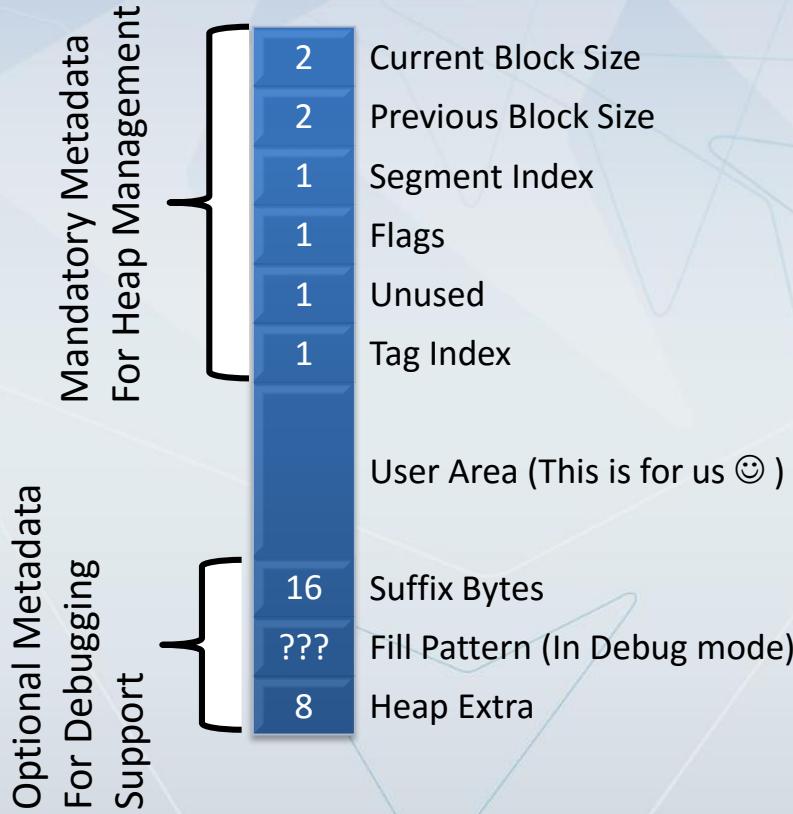


The Back End Allocator Data Structure

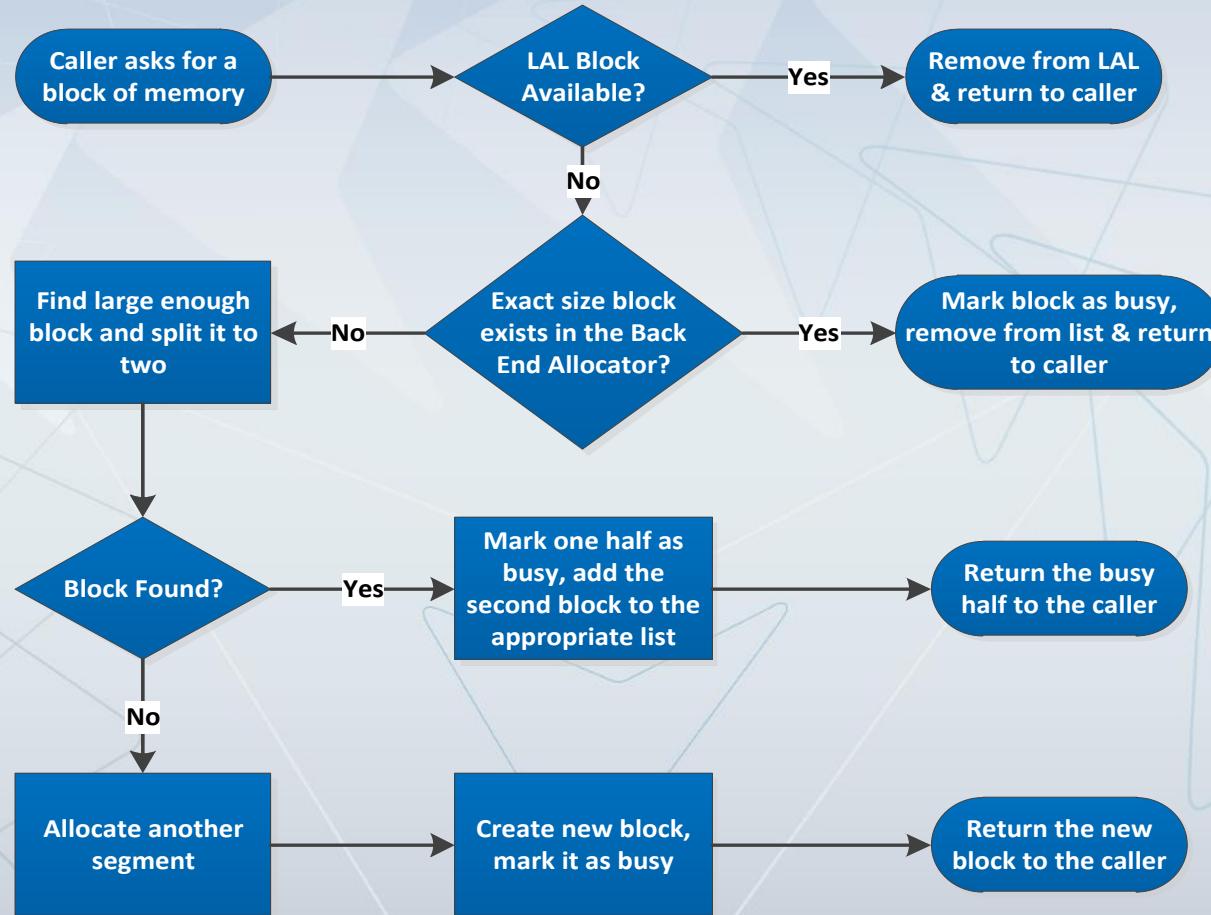
- The back-end maintains one varied block size free list and 126 fixed size look aside free lists



Each Block Data Structure



Heap Allocation Algorithm



Summary

- The Memory Manager provides various services to applications and system components
- Virtual memory is translated to physical memory while using page files as backup
- Sharing pages is implemented to minimize RAM usage and share data between processes
- User mode applications have plenty of memory related API and services

Module 6

SECURITY

459

Windows Internals



Agenda

- Windows Security Features
- Security Components
- User Access Control (UAC)
- Access Tokens
- Protecting Objects
- Privileges
- New Security features in NT 10.X
- Summary

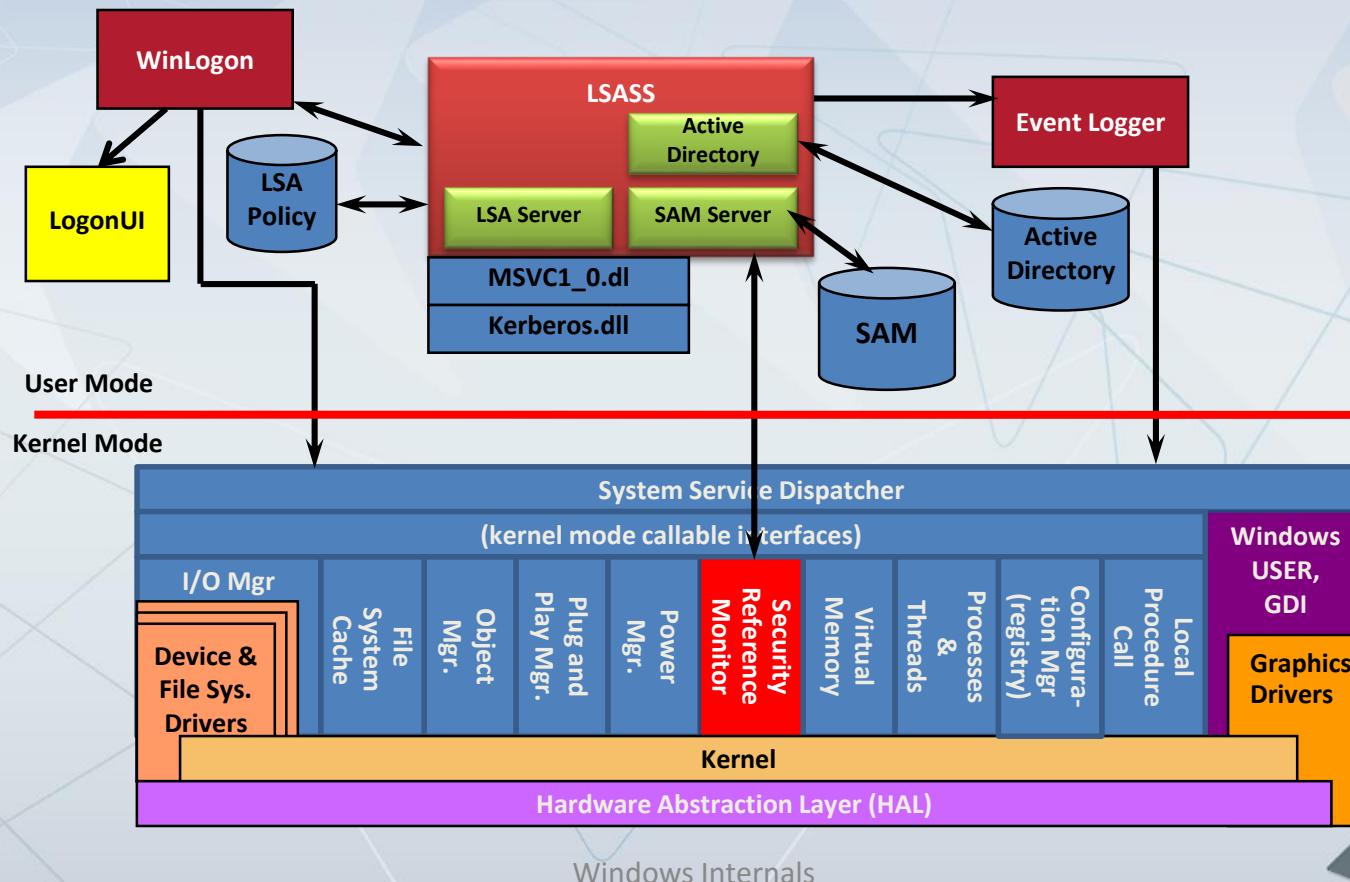
Windows Security Features

- Design for C2 level security of the DoD
- Permissions can be applied to all shareable resources
 - Including NTFS file system (but not FAT)
- NTFS encryption protects data from unauthorized physical access
- Support for Kerberos authentication
- Support for Digital Certificates
- Cryptographic API for encryption and hashing

Authentication and Authorization

- Authentication
 - Be sure that the user is who she claim she is
 - It is also about identity
- Authorization
 - Allow and deny read/write/control access to information, settings and resources
- There are other aspects such as integrity, Denial of Service, Cryptography , Intrusion Detection, Malware (Viruses , Worms) and more

Security Components Overview



Security System Components

- Security Reference Monitor (SRM)
 - Part of the Executive, providing access token definition, access checks on objects, manipulating privileges and auditing messages
- Local security authority subsystem (Lsass)
 - Running the image \Windows\System32\Lsass.exe
 - Responsible for local security policy, user authentication and sending events to the Event Log
 - Most functionality implemented by the Local security authority service (Lsass.dll)
- Lsass policy database
 - Local system security policy stored under HKLM\SECURITY

Security System Components

- Security Account Manager (SAM) service
 - A set of functions to manage the users and groups on the local system
 - Implemented in \Windows\System32\SamSrv.dll (runs under Lsass.Exe)
- SAM database
 - On non-domain controller systems, contains local users and groups (along with passwords)
 - On domain controller systems, stores admin recovery account data
 - Stored in the registry under HKLM\SAM

Security System Components

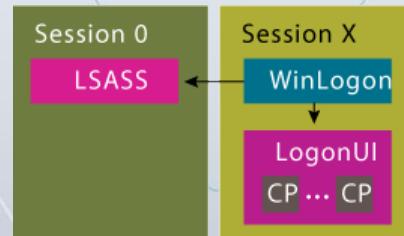
- Active Directory
 - A directory service storing objects data in a computer domain
 - Stored on computers designated as Domain Controllers
 - Runs as a service (\Windows\System32\Ntdsa.dll) under Lsass
- Authentication packages
 - DLLs that run in the context of Lsass or other processes that implement an authentication policy (e.g. returning whether a user may log in)
 - The returned information allows Lsass to generate a token

Security System Components

- Logon process
 - Running the image \Windows\System32\Winlogon.Exe
 - Responsible for responding to the SAS (Secure Attention Sequence, by default Ctrl+Alt+Del) and for managing interactive logon sessions (e.g. running the user's shell)

Authentication in Windows

- It begins with an instance of **WinLogon** process
 - In NT 6.x There is one WinLogon per session
- Usually at logon the user supply the user name and password
 - WinLogon supports custom credential provider
 - This enables scenarios such as Smart Card, Biometric or even face recognition based logon



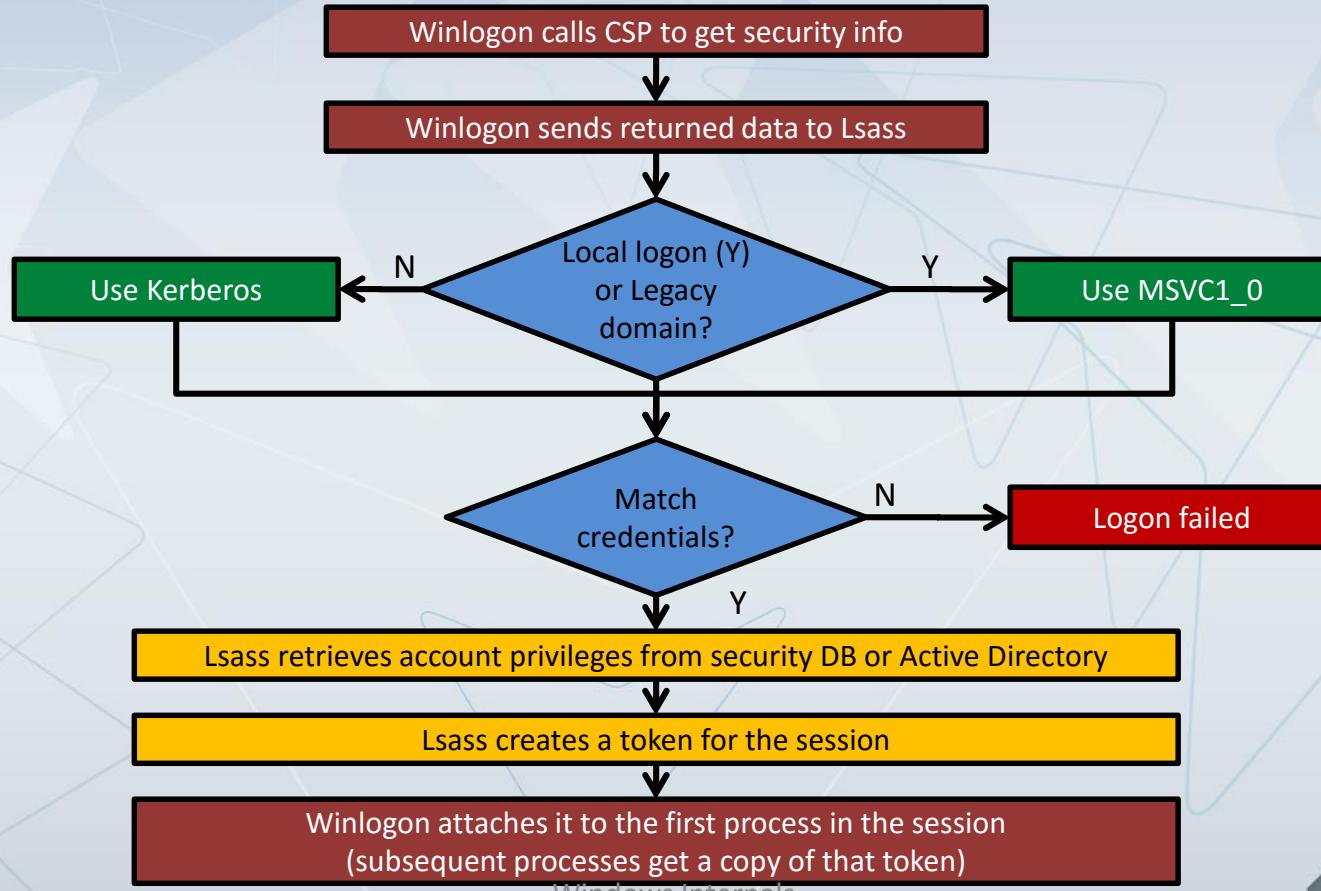
Credential Providers

- Prior to Windows Vista (< NT 6.0)
 - WinLogon loaded the Graphics Identification and Authentication (GINA) DLL to display a logon UI
 - Custom GINA hard to write
 - Only one can be configured
- Windows Vista (≥ 6.0) and later
 - The LogonUI.Exe process to display logon UI
 - Loads logon providers from registry info
 - Microsoft provides Interactive (authui.dll) and Smartcard (Smart-cardcredentialprovider.dll) providers
- Windows NT 10.X provides “[Windows Hello](#)”

Why Is Logon Secure?

- Winlogon creates the Winlogon desktop which is the only one visible
- Pressing SAS takes to that desktop
- Keyboard handler in the kernel disables hooks when SAS is detected
- Deregistering the SAS is only possible by the thread that registered it

Logon Sequence



Non-Interactive Logon

- A Network logon
 - Also called Non-Interactive logon
- Non-interactive authentication can only be used after an interactive authentication has taken place.
- In non-interactive authentication, the user does not input logon data,
 - Instead, previously established credentials are used
- The local LSA uses [SSPI](#), a standard security provider interface to pass the current logon information
- This enables Active Directory and Workgroup non-interactive logons

Additional Logon

- The [LogonUser](#) function attempts to log a user on to the local computer
 - The result is a primary access token
 - Use the [CreateProcessAsUser](#) to create a process for this user
 - Use [ImpersonateLoggedOnUser](#) to impersonate a thread
 - Call [RevertToSelf](#) to revert to self
- You can combine LogonUser & CreateProcessAsUser with [CreateProcessWithLogonW](#)

SID – Security Identifier

- A SID is a unique value used to identify a trustee
 - A trustee is one of User, Group, Domain, Well Known Group, Computer
- SIDs are issued by authority such as Domain Controller, or LSA
- Windows uses SIDs in:
 - Security Descriptors to identify the owner on an object
 - Access Control Entry, to identify the target trustee
 - Access Token, to identify the user and the groups

S-1-5-21-2280383352-4165795427-4139450486-1001

Well Known SIDs

- Windows defines some built in groups
- Examples

SID	Group	Description
S-1-1-0	Everyone	All users
S-1-2-0	Local	Users who log on physically
S-1-5-18	Local System	Local System account
S-1-5-20	Network Service	Network Service account
S-1-5-19	Local Service	Local Service account
S-1-5-32-544	Administrators	Administrators group

Authorization - The Access Token

- The Token is used for authorization
 - To check access to securable object
 - To perform a system task that requires privileges
- The token contains (partial list):
 - The SID for the user's account
 - SIDs for each group of which the user is a member
 - A logon SID that identifies the current logon session
 - A list of the privileges held by the user and the user's groups
 - The SID for the primary group
 - The default DACL
 - Whether the token is a primary or impersonation

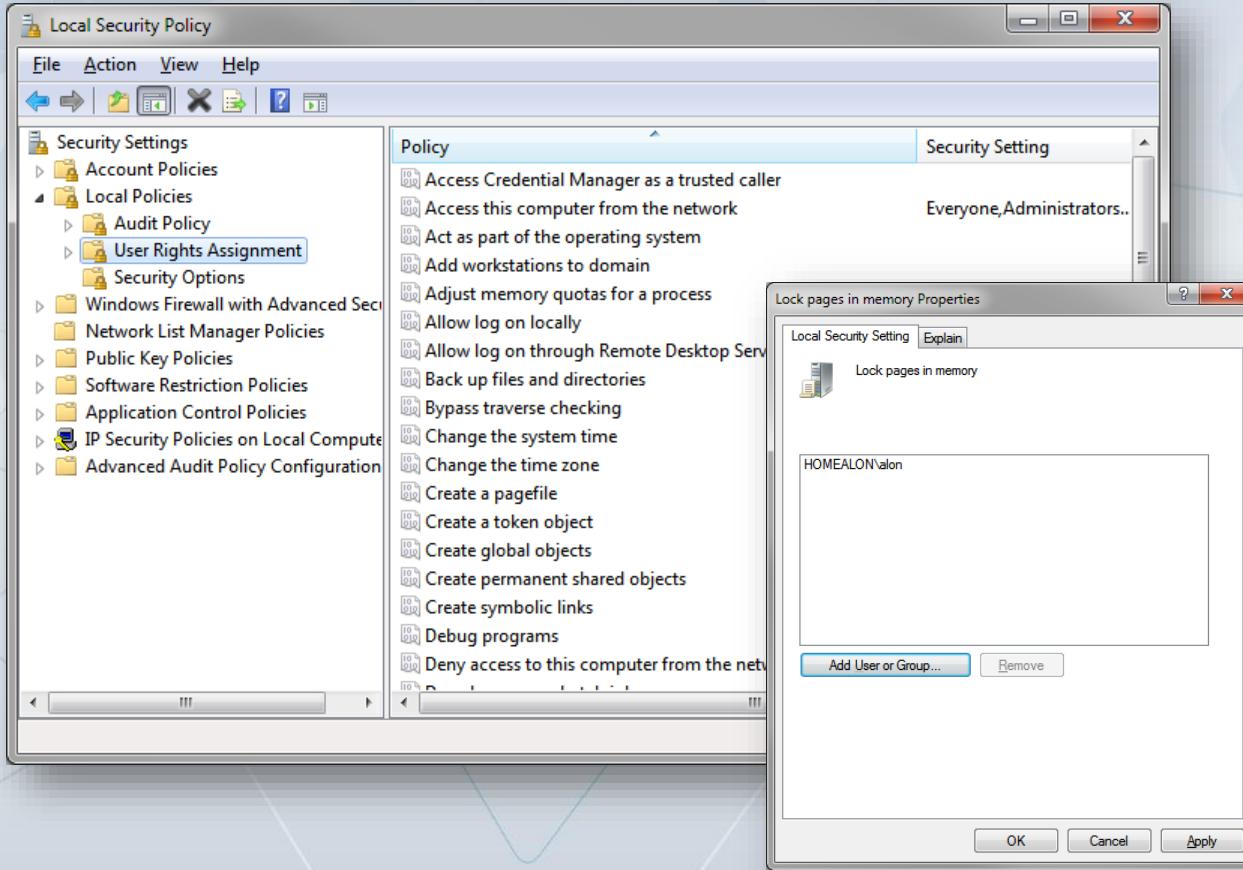
Using the Access Token

- Every process has a primary token
 - This is usually a copy of the logon token
- Each thread has a copy of the primary token
 - A thread can have another token – the impersonation token
 - Impersonation can be used in a server application to change the thread security context to match the client context
- To work with a token call: [OpenProcessToken](#)

Privileges

- The right to perform a system-related operation
 - Setting System times, add device, create page file,...
- For each user the system has a privileges db
 - The Access Token holds the list of privileges
 - Privilege can be enabled or disabled
- When trying to perform a privileged operation
 - The system checks the user's access token
 - The token should have the necessary privileges
 - And the privileges have to be enabled

Setting Privileges DB



User Account Control

- UAC is a security feature of \geq NT 6.x
 - It runs user applications in a security safe environment
 - It has changed the ecosystem to support running as a standard user
- UAC is based on a restricted filtered token
 - The user is not part of the Administrator group
 - The user has much less privileges
- Even if you add a privilege, it may not be there!
- We will talk more about UAC later

UAC Filtered Token

ProcExp + standard user notepad.exe + admin notepad.exe

DEMO

Debugging Privileges

Watch 1

Name	Value	Type
\$User	{...}	\$user register
Process	{...}	TOKEN
Name	HOMEALO	User Name
User SID	S-1-5-21-2	SID
Session Id	1	DWORD
Login Id	000566c7-	LUID
Impersonation Level	N/A (not in SECURITY_IMPERSONATION)	
Privileges	{...}	TOKEN_PRIVILEGES
SeLockMemoryPrivilege	Enabled	PRIVILEGE
SeIncreaseQuotaPrivilege	Disabled	PRIVILEGE
SeSecurityPrivilege	Disabled	PRIVILEGE
SeTakeOwnershipPrivilege	Disabled	PRIVILEGE
SeLoadDriverPrivilege	Disabled	PRIVILEGE
SeSystemProfilePrivilege	Disabled	PRIVILEGE
SeSystemtimePrivilege	Disabled	PRIVILEGE
SeProfileSingleProcessPrivilege	Disabled	PRIVILEGE
SeIncreaseBasePriorityPrivilege	Disabled	PRIVILEGE
SeCreatePagefilePrivilege	Disabled	PRIVILEGE
SeBackupPrivilege	Disabled	PRIVILEGE
SeRestorePrivilege	Disabled	PRIVILEGE
SeShutdownPrivilege	Disabled	PRIVILEGE
SeDebugPrivilege	Enabled	PRIVILEGE
SeSystemEnvironmentPrivilege	Disabled	PRIVILEGE
SeChangeNotifyPrivilege	Enabled by	PRIVILEGE
SeRemoteShutdownPrivilege	Disabled	PRIVILEGE
SeUndockPrivilege	Disabled	PRIVILEGE
SeManageVolumePrivilege	Disabled	PRIVILEGE
SeImpersonatePrivilege	Enabled by	PRIVILEGE
SeCreateGlobalPrivilege	Enabled by	PRIVILEGE
SeIncreaseWorkingSetPrivilege	Disabled	PRIVILEGE
SeTimeZonePrivilege	Disabled	PRIVILEGE
SeCreateSymbolicLinkPrivilege	Disabled	PRIVILEGE
Groups	{...}	TOKEN_GROUPS
Thread	No Token.	TOKEN

notepad.exe:11176 Properties

User: MEALON\alon

Image Performance Performance Graph Threads TCP/IP Security Environment Strings

Group Flags

BUILTIN\Administrators	Owner
BUILTIN\IIS_IUSRS	Mandatory
BUILTIN\Users	Mandatory
Everyone	Mandatory
HOMEALON\HelpLibraryUpdaters	Mandatory
HOMEALON\HomeUsers	Mandatory

Privilege Flags

SeBackupPrivilege	Disabled
SeChangeNotifyPrivilege	Default Enabled
SeCreateGlobalPrivilege	Default Enabled
SeCreatePagefilePrivilege	Disabled
SeCreateSymbolicLinkPrivilege	Disabled
SeDebugPrivilege	Disabled
SeImpersonatePrivilege	Default Enabled
SeIncreaseBasePriorityPrivilege	Disabled
SeIncreaseQuotaPrivilege	Disabled
SeIncreaseWorkingSetPrivilege	Disabled
SeLoadDriverPrivilege	Disabled
SeManageVolumePrivilege	Disabled
SeProfileSingleProcessPrivilege	Disabled
SeRemoteShutdownPrivilege	Disabled
SeRestorePrivilege	Disabled
SeSecurityPrivilege	Disabled
SeShutdownPrivilege	Disabled
SeSystemEnvironmentPrivilege	Disabled
SeSystemProfilePrivilege	Disabled
SeSystemtimePrivilege	Disabled
SeTakeOwnershipPrivilege	Disabled
SeTimeZonePrivilege	Disabled
SeUndockPrivilege	Disabled

OK Cancel

Securable Object Access Control

- A securable object is an object that can have a [security descriptor](#)
- All named Windows objects are securable
 - Some unnamed objects, such as [process](#) and thread objects, can have security descriptors too
 - You can pass [Security Attribute](#) that has a [Security Descriptor](#) when you create objects such as [file](#) & [process](#)
- Each type of object defines its own set of specific [access rights](#)
- Each type of object defines its [mapping of generic access rights](#)

Security Descriptor

- The security DB of an object
 - The SID of the owner and the primary group of the object
 - A list of access rights that tell who can do what with the object → DACL = List of ACEs
 - An information that tells the system the type of action that generates security audit event
- The SID is kept in the object
 - For example each NTFS file has a security NTFS stream
- The SID may contain inherited ACEs
 - This enables security setting of a container such as a directory to propagate to the child tree

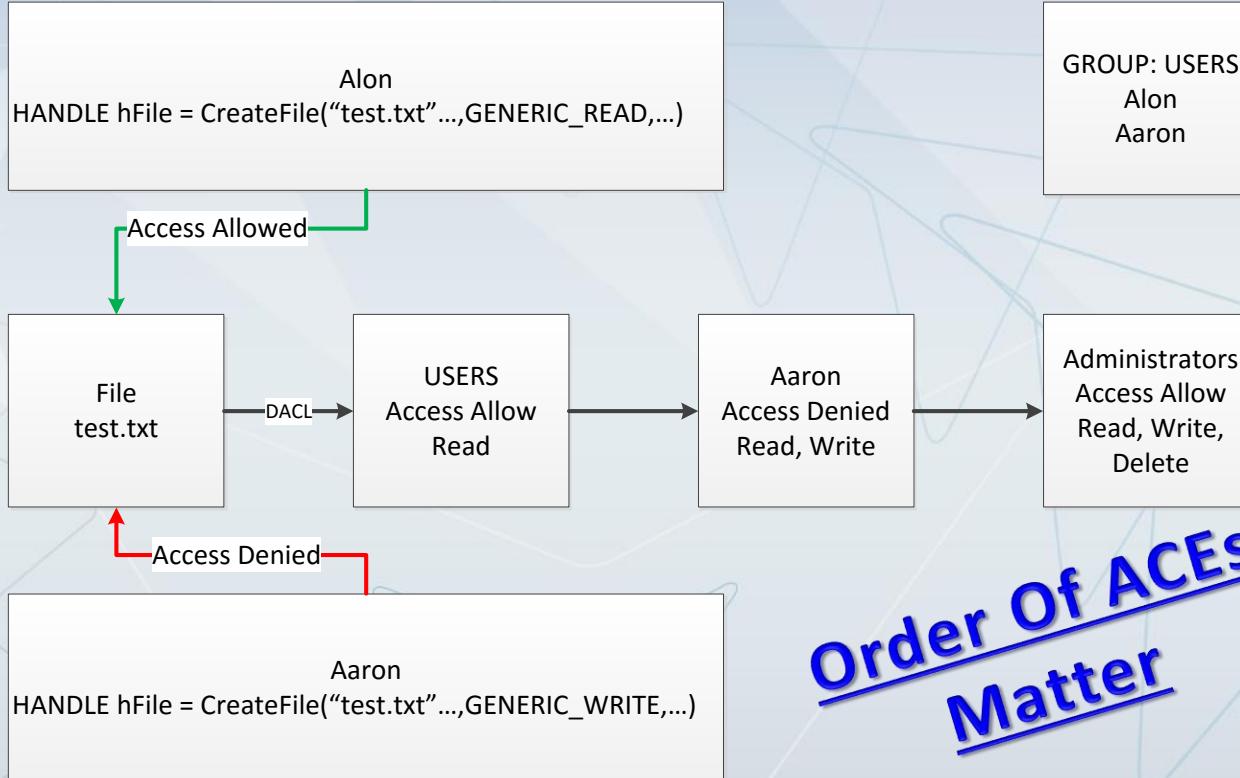
Access Mask

- Each Access Control Entry (ACE) has:
 - The target User or Group SID
 - The type of ACE (Allowed or Denied)
 - The Access Mask
 - Checking for access rights
 - a bitwise mask comparison process

AS	<u>Access SACL</u>
GA	Generic All
GE	Generic Execute
GW	Generic Write
GR	Generic Read



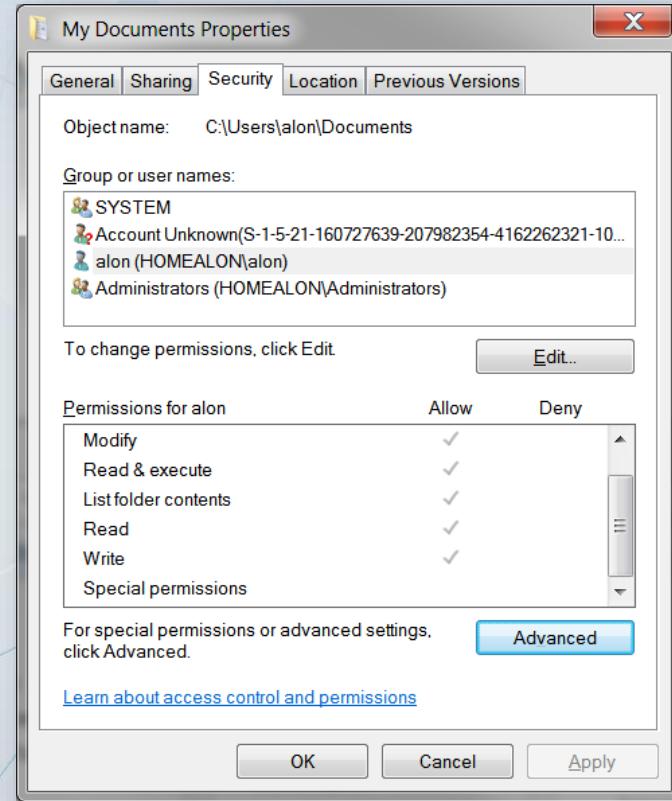
Accessing Securable Object



**Order Of ACEs
Matter**

File System Security Settings

DEMO



Object Handle Access Mask

- Most of Win32 Object handle creation APIs ask for access mask (dwDesiredAccess)
 - [OpenProcess](#), [CreateFile](#), [OpenThread](#), [RegCreateKeyEx](#), ...
- You should request only the access rights that you need
 - This prevents using the handle in an unintended way
- Most of the time use the generic access rights
 - This is simpler than specifying all the corresponding standard and specific rights
- The MAXIMUM_ALLOWED can be used to request that the object be opened with all the access rights that are valid for the caller

Object Handle Access Mask

- Each handle in the process handle table has a pointer to the object, some flags and an access mask
- The Object Handle Access Mask is set when the handle is created
 - Use DuplicateHandle to get another handle to an object with different access rights
- Since the access mask is part of the handle:
 - There is no farther DACL check when accessing the object
 - Impersonated client can access an object using any opened handle, even if this client can't create such an handle
 - If the handle has been opened prior the impersonation

Security Descriptor Definition Language

- A text based language that enables us to define a security descriptor information
 - It also provides a way to store and transfer security information (serialization)
 - [ConvertSecurityDescriptorToStringSecurityDescriptor](#)
 - [ConvertStringSecurityDescriptorToSecurityDescriptor](#)
 - [ObjectSecurity.GetSecurityDescriptorSddlForm](#)
 - [ObjectSecurity.SetSecurityDescriptorSddlForm](#)
- Example:
 - "0:AOG:DAD:(A;;RPWPCCDCLCSWRCDWOGA;;;S-1-0-0)"
- To learn more about SDDL language look [here](#)

Windows Authorization API

- Win32 has a large set of APIs, data structures and enums that let you manipulate and read securable objects
- Common use of these API are:
 - [Create object with SID](#) that contains DACL
 - Create and manipulate handle with specific Access Mask
 - Impersonate and Revert thread identity
 - Manipulate objects and container security settings

Windows Security in .NET

- Two namespaces provide Windows access control abilities in the .NET framework
 - [System.Security.Principal](#)
 - [System.Security.AccessControl](#)
- To deal with securable object you have two options:
 - The [ObjectSecurity](#) class (and its derived chain)
 - For Win32 objects that has a .NET wrapper
 - The [GenericSecurityDescriptor](#) (and its derived chain)
 - For Win32 objects that need to be interop to.
- Look at the [.NET Security Workshop Access Control](#)

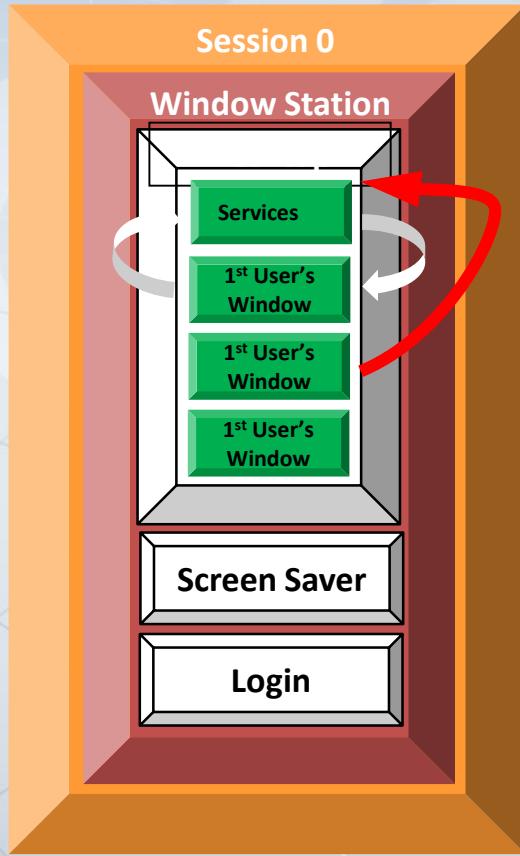
Session Zero Isolation

- Prior to Windows NT 6.x users applications and services ran in the same session
 - Sessions were introduced in Windows 2000 to support terminal service
- The main reason for sessions is isolation
 - Since isolation is a security feature, in Windows NT 6.x services run in Session zero, while users runs on session 1, 2, ...
 - Terminal services & Switch User
- **Session Zero Isolation makes an headache**
 - Can't send Windows Messages, Can't Display UI, Kernel Object are private to session, and more...

System Processes

- Session 0 is dedicated to system processes and services only
 - First interactive user session is 1
- When a user logs on
 - SMSS.EXE starts a new instance of itself to configure the new session
 - Creates the Windows subsystem process (CSRSS.EXE) and WinLogon.Exe
- Interactive Services
 - Work as expected

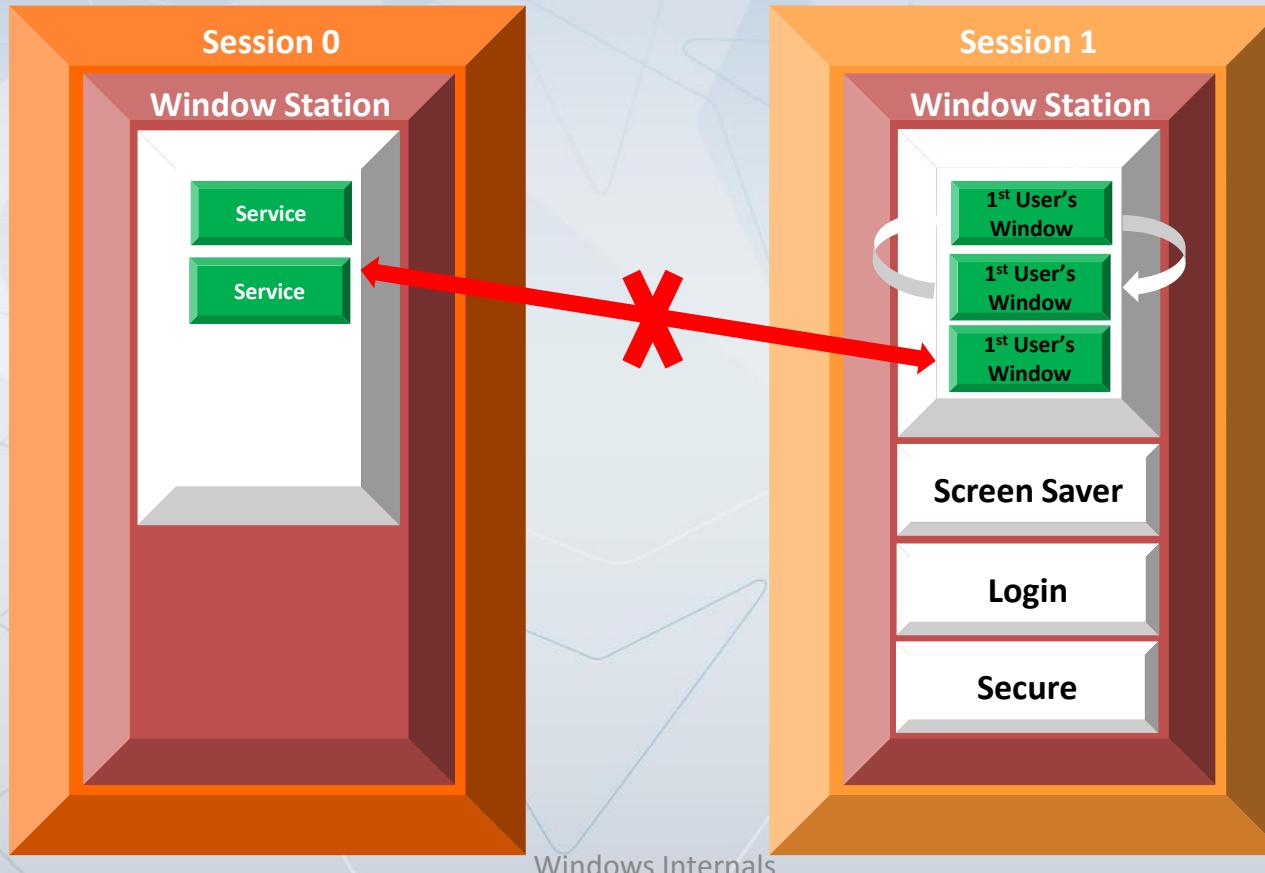
Sessions in XP/W2K/WS03



Shatter Attack

Windows Internals

Sessions in Vista/7/8



Integrity Levels

- The purpose of the Windows integrity mechanism is to restrict the access permissions of applications that are running under the same user account and that are less trustworthy
- There are four integrity levels:

Level	Typical process
System	Services
High	Elevated user apps
Medium	Default: Normal user apps
Low	IE Protected Mode

Object Integrity Labels

- Every securable object has one
- Includes Level and Policy
- Policies can include:
 - No-Write-Up - Lower IL can't write to object
 - No-Read-Up: -Lower IL can't read object
 - No-Execute-Up - Lower IL can't execute object
- Defaults:
 - Objects: Medium + No-Write-Up
 - Processes: No-Write-Up + No-Read-Up
- The integrity labels are stored as an ACE with a special SID in the SACL on the object

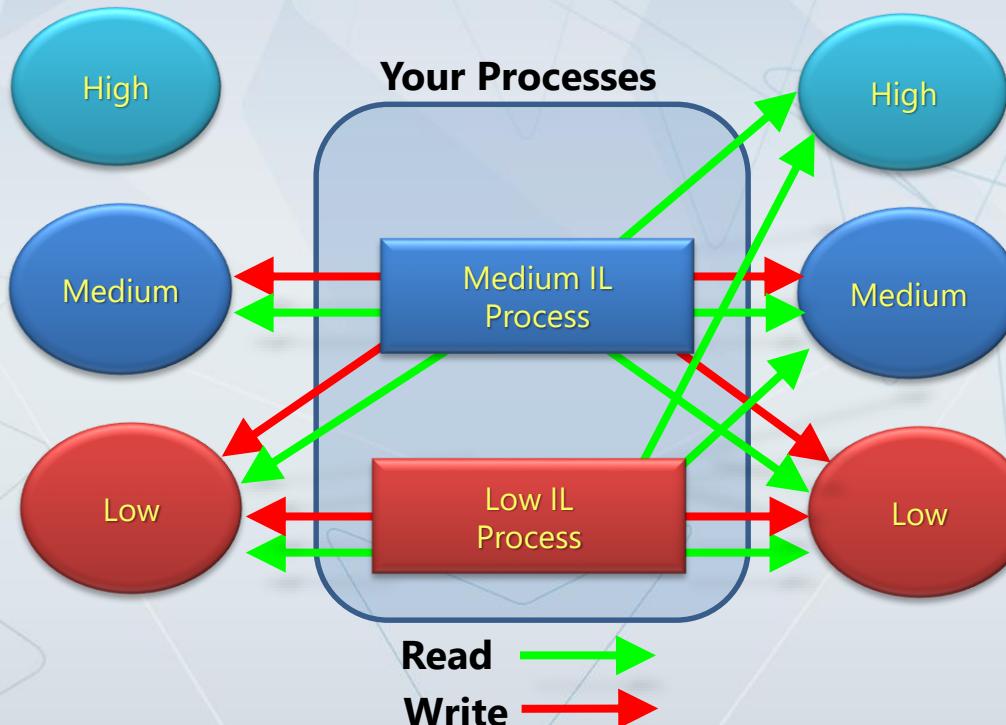
Complicated? Let's See Again

- Object can have an integrity label
 - Stored in its Security Descriptor
- Processes run at an integrity level (IL)
 - Stored in its Access Token
- Process cannot access object if its IL is lower than the object's label
 - Part of the access check

Object vs. Process Access

Other Processes

Objects



Mandatory Label Special SIDs

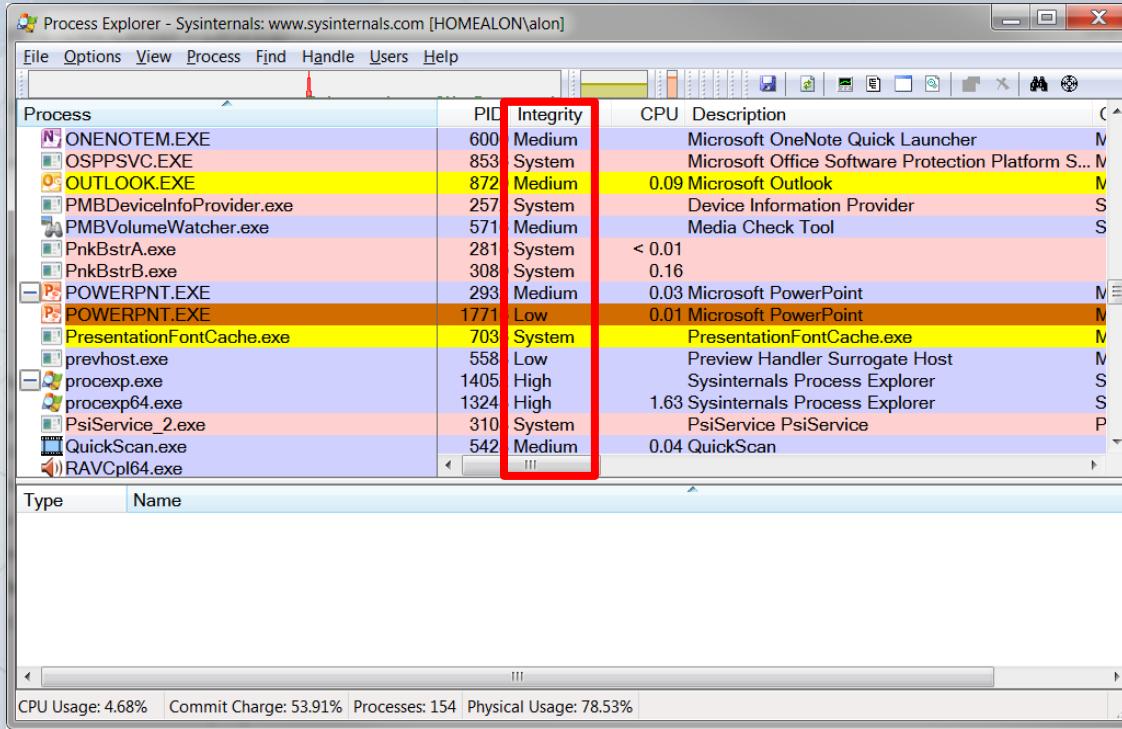
Integrity Level	Mandatory Label	Description
S-1-16-0	Untrusted process	Used for anonymous processes in some cases
S-1-16-4096	Low Level	Used for Internet Explorer
S-1-16-8192	Medium Level	Default for standard users
S-1-16-12288	High Level	Used for administrators
S-1-16-16384	System Level	Used for system processes and services
S-1-16-20480	Protected process	Used for certain protected processes, such as DRM processes

SDDL for Mandatory Labels

```
#define SDDL_MANDATORY_LABEL TEXT("ML") // Integrity label
//The SDDL strings for the mandatory label policy flags,
//which are in the access mask, are the following:
#define SDDL_NO_WRITE_UP           TEXT("NW")
#define SDDL_NO_READ_UP            TEXT("NR")
#define SDDL_NO_EXECUTE_UP          TEXT("NX")
//The SDDL SID strings for the integrity levels are the
//following:
#define SDDL_ML_LOW                TEXT("LW")
#define SDDL_ML_MEDIUM              TEXT("ME")
#define SDDL_ML_HIGH                TEXT("HI")
#define SDDL_ML_SYSTEM               TEXT("SI")
```

- Example: "S:(ML;;NW;;;LW)"

Process Explorer to the Rescue



Sharing Kernel Objects

- Sharing Kernel Object Between a service and user application is a bit tricky
 - We need to create the object in the Global namespace
 - We need to set the Integrity Level Label to Medium
 - We need to make sure that the object can be accessed by the client (DACL)
- The problem:
 - Any wrong setting will block the access
 - We will see how to do it in the following exercise

User Access Control (UAC)

- Goals
 - Running applications with standard user rights
 - Not as administrators
 - Allow applications to elevate to administrator rights when needed
 - Allow legacy applications to run with standard user rights even when assuming administrative rights

Running as a Standard User

- Windows Vista adds
 - Privileges that allow more granularity for various activities
 - E.g. changing time zone as opposed to changing time
 - More configurations options
 - e.g. wireless settings, power options, installing critical updates
 - Virtualization of file and registry to allow legacy applications to access global elements (e.g. **%ProgramFiles%** folder) without really compromising the system
- Windows 7 adds two more levels for UAC

Virtualization

- Applications should not write user-specific data in global locations in the file system or in the HKLM\Software registry hive
- When a legacy application tries to write to a global location, an access denied is returned
 - Redirected to a per user area
 - E.g. C:\Users\<username>\AppData\Local\VirtualStore
- When reading data, first checks the per user area, and if not found
 - global area
- “Legacy” means
 - 32 bit, runs with standard user rights and has no manifest file indicating it’s a Vista or later application

Administrator Approval Mode

- When a user logs in
 - If is a member of any administrators group
 - Two tokens are created
 - Otherwise
 - A standard user token is created
- When a process needs to run elevated
 - An administrator approves by clicking (AAM, Consent elevation)
 - A non-administrator needs to supply username and password of an administrator (Over The Shoulder elevation, OTS)



Running Elevated

- How to run a process with elevated credentials?
 - Right-click on the file in Windows Explorer and select “Run as Administrator”
 - Call the [ShellExecuteEx](#) API with the “runas” verb
 - Add a manifest file requesting administrative rights

```
<trustInfo xmlns="urn:schema-microsoft-com:asm.v3">
    <security>
        <requestedPrivileges>
            <requestedExecutionLevel Level="requireAdministrator"
uiAccess="false"/>
        </requestedPrivileges>
    </security>
</trustInfo>
```

The UAC Manifest

```
>mt.exe -inputresource:App.EXE;#1 -out:extracted.manifest
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<assembly xmlns="urn:schemas-microsoft-com:asm.v1" manifestVersion="1.0">
  <trustInfo xmlns="urn:schemas-microsoft-com:asm.v3">
    <security>
      <requestedPrivileges>
        <requestedExecutionLevel level="asInvoker"
          uiAccess="false"></requestedExecutionLevel>
      </requestedPrivileges>
    </security>
  </trustInfo>
</assembly>
```

asInvoker	Launch with the same token as the parent process
highestAvailable	Launch with the highest token that the user has
requireAdministrator	Highest token of the User provided User is a member of Administrators group

Executing an Elevated Process

- When executing an image requesting administrative rights
 - The Application Information Service (AIS, AppInfo.Dll) launches Consent.Exe
 - Captures the screen into a bitmap
 - Applies some fade effect
 - Switches to a desktop only accessible to the LocalSystem account
 - Displays the appropriate elevation dialog box
 - Signed by Microsoft or a Windows component – blue stripe
 - Signed by other than Microsoft – gray stripe
 - Unsigned – orange stripe
 - If the user declines the elevation, an access denied is returned to the process creator
 - Otherwise, AIS calls [CreateProcessAsUser](#) and redirects the original process to be the parent of the new process

User Interface Privilege Isolation

- Windows Vista introduced the concept of Integrity Levels
- By default, a process cannot send a window message to another process with a higher integrity level
- This mechanism is called User Interface Privilege Isolation (UIPI)
- Running as Standard User sets integrity level to Normal
- Running as Administrator sets integrity level to High
- Use the [ChangeWindowMessageFilter](#) & [ChangeWindowMessageFilterEx](#) to allow incoming messages
- Use UAC manifest: uiAccess to bypasses

Object Namespace

- protects named objects from unauthorized access
 - [CreatePrivateNamespace](#)
 - The lpAliasPrefix parameter of serves as the name of the namespace
 - The function requires an isolation boundary
 - use the [CreateBoundaryDescriptor](#) and [AddSIDToBoundaryDescriptor](#) functions
 - each namespace is uniquely identified by its name and boundaries
 - Two namespaces with the same name but different boundaries are different
- Sample: Microsoft SDKs\...\Samples\winbase\services\pvtnamespace

Other Security Mechanisms

- [Windows Firewall](#)
- [Nx \(DEP\)](#)
- [/GS](#)
- [Safe CRT functions](#)
- [ASLR](#)
- [BitLocker](#)
- [Crypto API](#)
- [Protected Process Light \(>= NT 6.3\)](#)
- [Virtual Secure Mode \(VSM\)\(NT 10.X\)](#)
- [Many more...](#)
- [Windows 10 Security Features](#)

Summary

- Security considerations must be factored into the design of a system
- All kernel based objects can be protected
- Windows uses a Security Descriptor to determine who can do what with an object

I/O SYSTEM

Agenda

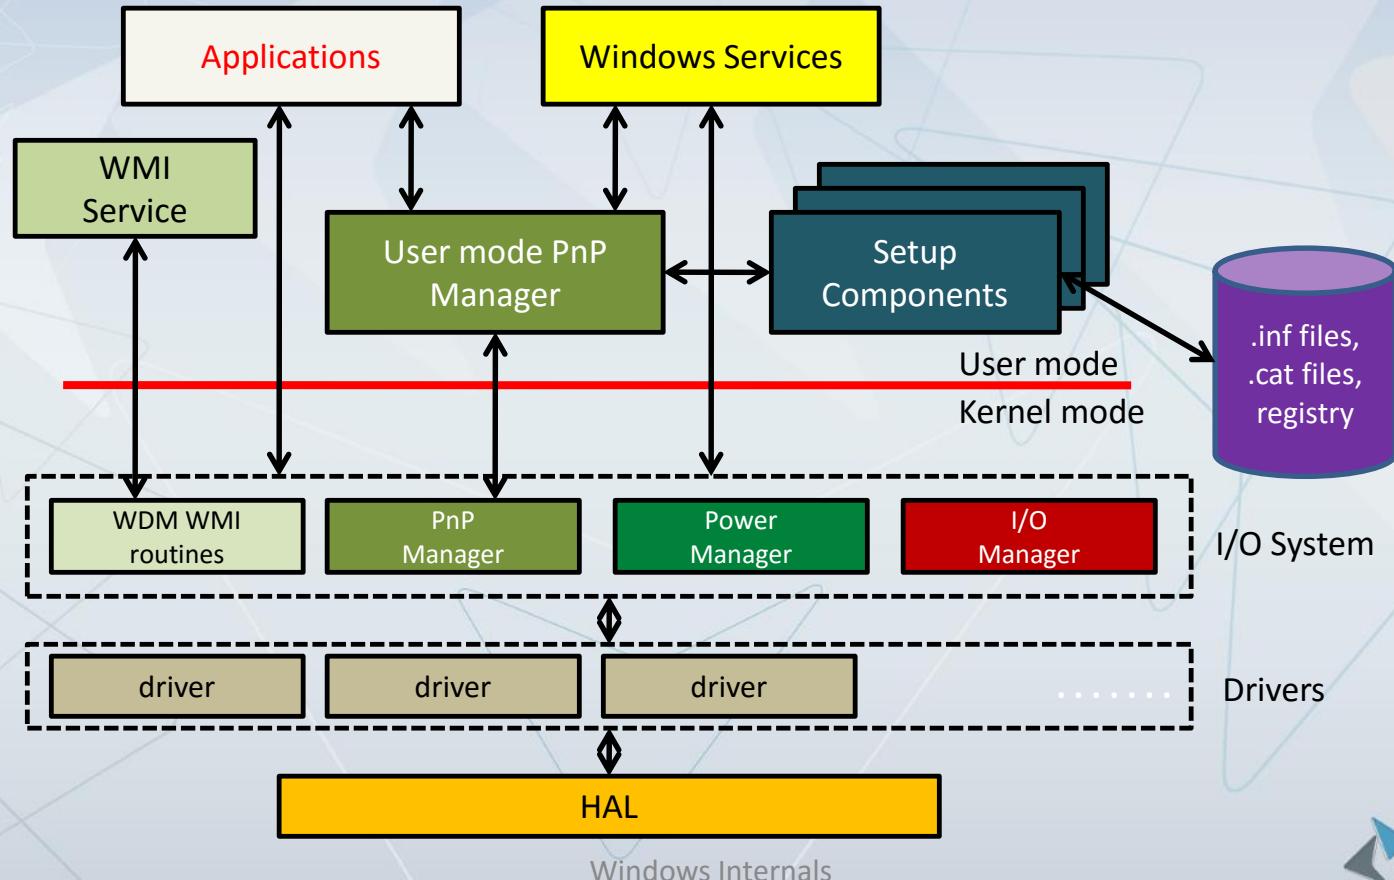
- The I/O System
- Device Drivers
- The Windows Driver Model (WDM)
- The Windows Driver Foundation (WDF)
- The Universal Driver Model
- Plug & Play
- IRP Processing
- I/O System Objects
- Power Management
- Driver Verifier

I/O Subsystem Design Features

- As many "setup" functions as possible are performed before driver is invoked
 - Reduced security risks - all security tests, buffer access methods, etc. are done by kernel-supplied code
 - Reduced risk of system crashes
- Many kernel mode functions needed by drivers are performed by system-supplied routines
 - Driver only decides which to call, and in what order
 - Drivers are easier to write, and more reliable
- I/O subsystem provides rich "feature set" for drivers
 - Hides many architectural and platform differences



I/O System Components



General Driver Types

- Function driver
 - Manages a hardware device
- Bus driver
 - Manages a bus (PCI, USB, Firewire, etc.)
- Filter driver(s)
 - Sits on top of a function driver (upper filter) or a bus driver (lower filter)
 - Allows interception of requests

More Device Drivers Types

- Virtual Device Driver
 - A user mode driver used to allow (some) access to hardware from DOS/Win16
- Kernel Device Driver
 - Video Device Driver
 - Translates GDI, DirectX and OpenGL commands for a specific video card
 - File System Driver
 - Plug & Play Device Driver
 - Handles Plug & Play requests from the Plug & Play Manager
 - WDM Device Driver
 - A Plug & Play device driver, which also handles Power Management and optionally WMI
 - Legacy Device Driver
 - Non Plug & Play Device Driver (usually NT 4 driver)
 - Software Driver

Kernel Device Driver

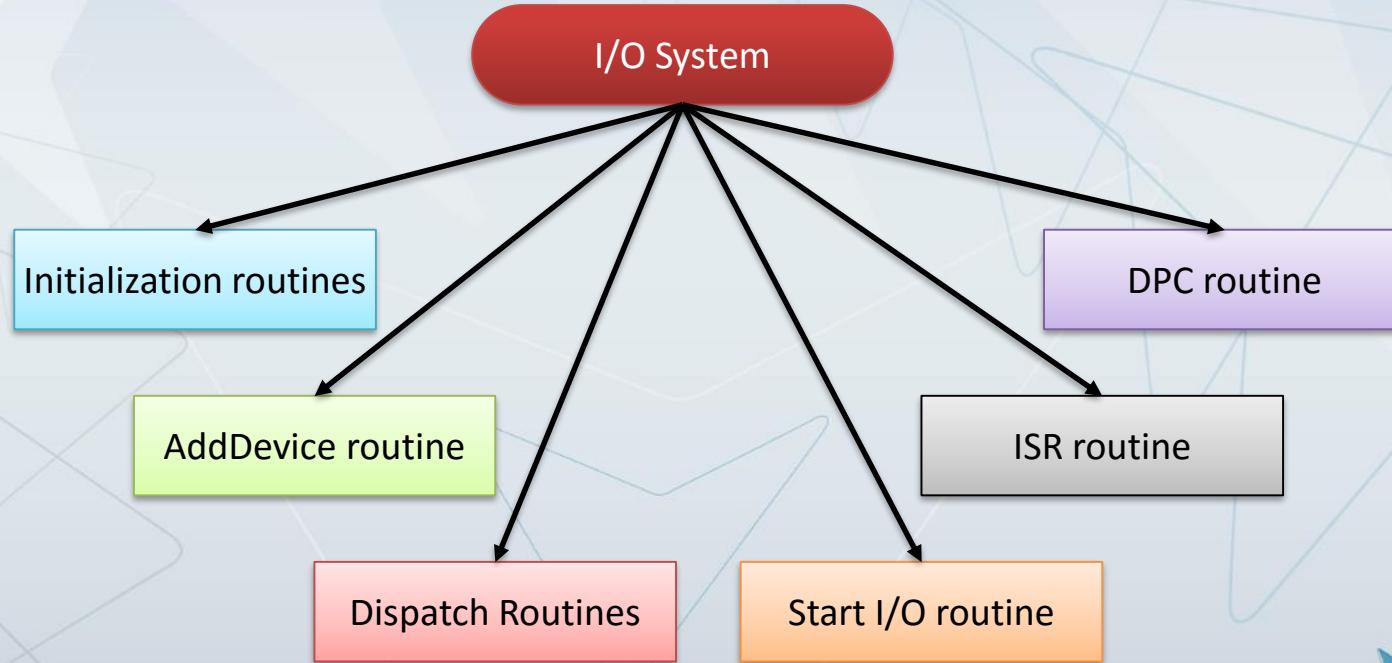
- The only one that can touch hardware, handle interrupts, etc.
 - Has a SYS extension
 - Its routines always run in kernel mode
 - Always uses the kernel mode stack
 - Limited in size: 12KB (32 bit), 24KB (64 bit)
 - No documented way to enlarge it
 - Unhandled exceptions will crash the system
 - Producing the infamous “Blue Screen of Death”

Kernel Device Driver

- Usually invoked by a user mode code ([ReadFile](#), [WriteFile](#), [DeviceIoControl](#))
- Fully interruptible, but not always pre-emptible
- System handles all device independent aspects of I/O
 - No need for assembly
- Layered model

Anatomy of a Driver

- A driver exports functionality, callable by the I/O system

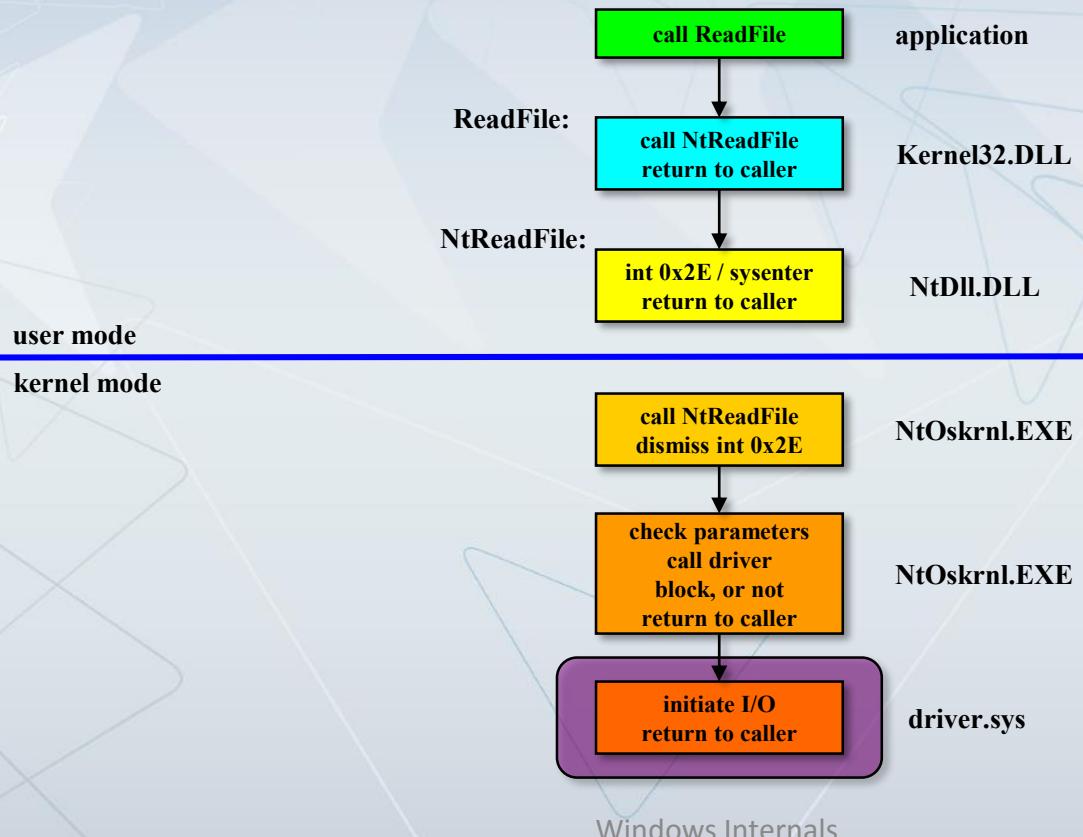


Driver and Device Objects

- Drivers are represented in memory using a DRIVER_OBJECT structure
 - Created by the I/O system
 - Provided to the driver in the [DriverEntry](#) function
 - Holds all exported functions
- Device objects are created by the driver on a per-device basis
 - Represented by the [DEVICE_OBJECT](#) structure
 - Typically created in the Driver's [AddDevice](#) routine
 - Several can be associated with a single driver object
- I/O system is device-centric, not driver-centric



Invoking a Driver



Accessing Devices

- A client that wants to communicate with a device, must open a handle to the device
 - [CreateFile](#) or [CreateFile2](#) from user mode
 - The [System.IO.FileStream](#) class in .NET
 - [ZwCreateFile](#) from kernel mode
- CreateFile accepts a "filename" which is actually a device symbolic link
 - "file" being just one specific case
 - For devices, the name should have the format \\.\name
 - Cannot access non-local device
 - Must use double backslashes "\\\\".\\name" in C/C++ or in C++ 11:
 - R"(\\".\\name)"

Asynchronous I/O

- The I/O manager supports an asynchronous model
 - Client initiates request, may not block, and get a notification later
- Device drivers must be written with asynchrony in mind
 - Should start an operation, mark the IRP as pending and return immediately
- The I/O manager supports several ways of receiving a notification when the operation completes
- To use I/O asynchronously, [CreateFile](#) must be called with the FILE_FLAG_OVERLAPPED flag
- Other I/O functions must provide a non-null OVERLAPPED structure pointer

What is WDM?

- A model for writing device drivers
 - Source compatibility between Windows 98/ME and Windows 2000/XP
 - Standard handling for Plug & Play (P&P), Power Management and Windows Management Instrumentation (WMI)
 - Supports a wide range of buses (PCI, USB, IEEE1394 and more)
 - Extensible to support future buses
 - Supports a wide range of device classes (HID, Scanners, Modems, etc.)
- Not included in WDM
 - File system drivers
 - Video drivers

What is WDF?

- A framework and a set of tools for the development of Windows device drivers
- Framework(s)
 - Kernel Mode Driver Framework (KMDF) – developing drivers running in kernel mode
 - User Mode Driver Framework (UMDF) – developing drivers running in user mode
- KMDF is a replacement for the Windows Driver Model (WDM)
- UMDF allows creation of certain drivers running in user mode

Why WDF?

- What's wrong with WDM?
 - Well, nothing... but...
- WDM
 - Originally created to get source-level compatibility between Windows 2K/XP and Windows 98/ME
 - Very flexible but complex
 - Poor DDI design
 - Very little default behavior support
 - Plug & Play and power management requires a lot of hard-to-write boilerplate code
- Most driver types must run in kernel mode
 - Potentially risking a blue screen
 - Harder to write and debug

Windows Driver Foundation

- User mode and kernel mode drivers
 - Same object model and concepts
- Consistent object model
 - Properties, methods and events
- Much boilerplate code and default handling already part of the library
 - Drivers only need to register for interesting events
- Object hierarchy simplifying object lifetime management
- Versionable with side by side support
- Simple but not simpler, and extensible

Universal Windows Drivers (NT 10.X)

- Enable developers to create a single driver that runs across multiple different device types
 - Desktop, Mobile, IoT Core, Windows Server 2016, Xbox One, HoloLens
 - The driver model equivalent to the UWP user mode API
 - A Universal Windows driver calls only device driver interfaces (DDIs) that are part of UWP
 - These DDIs are marked as Universal on the corresponding MSDN reference pages
 - Use the WDM or WDF driver model
 - New tools to package and install universal driver on the various platforms

What is Plug & Play?

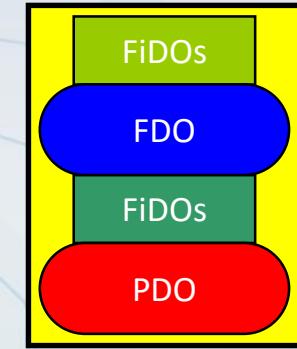
- Automatic and dynamic recognition of installed hardware
 - Hardware detected at initial system installation
 - Recognition of PnP hardware changes between boots
 - Run-time response to PnP hardware changes
- Dynamic loading and unloading of drivers in response to hardware insertion or removal
- Hardware resource allocation and reallocation
 - PnP manager may reconfigure resources at run-time in response to new hardware requesting resources that are already in use

Device Enumeration

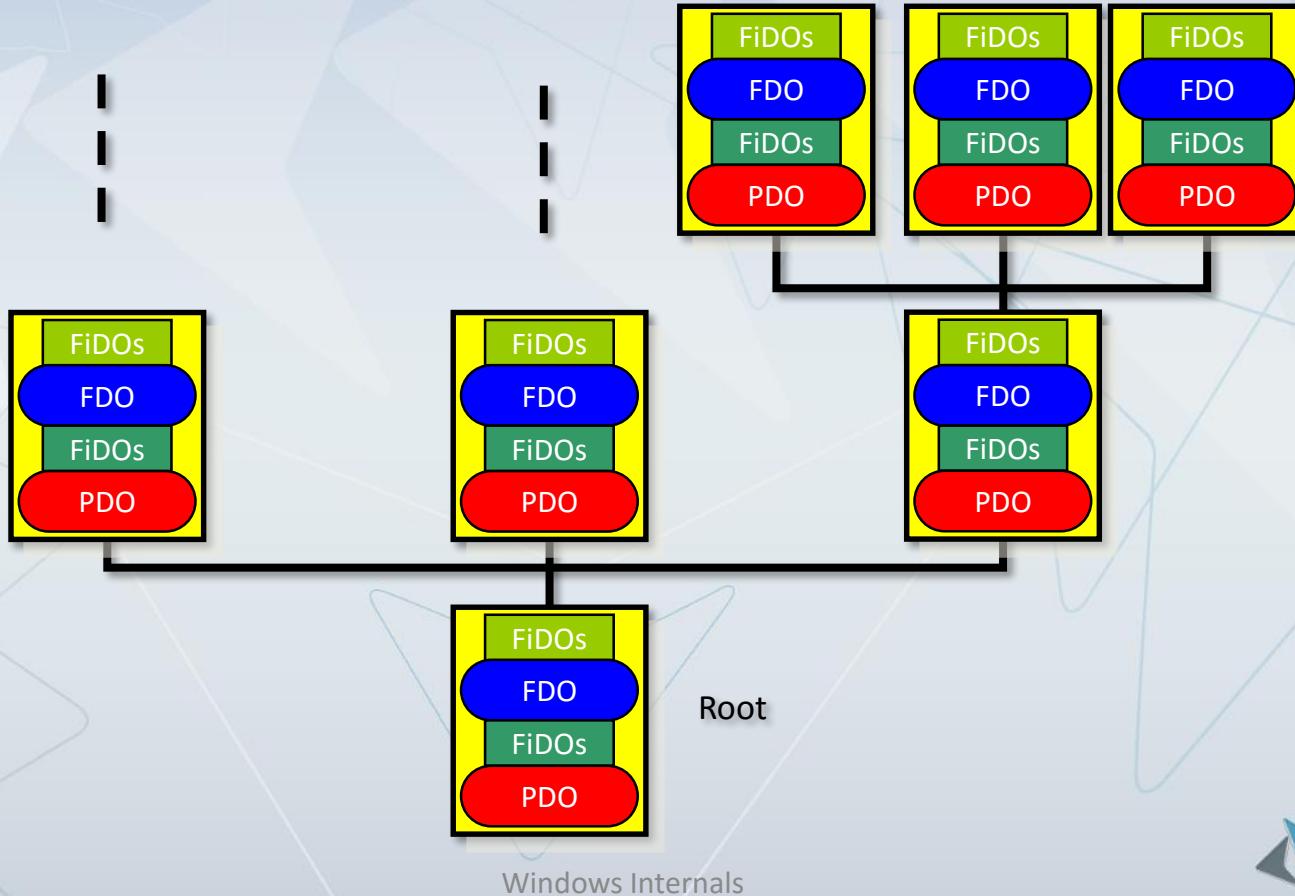
- Upon boot, the P&P Manager performs enumeration of buses and devices
 - Starts from an imaginary “Root” device
 - Scans the system recursively to walk the device tree
 - Creates a PDO (Physical Device Object) for each physical device
 - Loads lower filter drivers (if exist)
 - They create their FIDs (Filter Device Object)
 - Loads “the driver”
 - It should create the FDO (Functional Device Object)
 - Loads upper filter drivers (if exist)
 - They create their FIDs

Device Node (“DevNode”)

- Represents a stack of devices
 - PDO: Physical Device Object
 - Created by the bus driver
 - FiDO: Filter Device Object
 - Optional lower/upper device objects
 - FDO: Functional Device Object
 - The “actual” WDM driver created device object



Device Enumeration Tree



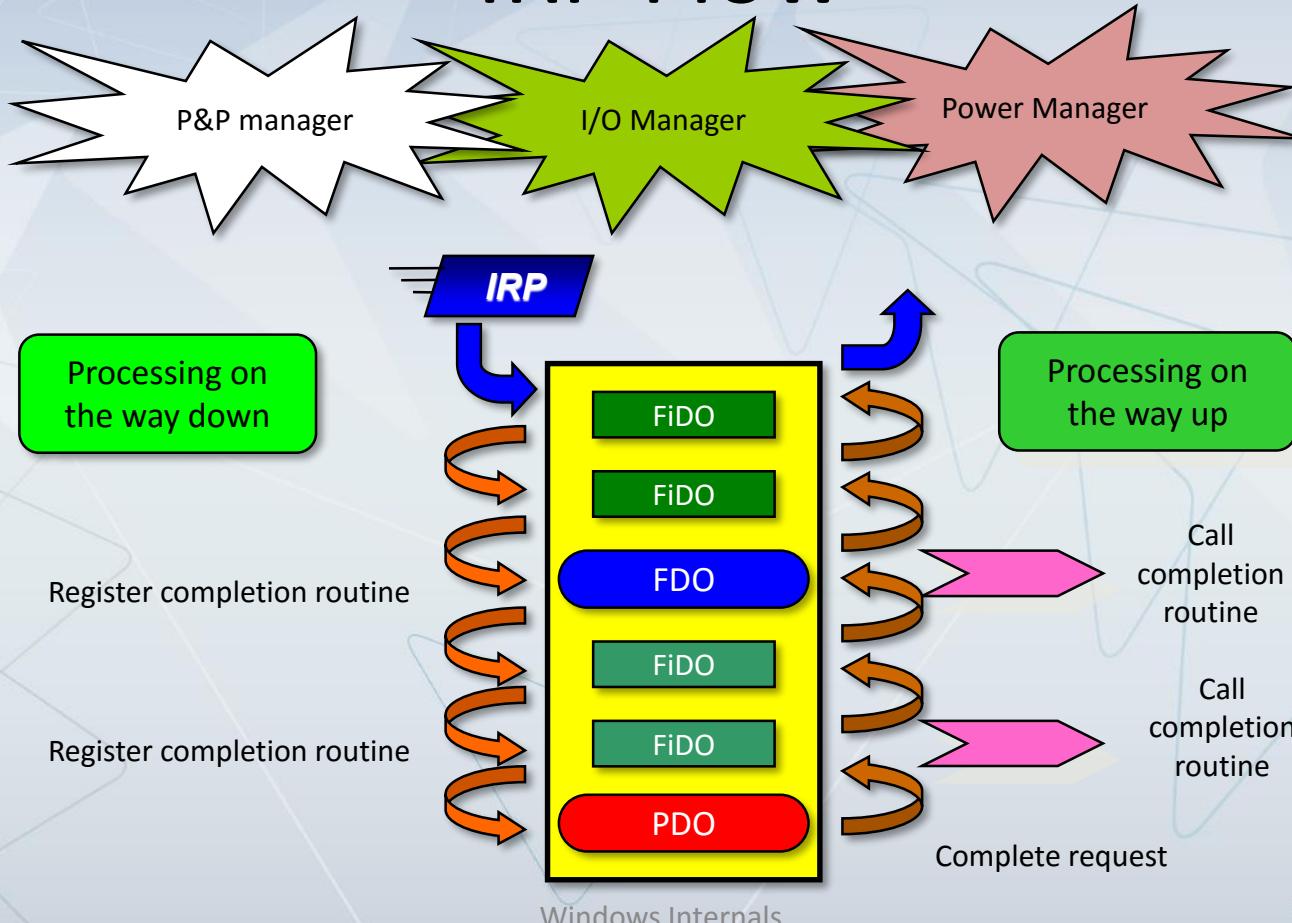
Important Registry Keys

- “Hardware” (Instance) keys
 - HKLM\System\CurrentControlSet\Enum
 - Information about a single device
- “Class” keys
 - HKLM\System\ CurrentControlSet\Control\Class
 - Information about all devices of same type
- “Software” (Service) key
 - HKLM\System\ CurrentControlSet \Services\drivername
 - Information about a specific driver

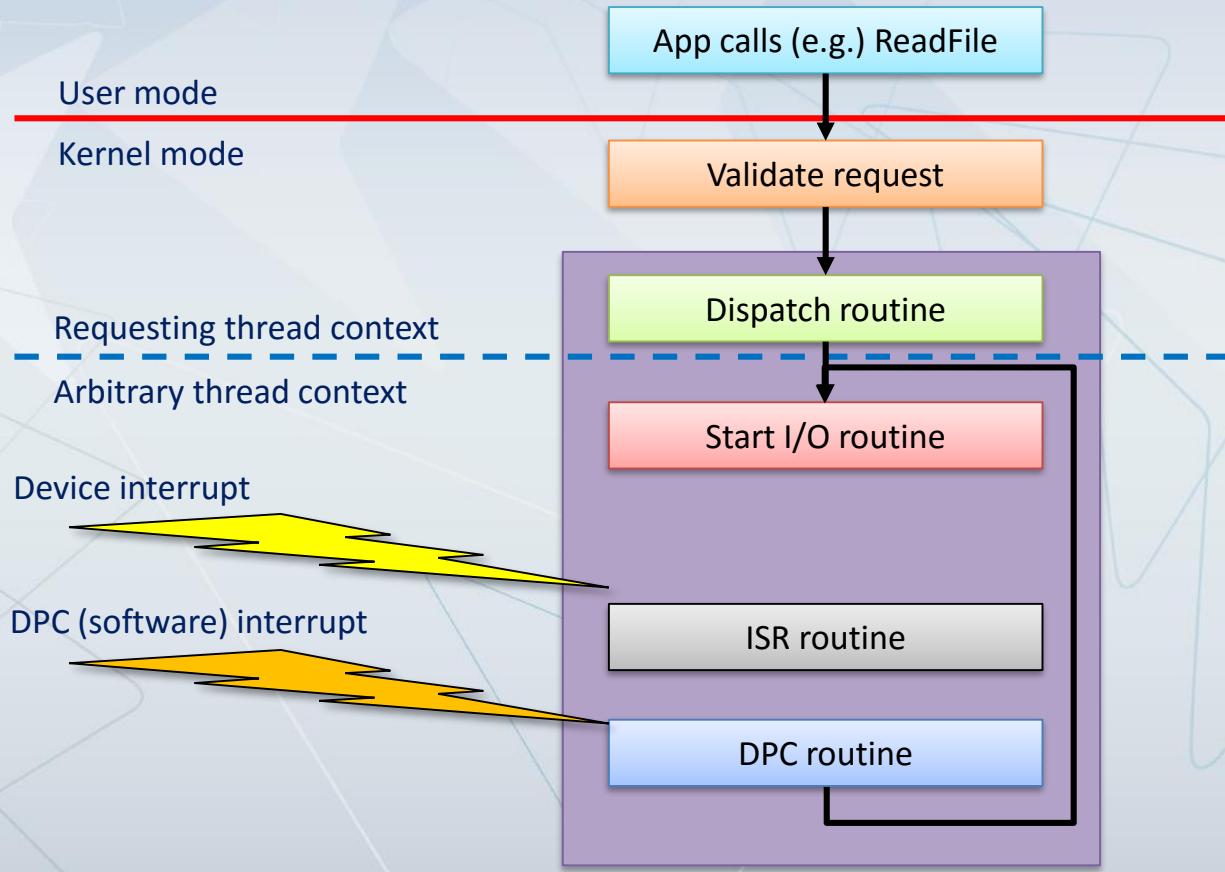
What is an IRP?

- A structure representing some request
 - The IRP structure is defined in `<wdm.h>`
 - Contains all details needed to handle the request (codes, buffers, sizes, etc.)
 - Always allocated from non-paged pool
 - Accompanied by a set of structures of type `IO_STACK_LOCATION`
 - Number of structures is the number of the devices in this DevNode
 - Complements the data in the IRP
 - Broadly speaking, the data needed for handling the request is “split” between the actual IRP object and the “current” I/O stack location

IRP Flow



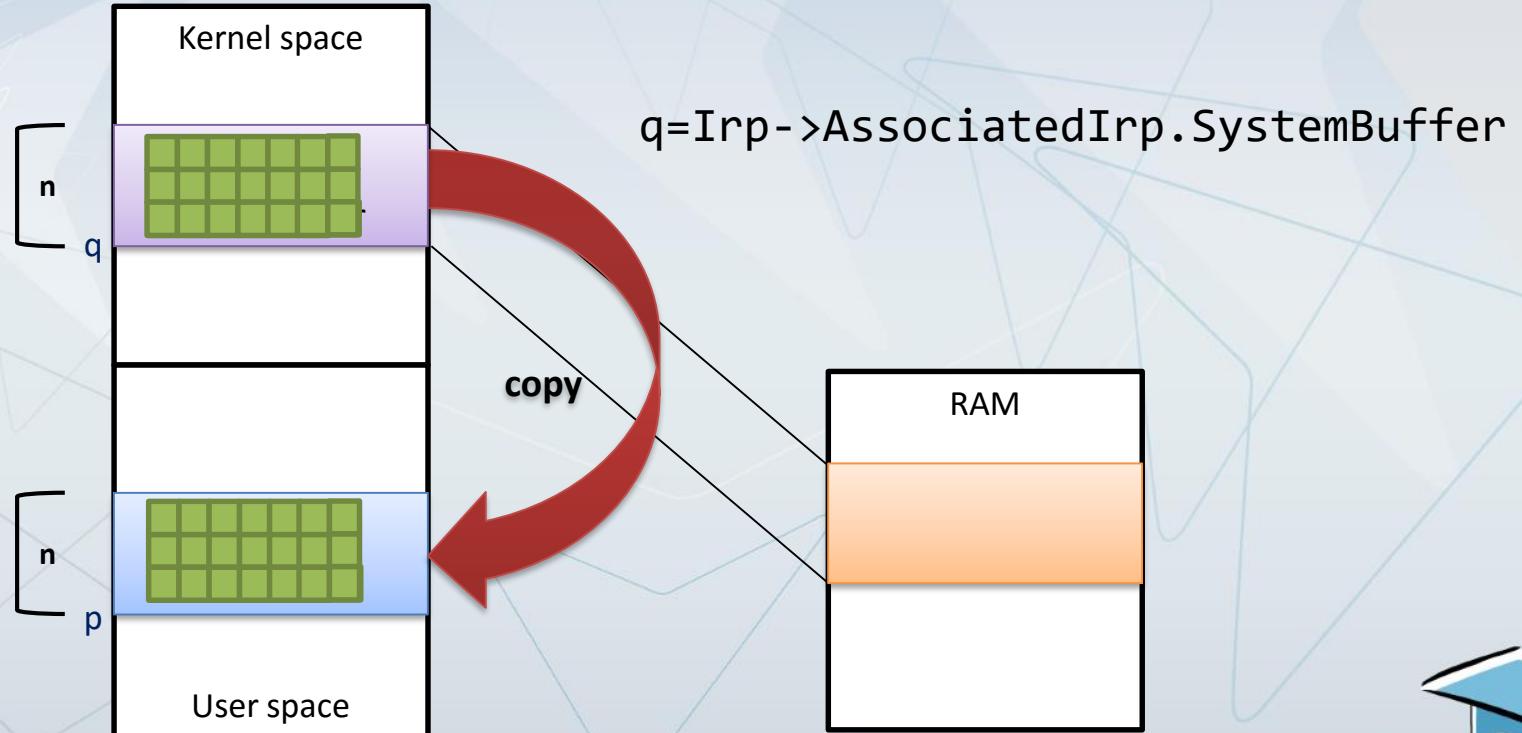
Typical IRP Processing



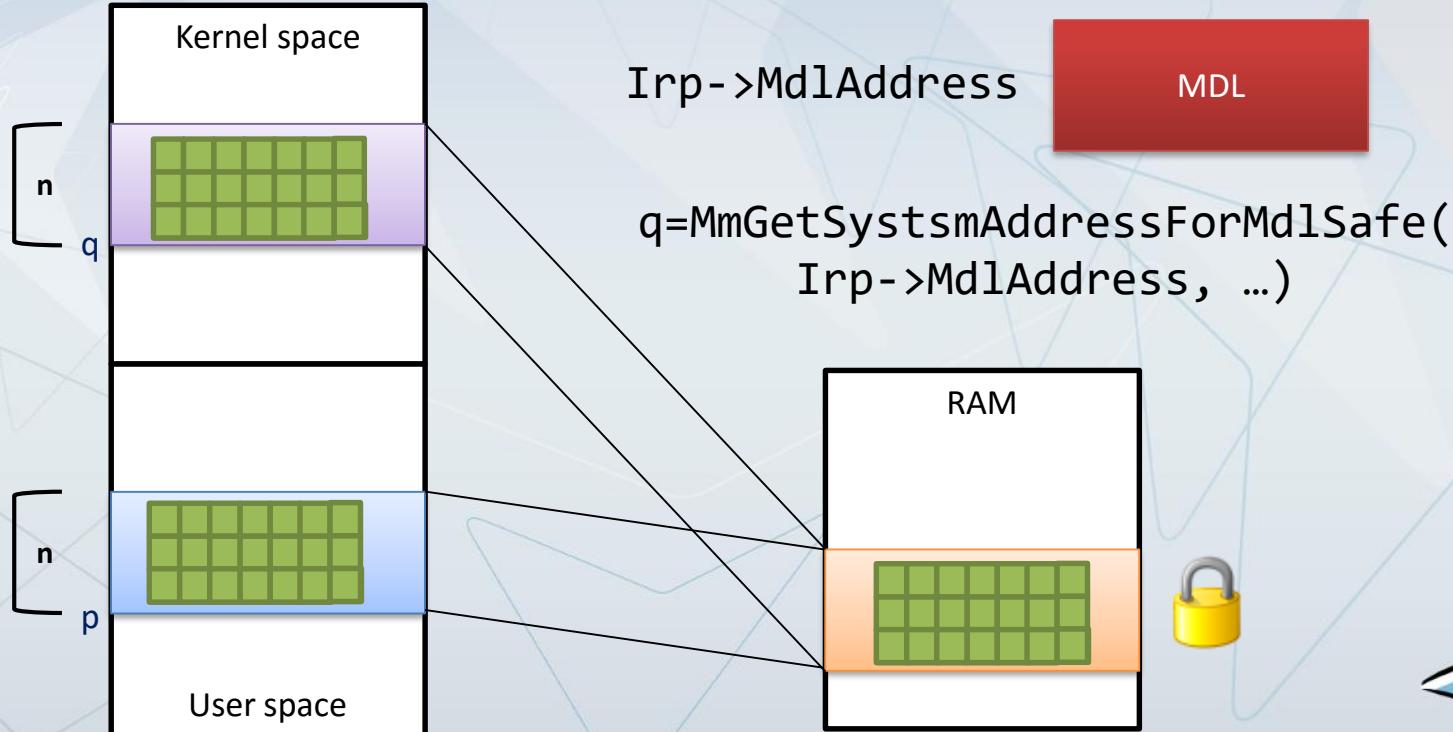
Referencing User Buffers

- Buffers provided in user space are not generally accessible from an arbitrary thread context and/or high IRQL ($>=2$)
- The I/O system provides ways to mitigate that
- Buffered I/O
 - Transfer is to and from an intermediate buffer in system address space
 - I/O Manager does all of the setup work
- Direct I/O
 - Transfer is to or from user's physical pages
 - I/O Manager does most of the setup work
- Neither I/O
 - No help from the I/O manager

Buffered I/O



Direct I/O



I/O System Objects

- Driver object ([DRIVER OBJECT](#))
 - Represents the driver for the I/O system
 - Includes dispatch routine pointers, [AddDevice](#) routine pointer, Unload routine pointer, etc.
 - Created by the kernel, passed to the driver and filled by it
- Device object ([DEVICE OBJECT](#))
 - Defines a specific device (usually hardware)
 - Associated with File objects
 - Allows for driver-defined extensions
 - Provides a DPC object for after Interrupt processing
 - Created by the driver using [IoCreateDevice](#)
 - Several may be created

I/O System Objects

- Interrupt object
 - Completely opaque
 - Represents an interrupt source
 - Associates an interrupt vector with an Interrupt Service Routine (ISR)
 - Created by calling [IoConnectInterrupt](#)
- Adapter object
 - Defines characteristics of a DMA controller
 - Either system DMA or bus master controller
 - Supports arbitration for access to DMA related resources
 - Created with [IoGetDmaAdapter](#)

Asynchronous Procedure Call (APC)

- An object representing a callback function to be called by a specific thread
- Example: after an interrupt may need to execute code in a particular (non-arbitrary) thread
 - e.g. reporting I/O completion
- APC forces a specified thread to run a specified routine in that thread
 - Makes the thread ready to run if it was waiting
 - At the end of the APC routine, the thread might go back to waiting (if it was in that state earlier)
- APC takes precedence over "normal" thread code
 - Only within that thread (does not affect its priority)
- Kernel mode API is exported but not documented

APC Types

- Special kernel APCs
 - Run in kernel mode, at IRQL APC_LEVEL (1)
 - Always deliverable
 - Used for I/O completion report from "arbitrary thread context"
- User mode APCs
 - Used for I/O completion callback routines (see [ReadFileEx](#), [WriteFileEx](#))
 - Only deliverable when thread goes into "alertable state" (see docs for e.g. [SleepEx](#), [WaitForSingleObjectEx](#), [MsgWaitForMultipleObjectsEx](#))
 - Can be explicitly queued with [QueueUserApc](#) (Win32)

I/O Prioritization

- To support I/O Prioritization in the Kernel:
 - The I/O manager has IRP queue for each priority
 - There are mechanisms to prevent I/O starvations
 - In case of a detection of a starvation, the `PsBoostThreadIo` and `IoBoostThreadIoPriority` are called
 - IRP has a priority field
 - This is a hint, not all device driver use it
 - However Idle class I/O is managed by the system storage class driver, so it is automatically applied to all storage devices

Power Management

- Why manage power?
 - Increase battery life in portable systems
 - Biggest savings in servers
- Energy Star requirements
 - When idle, computer can consume only 15 watts
- Why implement Energy Star?
 - Requirement for selling to US government
 - Many Europeans and Asian governments have similar requirements

Power Management Until Win2K

- Mainly managed by hardware
- Advanced Power Management (APM) BIOS
 - On x86 machines runs “hidden”
 - Unable to hook
 - Can interrupt the CPU at any time
- Problems
 - Scheme is x86-specific
 - Optimized for DOS – not NT
 - Does not allow devices to wake the system

WDM Power States

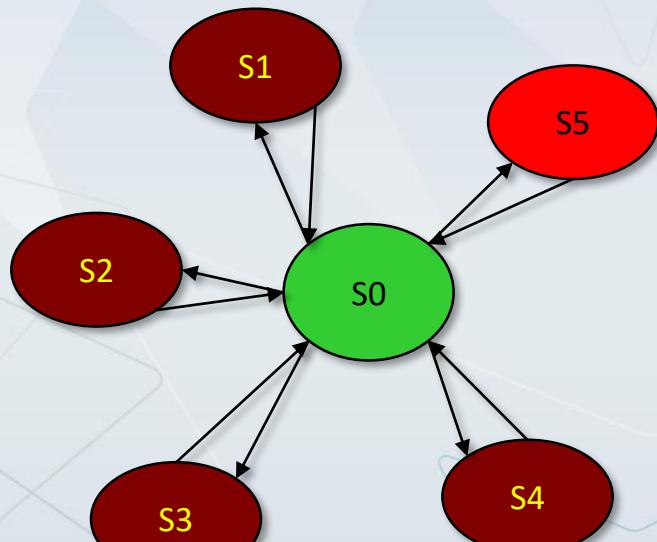
- System power states
 - Working (S0)
 - System is generally operational and running at full speed
 - Sleeping (S1-S3)
 - System appears off and is not performing any computational tasks. CPU is off. Memory remains powered up
 - Hibernate (S4)
 - Physical memory is written to disk. CPU is off
 - Off (S5)
 - No power to the system
- Waking
 - Transition from sleep state to Working state

WDM Power States

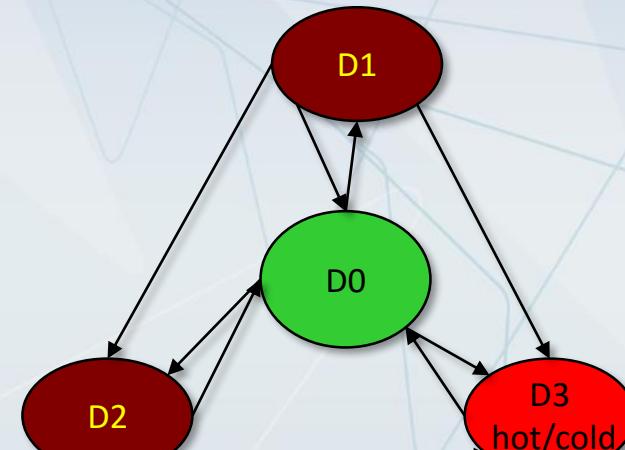
- D0 - Mandatory, Highest level of power consumption and performance
- D1
 - Less power consumption than D0 but more than D2
 - Restore latency shorter than D2
- D2
 - Less power consumption than D1
 - Restore latency longer than D1
- D3 (Hot & Cold) - Mandatory
 - No power to device
 - Longest restore time
 - The D3hot sub-state requires the device to remain accessible on its parent bus so that it can respond to bus-specific software commands
 - This requirement, and the power used to meet it, are removed in D3cold

Power Transitions

System Power States



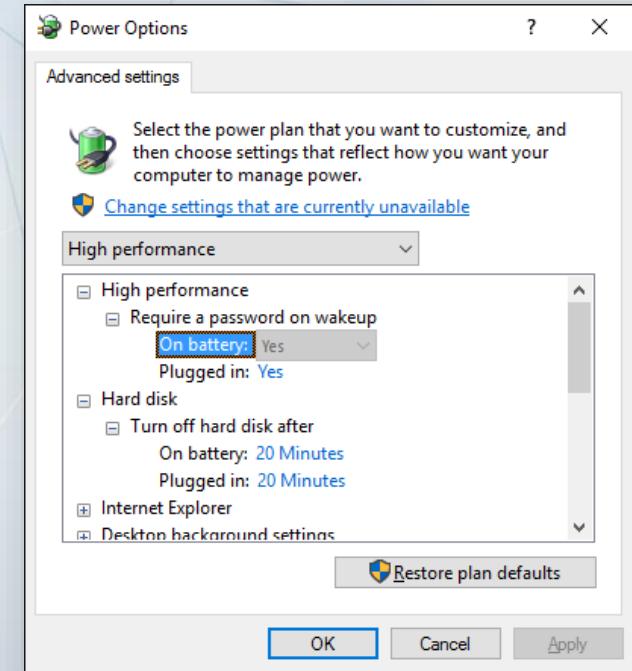
Device Power States



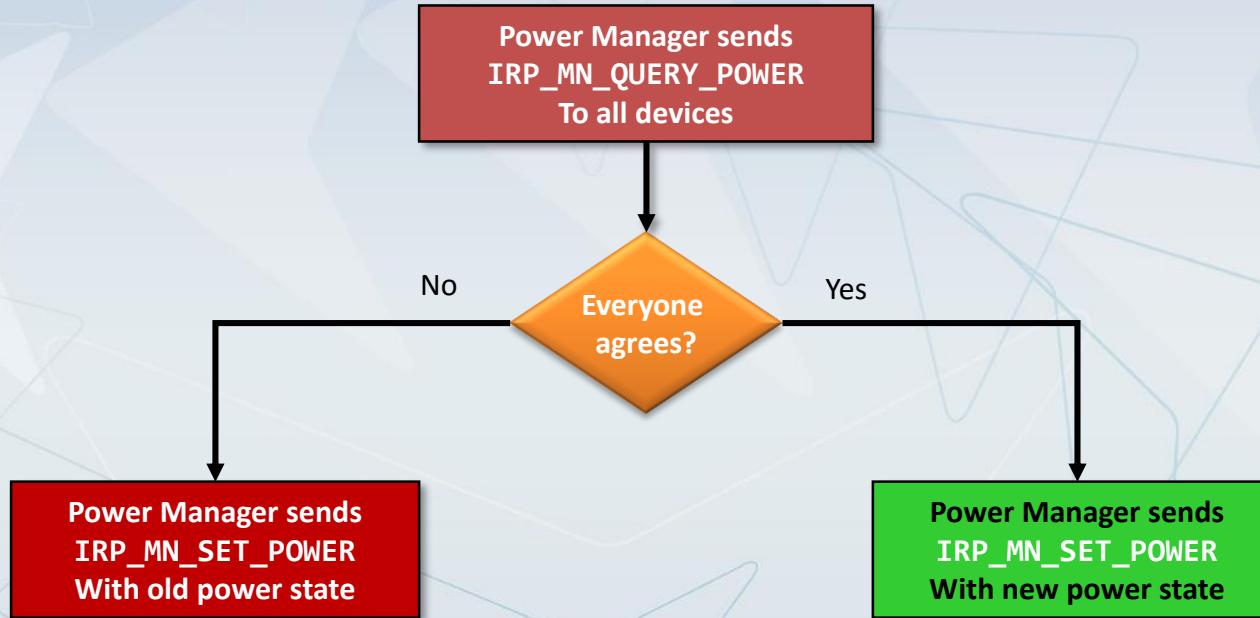
Windows Internals

Power Control Panel

- Specifies policy to the Power Manager
- ACPI machines can program the power button and sleep button (lid switch on laptops)
 - Shutdown, standby, hibernate, off
 - None (Sleep button only)
- Power Manager battery monitor alarm
 - Place machine in standby, hibernate or off state
 - Notification messages
 - Run program



How The System Powers Down



Installing Drivers

- Use an INF file
 - Required for WDM drivers
 - Format is new, but loosely based on Windows 95 INF file format
- When a new type of hardware is detected by the PnP Manager, the system device installer searches the `%SystemRoot%\INF` for a suitable INF file
 - If not found, a dialog box will pop up requesting the user to provide one
- INF file is not “run” from beginning to end

Driver Verifier

- A tool that allows monitoring device drivers activities and operations
- Different UI in Windows 2000 and XP
- Can be operated from the command line as well
 - Can even change some settings without a restart
- Does not require any special code or awareness from the driver writer
- Can monitor any driver

Driver Verifier Options

- Automatic Checks
 - Improper IRQL usage, improper Spin Lock acquisition or release, driver resources
- Special memory pool
 - The driver's allocations will be made from a special pool and monitored for overruns, underruns and illegal usage
- Forcing IRQL checking
 - Forces paging of all paged driver code/data, forcing checks of correct behavior in IRQL and Spin Lock usage
- Low resource simulation
 - Causes random failure in memory allocations
 - Checks driver's ability to cope

Driver Verifier Options

- Memory Pool Tracking
 - Checks whether the driver has freed all its allocated memory when unloaded
- I/O Verification
 - IRPs are allocated from a special pool, monitors driver's I/O handling
- DMA Verification
 - Monitors DMA buffer and map registers usage
- Deadlock Detection
 - Monitors usage of spin locks, mutexes and fast mutexes to detect possible deadlocks
- SCSI Verification
 - Monitors SCSI routine usage, excessive delays, etc.

Summary

- The I/O system provides a consistent model of flow for use by applications
- Device drivers are used to implement the device specific functionality
 - Integrate with the rest of the system
- The plug & play manager controls automatic recognition and loading of drivers

Module 8

NETWORKING

562

Windows Internals



Agenda

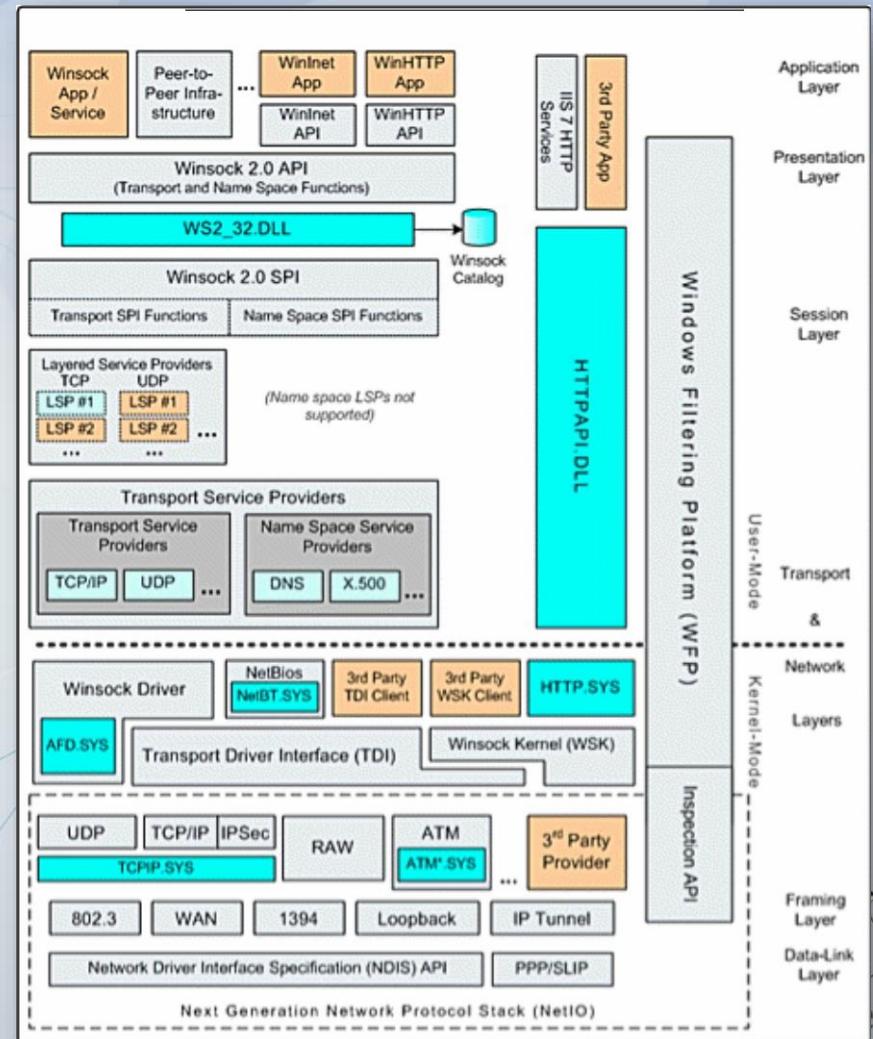
- Introduction to Windows Networking
- Networking Components
- Networking APIs
- Summary

Windows Networking

- Network support is integrated with the I/O system and the Windows API
- The network stack consists of services, APIs, protocol and network device drivers
- Broadly implements the classic OSI (Object System Interconnection) networking model

Networking Components

- Networking APIs
 - Protocol-independent way for apps to communicate across a network
 - Socket (TCP/UDP), Named pipe, mailslot, File sharing, RDP, etc.
- Transports Driver Interface
 - NDIS protocol drivers
 - Process IRPs coming from TDI clients
 - Can add protocol specific headers (e.g. TCP, UDP)



Networking Components

- NDIS library (`ndis.sys`)
 - Kernel mode DLL that encapsulates network adapter drivers
- NDIS miniport drivers
 - Kernel mode driver that interfaces with a specific hardware network card
 - They are wrapped by the NDIS library so they are easier to write
- Winsock Kernel (WSK) – transport independent kernel-mode networking API
 - Replaces the old TDI

Networking Components

- The Windows Filtering Platform
 - A set of API and system services that provide a platform for creating network filtering applications
 - Replaces the TDI Filter interface
 - Windows Firewall with Advanced Security (WFAS) is implemented using WFP
 - Provides a packet filtering infrastructure for both IP 4/6
 - Allows for data filtering, modification, and re-injection
 - Many other features



Networking APIs

- Windows Sockets (Winsock) and Winsock Kernel (WSK)
- Remote Procedure Calls (RPC)
- Web Access APIs
 - WinInet – FTP, HTTP
 - WinHttp - HTTP
- Named pipes and Mailslots
- Distributed Component Object Model (DCOM)
- Message queuing (MSMQ)
- Peer-to-Peer

Network Redirectors

- Windows supports multiple redirectors for accessing remote resources such as shared file
- Multiple Provider Router (MPR)
 - A DLL (`\Windows\System32\Mpr.dll`) which determines which network to access when using the Windows WNet API for browsing remote file systems
- Multiple UNC Provider (MUP)
 - A driver (`\Windows\System32\Drivers\Mup.sys`) which determines which network to access when using the Windows I/O API

Protocol Drivers

- Kernel mode drivers responsible for translating requests into the underlying protocol
- Separating APIs from protocols allows using more than one protocol to implement an API
- In Windows NT 6.X Microsoft rewrite the whole network stack around TCP/IP (IP V4 and IP V6)
- \Windows\System32\Drivers\Tcpip.sys
 - TCP/UDP/IP/ARP/ICMP and IGMP

NDIS Drivers

- NDIS (Network Driver Interface Specification) was developed to allow protocol drivers to communicate with network adapter drivers in a device independent way
- Network drivers that comply with the specification are NDIS drivers (or miniport drivers)
- Windows supplies Ndis.sys, a kernel mode helper library for communicating with and implementing an NDIS miniport driver

Summary

- Networking support is integrated with the rest of the system
- Various networking APIs exist to facilitate different programming models
- Windows provides extensibility and flexibility to further grow the networking system

APPENDIX

573

Windows Internals



Appendix A

INTRODUCTION TO UNIVERSAL WINDOWS PLATFORM

The Windows Runtime (>= NT 6.2)

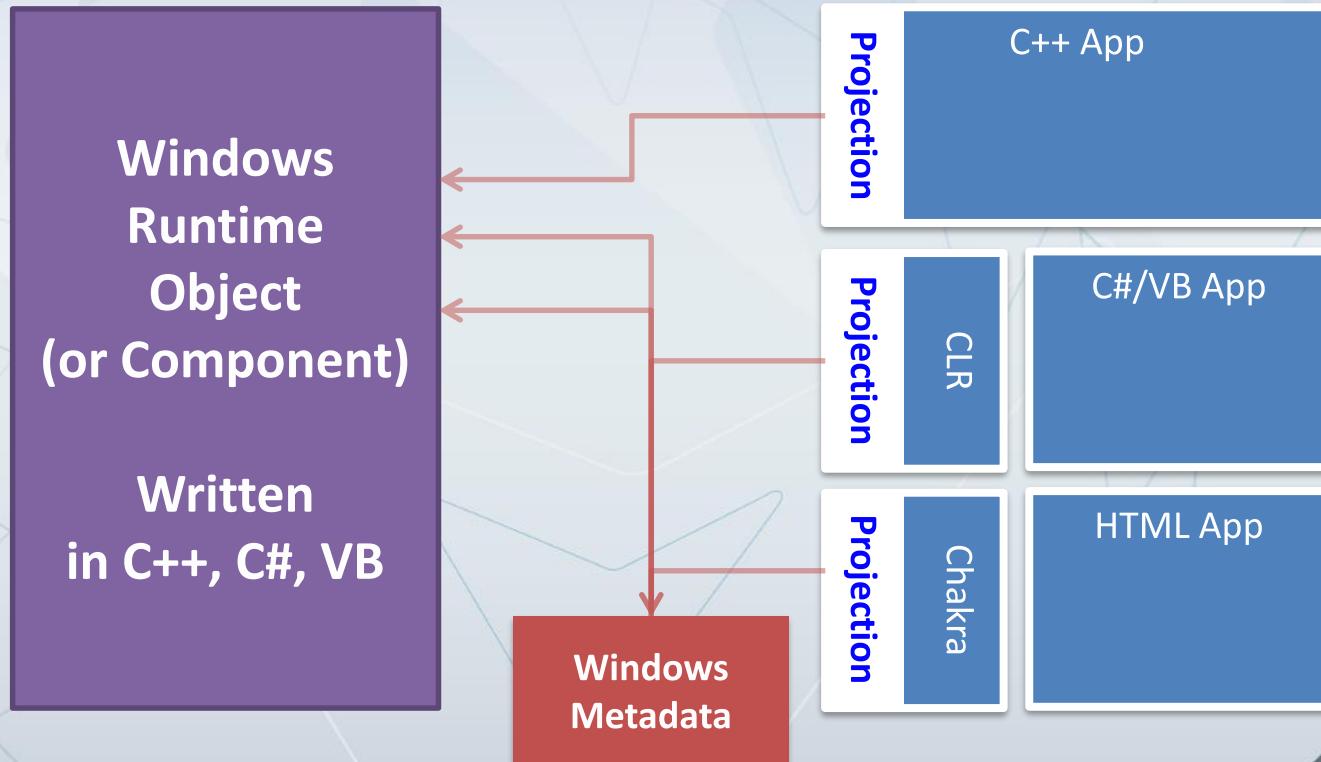
- The new runtime for Windows Store applications
- Based on an enhanced version of COM
 - Reference counted objects
 - Metadata with same format as CLR's metadata
- Projected to other environments
 - C++/CX – most direct projection
 - .NET platform (any language)
 - JavaScript
- Subset of Win32 and .NET work with Store apps

WinRT Objects

- COM objects implementing `IUnknown` and a new interface, `IIInspectable`

```
class IIInspectable : public IUnknown {  
public:  
    virtual HRESULT __stdcall GetIids(  
        ULONG *iidCount,  
        IID **iids) = 0;  
  
    virtual HRESULT __stdcall GetRuntimeClassName(  
        HSTRING *className) = 0;  
  
    virtual HRESULT __stdcall GetTrustLevel(  
        TrustLevel *trustLevel) = 0;  
};
```

Language Projections



VC++ Projections for WinRT

C++/CX:
Component Extensions

Language extensions

Well aligned with C++/CLI

Consumption: ^, ref new

Authoring: ref/value/interface
class, property, delegate,
event, generic interfaces

WRL:
Windows Runtime
Library

Library-like syntax

No exceptions (HRESULT model)

Consumption: ComPtr<>, make<>

Authoring: Lots of MIDL and
attributes

C++/CX Extensions

Key Bindings	Feature	Summary
1. Data Types	<code>ref class</code>	Reference type
	<code>value class</code>	Value type
	<code>interface class</code>	Interface
	<code>property</code>	Property with get/set
	<code>event</code>	“Delegate property” with add/remove/raise
	<code>delegate</code>	Type-safe function pointer
	<code>generic</code>	Type-safe generics
2. Allocation	<code>ref new</code>	Reference-counted allocation
3. Pointer & Reference	<code>^</code>	Strong pointer (“hat” or “handle”)
	<code>%</code>	Strong reference



Application Binary Interface (ABI)

- WinRT boundary
- Components adhering to the ABI are consumable in other languages
 - Simpler with the provided projections
- Public surface must conform to WinRT types
 - Implementation can use any type

WinRT Components

- DLLs with metadata (.winmd files)
- Cannot be distributed stand alone through the store
- Cannot be shared between applications
- Only sealed classes are allowed
- Only WinRT types can be passed across the interop boundary
- Can use the resulting library from C++, C# or JavaScript Store apps

Asynchrony in WinRT

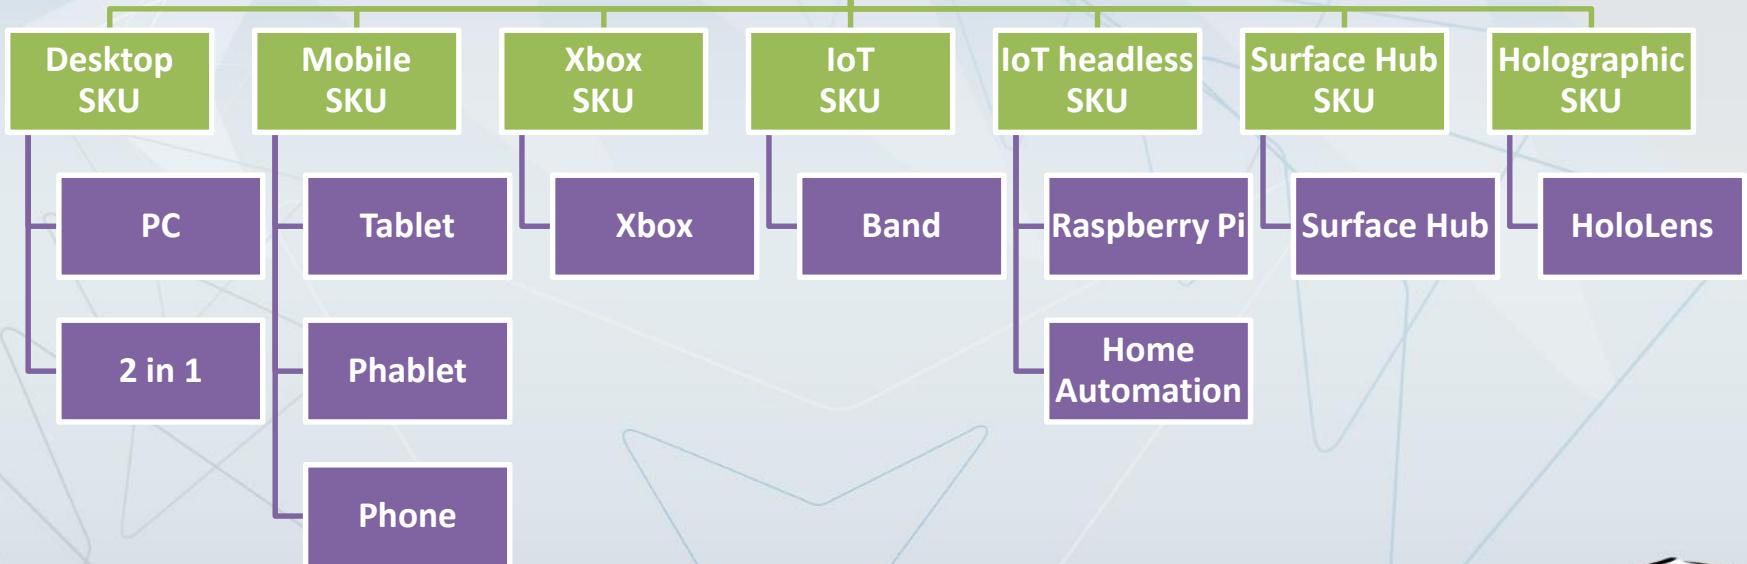
- Many WinRT APIs are asynchronous in nature
 - Call to start an operation
 - Returns immediately
 - Register to be notified when result is available
 - Handle the result
- Basic rule: if something may consume more than 50msec, make it asynchronous
- Some APIs provide asynchronous access only

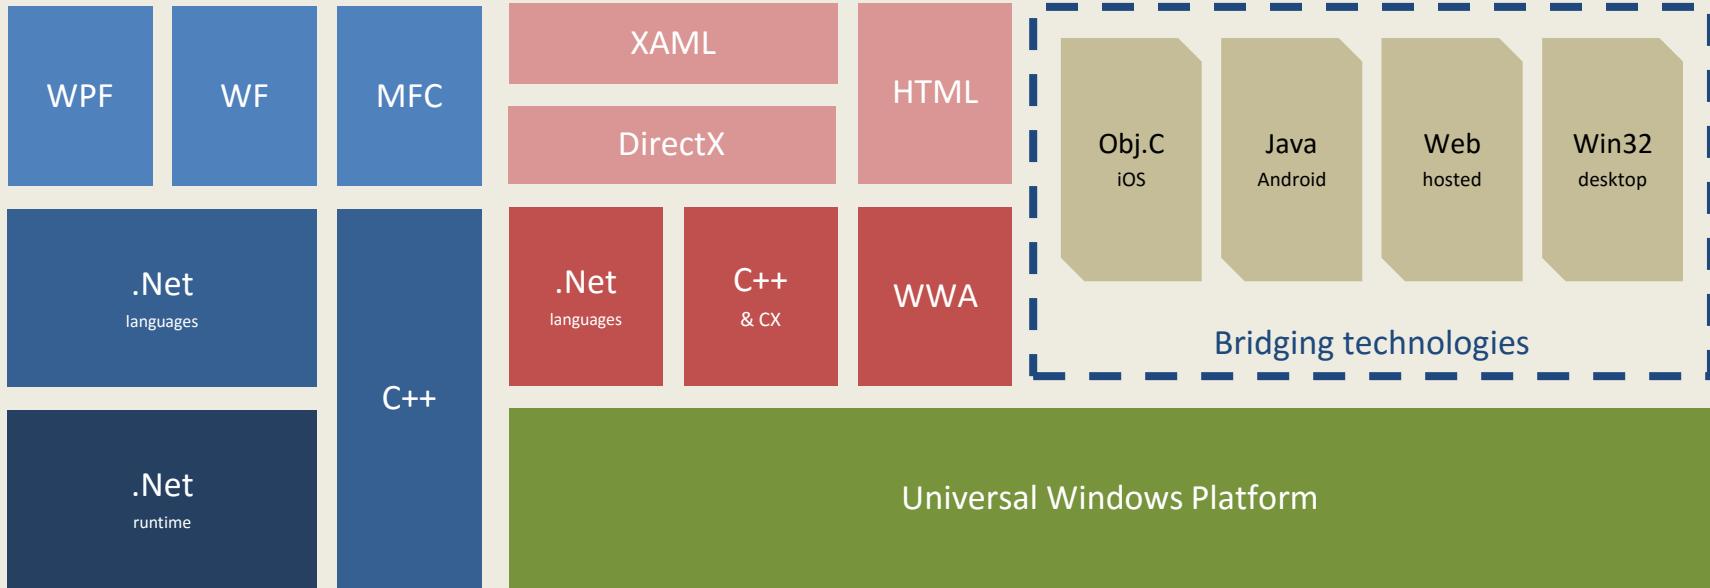
Universal Windows Platform

- One Operating System
 - One Windows core for all devices
- One App Platform
 - Apps run across every family
- One Dev Center
 - Single submission flow and dashboard
- One Store
 - Global reach, local monetization
 - Consumers, Business & Education



One Windows





Windows 10
operating system

Windows 10 IoT Core

- Use the same API, SDK, Visual Studio to develop Universal App
 - Extend UWP with device specific API
 - Run your desktop app on the PI
 - Extend it to control and to be controlled by hardware devices
 - You can Use [Arduino Wiring Pi](#) on Windows 10 IoT devices
 - C++ library that mimic the Arduino library

Setup Raspberry Pi 2 Device

IoT Dashboard

- My devices
- Set up a new device**
- Try some samples

Set up a new device

First, let's get Windows 10 IoT Core on your device.

Device type

Raspberry Pi 2 Windows 10 IoT Core for Raspberry Pi 2

Insert an SD card into your computer.
Note that this process will erase all content on your card.

[View the list of recommended SD cards](#)

Drive

I: 62Gb [Mass Storage Device USB Device]

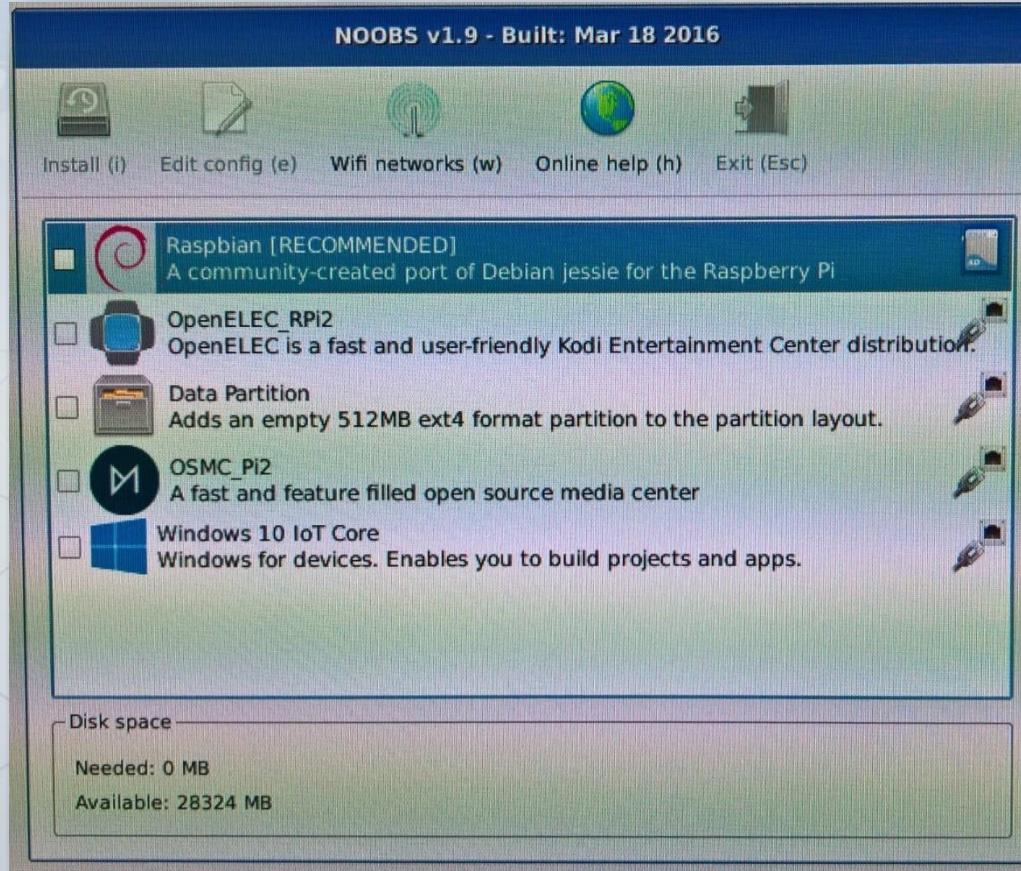
I accept the software license terms

Download and install

Flashing your SD card
18 MB downloading - 3%

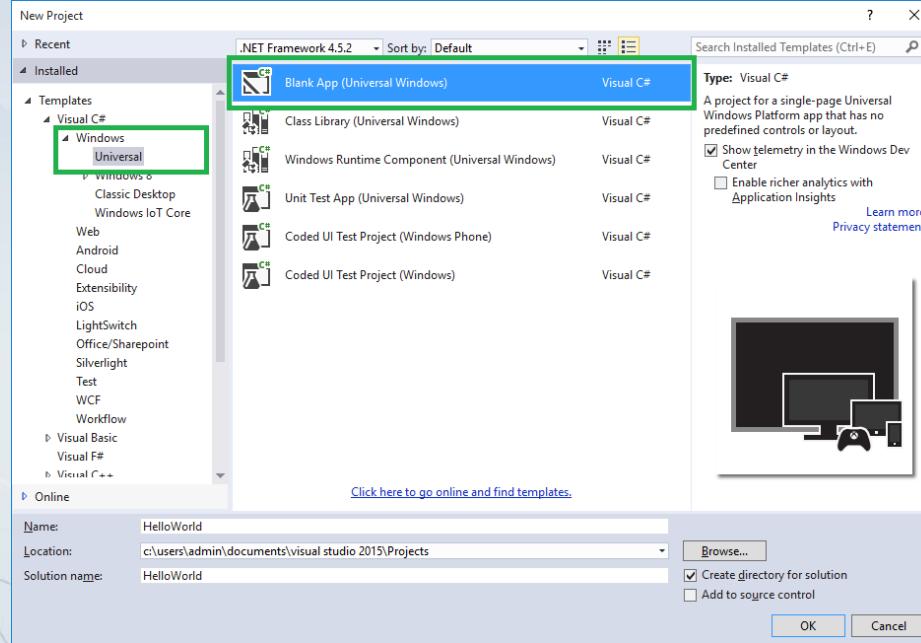
Cancel

NOOBS OS Selection Menu



Hello World – IoT Windows Universal App

- Start Coding



RPi 2 & 3 GPIO



3.3V PWR	1
I2C1 SDA	3
I2C1 SCL	5
GPIO 4	7
GND	9
GPIO 17	11
GPIO 27	13
GPIO 22	15
3.3V PWR	17
SPI0 MOSI	19
SPI0 MISO	21
SPI0 SCLK	23
GND	25
Reserved	27
GPIO 5	29
GPIO 6	31
GPIO 13	33
GPIO 19	35
GPIO 26	37
GND	39
2 5V PWR	2
4 5V PWR	4
6 GND	6
8 UART0 TX	8
10 UART0 RX	10
12 GPIO 18	12
14 GND	14
16 GPIO 23	16
18 GPIO 24	18
20 GND	20
22 GPIO 25	22
24 SPI0 CS0	24
26 SPI0 CS1	26
28 Reserved	28
30 GND	30
32 GPIO 12	32
34 GND	34
36 GPIO 16	36
38 GPIO 20	38
40 GPIO 21	40

Programming + the GPIO

```
using Windows.Devices.Gpio;

public void GPIO()
{
    // Get the default GPIO controller on the system
    GpioController gpio = GpioController.GetDefault();
    if (gpio == null)
        return; // GPIO not available on this system

    // Open GPIO 5
    using (GpioPin pin = gpio.OpenPin(5))
    {
        // Latch HIGH value first. This ensures a default value when the pin is set as output
        pin.Write(GpioPinValue.High);

        // Set the IO direction as output
        pin.SetDriveMode(GpioPinDriveMode.Output);

    } // Close pin - will revert to its power-on state
}
```

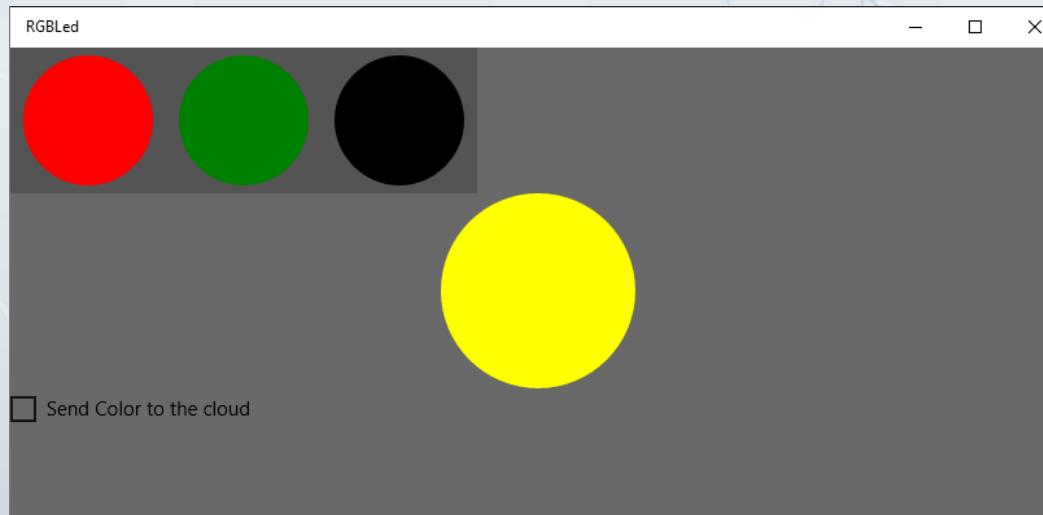
1 reference | Aparajita Dutta, 162 days ago | 1 author, 1 change

```
private void Timer_Tick(object sender, object e)
{
    if (pinValue == GpioPinValue.High)
    {
        pinValue = GpioPinValue.Low;
        pin.Write(pinValue);
        LED.Fill = redBrush;
    }
    else
    {
        pinValue = GpioPinValue.High;
        pin.Write(pinValue);
        LED.Fill = grayBrush;
    }
}
```



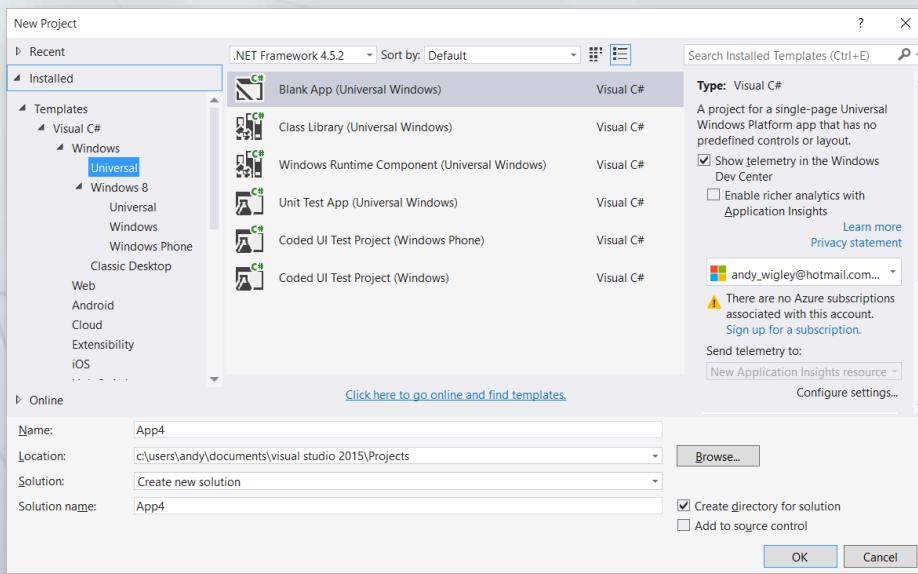
RGB Blinky Demo

- Using C++ and the WiringPi library
- Using Universal App to Control an RGB Led
- RPI GPIO Schema



Windows app

- A single app package
 - Running on any device
 - Testing for capabilities
 - Adjusting to devices



Windows App

Universal Windows Platform

Windows Core

Desktop
Device

Phone
Device

Xbox
Device



Adaptive design

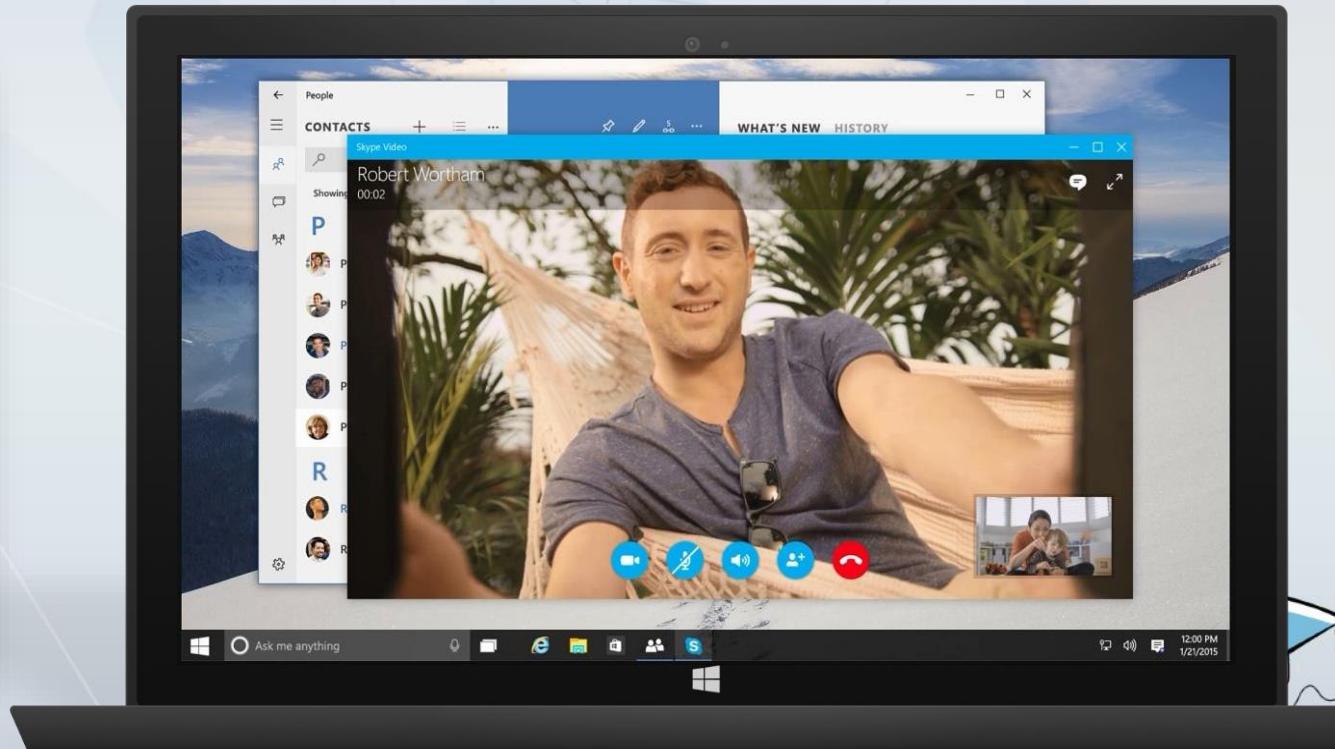
- Responsive design
 - Flexible layout responds to small changes
 - Many controls handle basic responsiveness
- Adaptive design
 - Smart layout adjusts to large changes
 - Features like visual states aid in this design
- Tailored design
 - A device-specific app can simplify design
 - Some devices have unique design languages

Adaptive design

Phone (portrait)

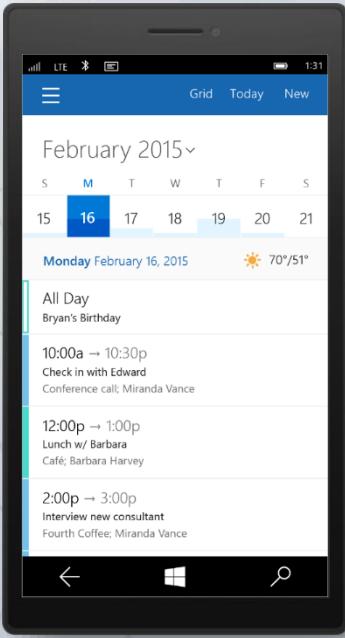


Tablet (landscape) / Desktop

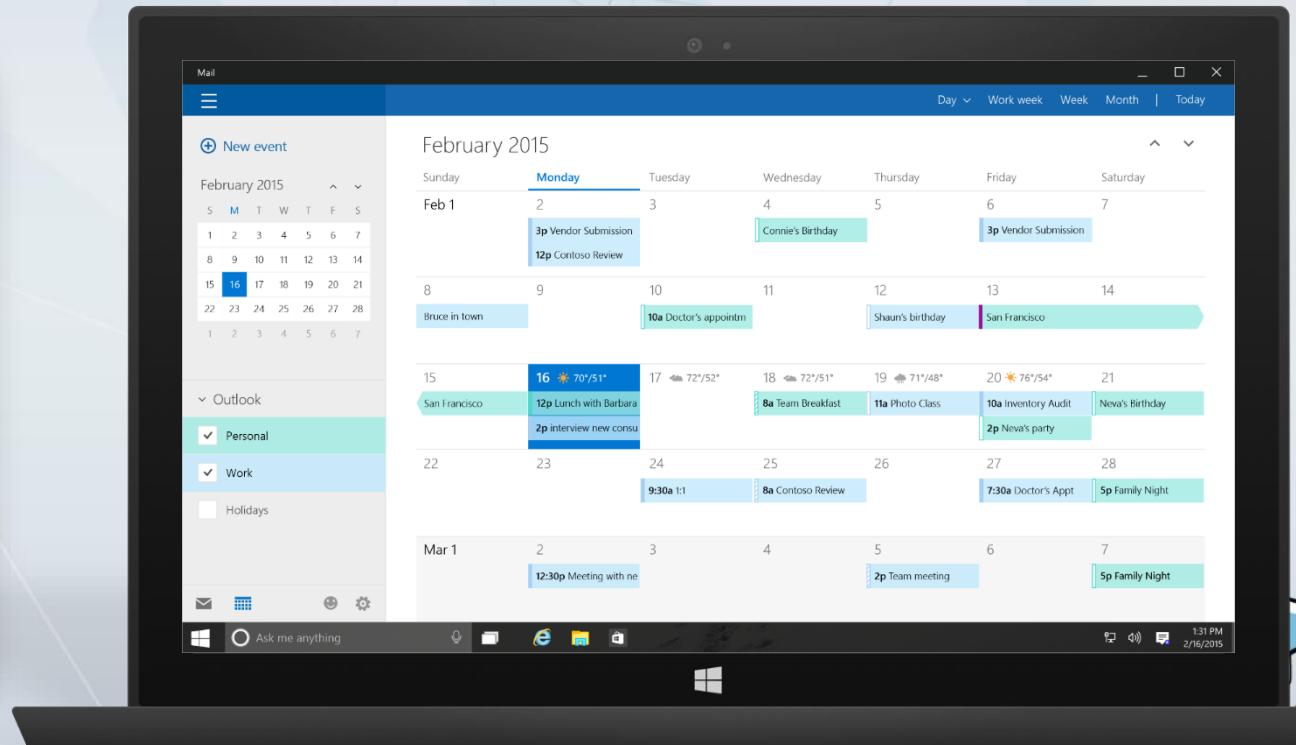


Tailored design

Phone (portrait)



Tablet (landscape) / Desktop



Continuum for
convertibles and
Phones



Adaptive code

- A compatible binary across devices
 - Universal API with device-specific implementation
- Light up our app with capabilities
 - Testing for capabilities and namespaces

Test capabilities at runtime

- Use Adaptive Code to light-up your app on specific devices

```
var api = "Windows.Phone.UI.Input.HardwareButtons";
if (Windows.Foundation.Metadata.ApiInformation.IsTypePresent(api))
{
    Windows.Phone.UI.Input.HardwareButtons.CameraPressed
        += CameraButtonPressed;
}
```

Demo

HELLO UWP

600

Windows Internals



Get Started

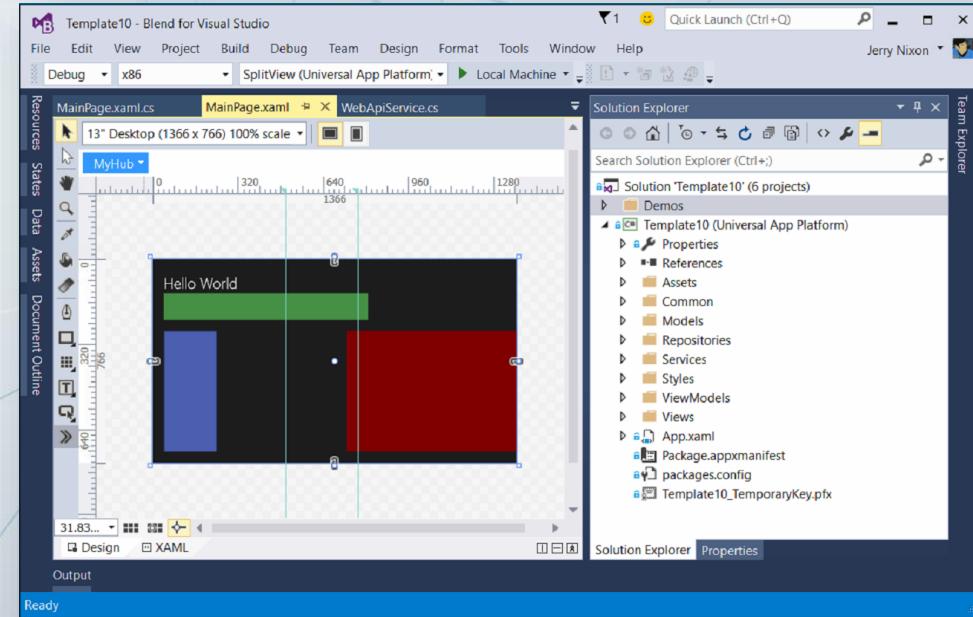
- This short lecture will provide the basic knowledge
- However this is defiantly not enough
 - Start here: <https://dev.windows.com/en-us/getstarted>
 - Download [visual studio Community 2015](#)
 - Follow [this](#) “Create a Hello World app” lab
 - Learn more about the [XAML Platform](#)
 - Learn C# - [this](#) is for absolute beginners
- Feel free to ask question – alonf@codevalue.net

Visual Studio IDE

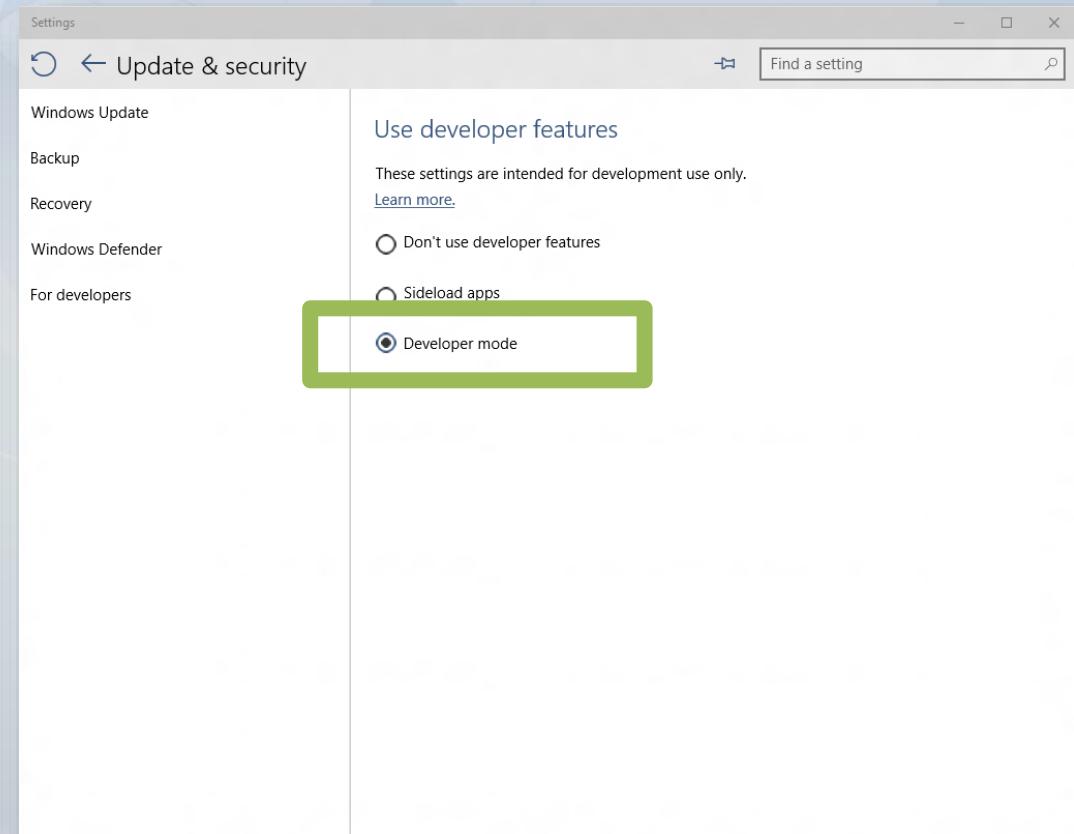
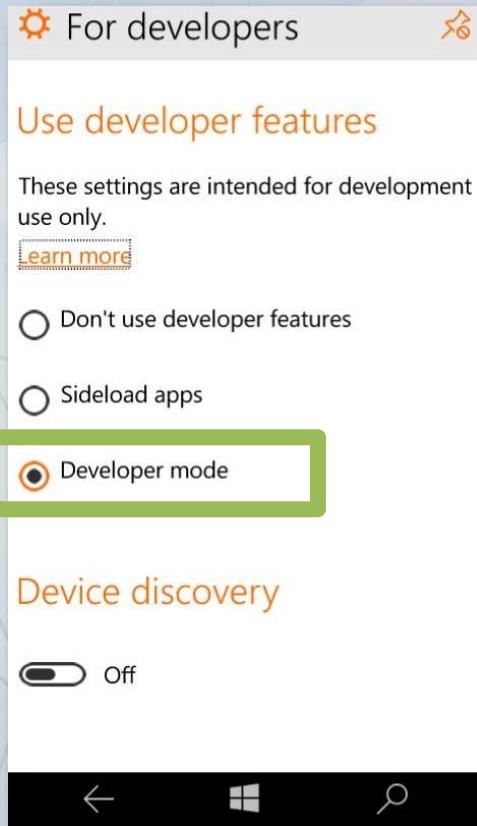
- Every project type
 - Desktop, Windows, Phone, Service, Web, Game, More...
- Every developer task
 - Code edit, Architecture design, UX design, Debug, Profile, Review, Test, More...
- Every development language
 - C++/CX, C#, Visual Basic, JavaScript, XAML, HTML, More...
- Visual Studio Online
 - Source repository, project management, bug tracking, More...

Blend for Visual Studio

- The XAML Developer's IDE
 - Always part of Visual Studio
 - Uses the Visual Studio shell
 - Full auto-complete & intellisense
- Validation
- Snippets
- Peek
 - File & solution management
 - Resource management
 - Data management
 - Animation
 - States



Developer unlock



.NET Native

- Next generation compiler in the cloud
 - Every Windows apps, only Windows app (right now)
- Apps use the standard C++ optimizer
 - As optimizer performance improves, so does .Net native
- Apps with .Net bootstrapper
 - Includes garbage collection
- There is no runtime
 - This is machine code

Real benefits with .Net Native

- 50% faster average startup time
- 14% less average memory usage

C#

- A multi-paradigm programming language encompassing strong typing, imperative, declarative, functional, generic, object-oriented and component-oriented programming disciplines
- Developed by Microsoft and approved as a standard by Ecma (ECMA-334) and ISO (ISO/IEC 23270:2006)
- C# is intended to be a simple, modern, general-purpose, object-oriented programming language
- The most recent version is C# 6.0, which was released on July 20, 2015

C# Basics

- The platform – Compiler, Jitter, CLR, GC, Libraries (references), .NET Native
- Types, Value types and reference types (boxing), Metadata (ILDASM)
- Class and Interfaces, explicit interface, polymorphism (abstract, override, new), partial
- Methods, ref and out parameters, params, lambda expression
- Exceptions, checked and unchecked, finally
- Control Flow (if-else, while, do-while, for, foreach, switch)
- Operator overloading
- Properties
- Delegates & Events
- Attributes
- Generic, constraints
- Enumerators (yield return)
- LINQ
- Task, Async/Await

https://en.wikipedia.org/wiki/Comparison_of_C_Sharp_and_Java

UI Designer and XAML Fundamentals

UI Designer and XAML Fundamentals

- Visual Studio
 - UI Designer
 - Toolbox
 - Document Outline
 - Properties View
 - Device Emulator
- What is XAML?
- Basic XAML
- Type Converters
- Markup Extensions
- Naming Elements
- XAML Rules
- Summary

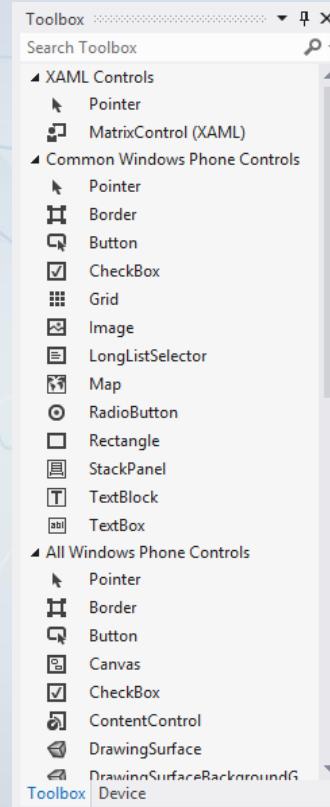
UI Designer

- Visual Studio UI Designer provides quick and easy way to produce rich UI interface
- Drag and drop UI controls from the toolbox into the designer
- The UI designer provides an easy way to select, resize, align, transform and more, using only the mouse
- The UI designer generates XAML and automatically updated when XAML changes



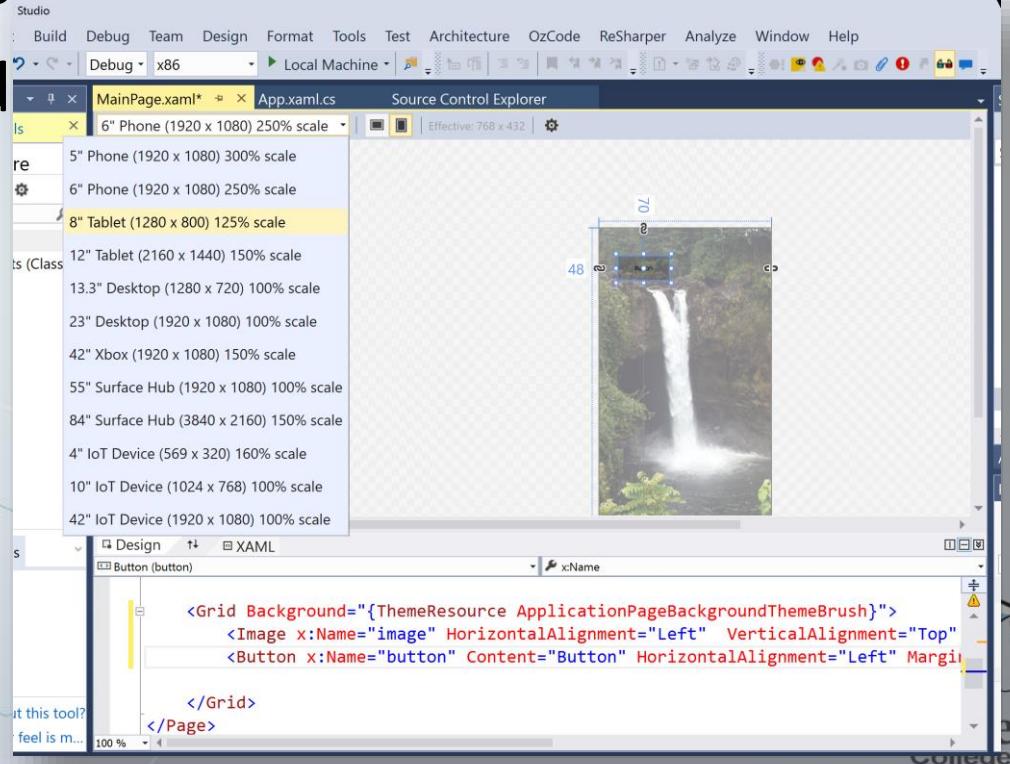
Toolbox

- The toolbox groups UI controls available at design time
- 3rd party and custom controls are also available from the toolbox
- Simple drag and drop a control from the toolbox to the designer canvas
- Element from the toolbox can be also drag and drop directly into XAML
- Controls can be searched within the search box and can be sorted
- Additional groups can be created



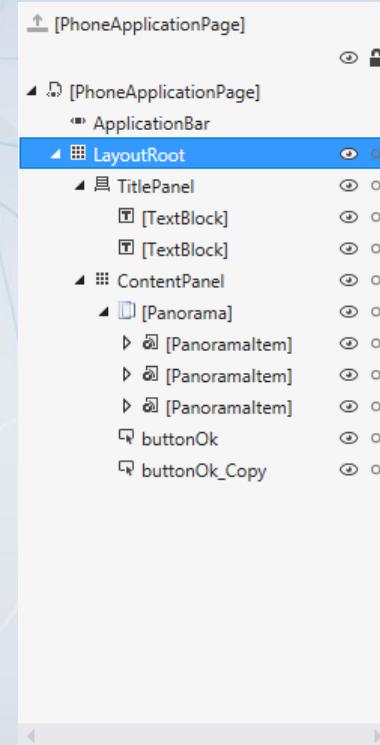
Device View

- The device view provides an easy way to change page layout properties
- Page orientation: Land
- screen resolutions:



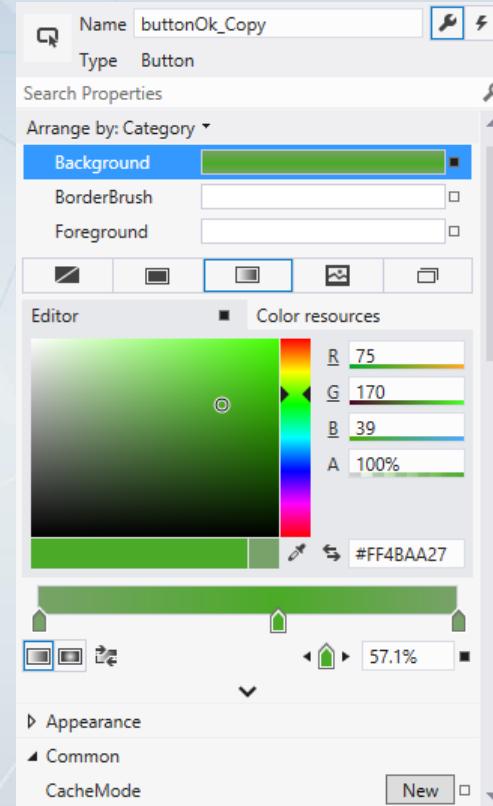
Document Outline

- The document outline view displays the logical tree of UI elements, each element name or type and an icon of the element's type
- Each element in the hierarchy can be design-time locked and hidden
- The document outline is very useful to navigate between UI elements, especially in complex XAML files



Properties View

- The properties view provides an easy a rapid way to:
 - Search for an element's property by its name
 - Set simple property value using plain text
 - Set complex property value using designers
 - Easily select color brushes, styles, font size, transformations and more
 - Register element's events
 - Arrange properties in groups
 - Create data bindings
 - Reset to default values



What is XAML?

- XML based language
- Enable separation of UI and behavior (code)
- Windows Phone related tools emit XAML
- XAML allows
 - Creation of objects
 - Setting of properties
 - Connection to events
 - Custom behaviors
- XAML cannot call methods directly

XAML vs. Code

- Anything that can be done in XAML can be done in code
 - But not vice versa
- XAML is usually shorter and more concise than the equivalent code
 - Thanks to type converters and markup extensions
- XAML should be used for initial UI
- Code will handle events and change items dynamically

Simple XAML Example

```
<Button xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"  
        Content="OK" />
```



```
Windows.UI.Xaml.Controls.Button b = new  
Windows.UI.Xaml.Controls.Button();  
b.Content = "OK";
```

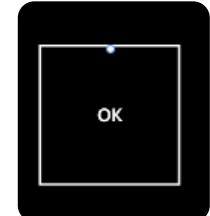
- Visual Studio UI designer generates XAML on each control picked from the toolbox
- XAML Can be visually viewed in the UI designer

XAML Namespaces

- The default XAML namespace is assigned a value that is mapped to some of the runtime namespaces contain UI elements
- Other XAML namespaces may be mapped to custom namespaces and other runtime namespaces
- The “x” namespace is mapped to a special namespace, contains XAML parser specific types
- XAML namespace can be defined on each element level

XAML Example

```
<!--ContentPanel - place additional content here-->
<Grid x:Name="ContentPanel" Grid.Row="1" Margin="12,0,12,0">
    <Button x:Name="button0k"
        Width="200"
        Height="200"
        Content="OK"
        Click="button0k_Click" />
</Grid>
```



Elements and Attributes

- Elements with type names only designate object creation (via the default constructor)
- Attributes indicate property or event values
 - Event values are event handlers (methods) names
- Complex properties are designated using a **<Type.Property>** element

XAML Example

```
<!--ContentPanel - place additional content here-->
<Grid x:Name="ContentPanel" Grid.Row="1" Margin="12,0,12,0">
    <Button x:Name="buttonOk"
        Width="200"
        Height="200"
        Content="OK"
        Click="buttonOk_Click" >
        <Button.Background>
            <LinearGradientBrush EndPoint="0.5,1"
                StartPoint="0.5,0">
                <GradientStop Color="#FFB2D9FF" Offset="0.004"/>
                <GradientStop Color="#FFB0D8FF" Offset="1"/>
                <GradientStop Color="#FF0A85FF" Offset="0.571"/>
            </LinearGradientBrush>
        </Button.Background>
    </Button>
</Grid>
```



XAML And Code Behind

- A root element, usually **Page** or **UserControl** classes, can have code behind file
- The name of the code behind file is correlated to the XAML file name
- For example: **MainPage.xaml** and **MainPage.xaml.cs**

```
<Page
    x:Class="UWPDemo.MainPage"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:UWPDemo"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d">
    <Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
        <Image x:Name="image" HorizontalAlignment="Left" VerticalAlignment="Top" Stretch="Fill" Opacity="0.5"/>
        <Button x:Name="button" Content="Button" HorizontalAlignment="Left" Height="61" Width="127"/>
    </Grid>
</Page>
```

Child Elements

- Child elements (that are not property elements) can be one of
 - The **Content** property of the object
 - A property adorned with the attribute
Windows.UI.Xaml.Controls.ContentProperty
 - Collection items
 - The object implements **IList** or **IDictionary**
 - A value that can be type-converted

Content Property

- A single property that is designated with the **ContentProperty** attribute on the type
- Allows shortening the markup

```
<Button Content="OK" >  
</Button>
```



```
<Button>  
    OK  
</Button>
```

```
<Button>  
    <Button.Content>  
        <Rectangle  
Fill="Blue"/>  
    </Button.Content>  
</Button>
```



```
<Button>  
    <Rectangle Fill="Blue"/>  
</Button>
```

Collection Items

- List (IList)

```
<ListBox>
<ListBox.Items>
<ListBoxItem Content="Item 1"/>
<ListBoxItem Content="Item 2"/>
</ListBox.Items>
</ListBox>
```

```
<ResourceDictionary>
<SolidColorBrush x:Key="br1" Color="Aqua" />
<Rectangle x:Key="rc1" Fill="Brown" />
</ResourceDictionary>
```

Summary of XAML Rules

- XML Element – create a new instance
- XML attribute – set a property or register an event
 - Type converter may execute
- **Type.Property** – set a “complex” property
- **ContentProperty** attribute – no need to specify **Type.Property**
- Property of type **IList** or **IDictionary**
 - Add child elements (XAML calls appropriate **Add** method)
 - Need a **x:Key** in case of a dictionary

Parsing and Using XAML

- Visual Studio compiles the XAML file to XBF (Binary XAML format) and embeds it as a resource
- The XBF file is parsed at runtime and the object tree created by the **InitializeComponent** method of the parent's element class

```
[GeneratedCodeAttribute("Microsoft.Windows.UI.Xaml.Build.Tasks", " 14.0.0.0")]
[global::System.Diagnostics.DebuggerNonUserCodeAttribute()]
public void InitializeComponent()
{
    if (_contentLoaded)
        return;
    _contentLoaded = true;

    Uri resourceLocator = new Uri("ms-appx:///MainPage.xaml");
    Application.LoadComponent(this, resourceLocator,
                             ComponentResourceLocation.Application);
}
```

Naming Elements

- Elements can be named using the **x:Name** XAML attribute
- The code-behind file will contain a field with that name
- Elements deriving from **FrameworkElement** contain a **Name** property that can be used in code to locate elements
 - **x:Name** and **Name** cannot be set on the same element

XAML Keywords

Keyword	Valid on	Meaning
x:Class	Root element	The class that derives from the element type
x:ClassModifier	Root element, must be used with x:Class	The class visibility (public by default)
x:FieldModifier	Element, must be used with x:Name	Visibility of the field created behind the element
x:Key	Element that its parent implements IDictionary	Key in a dictionary
x:Name	Element	The element's name, used for a field name for that element
X:Uid	Element	Identifies elements that should use localized resources

Mapping custom types to XAML namespaces

- You can define your own custom types in C# and then reference your custom types in XAML markup
- To use XAML for custom types - those that come from libraries other than the Windows Runtime core libraries:
 - You must declare and map a XAML namespace with a prefix
 - Use that prefix in element usages to reference the types that were defined in your library
 - You declare prefix mappings as `xmlns` attributes
- For example:
 - the attribute syntax to map a prefix `myTypes` to the namespace `myCompany.myTypes` is
 - `xmlns:myTypes="using:myCompany.myTypes"`
 - The representative element usage is: `<myTypes:CustomButton/>`

XAML Markup Extensions

- Represent some kind of "shortcut" that enables a XAML file to access a value or behavior that isn't simply declaring elements based on backing types
- In XAML attribute syntax, curly braces "{}" and "{}" indicate a XAML markup extension usage
- A XAML parser calls code that provides behavior for that particular markup extension
 - That code provides an alternate object or behavior result that the XAML parser needs
- Examples:
 - {x:Bind} {Binding} {StaticResource} {ThemeResource} {TemplateBinding} {RelativeSource} {CustomResource} {x:Null}

Markup Extension Example

```
<Canvas.Resources>
    <Style TargetType="Border" x:Key="PageBackground">
        <Setter Property="BorderBrush" Value="Blue"/>
        <Setter Property="BorderThickness" Value="5"/>
    </Style>
</Canvas.Resources>
...
<Border Style="{StaticResource PageBackground}">
...
</Border>
```

XAML and .NET Events

- XAML has a syntax for attaching event handlers to objects in the markup
- You specify the name of the event as an attribute name on the object where the event is handled
 - For the attribute value, you specify the name of an event-handler function that you define in code
- The XAML processor uses this name to create a delegate representation in the loaded object tree, and adds the specified handler to an internal handler list

```
<Button Click="showUpdatesButton_Click">Show updates</Button>
```

Summary

- XAML is mainly used to create a Windows app user interface
- It declaratively allows object creation, property and event assignment
- A code-behind file will usually contain the procedural logic
- Sharing with designers is easier
- Tools such as Expression Blend generate XAML that is immediately usable

BASIC CONCEPTS

Agenda

- Logical and Visual Trees
- Dependency Properties
- Attached Properties
- Routed Events
- Attached Events
- Resources
- Summary

Logical Tree

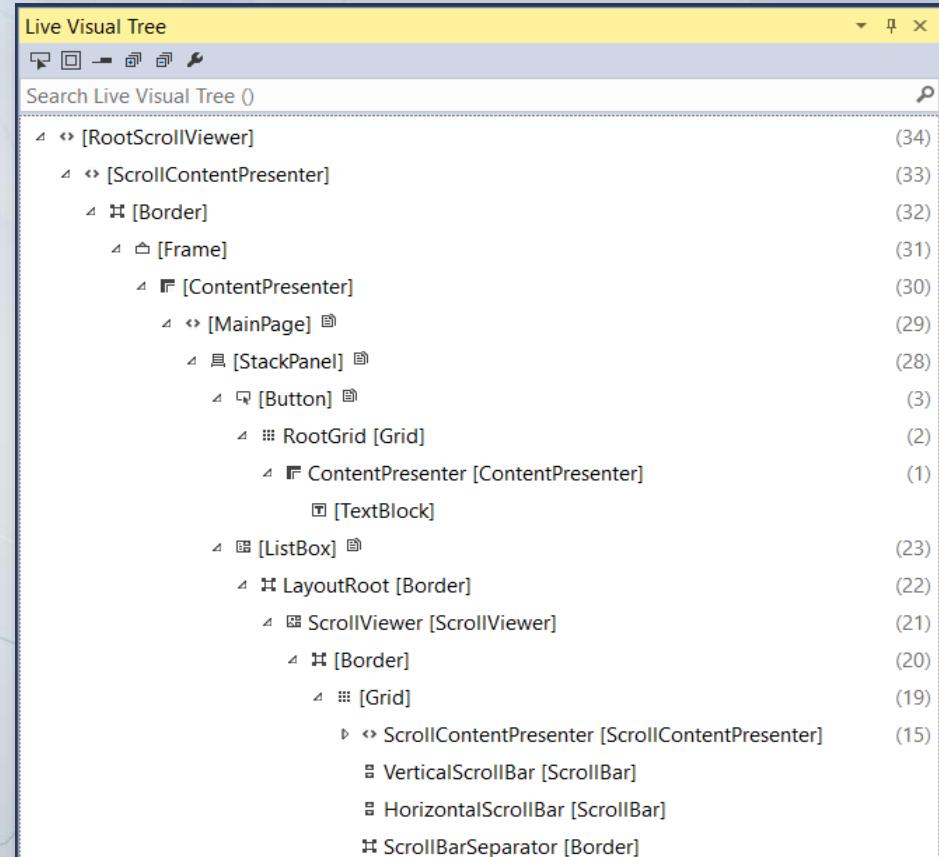
- A tree of elements/controls making up the user interface

```
<Page x:Class="UWPDemo.MainPage"  
      xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"  
      xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml">  
    <StackPanel>  
      <Button Content="OK" Background="Red" Margin="6"/>  
      <ListBox Margin="6" >  
        <ListBoxItem Margin="2">  
          <Border BorderThickness="4">  
            <TextBlock Text="Hello" FontSize="20" />  
          </Border>  
        </ListBoxItem>  
        <ListBoxItem Content="Item 2" />  
      </ListBox>  
    </StackPanel>  
</Page>
```



Visual Tree

- The actual visual objects making up a logical tree
 - Generated from the control templates



Traversing the Trees

- The visual trees can be examined using the static classes

`Windows.UI.Xaml.Media.VisualTreeHelper`

- `VisualTreeHelper`
 - `GetParent`, `GetChild`, `GetChildrenCount`

Dependency Properties

- Dependency properties are the workhorse of any XAML based technology
 - Provide the basis for many of UWP's features
 - E.g. property changed, data binding, animations, styles, default value
 - Value may be provided by many entities, each with its own priority
 - Highest active level wins out
- Must be declare in a class derived from DependencyObject
- Must be “registered” by calling the DependencyProperty.Register method
- Most UWP properties are dependency
- UWP can handle Dependency Properties in a new way

Dependency Property Example

```
// IsSpinningProperty is the dependency property identifier
// no need for info in the last PropertyMetadata parameter, so we pass null
    public static readonly DependencyProperty IsSpinningProperty =
        DependencyProperty.Register(
            "IsSpinning", typeof(Boolean),
            typeof(MainPage), null
        );
// The property wrapper, so that callers can use this property through a simple
// MainPageInstance.IsSpinning usage rather than requiring property system APIs
    public bool IsSpinning
    {
        get { return (bool)GetValue(IsSpinningProperty); }
        set { SetValue(IsSpinningProperty, value); }
    }
```

Dependency Property Precedence

- Animation
- Local value (also resource or data binding)
- Templated Property (Control or Data template)
- Style setters
- Default value from property metadata



Attached Properties

- Special kind of dependency properties
- May be “attached” to objects of different types than the declaring type
 - Declared with the static [**DependencyProperty.RegisterAttached**](#) method
- Allows “context” properties
 - E.g. **Canvas.Left** for elements that happen to be in a Canvas element
 - Can be set on any object
- XAML
 - An attribute with **Type.Property** syntax is used
- In code
 - The type exposes a **SetXxx** and a **GetXxx** with the element reference

Attached Properties Example

- XAML

```
<Canvas>
    <Button x:Name="cmdOK" Canvas.Left="30" Canvas.Top="20"
            Content="OK" Padding="10" FontSize="26">
    </Button>
</Canvas>
```

- Code

```
Canvas.SetLeft(cmdOK, 30);
Canvas.SetTop(cmdOK, 20);
```



```
cmdOK.SetValue(Canvas.LeftProperty, 30);
cmdOK.SetValue(Canvas.TopProperty, 30);
```

DependencyObject

- Represents an object that participates in the dependency property system
- Base class of many important UI-related classes: [**UIElement**](#), [**Geometry**](#), [**FrameworkTemplate**](#), [**Style**](#), and [**ResourceDictionary**](#)
- The dependency property system's primary function is to compute the values of properties, and to provide system notification about values that have changed
- Dependency & Attach property hosting support for the existing Windows Runtime dependency properties and custom properties
- The [**Dispatcher**](#) property for advanced threading scenarios

Routed Events

- UWP events are not implemented via the default .NET event implementation
- A routed event is an event that is potentially passed on (*routed*) from a child object to each of its successive parent objects
- The Windows Runtime supports the concept of a routed event for a set of events that are present on most UI elements
 - These events are for input and user interaction scenarios
 - They are implemented on the [**UIElement**](#) base class
 - For example:
 - [**DoubleTapped**](#)
 - [**KeyDown**](#)
 - [**KeyUp**](#)

Binary Resources

- The resources that can be used by any .NET application
 - Usually store bitmaps, icons, etc.
- Packaging
 - Embedded inside an Assembly
 - Loose files that are known to the application at compile time
 - Loose files that are not known to the application at compile time
- May be localized

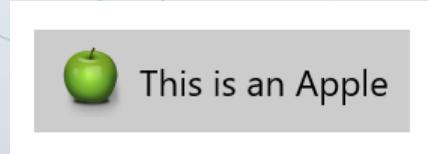
Defining Binary Resources

- Binary resources can be embedded using Visual Studio by adding the resource and selecting an action type
 - **Resource**
 - Embedded inside the assembly
 - **Content**
 - Remains as a loose file, but an **AssemblyAssociatedContentFile** attribute is added with the relative path of the file
- Don't use the **Embedded Resource** action! (UWP can't use such resources)

Accessing Binary Resources

- Binary resources that are added to the project as Content or Resource can be referenced easily, even if loose files are involved (as long as they are in the

```
<Button VerticalAlignment="Center" HorizontalAlignment="Center" >
    <StackPanel Orientation="Horizontal" >
        <Image Source="Assets/apple.png" Width="48" Height="48" Margin="4"/>
        <TextBlock Text="This is an Apple" FontSize="24"
            VerticalAlignment="Center" Margin="4"/>
    </StackPanel>
</Button>
```



Logical Resources

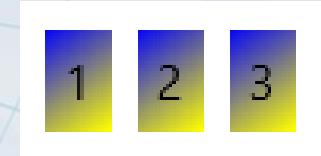
- Arbitrary named .NET objects, stored in the **Resources** collection property of an element
 - Typically for sharing the resource among child objects
- The **FrameworkElement** type define a **Resources** property (of type **ResourceDictionary**)

Creating and Using Resources

- Add a **Resources** property to some element
 - Usually a **Page** or the **Application**
 - Any child element can reference those resources
- Add the objects with a **x:Key** attribute (must be unique in this resource dictionary)
- Use the **StaticResource** markup extension with the resource key name

Resources Example

```
<StackPanel Margin="4" Orientation="Horizontal">
    <Button Margin="4" Content="1" Padding="4">
        <Button.Background>
            <LinearGradientBrush>
                <GradientStop Offset="0" Color="Blue" />
                <GradientStop Offset="1" Color="Yellow" />
            </LinearGradientBrush>
        </Button.Background>
    </Button>
    <Button Margin="4" Content="2" Padding="4">
        <Button.Background>
            <LinearGradientBrush>
                <GradientStop Offset="0" Color="Blue" />
                <GradientStop Offset="1" Color="Yellow" />
            </LinearGradientBrush>
        </Button.Background>
    </Button>
    <StackPanel Margin="4" Orientation="Horizontal">
        <StackPanel.Resources>
            <LinearGradientBrush x:Key="back">
                <GradientStop Offset="0" Color="Blue" />
                <GradientStop Offset="1" Color="Yellow" />
            </LinearGradientBrush>
        </StackPanel.Resources>
        <Button Margin="4" Content="1" Padding="4" Background="{StaticResource back}" />
        <Button Margin="4" Content="2" Padding="4" Background="{StaticResource back}" />
        <Button Margin="4" Content="3" Padding="4" Background="{StaticResource back}" />
    </StackPanel>
</StackPanel>
```



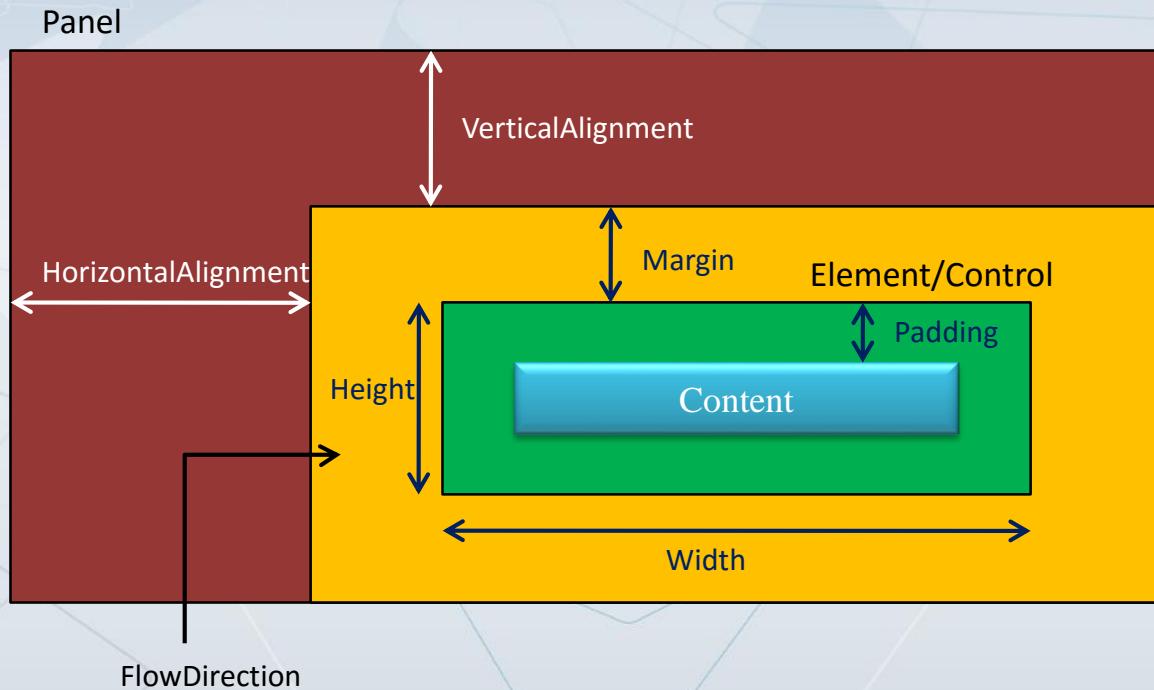
Layout

- Layout is the arranging of user interface elements within some container
- Older technologies (e.g. Windows Forms) mostly used exact position and sizes
 - Limited in flexibility and adaptability
- UWP provides several layout panels that can control dynamically size and placement of elements
- Elements may advertise their size and position needs
- Two layout kind:
 - Static layout: Explicit pixel sizes and positions (Canvas)
 - Fluid layout: shrink, grow and reflow to adapt the visual space available

Size and Position of Elements

- Element sizing and positioning is determined by the element itself and its logical parent
- A child element may request various settings
- The parent panel does not have to comply

Element Layout Properties



Element Size

- **Width** and **Height** properties (from `FrameworkElement`)
 - Control the exact size of the element
 - Default value is **Double.NaN**
 - Meaning: be as large as it needs to be
 - usually a bad idea to use these properties
 - Prevents smart resizing by panel
 - Reasonable only in a Canvas
- **MinWidth**, **MinHeight**, **MaxWidth**, **MaxHeight** properties
 - Defaults are 0 (**MinWidth**, **MinHeight**),
Double.PositiveInfinity (“Infinity” in XAML)
(**MaxWidth**, **MaxHeight**)

More Size Properties

- Read only properties (UIElement)
 - **DesiredSize**
 - Set indirectly by elements to report their desired size to their parent
 - Internally used by panels
 - **RenderSize**
 - The actual size the element is rendered with
 - **ActualWidth, ActualHeight**
 - Just the components of RenderSize (**RenderSize.Width**, **RenderSize.Height**)

Margin and Padding

- Both of type **Thickness** (value type)
 - Maintains the properties **Left, Top, Right, Bottom** indicating distance from the corresponding edge
- **Margin** (from FrameworkElement)
 - The amount of space to add around the element
- **Padding** (from Control, and Border)
 - The amount of space to add around the content of the control
- In XAML, can supply one, two or four numbers

Visibility

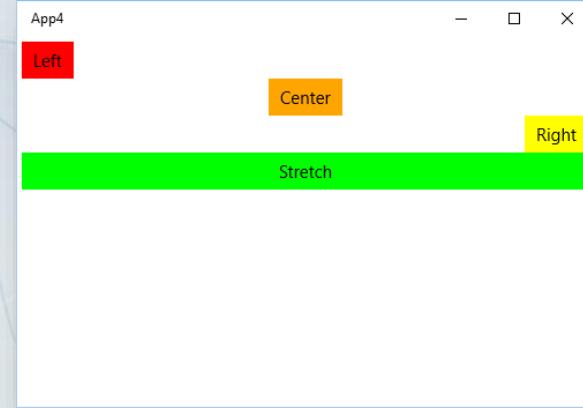
- Visibility of elements is determined by the **Visibility** property (from `UIElement`) of the **Windows.UI.Xaml.Visibility** enumeration
 - **Visible**
 - The element is rendered and participates in layout
 - **Collapsed**
 - The element is invisible and does not participate in layout

Element Positioning

- Element position is determined by the containing panel policy
 - Simple X,Y is not the usual case
- Elements can indicate their position constraints or preferences to their containing panel

Alignment

- Alignment indicates what should be done with any extra space given to an element
- **HorizontalAlignment**
 - **Left, Right, Center, Stretch**
- **VerticalAlignment**



```
<StackPanel Margin="4">
    <Button HorizontalAlignment="Left" Background="Red">Left</Button>
    <Button HorizontalAlignment="Center" Background="Orange">Center</Button>
    <Button HorizontalAlignment="Right" Background="Yellow">Right</Button>
    <Button HorizontalAlignment="Stretch" Background="Lime">Stretch</Button>
</StackPanel>
```

Content Alignment

- Similar to element alignment
- What to do with extra space when the content is smaller than its control
- HorizontalContentAlignment
- VerticalContentAlignment



```
<StackPanel Margin="4" >
    <Button HorizontalAlignment="Stretch" HorizontalContentAlignment="Left" Background="Red">Left</Button>
    <Button HorizontalAlignment="Stretch" HorizontalContentAlignment="Center" Background="Orange">Center</Button>
    <Button HorizontalAlignment="Stretch" HorizontalContentAlignment="Right" Background="Yellow">Right</Button>
    <Button HorizontalAlignment="Stretch" HorizontalContentAlignment="Stretch" Background="Lime">Stretch</Button>
</StackPanel>
```

Flow Direction

- The **FlowDirection** property indicates the flow of layout
 - **LeftToRight** (the default)
 - **RightToLeft**
- The **RightToLeft** setting reverses the meaning of “left” and “right”

Layout Panels

- Layout panels derive from the abstract ***Windows.UI.Xaml.Controls.Panel*** class
- Maintain a **Children** property of type **UIElementCollection** (its **ContentProperty**)
- Each child element can be a panel as well
 - Allows creation of complex and adaptive user interfaces
- UWP provides several built in panels
 - Custom layout panels can be created as well

UWP Layout Panels

- Main layout panels
 - **Canvas**
 - Arranges children in a 2D coordinate system
 - **StackPanel**
 - Arranges children in a horizontal or vertical “stack”
 - **VariableSizedWrapGrid**
 - Provides a grid-style layout panel where each tile/cell can be variable size based on content
 - **Grid**
 - Arranges children in a flexible grid
 - **RelativePanel**
 - Defines an area within which you can position and align child objects in relation to each other or the parent panel

Canvas

- Doesn't support fluid UI
 - you control all aspects of positioning and sizing child elements
 - You typically use it for special cases like creating graphics or to define small static areas of a larger adaptive UI
 - You can use code or visual states to reposition elements at runtime
- Elements are positioned absolutely using Canvas.Top and Canvas.Left attached properties
- Layering can be explicitly specified using the Canvas.ZIndex
- If an element's size is not set explicitly, it sizes to its content
- Child content is not visually clipped if larger than the panel
- Child content is not constrained by the bounds of the panel

Canvas Example

```
<Page
    x:Class="CanvasDemo.MainPage"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:CanvasDemo"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d">

    <Canvas x:Name="_canvas" PointerPressed="OnPress" Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
        <Rectangle Fill="Blue" Width="100" Height="100" Canvas.Left="100" Canvas.Top="100"/>
    </Canvas>
</Page>
```

```
private void OnPress(object sender, PointerRoutedEventArgs e)
{
    var ellipse = new Ellipse
    {
        Width = Radius, Height = Radius,
        Fill = new SolidColorBrush(Colors.BlueViolet),
    };
    _canvas.Children.Add(ellipse);
    var position = e.GetCurrentPoint(_canvas).Position;
    Canvas.SetLeft(ellipse, position.X - Radius/2);
    Canvas.SetTop(ellipse, position.Y - Radius/2);
}
```

StackPanel

- Stacks its elements in a vertical or horizontal “stack”
- **Orientation** property
 - **Vertical** (default) or **Horizontal**
- Alignment is ignored in the direction of stacking
- In the direction specified by property, an element sizes t



```
<Page
    x:Class="StackPannelDemo.MainPage"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:StackPannelDemo"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d"
    Width="400"
    Height="200">

    <Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
        <Grid.ColumnDefinitions>
            <ColumnDefinition/>
            <ColumnDefinition/>
        </Grid.ColumnDefinitions>
        <StackPanel Orientation="Horizontal" Grid.Column="0" >
            <Button Margin="4" Background="Red" VerticalAlignment="Stretch">One</Button>
            <Button Margin="4" Background="Orange" VerticalAlignment="Stretch">Two</Button>
            <Button Margin="4" Background="LightGreen" VerticalAlignment="Stretch">Three</Button>
        </StackPanel>
        <StackPanel Orientation="Vertical" Grid.Column="1">
            <Button Margin="4" Background="Red" HorizontalAlignment="Stretch">One</Button>
            <Button Margin="4" Background="Orange" HorizontalAlignment="Stretch">Two</Button>
            <Button Margin="4" Background="LightGreen" HorizontalAlignment="Stretch">Three</Button>
        </StackPanel>
    </Grid>
</Page>
```



The Grid

- The most versatile and useful panel
- Usually used as the top-level panel
 - Visual Studio and Expression Blend windows/user controls start this way
- supports fluid resizing of child elements
 - You can use code or visual states to reposition and reflow elements
- Arranges its children in a multi-row and multi-column way
 - Their sizes and number can be manipulated in interesting ways
 - Somewhat similar to an HTML table
- Child content is visually clipped if larger than the panel

Creating a Grid

- For rows
 - Set the **RowDefinitions** property
 - Add a **RowDefinition** object for each row
 - Set any special properties
- For columns
 - Set the **ColumnDefinitions** property
 - Add a **ColumnDefinition** object for each column
 - Set any special properties
- For each element
 - Set the **Grid.Row** and **Grid.Column** attached properties (default is 0, 0)

Sizing Rows and Columns

- By default, all rows are of equal height and all columns are of equal width
 - Can change the height of a row using the **RowDefinition.Height** property
 - Can change the width of a column using the **ColumnDefinition.Width** property
 - Each one of type **GridLength**
 - The unit is controlled by the **GridUnitType** property
 - **Auto** – size as required by content
 - **Pixel** – (double value) size is the number specified
 - **Star** – size is a weighted proportional (default)
 - “*”, “2*”, etc. in XAML
 - Spanning
 - A row may span more than one column and vice versa
 - Can be set by the **Grid.RowSpan** and **Grid.ColumnSpan** attached properties
 - Default for both is 1

Grid Sample

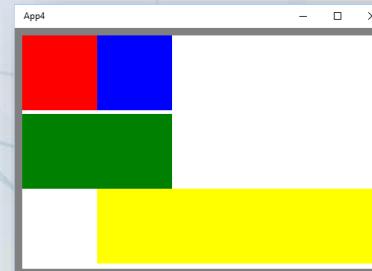
Column_1	Auto	The column will size to its content.
Column_2	*	After the Auto columns are calculated, the column gets part of the remaining width. Column_2 will be one-half as wide as Column_4.
Column_3	44	The column will be 44 pixels wide.
Column_4	*2	After the Auto columns are calculated, the column gets part of the remaining width. Column_4 will be twice as wide as



```
<Grid>
    <Grid.ColumnDefinitions>
        <ColumnDefinition Width="Auto"/>
        <ColumnDefinition/>
        <ColumnDefinition Width="44"/>
        <ColumnDefinition Width="2*"/>
    </Grid.ColumnDefinitions>
    <TextBlock Text="Column 1 sizes to its content." FontSize="24"/>
</Grid>
```

RelativePanel

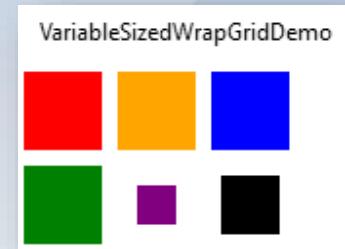
- A layout container that is useful for creating UIs that do not have a clear linear pattern
- Elements are arranged in relation to the edge or center of the panel, and in relation to each other
- Child content is visually clipped if larger than the panel



```
<RelativePanel BorderBrush="Gray" BorderThickness="10">
    <Rectangle x:Name="RedRect" Fill="Red" MinHeight="100" MinWidth="100"/>
    <Rectangle x:Name="BlueRect" Fill="Blue" MinHeight="100" MinWidth="100" RelativePanel.RightOf="RedRect"/>
    <!--Width is not set on the green and yellow rectangles.
        It's determined by the RelativePanel properties. -->
    <Rectangle x:Name="GreenRect" Fill="Green" MinHeight="100" Margin="0,5,0,0" RelativePanel.Below="RedRect"
        RelativePanel.AlignLeftWith="RedRect" RelativePanel.AlignRightWith="BlueRect"/>
    <Rectangle Fill="Yellow" MinHeight="100" RelativePanel.Below="GreenRect"
        RelativePanel.AlignLeftWith="BlueRect" RelativePanel.AlignRightWithPanel="True"/>
</RelativePanel>
```

VariableSizedWrapGrid

- Elements are arranged in rows or columns
- Elements are automatically wrap to a new row or column
 - When the MaximumRowsOrColumns value is reached
- Whether elements are arranged in rows or columns is specified by the Orientation property
- Elements can span multiple rows and columns using:
 - VariableSizedWrapGrid.RowSpan and VariableSizedWrapGrid.ColumnSpan
- Elements are sized as specified by the ItemHeight and ItemWidth properties
 - If these properties are not set, the item in the first cell sizes to its content, and all other cells inherit this size
- Child content is visually clipped if larger than the panel



VariableSizedWrapGrid Example

<Page

```
x:Class="VariableSizedWrapGridDemo.MainPage"
xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
xmlns:local="using:VariableSizedWrapGridDemo"
xmlns:d="http://schemas.microsoft.com/expressionblend/2008"
xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
mc:Ignorable="d">

<VariableSizedWrapGrid MaximumRowsOrColumns="3" Orientation="Horizontal"
                        Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
    <Rectangle Margin="4" Fill="Red" Height="40" Width="40"></Rectangle>
    <Rectangle Margin="4" Fill="Orange" Height="40" Width="40"></Rectangle>
    <Rectangle Margin="4" Fill="Blue" Height="80" Width="80"></Rectangle>
    <Rectangle Margin="4" Fill="Green" Height="100" Width="140"></Rectangle>
    <Rectangle Margin="4" Fill="Purple" Height="20" Width="20"></Rectangle>
    <Rectangle Margin="4" Fill="Black" Height="30" Width="30"></Rectangle>

</VariableSizedWrapGrid>
</Page>
```

VariableSizedWrapGridDemo



What is a Control?

- Controls are elements capable of receiving focus and handling input
 - You add a control to your app UI.
 - You set properties on the control, such as width, height, or foreground color
 - You hook up some code to the control so that it does something
- Many controls are available “out of the box”
- Custom controls can be created
 - User controls that wrap one or more controls and expose higher level properties
 - Custom controls that derive from an existing control and extend its functionality

Example: Static Text

- The **TextBlock** Element
 - The **Text** property
 - Font related properties
 - **FontSize**, **FontFamily**, etc.
 - **TextAlignment**, **TextTrimming**, **TextDecorations**
- An optional collection of “inlines”
 - Replacement for the **Text** property
 - Inheriting from the abstract **Inline** class
 - **Hyperlink**, **Bold**, **Run**, **Span**, **Underline**, **Italic**, **LineBreak**

Example: TextBlock

```
<TextBlock FontSize="16" Margin="4">
    <Run Text="Hello" />
    <Bold>
        Hello
        <Italic>Hello</Italic>
    </Bold>
    <LineBreak />
    <Hyperlink>Go to my web site</Hyperlink>
    <LineBreak /><LineBreak />
    <Underline>This line is underlined</Underline>
    <LineBreak />
    <Span FontSize="20">
        Hello in Bigger Font
    </Span>
</TextBlock>
```

UWPDemo

Hello **Hello** **Hello**
[Go to my web site](#)

This line is underlined
Hello in Bigger Font

Other Controls

- Here you can find the full [control list](#)

Common XAML Controls	
	Pointer
	Border
	Button
	CheckBox
	ComboBox
	DatePicker
	FlipView
	Flyout
	Grid
	GridView
	Hub
	Image
	ListView
	MenuFlyout
	Pivot
	PivotItem
	RadioButton
	Rectangle
	RelativePanel
	StackPanel
	SplitView
	TextBlock
	TextBox
	TimePicker

DATA BINDING

What is Data Binding?

- A way for your app's UI to display data, and optionally to stay in sync with that data
- Data binding allows you to separate the concern of data from the concern of UI, for better:
 - Readability, Testability, and Maintainability
- You can choose to use either the {x:Bind} or {Binding} markup extension
- {x:Bind} is new for Windows 10 and it has better performance
 - {Binding} has more features
- Data binding means tying two arbitrary objects
- Typical scenario is a non-visual object (or collection) to a visual element
 - Any changes to the non-visual object are reflected in the visual element (and optionally vice versa)

Data Binding Concepts

- Source
 - The data object to bind to
- Property Path
 - The property on the source object to use
 - May be a nested property, an array element or an indexer
- Target
 - The target object to modify
 - Must be a dependency property
- Binding Mode
 - Typically one way or two way (target update source)

Using Data Binding

- Typically done in XAML using the {Binding} markup extension
 - The Binding class is the workhorse behind the scenes
 - Set on the target property
- {x:Bind} – a new binding markup extension
 - Provides compile-time syntax validation
 - Provides better performance
 - Where Binding used reflection at runtime to handle binding, x:Bind is able to produce strongly typed code at compile-time to handle the bindings
 - The defaults is Mode=OneTime whereas Binding will default to Mode=OneWay
 - It requires defining the type of the variable using x:DataType

Binding Direction

- The **Binding** object allows specifying how the target / source properties are updated
- **Mode** property (of type enum **BindingMode**)

Binding Mode	Meaning
OneWay	The target property is updated by source property changes
TwoWay	OneWay + the source property is updated by changes of the target property
OneTime	Target is updated by the source the first time they are bound

Binding to Objects

- **Source**
 - Reference to the source object
- **RelativeSource**
 - Sets the source based on the “location” of the target in the layout tree
 - Useful in data and control templates
- **DataContext**
 - Used by default if **Source** or **RelativeSource** is not specified
 - Searches up the element tree if not found on target element

Change Notifications

- An object must notify when one of its properties changes
 - By defining the property as a dependency property
 - Or by implementing the **INotifyPropertyChanged** interface
 - Raise the **PropertyChanged** event

The DataContext

- Sometimes many elements bind to the same object
 - Perhaps with different properties
- The object may be specified as the **DataContext** property on any common parent element
- Whenever the **Source** or **RelativeSource** properties are not specified in the **Binding**, a data context object is searched up the element hierarchy
 - If found, becomes the binding source object
- Can be used programmatically without the need to create the source object in XAML

Bind Example

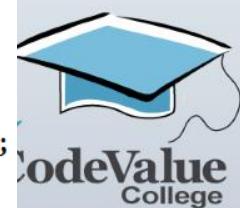
<Page

```
x:Class="DataBindingDemo.MainPage"
xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
xmlns:local="using:DataBindingDemo"
xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
mc:Ignorable="d"
FontSize="36">

<Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
    <Grid.ColumnDefinitions>
        <ColumnDefinition/>
        <ColumnDefinition/>
        <ColumnDefinition/>
    </Grid.ColumnDefinitions>
    <Grid.RowDefinitions>
        <RowDefinition Height="Auto"/>
        <RowDefinition/>
    </Grid.RowDefinitions>
    <TextBlock Grid.Row="0" Grid.ColumnSpan="3" HorizontalAlignment="Center" Text="{x:Bind Counter, Mode=OneWay}">0</TextBlock>
    <Button Grid.Row="1" Grid.Column="0" Click="OnUp" HorizontalAlignment="Center" VerticalAlignment="Top">+</Button>
    <Button Grid.Row="1" Grid.Column="1" Click="OnDown" HorizontalAlignment="Center" VerticalAlignment="Top">-</Button>
    <Button Grid.Row="1" Grid.Column="2" Click="OnReset" HorizontalAlignment="Center" VerticalAlignment="Top">R</Button>
</Grid>
</Page>
```

```
public sealed partial class MainPage : INotifyPropertyChanged
{
    private int _counter;
    1 reference
    public MainPage()
    {
        InitializeComponent();
        DataContext = this;
    }
    8 references
    public int Counter
    {
        get { return _counter; }
        set
        {
            if (_counter == value)
                return;
            _counter = value;
            OnPropertyChanged();
        }
    }
    1 reference
    private void OnUp(object sender, RoutedEventArgs e) { ++Counter; }
    1 reference
    private void OnDown(object sender, RoutedEventArgs e) { --Counter; }
    1 reference
    private void OnReset(object sender, RoutedEventArgs e) { Counter = 0; }

    public event PropertyChangedEventHandler PropertyChanged;
    1 reference
    private void OnPropertyChanged([CallerMemberName] string propertyName = null)
    {
        PropertyChanged?.Invoke(this, new PropertyChangedEventArgs(propertyName));
    }
}
```



Data Templates

- A data template is a piece of UI that describes how to display a source object
- Various elements have properties of type **DataTemplate** just for this
 - **ContentControl** has a **ContentTemplate** property
 - **ItemsControl** has a **ItemTemplate** property
 - **HeaderedContentControl** and **HeaderedItemsControl** have a **HeaderTemplate** property
- With data binding, the source object is automatically set to the current object that is being rendered

Value Converters

- A value converter can completely alter the way the source is interpreted into the target
- Often used to match source and target that are of incompatible types
 - E.g. show a red background when the price of a book is greater than 50
- A value converter implements the **IValueConverter** interface
- Create an instance of the converter as a resource and use in a binding expression

```
public interface IValueConverter {  
    object Convert(object value, Type targetType, object parameter, CultureInfo culture);  
    object ConvertBack(object value, Type targetType, object parameter,  
                      CultureInfo culture);  
}
```

Converter Parameters

- The “parameter” object is null by default
 - Can be set using the **ConverterParameter** property of the Binding object
 - Not “bindable” in itself!
- The culture info supplied is by default set to “en-US”
 - A type converter exists that can be used via the **ConverterCulture** property of the Binding object (e.g. “fr-CA”, “he-IL”)
- The return value can be **Binding.DoNothing**, which cancels the binding operation

Data Template Selectors

- A way to replace a data template completely
 - Cannot be done in XAML alone
- Derive a class from ***DataTemplateSelector***
- Override the **SelectTemplate** method
- Set the **ItemTemplateSelector** property on the **ItemsControl** element to an instance of the custom template selector

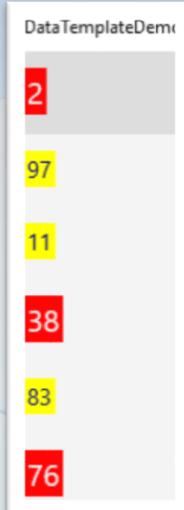
Data Template Selector Example

```
<Application.Resources>
    <local:AlternateTemplateSelector x:Key="Selector" EvenTemplate="evenTemplate" OddTemplate="oddTemplate" />
    <DataTemplate x:Key="evenTemplate">
        <Border Background="Red">
            <TextBlock Margin="2" Foreground="White" FontSize="20" Text="{Binding}"/>
        </Border>
    </DataTemplate>
    <DataTemplate x:Key="oddTemplate">
        <Border Background="Yellow">
            <TextBlock Margin="2" Foreground="DarkBlue" FontSize="15" Text="{Binding}"/>
        </Border>
    </DataTemplate>
</Application.Resources>
```

```
<Grid Background="{StaticResource ApplicationPageBackgroundBrush}">
    <ListBox Name="_list" ItemTemplateSelector="{StaticResource Selector}"
        HorizontalContentAlignment="Stretch"/>
</Grid>
```

```
public class AlternateTemplateSelector : DataTemplateSelector
{
    public string OddTemplate { get; set; }
    public string EvenTemplate { get; set; }
    protected override DataTemplate SelectTemplateCore(object item, DependencyObject container)
    {
        return (DataTemplate)(Application.Current.Resources[
            (int)item % 2 == 0 ? EvenTemplate : OddTemplate]);
    }
}
```

```
public MainPage()
{
    InitializeComponent();
    var data = new List<int>();
    Random r = new Random();
    for (int i = 1; i < 100; i++)
        data.Add(r.Next(1, 100));
    _list.ItemsSource = data;
}
```

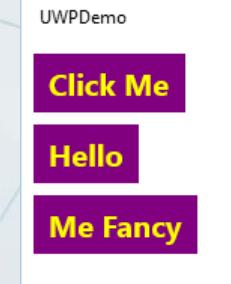


Styles

- A style (instance of the **Windows.UI.Xaml.Style** class) is a collection of dependency property setters (and triggers)
- Usually defined as a resource
- Can be applied to any element with the **Style** property
- Any specific property changed by the element that conflicts with the style takes precedence over the style setting

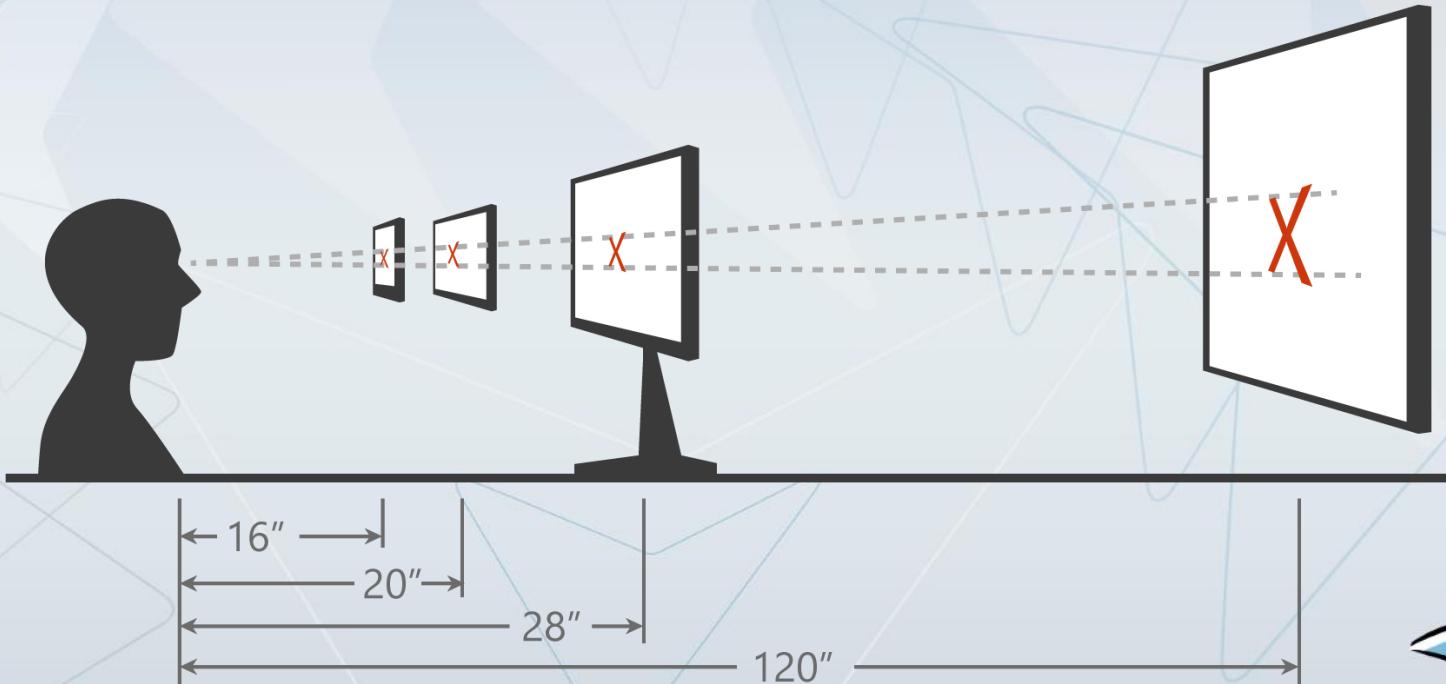
Style Example

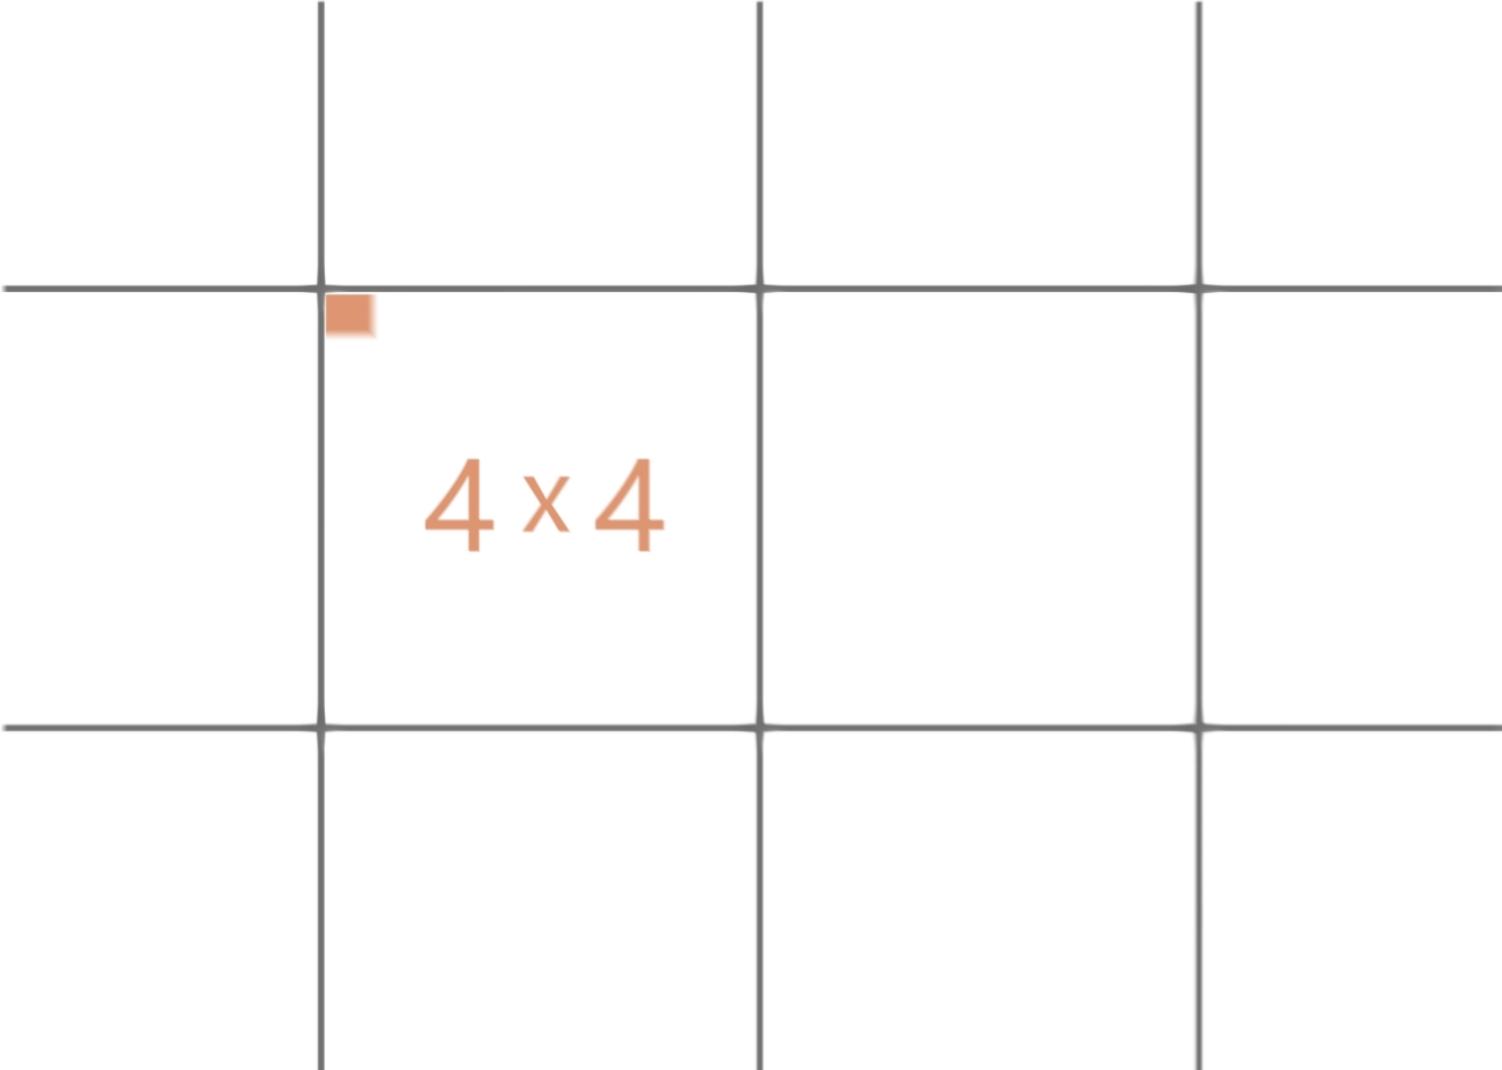
```
<Page x:Class="UWPDemo.MainPage"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    >
    <Page.Resources>
        <Style x:Key="FancyButton" TargetType="Button">
            <Setter Property="Button.Background" Value="Purple" />
            <Setter Property="Button.FontSize" Value="20" />
            <Setter Property="Button.FontWeight" Value="Bold" />
            <Setter Property="Button.Foreground" Value="Yellow" />
            <Setter Property="Button.Margin" Value="4" />
        </Style>
    </Page.Resources>
    <StackPanel Margin="4">
        <Button Content="Click Me" Style="{StaticResource FancyButton}" />
        <Button Content="Hello" Style="{StaticResource FancyButton}" />
        <Button Content="Me Fancy" Style="{StaticResource FancyButton}" />
    </StackPanel>
</Page>
```



ADAPTIVE UI

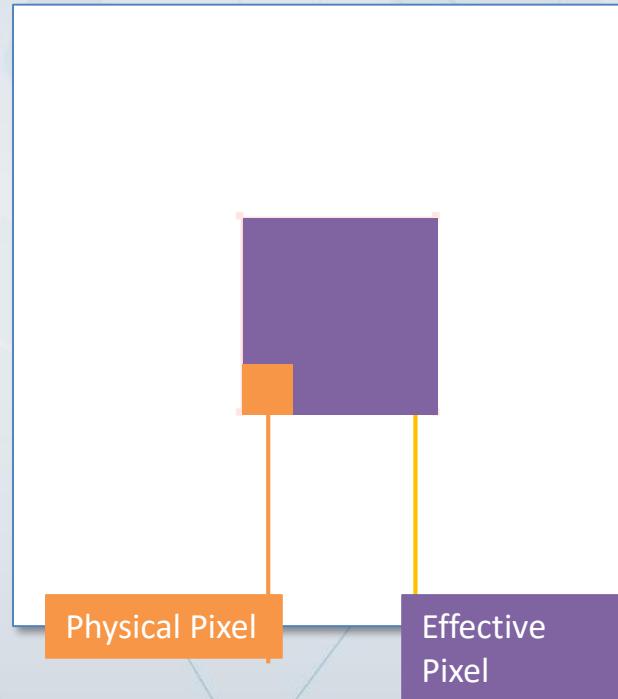
Scaling algorithm





4×4

Effective pixel



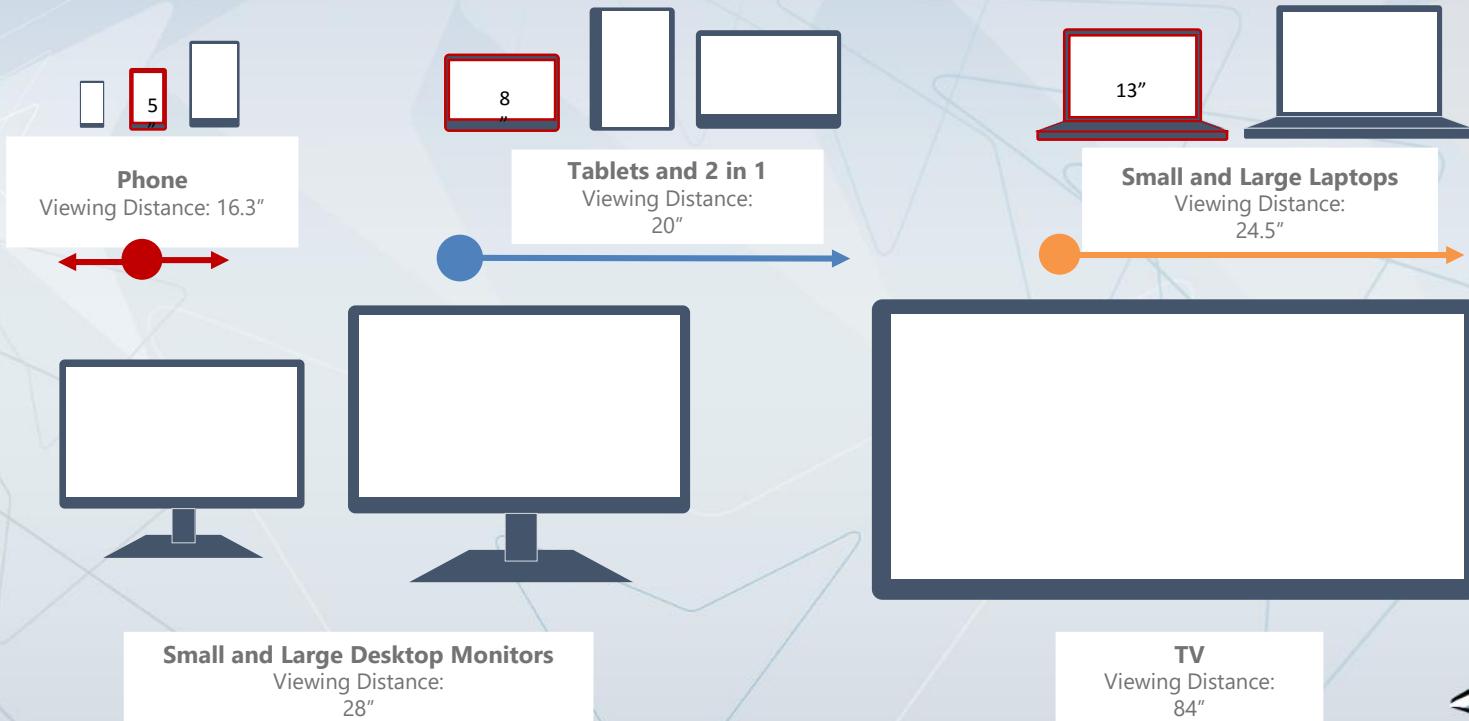
NDA Confidential

Ignore scale, resolution, & dpi
Design in Effective Pixels

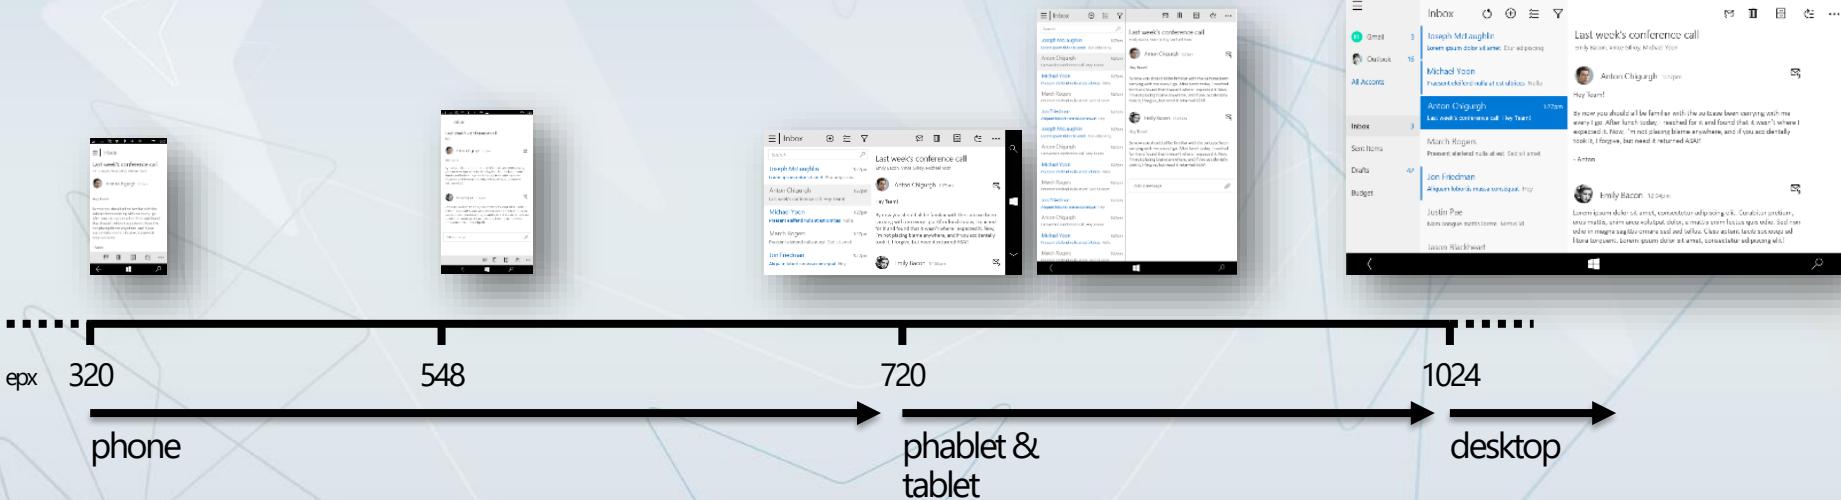
XAML is already in Effective Pixels



Planning your design



Snap points



NDA Confidential

Design Techniques for Adaptive UI

Use standard responsive/adaptive design techniques

1

Reposition

4

Reveal

2

Resize

5

Replace

3

Reflow

6

Re-architect

Adaptive design

- Build a page that adapts to different screen sizes and orientations
 - Use Visual States and Adaptive Triggers to change layout
 - Use RelativePanel to position blocks of content relative to peers, re-positioning in different visual states

Phone (portrait)



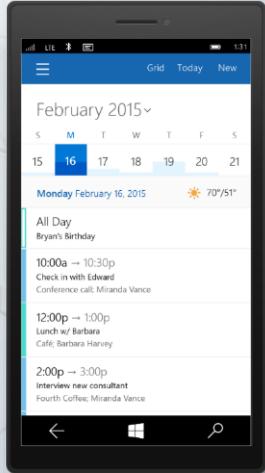
Tablet (landscape) / Desktop



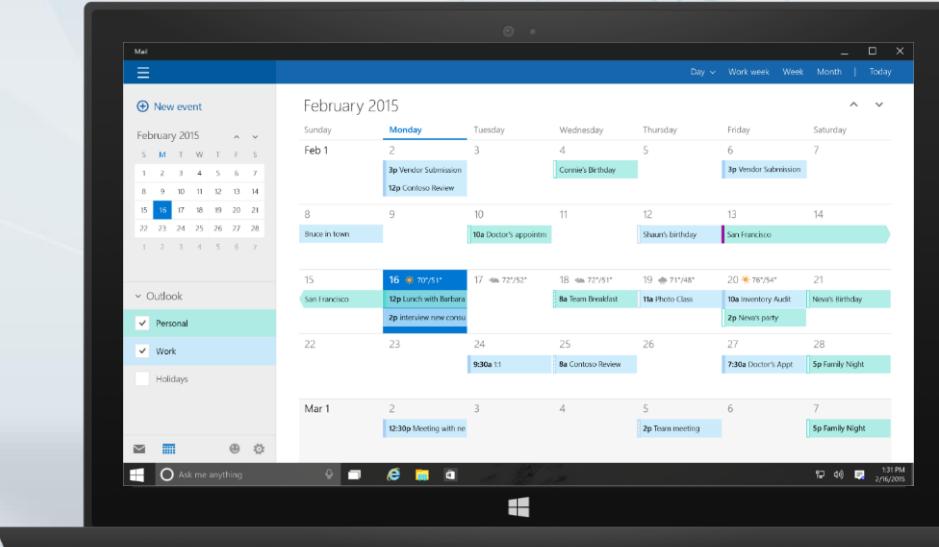
Tailored design

- Build unique experiences on different devices

Phone (portrait)



Tablet (landscape) / Desktop



Adaptive UI – The VisualStateManager

- Manages visual states and the logic for transitions between visual states for controls
- Each VisualState has a name that is representative of a UI state that can be changed by the user, or changed by control logic
- One way to handle the current VSM state is to handle the SizeChanged event and to command the VSM to GoToState
- However, VSM states can be triggered using the AdaptiveTrigger element
 - Represents a declarative rule that applies visual states based on window properties

VSM Sample

```
<Page x:Class="AdaptiveUI.MainPage"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:AdaptiveUI"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d">
    <Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
        <VisualStateManager.VisualStateGroups>
            <VisualStateGroup x:Name="VisualStateGroup">
                <VisualState x:Name="VisualStateMin320">
                    <VisualState.StateTriggers>
                        <AdaptiveTrigger MinWindowWidth="320"/>
                    </VisualState.StateTriggers>
                </VisualState>
            </VisualStateGroup>
        </VisualStateManager.VisualStateGroups>
    </Grid>
</Page>
```

VSM Sample

```
<VisualState.Setters>
    <Setter Target="Hero.(Grid.ColumnSpan)" Value="2" />
    <Setter Target="Hero.(Grid.RowSpan)" Value="1" />
    <Setter Target="Metadata.(Grid.Column)" Value="0" />
    <Setter Target="Metadata.(Grid.Row)" Value="1" />
    <Setter Target="Metadata.(Grid.ColumnSpan)" Value="2" />
    <Setter Target="Metadata.(Grid.RowSpan)" Value="1" />
    <Setter Target="Description.(RelativePanel.RightOf)" Value="" />
    <Setter Target="Description.(RelativePanel.Below)" Value="Username" />
    <Setter Target="Description.Margin " Value="0,12,0,0" />
    <Setter Target="ImageName.FontSize" Value="20" />
</VisualState.Setters>
</VisualState>
```

VSM Sample

```
<VisualState x:Name="VisualStateMin548">
    <VisualState.StateTriggers>
        <AdaptiveTrigger MinWindowWidth="548"/>
    </VisualState.StateTriggers>
    <VisualState.Setters>
        <Setter Target="Hero.(Grid.ColumnSpan)" Value="2" />
        <Setter Target="Hero.(Grid.RowSpan)" Value="1" />
        <Setter Target="Metadata.(Grid.Column)" Value="0" />
        <Setter Target="Metadata.(Grid.Row)" Value="1" />
        <Setter Target="Metadata.(Grid.ColumnSpan)" Value="2" />
        <Setter Target="Metadata.(Grid.RowSpan)" Value="1" />
        <Setter Target="Description.(RelativePanel.RightOf)" Value="Avatar" />
        <Setter Target="Description.(RelativePanel.Below)" Value="" />
        <Setter Target="Description.Margin " Value="12,0,0,0" />
        <Setter Target="ImageName.FontSize" Value="20" />
    </VisualState.Setters>
</VisualState>
```

VSM Sample

```
<VisualState x:Name="VisualStateMin1024">
    <VisualState.StateTriggers>
        <AdaptiveTrigger MinWindowWidth="1024"/>
    </VisualState.StateTriggers>
    <VisualState.Setters>
        <Setter Target="Hero.(Grid.ColumnSpan)" Value="1" />
        <Setter Target="Hero.(Grid.RowSpan)" Value="2" />
        <Setter Target="Metadata.(Grid.Column)" Value="1" />
        <Setter Target="Metadata.(Grid.Row)" Value="0" />
        <Setter Target="Metadata.(Grid.ColumnSpan)" Value="1" />
        <Setter Target="Metadata.(Grid.RowSpan)" Value="2" />
        <Setter Target="LeftCol.Width" Value="2*" />
        <Setter Target="RightCol.Width" Value="1*" />
        <Setter Target="Description.(RelativePanel.RightOf)" Value="Avatar" />
        <Setter Target="Description.(RelativePanel.Below)" Value="" />
        <Setter Target="Description.Margin " Value="12,0,0,0" />
        <Setter Target="ImageName.FontSize" Value="24" />
    </VisualState.Setters>
</VisualState>
```

VSM Sample

```
</VisualStateGroup>
</VisualStateManager.VisualStateGroups>
<Grid.RowDefinitions>
    <RowDefinition />
    <RowDefinition />
</Grid.RowDefinitions>
<Grid.ColumnDefinitions>
    <ColumnDefinition x:Name="LeftCol" />
    <ColumnDefinition x:Name="RightCol" />
</Grid.ColumnDefinitions>
<Image x:Name="Hero" Grid.Column="0" Source="Assets/airtime.jpg" Stretch="UniformToFill"
    HorizontalAlignment="Center" VerticalAlignment="Center" />
<ScrollViewer x:Name="Metadata" VerticalScrollBarVisibility="Auto">
    <RelativePanel Grid.Column="1" Background="LightBlue" Padding="12">
        <Image x:Name="Avatar" Source="Assets/avatar.jpg" Width="100" Height="100" HorizontalAlignment="Left" />
        <TextBlock x:Name="Username" RelativePanel.Below="Avatar"
            RelativePanel.AlignHorizontalCenterWith="Avatar" Text="phutureproof" />
```

VSM Sample

```
<StackPanel x:Name="Description">
    <TextBlock x:Name="ImageName" Foreground="White" FontWeight="Light" Text="Airtime" />
    <TextBlock Text="9/15/15" />
    <TextBlock Text="Lorem ipsum dolor sit amet, consectetur adipiscing elit.
        Nulla imperdiet pulvinar nunc. In et gravida ipsum.
        Morbi congue consequat ullamcorper. Integer "
        TextWrapping="WrapWholeWords" />
</StackPanel>
</RelativePanel>
</ScrollViewer>
</Grid>
</Page>
```



 phutureproof

Airtime

9/15/15

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nulla imperdiet pulvinar nunc. In et gravida ipsum. Morbi congue consequat ullamcorper. Integer ornare porta convallis. Pellentesque habitant morbi tristique senectus et netus malesuada fames ac turpis egestas. Vivamus sem nisi, ornare vel laoreet vel, accumsan facilisis tortor. Pellentesque ut nunc in leo vehicula pretium et at quam. Aliquam euismod id purus nec ultrices. Aliquam sed nisl at erat maximus finibus in sed urna. Nulla ullamcorper vehicula ex, in porta ante ullamcorper id. Phasellus a enim vitae odio ultricies semper. Suspendisse fermentum, erat in sodales accumsan, lacus urna aliquam nisi, sed ultricies dolor orci quis ligula.



 phutureproof

Airtime

9/15/15

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nulla imperdiet pulvinar nunc. In et gravida ipsum. Morbi congue consequat ullamcorper. Integer ornare porta convallis. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Vivamus sem nisi, ornare vel laoreet vel, accumsan facilisis tortor. Pellentesque ut nunc in leo vehicula pretium et at quam. Aliquam euismod id purus nec ultrices. Aliquam sed nisl at erat maximus finibus in sed urna. Nulla ullamcorper vehicula ex, in porta ante ullamcorper id. Phasellus a enim vitae odio ultricies semper. Suspendisse fermentum, erat in sodales accumsan, lacus urna aliquam nisi, sed ultricies dolor orci quis ligula.



 phutureproof

Airtime

9/15/15

The Application Object

- Every UWP application is represented by an instance of **Windows.UI.Xaml.Application**
 - A singleton
- The static property **Application.Current** returns that instance
- The Visual Studio wizard creates a XAML file for the application
 - Useful for application level resources
 - Also can set the **StartupUri** property to indicate which window to show on startup

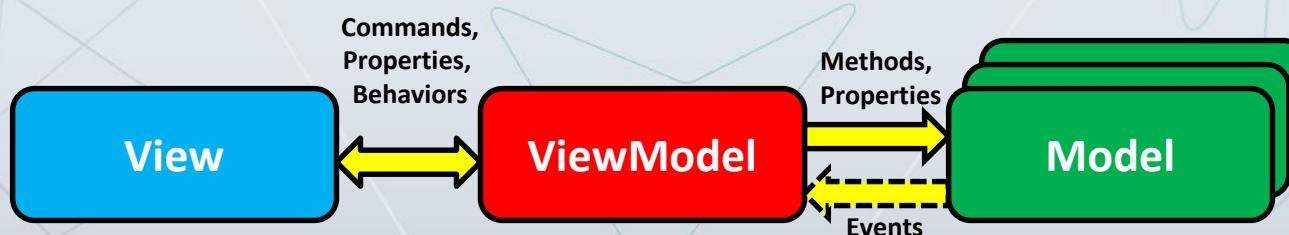
COMMANDS AND MVVM

Agenda

- Introduction to the MVVM Pattern
- Commands
- UWP Command Support
- Implementing Commands for MVVM
- Simple MVVM Framework
- Summary

The MVVM Pattern

- Model – View – ViewModel
- Based on similar principles of Model View Controller (MVC) and Model View Presenter (MVP)
- Natural pattern for XAML based applications
 - Data binding is key
- Enables developer-designer workflow
- Increases application testability



MVVM Participants

- Model
 - Business logic and data
 - May implement change notification for properties and collections
- View
 - Data display and user interactivity
 - Implemented as a **Page**, **UserControl**, **DataTemplate** or custom control
 - Has little or no code behind
- ViewModel
 - UI logic and data for the View
 - Abstracts the Model for View usage
 - Exposes commands (**ICommand**) to be used by the View
 - Implements change notifications
 - Maintains state for the View (communicates via data binding)

The View

- Provides the user interface and interaction
- The `DataContext` property points to the View Model
- Updated using property changes from the View Model
- Binds to commands (on `ICommandSource` elements) provided by the View Model

The ViewModel

- Exposes properties the View binds to
- Can be an adapter if some functionality missing from Model classes
- Exposes commands to be invoked by the view
- Maintains state for the View
- Implements change notifications (**IPropertyChanged**)
 - Uses **ObservableCollection<T>** that already implements **INotifyCollectionChanged**

The Model

- Responsible for business logic and data, e.g.
 - Data Transfer Objects (DTO)
 - POCOs (Plain Old CLR Objects)
 - Generated entity objects
 - Generated proxy objects
- May provide change notifications
- Provides validation if appropriate

Introduction to Commands

- Handling routed events and executing some code is fine for simple applications
- Sometimes the same code needs to execute from unrelated events (e.g. mouse click, keyboard shortcut, menu item, toolbar)
- Maintaining UI state (e.g. enabled/disabled) becomes difficult
- Higher level functionality, such as an undo / redo system is not possible
- Solution: use commands (the “Command” design pattern) with some support from UWP

The Command

- A command is an object implementing the **System.Windows.Input.ICommand** interface

```
public interface ICommand {  
    event EventHandler CanExecuteChanged;  
  
    bool CanExecute(object parameter);  
    void Execute(object parameter);  
}
```

Commands for MVVM

- MVVM frameworks typically provide a basic **ICommand** implementation that uses a delegate
 - Class typically called **DelegateCommand** or **RelayCommand**
- Other implementations possible
 - E.g. the **CompositeCommand** class from PRISM that holds a list of commands
- Using commands in MVVM
 - Some controls expose a **Command** and **CommandParameter** properties that can be bound to a command exposed by the ViewModel
 - **ButtonBase**, **Hyperlink** and **MenuItem** expose these
 - Technically part of the **ICommandSource** interface

Wiring the View and the View Model

- The View's `DataContext` must be set to its supporting `ViewModel`
- Some options
 - The View can create an instance of the right VM (even in XAML)
 - The ViewModel can be injected using some dependency injection technique (e.g. Unity or MEF)
 - Use some global ViewModel locator object
 - A Main VM can be set explicitly on the main View, and other VMs can be exposed as properties, which will be bound by child views

Summary

- Commands allow high level segregation of tasks
- The MVVM pattern is common in WPF to separate logic from UI and increase testability

Asynchronous Operations

- UWP, like other UI technologies, is mostly single threaded
- Synchronous operations are easy to use, but long operations may freeze the UI (“Not responding”)
 - Unacceptable user experience
- Asynchronous operations on the UI thread can be scheduled using the Dispatcher object responsible for that UI thread
 - Accessible using **Dispatcher.CurrentDispatcher** static property if on the UI thread
 - Or by calling **DispatcherObject.Dispatcher** instance property (which is inherited by all UWP elements)

Using the Dispatcher

- An operation can be scheduled using the **Invoke** or **BeginInvoke** methods
 - Accept at least a delegate and an optional **DispatcherPriority**
 - .NET 4.5 adds **InvokeAsync** variants
- **DispatcherPriority** enumeration
 - Indicates the priority to process the operation
 - Lower priorities can assume higher priority operations have already completed
 - Default is **DispatcherPriority.Normal**

Updating the UI

- When doing work on a background thread (either explicitly created or using the thread pool or task), some result may need to update UWP elements
 - Accessing the object directly will cause an exception
- Need to marshal the required operation to the UI thread using the Dispatcher
 - Call **Invoke** (for synchronous invocation) or **BeginInvoke/InvokeAsync** (for asynchronous invocation)
 - Specify a priority for the update operation (usually **DispatcherPriority.Normal**)

Async Patterns in .NET

- Asynchronous Programming Model (APM)
 - Existed since .NET 1.0
- Event Asynchronous Pattern (EAP)
 - Introduced in .NET 2.0
- Task Asynchronous Pattern (TAP)
 - New to .NET 4.5 and C# 5.0

Asynchrony with C# 5.0

- Synchronous

```
private void OnGetData(object sender, RoutedEventArgs e) {  
    _cmdGet.IsEnabled = false;  
    var wc = new WebClient();  
    _text.Text = wc.DownloadString(new Uri("http://msdn.microsoft.com"));  
    _cmdGet.IsEnabled = true;  
}
```

- Asynchronous with C# 5.0

```
private async void OnGetData(object sender, RoutedEventArgs e) {  
    _cmdGet.IsEnabled = false;  
    var wc = new WebClient();  
    _text.Text = await wc.DownloadStringTaskAsync("http://msdn.microsoft.com");  
    _cmdGet.IsEnabled = true;  
}
```

Asynchronous methods

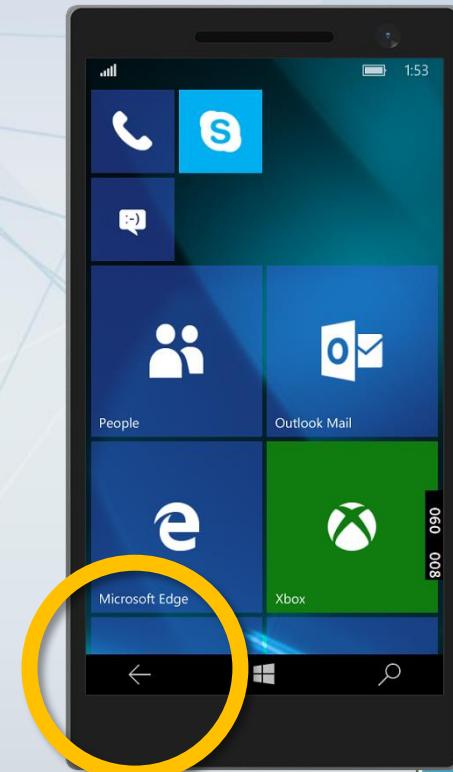
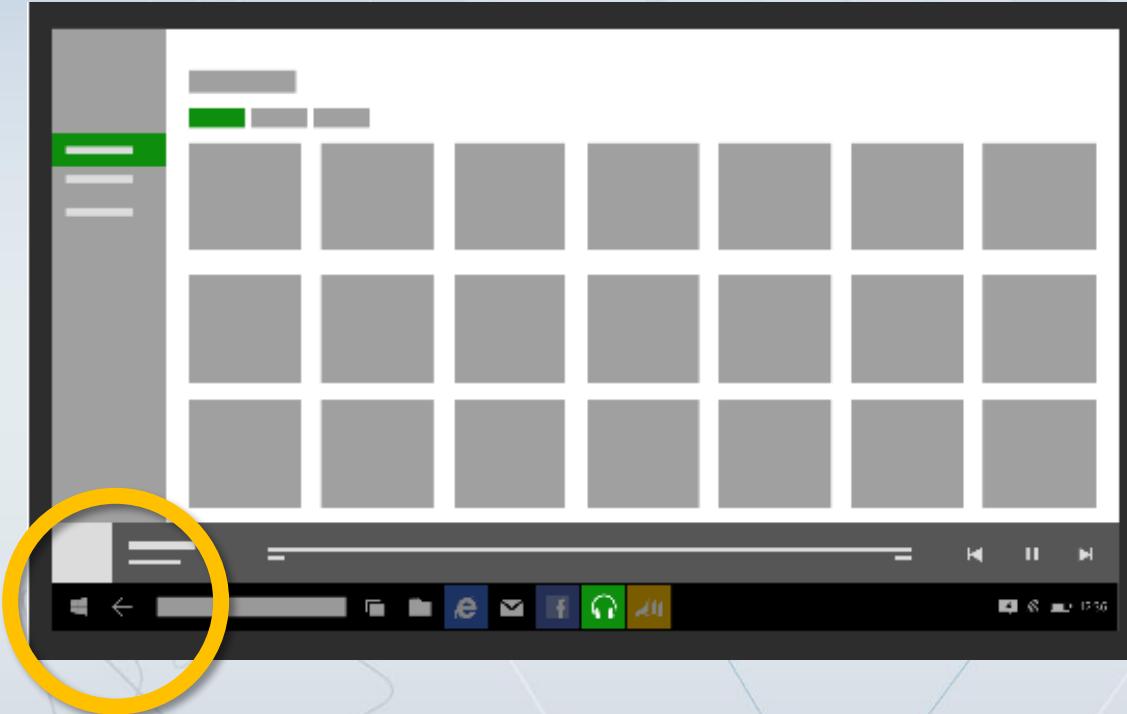
- Marked with the `async` modifier
- Must return `void`, `Task` or `Task<T>`
 - Avoid returning `void`
 - Does not allow callers to be notified of completion
- Use `await` to cooperatively yield control
- Are resumed when awaited operation completes
- Can await anything that implements the “awaiter pattern”
- Execute in the synchronization context of their caller
 - Depends on the implementing awriter
- Allow composition using regular programming constructs

HANDLING PAGE NAVIGATION

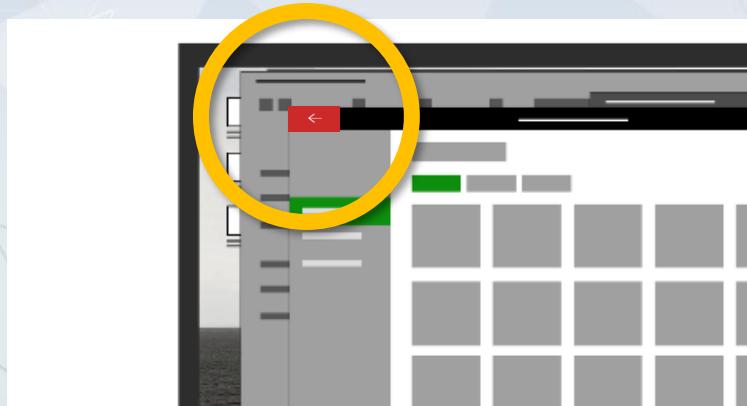
The Navigation Model

- The UWP app navigation model consists of:
 - Launching, activating, resuming, suspending, and terminating apps
 - Moving between apps
 - Moving between pages in an app
 - Moving between views within an app
- The UWP navigation model uses **Frame** and **Page** objects
 - The frame hosts the pages and keeps the navigation history
 - You can also pass data between pages as you navigate

Shell-drawn back button for Mobile and Tablet

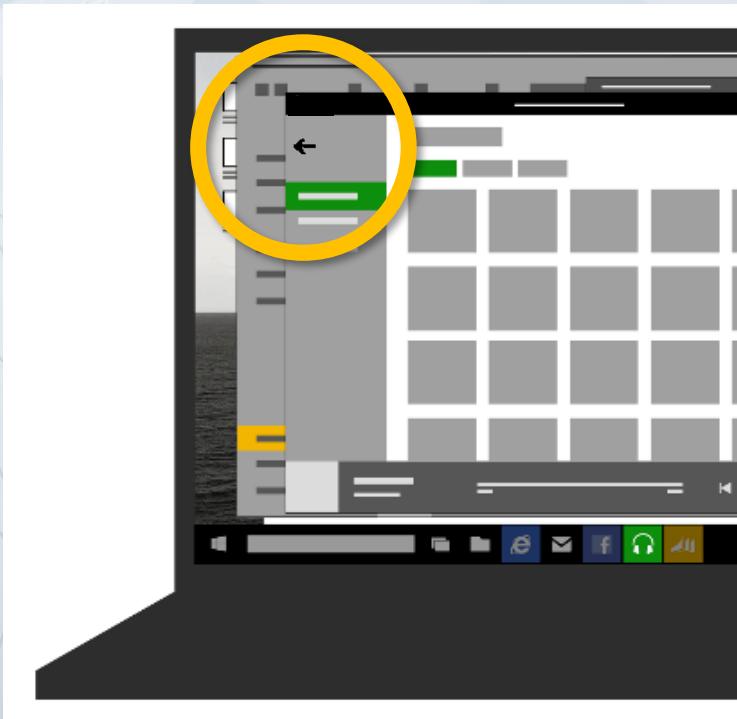


Desktop, Windowed mode: Opt-in, shell-drawn back button on Title Bar



```
if (Frame.CanGoBack)
{
    // Setting this visible is ignored on Mobile and when in tablet mode!
    Windows.UI.Core.SystemNavigationManager.GetForCurrentView().AppViewBackButtonVisibility =
        AppViewBackButtonVisibility.Visible;
}
```

Desktop, Windowed mode: Or provide your own on-canvas Back Button



Navigation in Depth

- Refer to [this](#) location for more navigation details

More Options – How to guides for UWP

Topic	Description
Accessibility	Create Universal Windows apps that are accessible to the widest possible range of device types and user audiences, including people who have impairments or disabilities.
App data and settings	Find out how and where a Universal Windows app stores both app data and user data. Learn from scenarios that describe the different storage options and the situations where each provides the best experience.
App-to-app communication	Learn how Universal Windows apps (including Windows web apps) can launch other apps and exchange data and files. Complex tasks that would normally require a user to manage multiple apps can now be handled seamlessly.
Audio, video, and camera	Capture photos and videos from a capture device, such as a webcam, and render audio streams in an app.
Contacts and calendar	Let your users access their Windows contacts and calendar appointments from your app, so they can share content, email, and calendar info, or send messages without having to switch between Windows apps.
Controls, layouts, and text	Learn about user interface components, like text, buttons, lists, and windows that are available and the various options and layouts available to support a fully-interactive user experience.
Custom user interactions	Find out about the user interaction platform, the input sources (including touch, touchpad, mouse, pen/stylus, and keyboard), modes (touch keyboard, mouse wheel, pen eraser, and so on), and user interactions supported by Universal Windows apps.
Data binding	Synchronize the UI elements of your Universal Windows app with different sources of data, including databases, files, and internal objects, to provide a data-driven user experience.
Debugging, testing, and performance	Learn about the testing and debugging cycle and how to use the related tools provided with Microsoft Visual Studio or as separate downloads. Make sure your Universal Windows app delivers the experience you intend and is ready for publication to the Windows Store.

More Options – How to guides for UWP

Topic	Description
<u>Devices, sensors, and power</u>	Integrate different devices like printers, cameras, and sensors into your Universal Windows app to provide a robust and flexible connected-device experience for your users.
<u>Files, folders, and libraries</u>	Learn how to read and write text and other data formats in files, and manage files and folders. Also find info about reading and writing app settings, about file and folder pickers, and about special, "sandboxed" locations such as the Video/Music library.
<u>Games</u>	Understand the basics of creating games on the new Universal Windows Platform (UWP), and how to provide visuals that match the ambition of your game with Microsoft DirectX 12.
<u>Globalization and localization</u>	Adapt your Universal Windows app for additional languages, markets, cultures, and regions. This info guides you through a set of best practices for development, and points you to more details about how to prepare your app for international markets.
<u>Graphics and animation</u>	Enhance your Universal Windows app with UI graphics and animations that keep users visually engaged and interested in the user experience.
<u>Launching, resuming, and background tasks</u>	Create background tasks and register for system-generated events to provide functionality even when your Universal Windows app is suspended or not running.
<u>Maps and location</u>	Learn how your Universal Windows app can tap into the Bing Maps service and produce accurate map visuals that now include aerial 3D imagery and street-level views.
<u>Monetize your app</u>	Create free apps, trials (both time-based and feature-based), paid apps, and in-app products, to give your customers the option to try your app for free and make purchase decisions during their experience with your app.
<u>Navigation</u>	Learn about the various options you have to support navigation between pages and content in your app.

More Options – How to guides for UWP

Topic	Description
Navigation	Learn about the various options you have to support navigation between pages and content in your app.
Networking and web services	Create a connected, or network-aware, Universal Windows app that can use available network connections to do things like fetch RSS feeds, engage in multiplayer games, or interact with nearby devices.
Packaging apps	Understand the app package that contains the files that constitute your Universal Windows app, and how you work with it to deploy, manage, and update your app through the Windows Store. Also learn about app capabilities, which must be declared in the app package manifest for access to specific resources.
Porting apps to Windows 10	Bring an existing app to the UWP where you can create a single app package that not only targets the Windows-based devices of your choosing, but also capitalizes on features and user experiences unique to each device type.
Security	Manage sensitive user info and help secure app data and resources while keeping the user experience intact. Features like basic password protection, roaming credentials, single sign-on, Microsoft account authentication, and cryptography are all at your disposal.
Threading and async programming	Use asynchronous programming to help your app stay responsive by allowing it to continue to run and respond to the UI while it completes other work that might take an extended amount of time.
Tiles, badges, and notifications	Design tiles (including secondary tiles and lock-screen apps), badges, and toast notifications that your Universal Windows app can use to communicate new, real-time info to users in the form of text, images, or both.
Windows Runtime components	Learn more about these self-contained objects that you can initialize and use from any language, including C#, Visual Basic, JavaScript, and C++. For example, you could create a Windows Runtime component in C++ that uses a third-party library to perform a computationally expensive operation, or simply reuse some Visual Basic or C# code in your Universal Windows app.
XAML platform	Get started with the basic concepts of the XAML programming language. Or, if you're already familiar with XAML, jump ahead and learn how to implement Windows Runtime features in XAML using Visual Studio to create a great Universal Windows app.

Appendix B

INTRODUCTION TO COM

Agenda

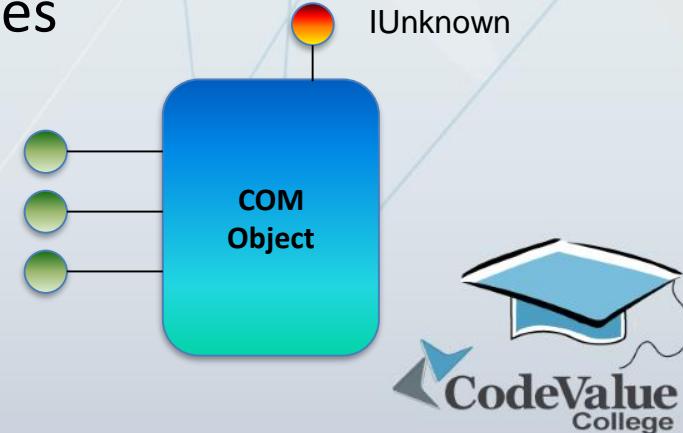
- What is the Component Object Model?
- Interfaces and IUnknown
- Other COM Entities
- The Active Template Library (ATL)
- Using ATL for Client Access
- Using ATL to Implement COM Classes
- Object Reuse
- Threading Models and Apartments
- Summary

Component Object Model

- Binary standard interface specification for objects
 - Based closely on C++ vtable mechanism for dispatch
 - Ideally language independent
 - Can use with C, C++, VB, .NET, etc.
- A set of rules and a supporting runtime for building component based systems
- Location transparency
 - Fast access to in-process objects
 - Seamless inter-process communication to out-of-process objects
- Interface based programming

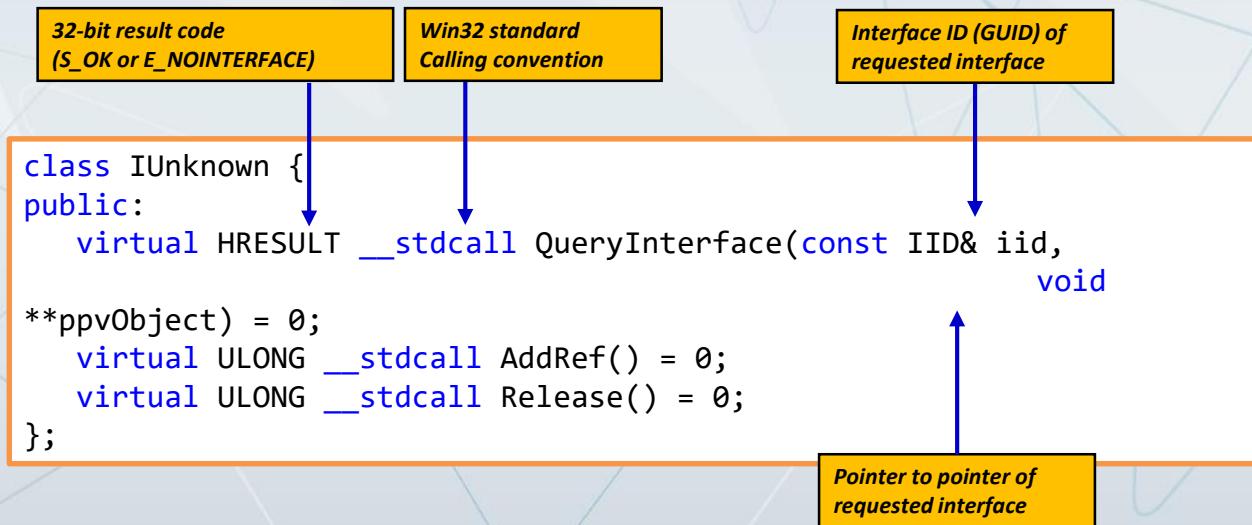
What is a COM Object?

- A COM object
 - Separates its interface(s) from its implementation
 - Provides multiple services through separate interfaces
- All COM objects must implement the **IUnknown** interface
 - Includes a `QueryInterface()` function that clients can use to acquire pointers to other interfaces
 - Also includes `AddRef()` and `Release()` functions for reference counting



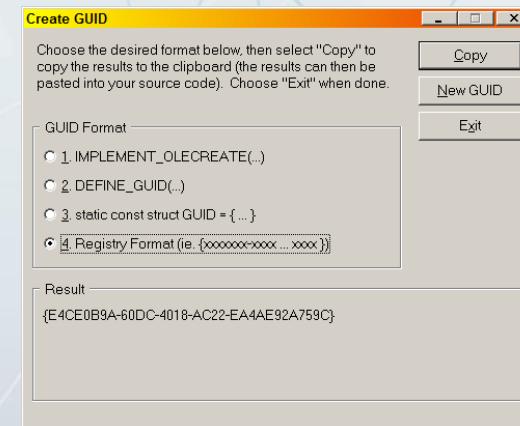
The IUnknown Interface

- Base interface
 - Must be implemented by all COM objects
 - All interfaces must extend **IUnknown**



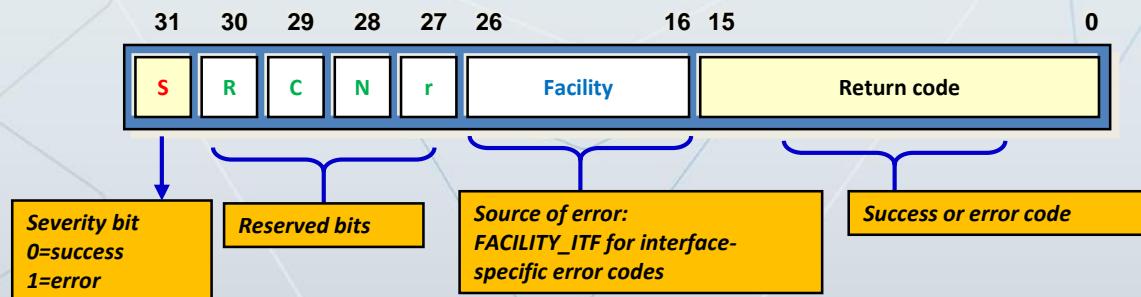
GUIDs

- Every entity in COM is identified by a GUID
 - 128 bit Global Unique Identifiers, statistically unique across space and time
- Used to identify interfaces, classes, type libraries and more
- Can generate with the Guidgen.exe tool
 - Calls **CoCreateGuid** internally
 - Several formats available



HRESULTS

- Interface functions should not throw Win32 or C++ exceptions
 - Not meaningful to all types of clients
 - Cannot cross process boundaries
- Interface methods should return an HRESULT
 - 32-bit value that contains success or error code
 - AddRef() and Release() are only exceptions to this rule



Common HRESULTs

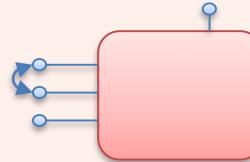
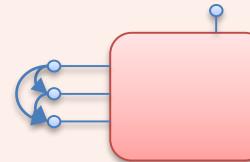
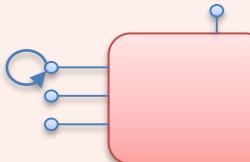
HRESULT	Meaning
S_OK	Standard success code (=0)
S_FALSE	Partial success in some sense (=1)
E_FAIL	Generic failure code
E_POINTER	Bad pointer supplied
E_UNEXPECTED	Unexpected call at this time
E_NOINTERFACE	The requested interface is not supported
E_NOTIMPL	Functionality not implemented (yet)
E_OUTOFMEMORY	Not enough memory to complete the operation
E_INVALIDARG	Argument is invalid

HRESULT Macros

Macro	Description
SUCCEEDED	Returns true if the severity bit is off
FAILED	Returns true if the severity bit is on
HRESULT_CODE	Returns the information field (e.g., error code) of the HRESULT
HRESULT_FACILITY	Returns the facility field of the HRESULT
HRESULT_SEVERITY	Returns the severity field of the HRESULT
MAKE_HRESULT	Constructs an HRESULT from individual fields
HRESULT_FROM_WIN32	Constructs an HRESULT from a Win32 error code

QueryInterface Rules

- **QueryInterface** must be

Symmetric	Transitive	Reflexive
		

- When **IUnknown** is requested, **QueryInterface** must return the same pointer value
 - Serves as the identity of the COM object

Reference Counting

- A COM object determines its own lifetime
 - Maintains a reference count
 - Incremented by a call to **AddRef()** or a successful call to **QueryInterface()**
 - Decrement by a call to **Release()**
 - When all reference counts drop to zero, deletes itself
- Clients of COM objects must follow the rules
 - Call **AddRef()** after duplicating an interface pointer
 - Call **Release()** when an interface is no longer required
 - Failure to follow these rules may lead to memory leaks or crashes



COM Servers

- One or more COM classes can be implemented as
 - A DLL (in-proc) server, or
 - An EXE (local) server
- Each class of COM object has an associated class object (class factory)
 - Usually implements the standard interface **IClassFactory**
 - Creates instance of the associated COM object
- To implement a DLL server
 - Implement the class object
 - Implement the exported DLL global functions (see later)
 - Generate a GUID to identify the COM object's class
 - Build the DLL
 - Register the COM object with the Registry (see later)

Activation APIs

- Client wants to create an instance of a COM class
 - Calls **CoGetClassObject** to get a class object (class factory)
 - Typically requesting the **IClassFactory** interface
 - Calls **IClassFactory::CreateInstance** to get an actual instance

```
HRESULT CoGetClassObject(  
    __in REFLSID rclsid,          // GUID of COM class  
    __in DWORD dwClsContext,      // CLSCTX_ALL  
    __in_opt COSERVERINFO *pServerInfo, // NULL for local machine  
    __in REFIID riid,            // class factory interface ID  
    __out LPVOID *ppv            // result  
);
```

CoGetClassObject Example

```
CoInitialize(0);

IClassFactory* pCF;
HRESULT hr = CoGetClassObject(CLSID_Account, CLSCTX_ALL, NULL,
    IID_IClassFactory, (void**)&pCF);
if(SUCCEEDED(hr)) {
    // acquired class factory, create object
    IAccount* pAcc;
    hr = pCF->CreateInstance(NULL, IID_IAccount, (void**)&pAcc);
    pCF->Release();

    if(SUCCEEDED(hr)) {
        // do something with pAcc
        pAcc->Deposit(100);
        // interface no longer needed
        pAcc->Release();
    }
}
```

Activation APIs

- The previous code works, but a bit tedious
- Most class factories implement **IClassFactory**
- Call **CoCreateInstance** as a convenience that always requests IClassFactory in its internal call to **CoGetClassObject**
 - Always uses the local machine

```
HRESULT CoCreateInstance(
    __in  REFCLSID rclsid,          // requested class ID
    __in  LPUNKNOWN pUnkOuter,      // for aggregation (see later)
    __in  DWORD dwClsContext,
    __in  REFIID riid,             // requested interface
    __out LPVOID *ppv              // result
);
```

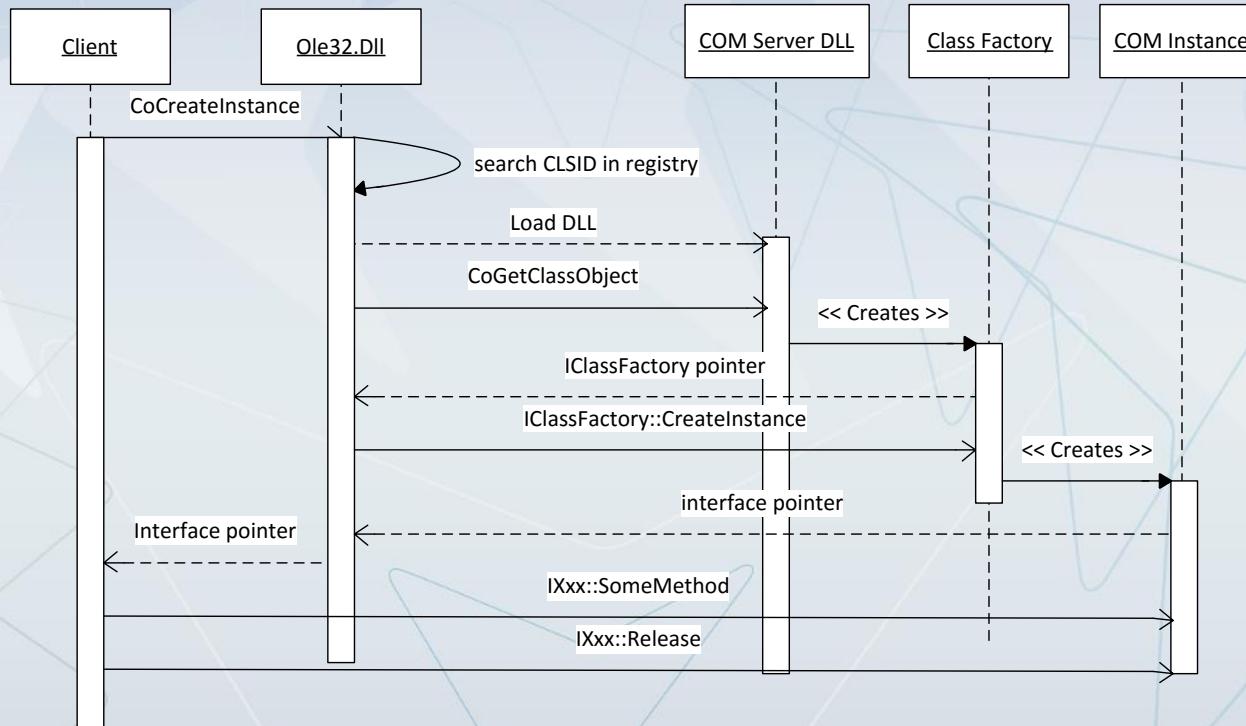
CoCreateInstance Example

```
CoInitialize(0);

IAccount* pAcc;
HRESULT hr = CoCreateInstance(CLSID_Account, NULL, CLSCTX_ALL,
    IID_IAccount, (void**)&pAcc);
if(SUCCEEDED(hr)) {
    // use account
    pAcc->Deposit(100);

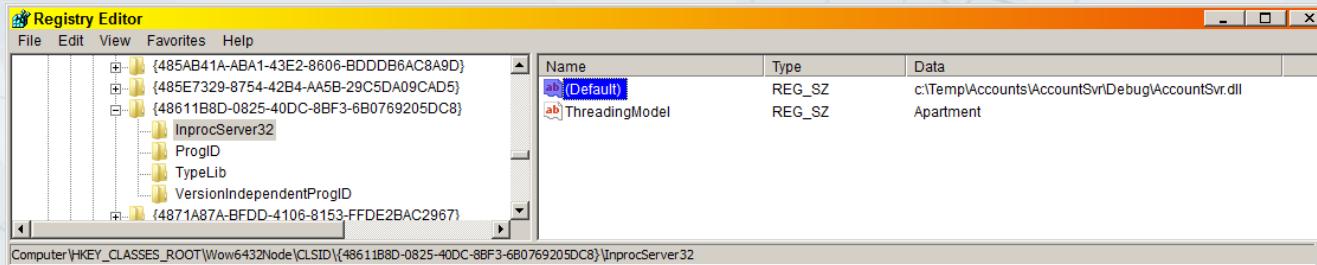
    // interface no longer needed
    pAcc->Release();
}
```

How a COM Object is Created

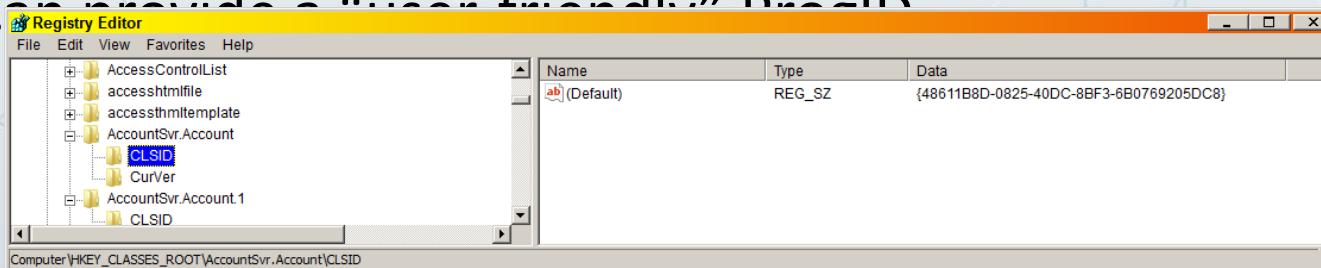


Registering the COM Class

- COM classes must be registered
 - Minimum is the CLSID to DLL/EXE mapping



- Companies provide a “user friendly” ProgID

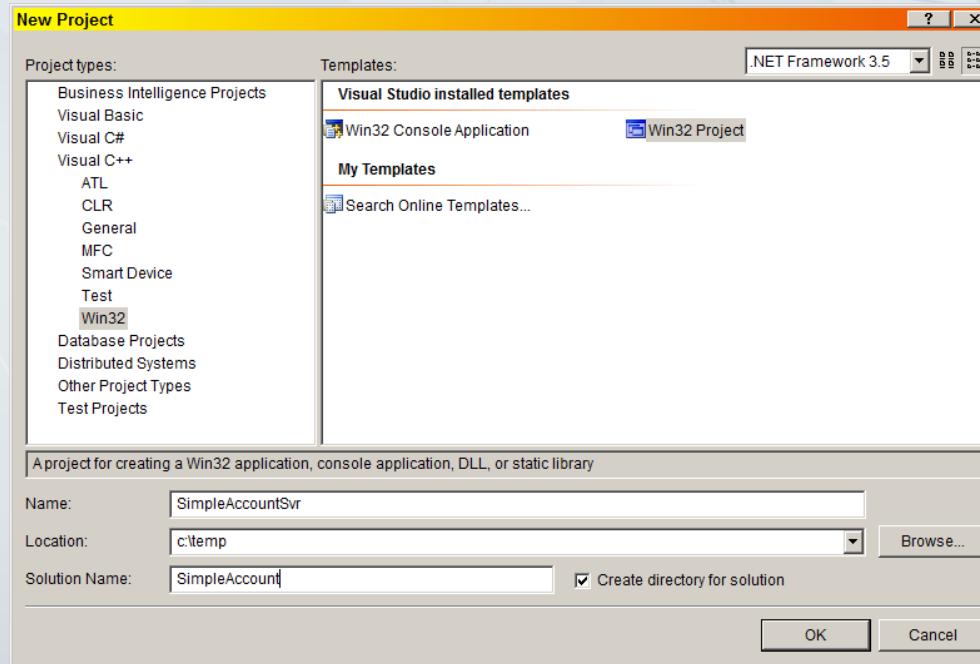


Implementing a COM Class in C++

- Create a Win32 DLL Project
- Define all relevant GUIDs
- Define required interfaces
- Implement the actual class
- Implement a Class Factory
- Implement the following global functions
 - **DIIGetClassObject** (must)
 - **DIIRegisterServer** (highly recommended)
 - **DIIUnregisterServer** (highly recommended)
 - **DIICanUnloadNow** (optional)

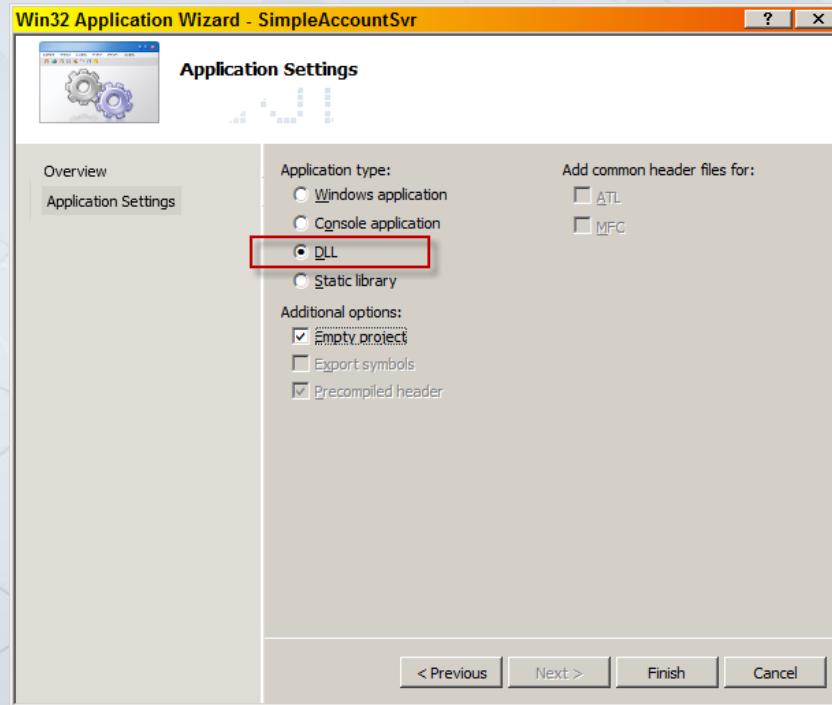
Create a Win32 DLL Project

- Open Visual Studio, Select File->New Project...
 - C++ / Win32 -> Win32 Project



Create a Win32 DLL Project

- Select an Application Type of DLL



Define Relevant GUIDs

- Use the GuidGen.Exe tool

```
// GUIDs.h

#pragma once

#include <InitGuid.h>

// {D28CD8CD-18DE-456b-B544-4D771F96B6F9}
DEFINE_GUID(CLSID_SimpleAccount,
            0xd28cd8cd, 0x18de, 0x456b, 0xb5, 0x44, 0x4d, 0x77, 0x1f,
            0x96, 0xb6, 0xf9);

// {B66AB2DF-D047-46ba-9AAD-F7BE32B1140F}
DEFINE_GUID(IID_IAccount,
            0xb66ab2df, 0xd047, 0x46ba, 0x9a, 0xad, 0xf7, 0xbe, 0x32,
            0xb1, 0x14, 0xf);
```

Define Interfaces

```
#include <Unknwn.h>

struct IAccount : IUnknown {
    STDMETHOD(Deposit)(double amount) = 0;
    STDMETHOD(Withdraw)(double amount) = 0;
    STDMETHOD(GetBalance)(double* balance) = 0;
};
```

- This is typically defined with the Interface Definition Language (IDL)

Implement the Actual COM Class

```
// SimpleAccount.h

#include "GUIDs.h"
#include "Interfaces.h"

class CSimpleAccount : public IAccount {
public:
    CSimpleAccount() : m_Balance(0), m_RefCount(1) {
    }

    // IUnknown
    STDMETHOD(QueryInterface)(REFIID riid, void** ppv);
    STDMETHOD_(ULONG, AddRef)();
    STDMETHOD_(ULONG, Release)();

    // IAccount
    STDMETHOD(Deposit)(double amount);
    STDMETHOD(Withdraw)(double amount);
    STDMETHOD(GetBalance)(double* balance);

private:
    double m_Balance;
    unsigned m_RefCount;
};
```

Implement the Actual COM Class

```
// SimpleAccount.cpp

#include "SimpleAccount.h"
#include "GUIDs.h"

// IUnknown

STDMETHODIMP CSimpleAccount::QueryInterface(REFIID riid, void** ppv) {
    if(riid == IID_IUnknown || riid == IID_IAccount) {
        AddRef();
        return *ppv = static_cast<IAccount*>(this), S_OK;
    }
    return E_NOINTERFACE;
}

STDMETHODIMP_(ULONG) CSimpleAccount::AddRef() {
    return ++m_RefCount;
}

STDMETHODIMP_(ULONG) CSimpleAccount::Release() {
    if(--m_RefCount == 0) {
        delete this;
        return 0;
    }
    return m_RefCount;
}
```

Implement the Actual COM Class

```
// IAccount

STDMETHODIMP CSimpleAccount::Deposit(double amount) {
    if(amount <= 0) return E_INVALIDARG;
    m_Balance += amount;
    return S_OK;
}

STDMETHODIMP CSimpleAccount::Withdraw(double amount) {
    if(amount <= 0) return E_INVALIDARG;
    if(m_Balance < amount) return E_UNEXPECTED;

    m_Balance -= amount;
    return S_OK;
}

STDMETHODIMP CSimpleAccount::GetBalance(double* balance) {
    if(balance == NULL) return E_POINTER;
    *balance = m_Balance;
    return S_OK;
}
```

Implement the Class Factory

```
class CSimpleAccountCF : public IClassFactory {  
public:  
    // IClassFactory methods  
    STDMETHOD(CreateInstance)(IUnknown *pUnkOuter, REFIID riid, void **ppvObject);  
    STDMETHOD(LockServer)(BOOL fLock);  
  
    // IUnknown methods  
    STDMETHOD(QueryInterface)(REFIID riid, void** ppv);  
    STDMETHOD_(ULONG, AddRef)() { return 2; }      // never dies  
    STDMETHOD_(ULONG, Release)() { return 1; }      // never dies  
};  
  
// implement IClassFactory  
  
STDMETHODIMP CSimpleAccountCF::CreateInstance(IUnknown* pOuter, REFIID riid, void** ppv) {  
    CSimpleAccount* pAcc = new CSimpleAccount;  
    HRESULT hr = pAcc->QueryInterface(riid, ppv);  
    pAcc->Release();  
    return hr;  
}  
  
STDMETHODIMP CSimpleAccountCF::LockServer(BOOL fLock) {  
    return S_OK; // don't keep track  
}
```

Implement Global Functions

```
HINSTANCE g_hInstD1l;

BOOL WINAPI DllMain(HINSTANCE hDll, DWORD reason, LPVOID) {
    switch(reason) {
        case DLL_PROCESS_ATTACH:
            g_hInstD1l = hDll;
            break;
    }
    return TRUE;
}

// maintain a singleton class factory
CSimpleAccountCF g_SimpleAccountCF;

STDAPI DllGetClassObject(REFCLSID rclsid, REFIID riid, void** ppv) {
    if(rclsid == CLSID_SimpleAccount)
        return g_SimpleAccountCF.QueryInterface(riid, ppv);
    return CLASS_E_CLASSNOTAVAILABLE;
}
```

; SimpleAccountSvr.def

LIBRARY "SimpleAccountSvr"

EXPORTS

DllGetClassObject

PRIVATE

DllRegisterServer

PRIVATE

DllUnregisterServer

PRIVATE

Implementing Global Functions

```
STDAPI DllRegisterServer() {
    HKEY hKey = NULL;
    DWORD code = RegCreateKeyExW(HKEY_CLASSES_ROOT,
        L"CLSID\\{D28CD8CD-18DE-456b-B544-4D771F96B6F9}\\InProcServer32",
        0, 0, 0, KEY_ALL_ACCESS, 0, &hKey, NULL);
    if(hKey == NULL) return HRESULT_FROM_WIN32(code);
    WCHAR path[MAX_PATH] = { 0 };
    GetModuleFileNameW(g_hInstDll, path, MAX_PATH);
    RegSetValueEx(hKey, L"", 0, REG_SZ, (const BYTE*)path,
        lstrlenW(path) * sizeof(WCHAR));
    RegCloseKey(hKey);
    return S_OK;
}

STDAPI DllUnregisterServer() {
    RegDeleteKeyW(HKEY_CLASSES_ROOT,
        L"CLSID\\{D28CD8CD-18DE-456b-B544-4D771F96B6F9}\\InProcServer32");
    return S_OK;
}
```

- Use `regsvr32.exe` with the DLL path to register
 - Calls `DllRegisterServer`
 - Use `/u` to unregister (calls `DllUnregisterServer`)

A Simple Client

```
#include "..\SimpleAccountSvr\GUIDs.h"
#include "..\SimpleAccountSvr\Interfaces.h"

void _tmain(int argc, _TCHAR* argv[]) {
    CoInitialize(0);

    IAccount* pAcc;
    HRESULT hr = CoCreateInstance(CLSID_SimpleAccount, NULL, CLSCTX_ALL,
        IID_IAccount, (void**)&pAcc);
    if(SUCCEEDED(hr)) {
        pAcc->Deposit(100);
        pAcc->Withdraw(60);
        double balance;
        pAcc->GetBalance(&balance);
        cout << "Current balance: " << balance << endl;

        // finished working with object
        pAcc->Release();
    }

    CoUninitialize();
}
```

Implementation Notes

- The “interesting” implementation is of the CSimpleAccount class
- Most other code is boilerplate and would probably be repeated in other projects
 - Class factory implementation
 - IUnknown implementation
 - Registration / unregistration code
 - DllGetClassObject implementation
- Enter the Active Template Library (ATL)

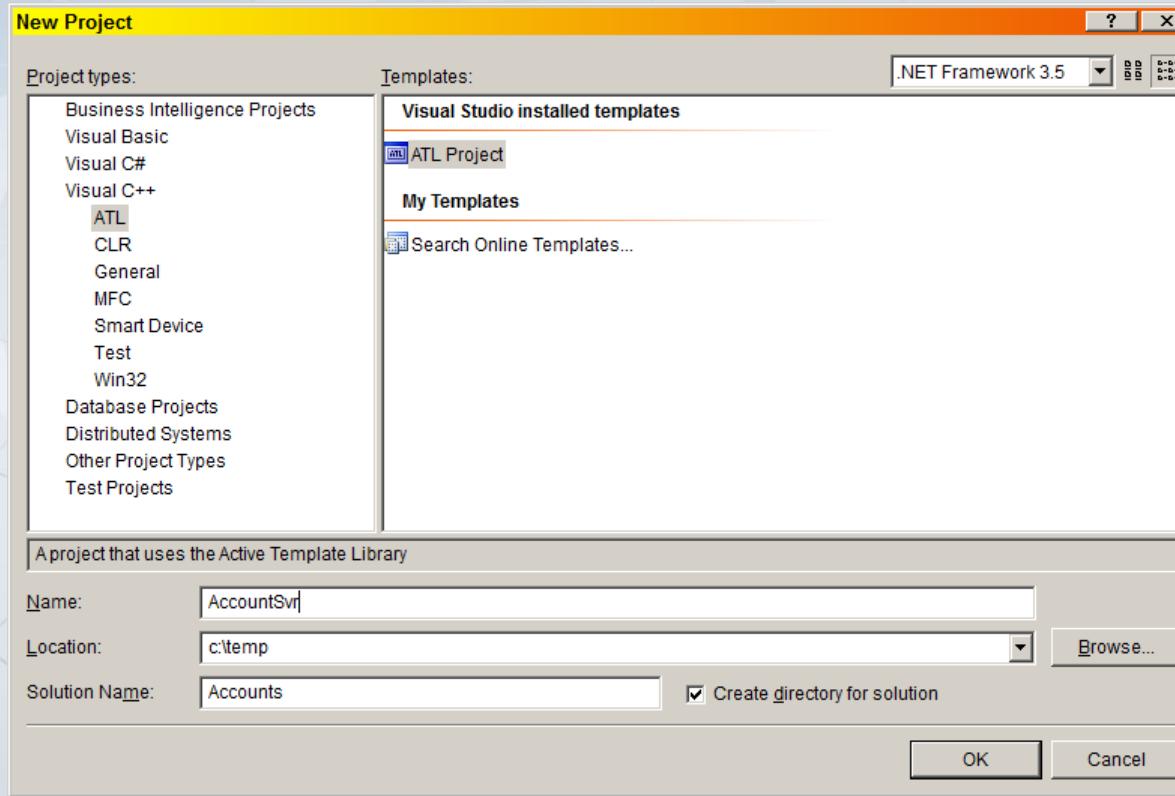
What is ATL?

- A collection of C++ template classes, macros and functions that make it easier to develop COM components
 - Highly optimized with no run time library
 - Takes care of “boilerplate” code such as class factories, IUnknown implementation and registration
- Provides client side support as well with smart pointers and other COM entity wrappers
 - BSTR, VARIANT, SAFEARRAY

Building a COM Component with ATL

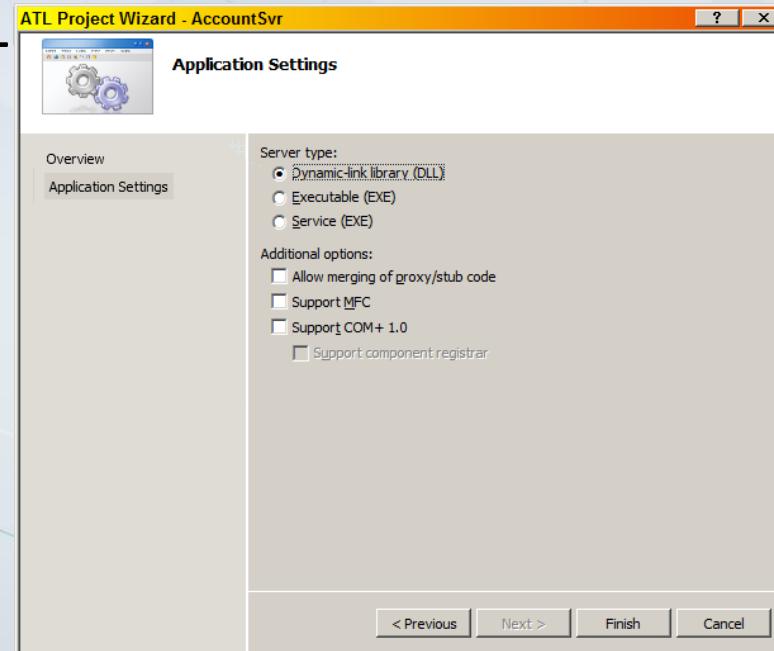
- Create an ATL project
- Add a COM class
 - Select interface type and other attributes
- Implement methods (can use a helping wizard)
- Build and you're done
 - No need to explicitly implement IUnknown
 - No need to implement a class factory
 - No need to implement registration code

Create an ATL Project



Create an ATL Project

- Select a server type
 - Typically a DL



New ATL Project Contents

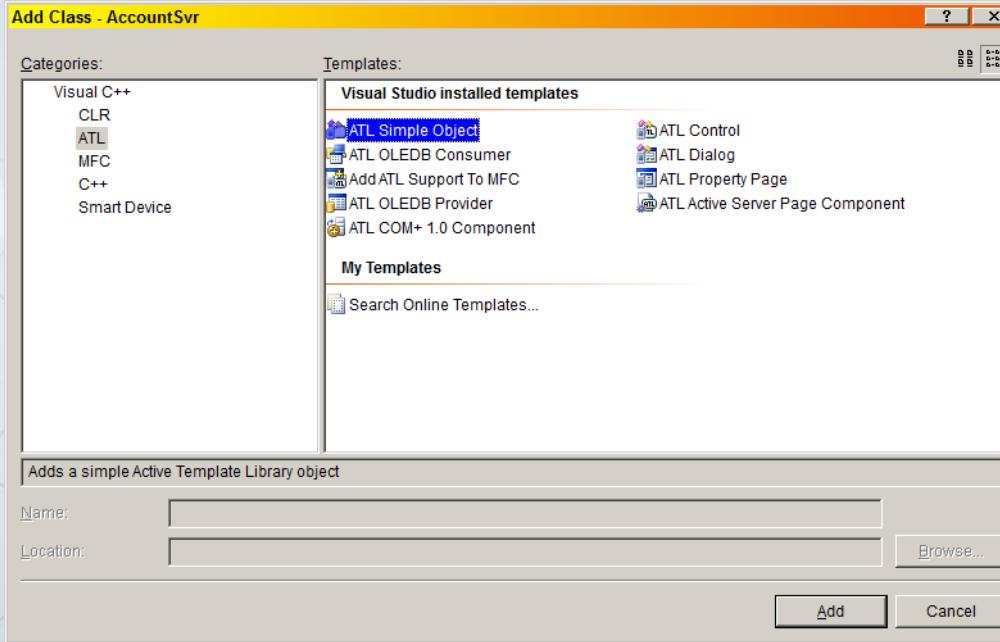
- Two projects are actually created
- One is the main project where our code resides
- The other (ending with PS) is a proxy/stub DLL
 - Code automatically generated based on interface definitions in the main project
 - Only needed if cross apartment / process communication may occur
 - Can do without it when working with the Type Library Marshalar (not covered in this course)
 - Typically unnecessary when working with Media Foundation as we'll implement MF interfaces and not define our own
- Can usually delete the PS project

Main Files in an ATL Project

- *ProjectName.IDL*
 - Interface, methods and COM class definitions (more on that later)
- *ProjectName.cpp*
 - Global functions implementation
- *dllmain.cpp*
 - DllMain implementation
- *ProjectName.def*
 - Module definition file with exports
- *ProjectName.rgs*
 - Project registry script (more on this later)

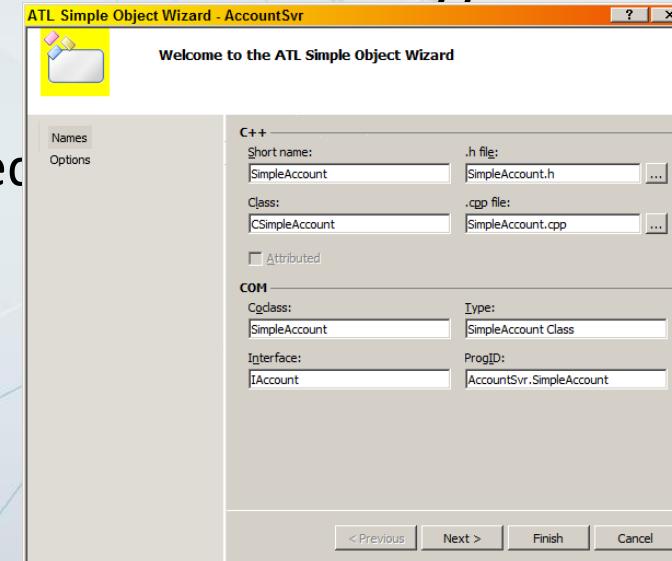
Adding a COM Class

- Right click project, select “Add Class”, select “ATL Simple Object”, then click Add



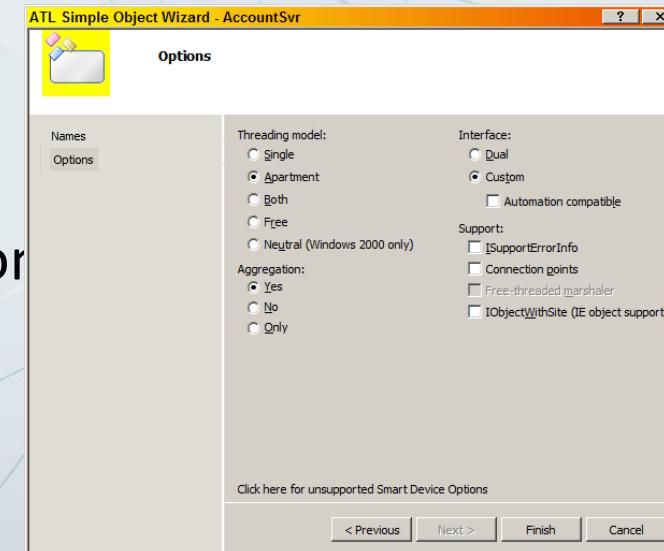
Configuring a COM Class

- Select a Simple name as a base name
- Select first interface to implement
(IUnknown is present automatically)
- C++ section
 - Affects the generated source code
- COM section
 - Reflected in the IDL and type library



Configuring a COM Class

- Custom interface inherits from **IUnknown**
- Dual interface inherits from **IDispatch**
 - Provides automation capabilities



Interface Definition Language (IDL)

- Used to define COM interfaces, methods, properties and other entities
- Has no implementation, just declarations
- Allows placements of attributes
- Compiled by the Microsoft IDL (MIDL) compiler
 - Generates appropriate header files
 - Generates a type library
- Used to generate a proxy/stub DLL (if needed)

A Typical IDL File

```
import "oaidl.idl";           ← Standard IDL imports
import "ocidl.idl";           ←

[                                ← IID
    object,
    uuid(C8084AF0-4875-4243-AFB4-1C194EC03A85),
    helpstring("IAccount Interface"),
    pointer_default(unique)
]
interface IAccount : IUnknown {
    [helpstring("method Deposit")] HRESULT Deposit([in] DOUBLE amount);
    [helpstring("method Withdraw")] HRESULT Withdraw([in] DOUBLE amount);
    [helpstring("method GetBalance")] HRESULT GetBalance([out,retval] DOUBLE* balance);
};

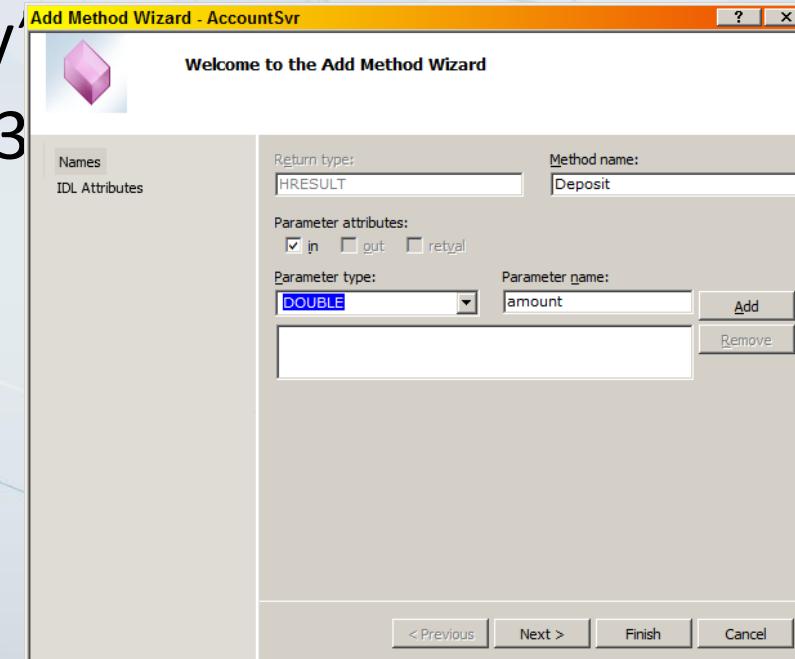
[                                ← Interface attributes
    uuid(52BAE9B-BD95-42A5-BBC5-57D220F606AC),
    version(1.0),
    helpstring("AccountSvr 1.0 Type Library")
]
library AccountSvrLib {          ← Type library declaration
    importlib("stdole2.tlb");
    [
        uuid(E29B6DC7-6A8C-4F25-9589-4AA563C27380), ← CLSID
        helpstring("SimpleAccount Class")
    ]
    coclass SimpleAccount {          ← A COM class
        [default] interface IAccount;
    };
};
```

Compiling an IDL File

- IDL files are compiled by the MIDL compiler
 - A C/C++ header file needed for server compilation (and optionally for the client too)
 - A type library file (TLB extension)
 - Provides a “universal header file” usable by any language (C++, VB, .NET languages, others)
 - Added to the final DLL/EXE as a resource in an ATL project
 - A GUIDs file
 - More files for creating a proxy/stub DLL

Adding Members to an Interface

- Right click the interface in “Class View” and select “Add Method” or “Add Property”
- Adds code in 3 locations
 - IDL file
 - H file
 - CPP file



A Typical Header File

```
class ATL_NO_VTABLE CSimpleAccount :  
public CComObjectRootEx<CComSingleThreadModel>, <-- Implements IUnknown  
public CComCoClass<CSimpleAccount, &CLSID_SimpleAccount>, <-- Define class factory  
public IAccount {  
public:  
    CSimpleAccount() {}  
  
DECLARE_REGISTRY_RESOURCEID(IDR_SIMPLEACCOUNT)  
  
BEGIN_COM_MAP(CSimpleAccount)  
    COM_INTERFACE_ENTRY(IAccount)  
END_COM_MAP()  
  
DECLARE_PROTECT_FINAL_CONSTRUCT()  
  
HRESULT FinalConstruct() {  
    return S_OK;  
}  
  
void FinalRelease() {}  
  
public:  
    STDMETHOD(Deposit)(DOUBLE amount);  
    STDMETHOD(Withdraw)(DOUBLE amount);  
    STDMETHOD(GetBalance)(DOUBLE* balance);  
};  
  
OBJECT_ENTRY_AUTO(__uuidof(SimpleAccount), CSimpleAccount) <-- Export class for external clients
```

Our custom interface

Registry script resource ID
For registering this as a COM Class in the registry

Map for implementing IUnknown::QueryInterface

Post construction

Interface methods

Export class for external clients

A Typical Registry Script File (RGS)

```
HKCR <-- HKEY_CLASSES_ROOT
{
    AccountSvr.SimpleAccount.1 = s 'SimpleAccount Class' <-- ProgID
    {
        CLSID = s '{E29B6DC7-6A8C-4F25-9589-4AA563C27380}' <-- CLSID
    }
    AccountSvr.SimpleAccount = s 'SimpleAccount Class'
    {
        CLSID = s '{E29B6DC7-6A8C-4F25-9589-4AA563C27380}'
        CurVer = s 'AccountSvr.SimpleAccount.1'
    }
    NoRemove CLSID
    {
        ForceRemove {E29B6DC7-6A8C-4F25-9589-4AA563C27380} = s 'SimpleAccount Class'
        {
            ProgID = s 'AccountSvr.SimpleAccount.1'
            VersionIndependentProgID = s 'AccountSvr.SimpleAccount'
            InprocServer32 = s '%MODULE%' <-- Path to DLL
            {
                val ThreadingModel = s 'Apartment'
            }
            'TypeLib' = s '{52BEAE9B-BD95-42A5-BBC5-57D220F606AC}' <-- Type library ID
        }
    }
}
```

ATL Smart Pointers

- Automate the acquiring, releasing and accessing of interface pointers
- **CComPtr<T>**
 - Basic smart pointer
- **CComQIPtr<T>**
 - Supports QueryInterface with a constructor or operator=
- Overload operators for convenient access to the underlying interface pointer (->, *, &)
- Destructor calls IUnknown::Release

Client Code with Smart Pointers

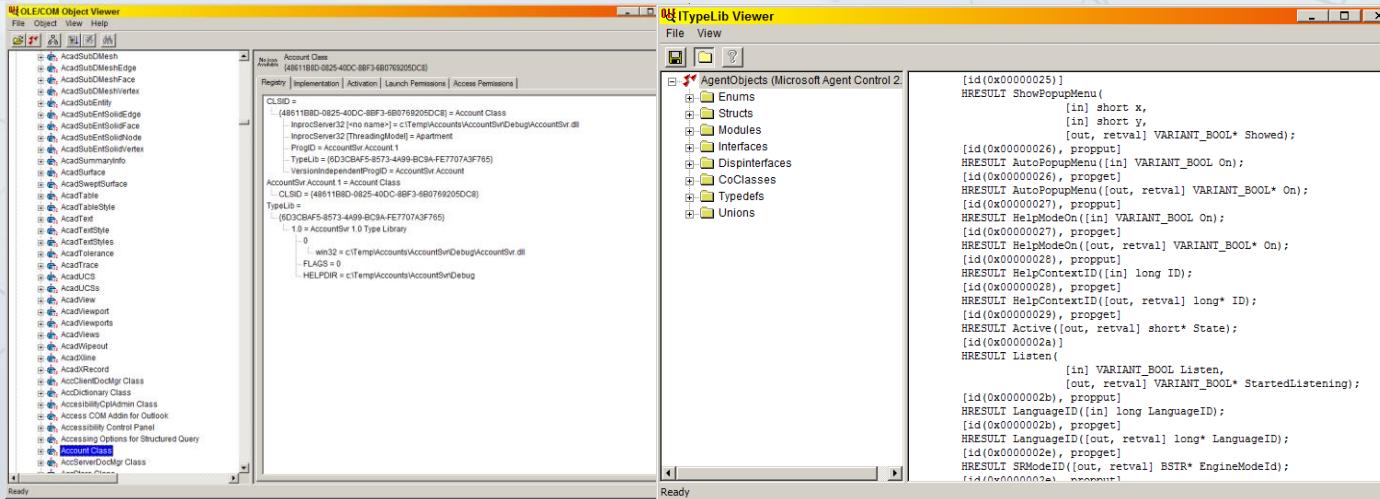
- #include <atlbase.h>

__uuidof(Account) can be used instead of CLSID_Account

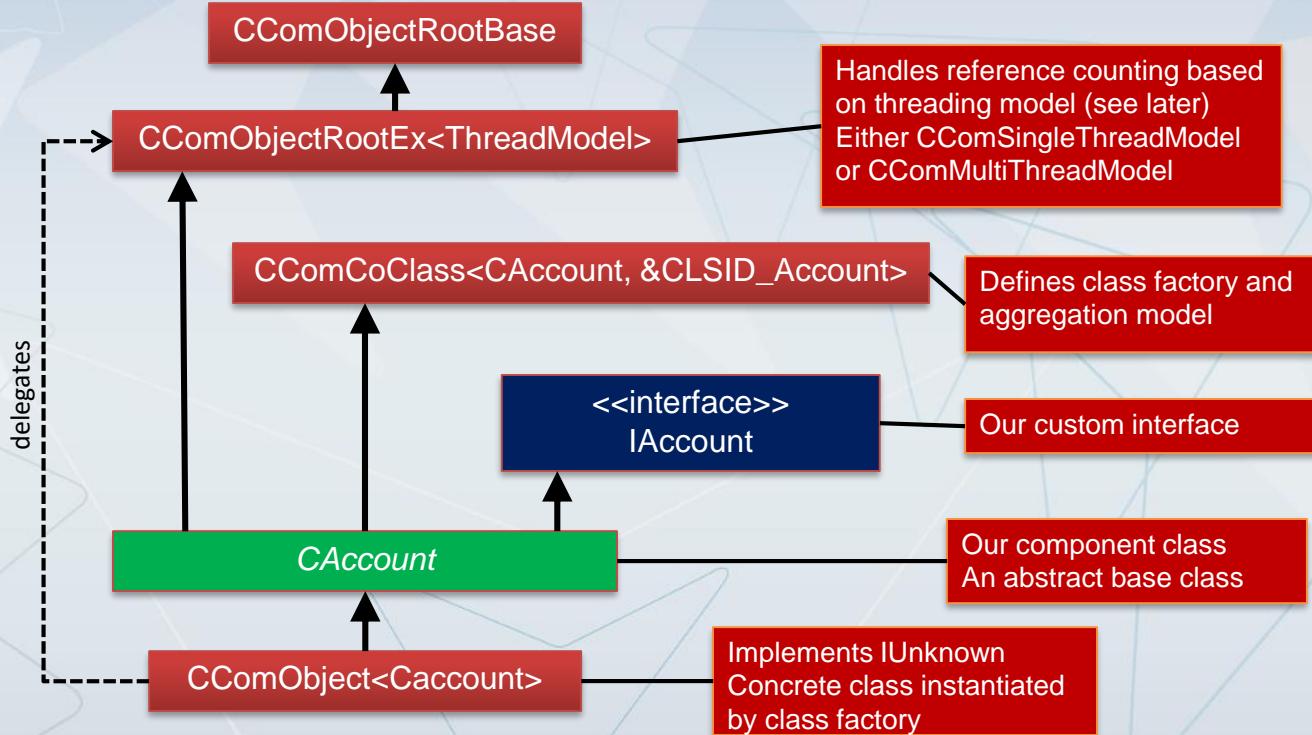
```
CComPtr<IAccount> spAcc;  
spAcc.CoCreateInstance(__uuidof(Account)); // calls CoCreateInstance  
if(spAcc) {  
    spAcc->Deposit(100);  
    spAcc->Withdraw(1000);  
    // no need to release interface pointer  
    // can explicitly release with spAcc.Release() or spAcc = NULL  
}
```

Looking At COM Components

- Can use the OLE/COM Object Viewer (oleview.exe)
 - View registry information conveniently
 - Can create instances
 - Can “decompile” type libraries back into IDL



ATL Class Relationships



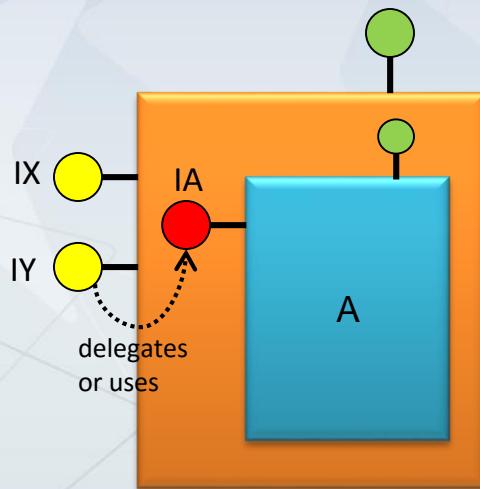
COM Object Reuse

- Traditional OOP reuse is through containment or inheritance
- Containment
 - COM supports containment easily
 - A COM object can be a client of another COM object
 - Can even implement same interface and provide delegation
- Inheritance
 - As a binary standard, COM cannot support true inheritance, as not enough common ground exists between languages
 - As an alternative, COM supports aggregation

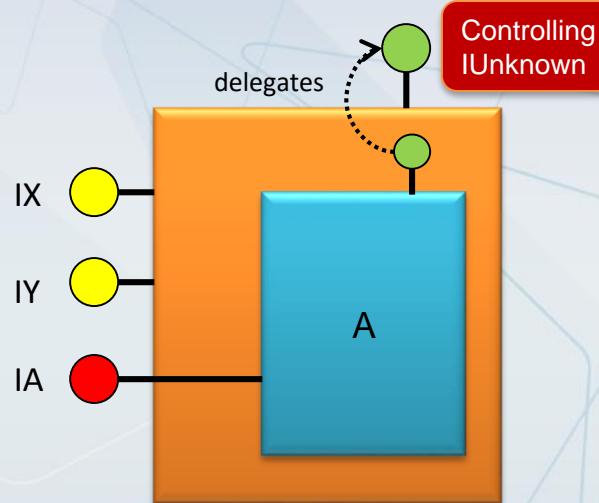


Containment vs. Aggregation

Containment



Aggregation



Aggregation Example

- The Account class aggregates a Calculator class

```
class ATL_NO_VTABLE CSimpleAccount :  
public CComObjectRootEx<CComSingleThreadModel>,  
public CComCoClass<CSimpleAccount, &CLSID_SimpleAccount>,  
public IAccount  
{  
public:  
  
DECLARE_REGISTRY_RESOURCEID(IDR_SIMPLEACCOUNT)  
BEGIN_COM_MAP(CSimpleAccount)  
    COM_INTERFACE_ENTRY(IAccount)  
    COM_INTERFACE_ENTRY_AGGREGATE(IID_ICalculator, m_spUnkCalc.p)  
END_COM_MAP()  
  
DECLARE_GET_CONTROLLING_UNKNOWN()  
DECLARE_PROTECT_FINAL_CONSTRUCT()  
  
HRESULT FinalConstruct() {  
    return m_spUnkCalc.CoCreateInstance(  
        CLSID_Calculator, GetControllingUnknown());  
}  
  
CComPtr<IUnknown> m_spUnkCalc;  
...  
};
```

```
CComPtr<IAccount> spAcc;  
spAcc.CoCreateInstance(__uuidof(SimpleAccount));  
if(spAcc) {  
    CComQIPtr<ICalculator> spCalc(spAcc);  
    ATLASSERT(spCalc);  
    // do something with spCalc...  
  
    // can go the other way around  
    CComQIPtr<IAccount> spAcc2(spCalc);  
    ATLASSERT(spAcc2);  
}
```

Processes & Threads

- Process
 - A management and containment object, providing resources to execute a program
 - Manages a private virtual address space
 - By default 2 GB (32 bit), 8 TB (64 bit)
- Thread
 - Entity scheduled by the kernel to execute code
 - Has a stack, priority, optional message queue
 - Has a scheduling state (running, ready, waiting)

COM & Threading

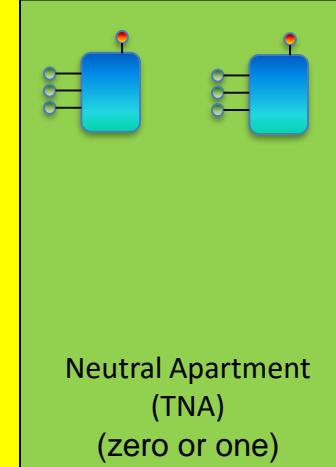
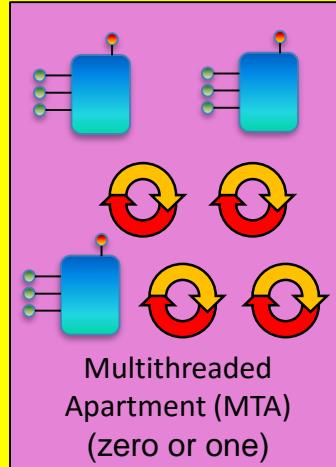
- Client code may have multiple threads
 - May create objects on many threads
 - May access an object concurrently from multiple threads
 - Object may call back to the client
- Objects may have multiple threads
 - Worker threads
 - May access object state
- Some objects are thread safe, some are not
- COM Apartments provide a solution

COM Apartments

- An “Apartment” groups objects with the same concurrency requirements
- A process may have any number of apartments
- Every object is associated with exactly one apartment
- Objects may be called by threads in their apartment only
- Cross apartment access requires a proxy

Process & Apartments

Process



- Apartments may be created explicitly, or implicitly by COM

Apartment Types

- Single Threaded Apartment (STA)
 - Only one thread can live in the apartment
 - Any number of STAs may exist in a process
 - Method calls from other apartments serialized by the thread's message queue
 - Implies thread affinity
 - Thread safety implicitly provided
- Multithreaded Apartment (MTA)
 - Any number of threads can reside there
 - Zero or one MTA per process
 - Any thread in the MTA can call any object in there
 - No thread affinity
 - No thread safety
- Thread Neutral Apartment (TNA)
 - No threads may live there
 - Zero or one TNA per process
 - No thread safety, no thread affinity
 - Mostly useful in COM+ scenarios

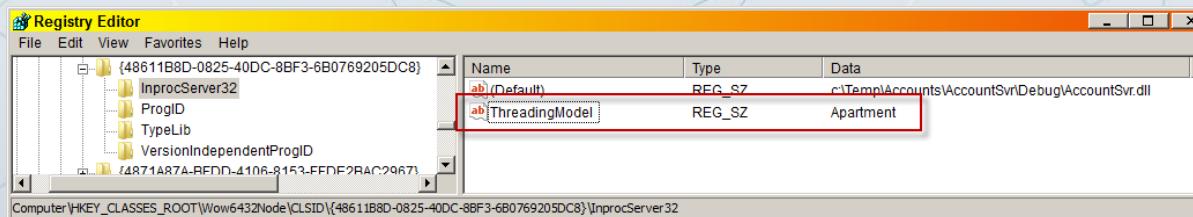
Entering an Apartment

- A thread that wants to work with COM must enter an apartment by calling **CoInitializeEx**
COINIT_APARTMENTTHREADED enters a new STA
COINIT_MULTITHREADED enters the lone MTA
`CoInitialize(0)` is equivalent to
`CoInitializeEx(0,COINIT_APARTMENTTHREADED)`
- Cannot switch apartments unless
 - Calls **CoUninitialize** to get out of its apartment

Objects & Apartments

- Objects may have their own concurrency requirements
- A COM class residing in a DLL indicates its threading model via the registry

Setting	Apartment Type	Notes
Single (or nothing)	Main STA	First STA in process (legacy setting)
Apartment	STA	
Free	MTA	
Both	STA or MTA	In the calling thread's apartment
Neutral	TNA	

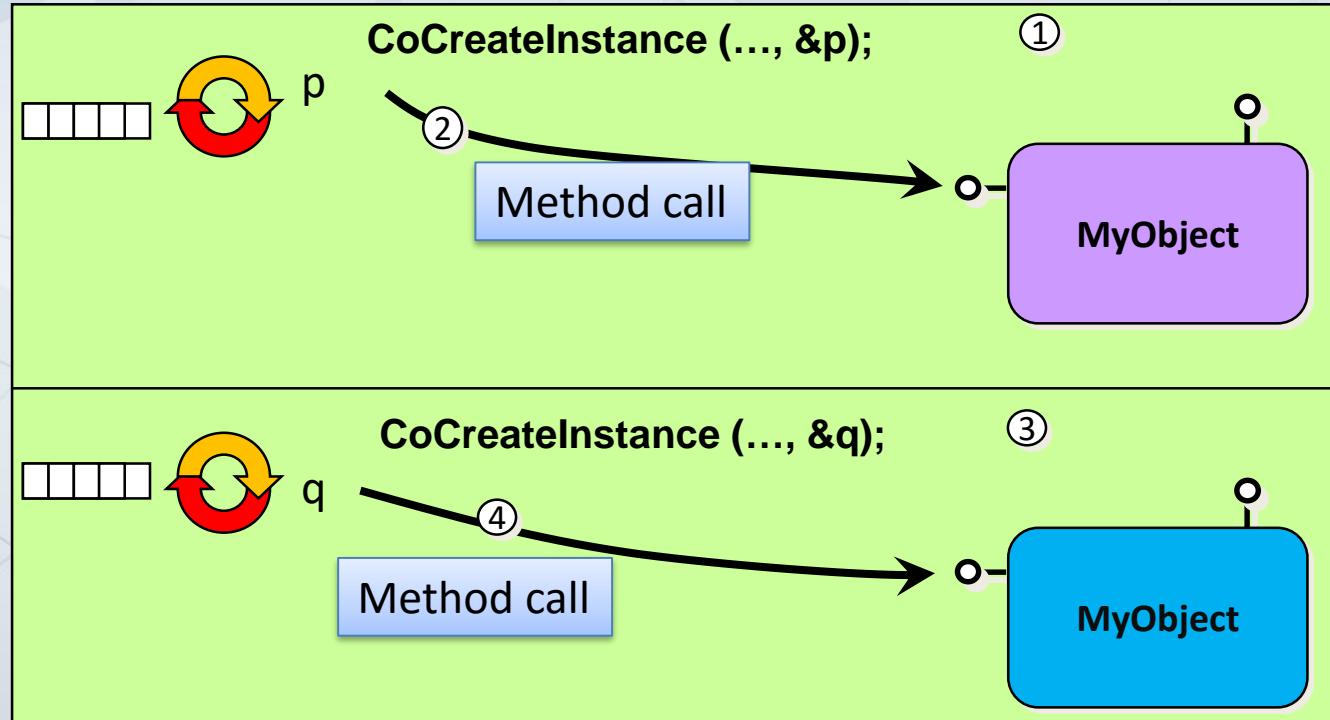


Apartments & Marshaling

- When a thread calls **CoCreateInstance**
 - If the threading model of the class is where that thread is located, the object is created in that apartment and the thread gets back a direct pointer
 - Otherwise, a proxy is created (based on the interface requested) that marshals calls correctly to the target apartment
- A **ThreadingModel=“Both”** setting means the object is always created in the creator’s apartment
- Explicit marshaling is sometimes needed

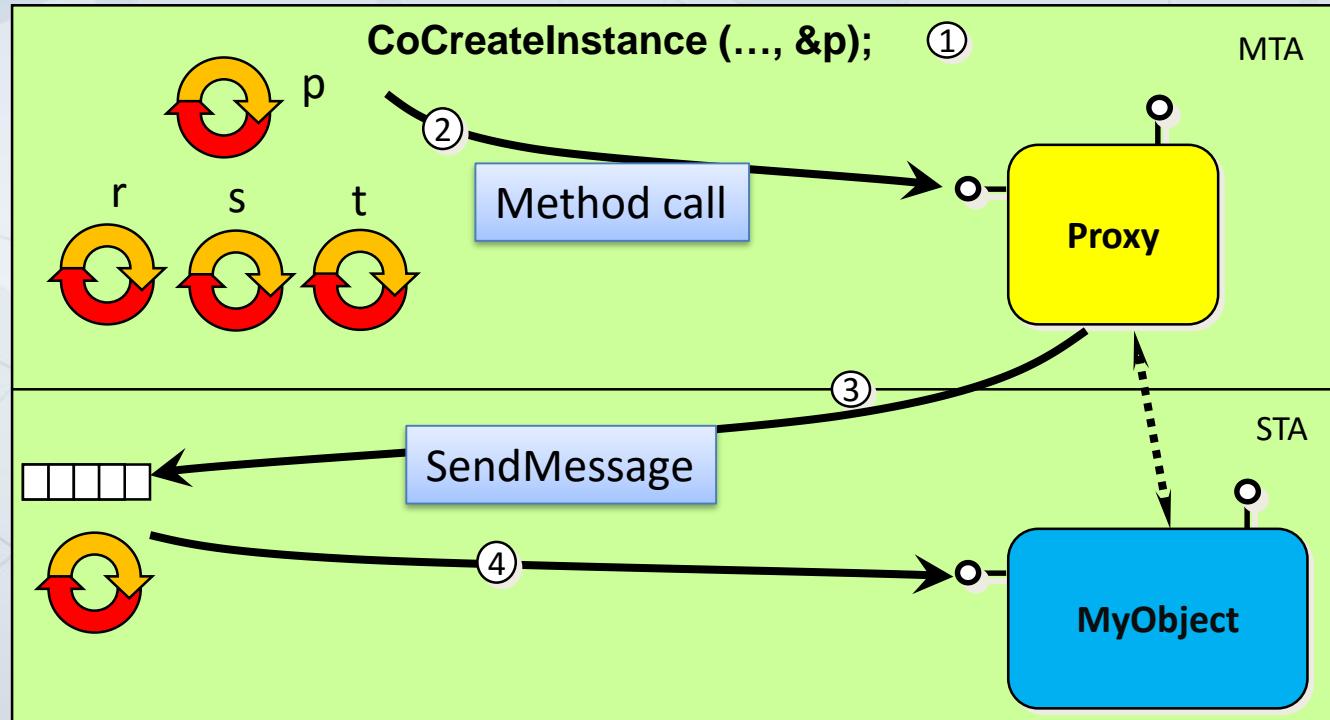
Apartments Example

- STA client / STA object



Apartments Example

- MTA client / STA object



Explicit Marshaling

- An interface pointer is marshaled by using the **CoMarshallInterface** and **CoUnmarshallInterface** functions
 - Usually happens behind the scenes
 - Won't be covered in this course
- Simpler alternatives exist
 - **CoMarshallInterThreadInterfaceInStream**,
CoGetInterfaceAndReleaseStream
 - Suitable for some scenarios
- A more convenient way to marshal interfaces is using the Global Interface Table (GIT)

The Global Interface Table (GIT)

- One per process
- Can marshal direct interface pointers or pointers to proxies
- Accessed with the **IGlobalInterfaceTable** interface
 - Class id is **CLSID_StdGlobalInterfaceTable**

```
interface IGlobalInterfaceTable : IUnknown {  
    HRESULT RegisterInterfaceInGlobal(  
        [in, iid_is(rIID)] IUnknown *pIUnk,  
        [in] REFIID riid,  
        [out] DWORD *pdwCookie );  
  
    HRESULT RevokeInterfaceFromGlobal(  
        [in] DWORD dwCookie );  
  
    HRESULT GetInterfaceFromGlobal(  
        [in] DWORD dwCookie,  
        [in] REFIID riid,  
        [out,iid_is(rIID)] void **ppv );  
}
```

Marshal an interface into the GIT
Get back a neutral cookie

Revoke registration

Unmarshal an interface into the target
apartment (proxy created if needed)

Using the GIT Example

```
CoInitialize(0);      // STA

CComPtr<IAccount> spAcc;
spAcc.CoCreateInstance(__uuidof(SimpleAccount));
if(spAcc) {
    CComPtr<IGlobalInterfaceTable> spGIT;
    spGIT.CoCreateInstance(CLSID_StdGlobalInterfaceTable);
    DWORD cookie;
    spGIT->RegisterInterfaceInGlobal(spAcc, __uuidof(IAccount), &cookie);

    CreateThread(0, 0, MyThread, (void*)cookie, 0, NULL);
}
```

```
DWORD WINAPI MyThread(PVOID param) {
    CoInitializeEx(0, COINIT_MULTITHREADED);      // MTA

    CComPtr<IAccount> spAcc;
    CComPtr<IGlobalInterfaceTable> spGIT;
    spGIT.CoCreateInstance(CLSID_StdGlobalInterfaceTable);
    spGIT->GetInterfaceFromGlobal((DWORD)param, __uuidof(IAccount), (void**)&spAcc);
    ATLASSERT(spAcc);
    // do something with spAcc...

    return 0;
}
```

The Free Threaded Marshalar (FTM)

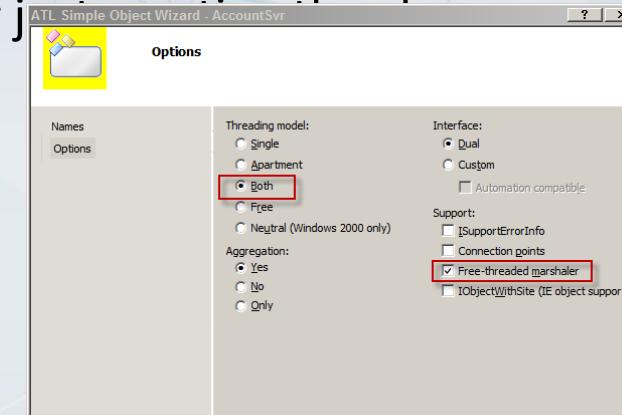
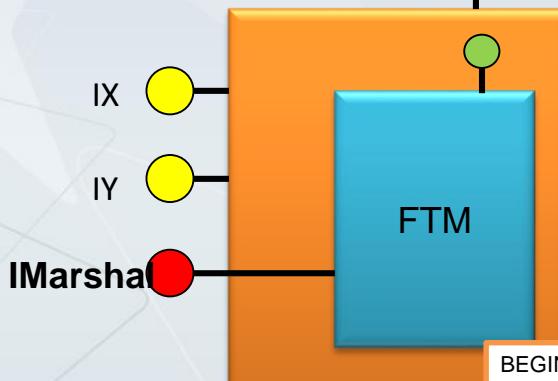
- COM classes marked “Both”
 - Can live in an STA or MTA
 - Never created with a proxy
 - But a proxy might be created if object accessed from a different apartment
- How can such an object make sure it never gets a proxy in-process?
 - Aggregate the Free Threaded Marshalar (FTM)
 - Implements the **IMarshal** interface and always returns a direct pointer when unmarshaling

Aggregating the FTM

- Can use

CoCreateFreeThreadedMarshalar

- ATL wizard can help if j



```
BEGIN_COM_MAP(CSomeClass)
...
    COM_INTERFACE_ENTRY_AGGREGATE(IID_IMarshal, m_pUnkMarshaler.p)
END_COM_MAP()

HRESULT FinalConstruct() {
    return CoCreateFreeThreadedMarshaler(
        GetControllingUnknown(), &m_pUnkMarshaler.p);
}
```

FTM Issues

- An object aggregating the FTM should not store pointers to other COM objects that do not use the FTM
 - Can lead to COM rule violation (such as the unhelpful **RPC_E_WRONGTHREAD** HRESULT)
- Solution
 - Use the GIT to get back cookies and store those
 - Cookies are apartment agnostic
 - Create an interface pointer on the fly whenever needed with the help of the GIT and the cookie

Summary

- COM is a binary standard for object communication
- Interface based programming
- Dynamic loading and location transparency
- ATL helps with building COM servers and clients
- The Apartments model helps in dealing with concurrency issues

Appendix C

INTRODUCTION TO WINDBG

WinDbg Overview

- WinDbg is a standalone GUI debugger
 - No need to install – simply xcopy
 - Used by Microsoft to debug Windows itself
 - User mode or kernel mode debugger
- UI windows
 - Command – most important window
 - Call Stack, Processes & Threads, Source, Locals, Watch, Registers, others
- Command window can do anything
 - Some shortcuts available through the GUI

WinDbg Workspaces

- WinDbg uses Workspace to keep the debugging session settings
 - This is somehow similar to VS modes
- There are several workspaces
 - A base workspace – no debugging session
 - Default user-mode workspace
 - Remote default workspace
 - Kernel-mode workspace (& Processor-specific) workspace
 - Named workspace
- For each debugging session (including dump analysis) a new workspace is created
- Workspaces are cascaded, i.e. when debugging user-mode app, the setting is accumulated.

Starting Process Debugging

- WinDbg can debug multiple processes at the same time
 - Create a new process
 - **.create <exe_path>**
 - Or **File->Open Executable...**
 - Attach to a running process
 - **.attach <process_id>**
 - Or **File->Attach to a Process...**
- By default, stops at **Initial Breakpoint**

Starting WinDbg from cmd.exe

windbg.exe <app_to_start>

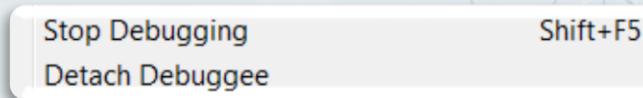
Windbg.exe -p <PID>

Windbg.exe -pn <exe_name>

Windbg.exe -z <crash_dump_file>

Stop Debugging

- Stop the debugging session by:
 - q command – quit
 - .detach
 - From the menu:



Debugging Information Windows

- Most of the WinDBG [windows](#) resemble the Visual Studio debugger main windows
 - Watch, Local, Registers, Disassembly, Process & Threads and Call Stack
 - Many of them provide more information
- The [Command Window](#) is the main debugger information window
- The [Scratch Pad](#) is like notepad
- The Command Browser is a GUI version of the Command Windows
- The Source windows display the source code

Debugger Commands

- The command prompt shows the ordinal process number : thread number
- To break long command execution use the Debug->Break menu, or **CTRL+BREAK**
- `.cls` command clears the command window text
- Debugger Command Programs
 - small application that consists of debugger commands and control flow tokens,
 - `if`, `.else`, `.elsif`, `.foreach`, `.for`, `.do`, `.break`, `.while`,
 - `...`
 - `.block { }`
 - You can execute a program from the command window or from a script

WinDbg Command Types

- Regular commands
 - Intrinsic to WinDbg
 - Have no prefix
 - Work on the debugged process
- Meta commands
 - Work on the debugger itself or the process of debugging itself
 - Prefixed with a dot (.)
- Extension commands (“bang” commands)
 - Supplied by extension (custom) DLLs
 - Prefixed with an exclamation mark (!)
 - Some extension DLLs are loaded automatically

Regular Commands

Command	Description
B[A C D E L P][<bps>]	Manipulate breakpoint(s)
D[type][<range>]	Dump memory
DT ...	Dump using type information
DV [<name>]	Dump local variables
E[type] <add> <vals>	Enter memory values
G[H N] [...] P T	Go, Step Over, Trace into
K KP KB	Stack Trace
LM[k l u v]	list modules
LN <expr>	list nearest symbols
Q	Quit
R [[<reg>[= <expr>]]]	view or set registers
S[<opts>] <range> <values>	Search memory
SX [...]	Event filter
U [<range>]	Unassemble
X [<*> module>!]<*> symbol>	View symbols

Meta (Dot) Commands

- Control the behavior of the debugger, or the flow of execution

Command	Description
.help	List all meta-command
.hh	Open the Debugger help viewer
.attach	Attach to process
.detach	Detach from process
.kill	End debugging session and close the target application
.beep	Speaker beep
.create	Create Process
.dump	Create dump file
.echotime	Display the current time
.ecxr	Display exception context
.eventlog	Display recent debug events

Extension Commands

- Extension commands are dll based debugger commands
 - They extend the debugger
 - They can be tailored to specific needs
 - You can develop your own
- Use .load to load extension Dll
- To use extension command
 - ! [module.]extension [arguments]

Some User-Mode Ext. Commands

Command	Description
!analyze	Displays information about the current exception
!help	List extension commands
!address	Displays the address space layout
!cpuid	Displays CPU version info for all CPUs
!error [err]	Display Win32 error
!exchain	Display exception chain
!for_each_frame <cmd>	Executes command for each frame in current thread
!for_each_local <cmd>	Executes command for each local variable
!gle	Displays last error & status for current thread
!obja <address>	Displays OBJECT_ATTRIBUTES[32 64]
!std_map <address>	Displays a std::map<>
![u]str <address>	Display string
!gflag	Sets or displays the global flags

Threads

Thread identifier	Description
~.	The current thread
~#	The thread that caused the current exception or debug event
~*	All threads in the process
~Number	The thread whose ordinal is Number.
~~[TID]	The thread whose thread ID is TID. (The brackets are required And you cannot add a space between the second tilde and the opening bracket.)
~[Expression]	The thread whose thread ID is the integer to which the numerical Expression resolves.

Process

Process identifier	Description
.	The current process
#	The process that caused the current exception or debug event
*	All processes
Number	The process whose ordinal is Number.
~[PID]	The process whose process ID is PID
[Expression]	The process whose process ID is the integer to which the numerical Expression resolves
Number s	Select process whose ordinal is number

Self Repeating Commands

- The j (Execute If-Else) command

```
bp `loop.cpp:143` "j (poi(i)>0n20) ''; 'gc' "
```

- The z (Execute While) command

```
r$t0 = 0
```

```
?? (wchar_t *)envp[@$t0]; r$t0 = @$t0 + 1; z (@$t0 < 10)
```

- The ~e (for each thread) command

```
~*e kb 2
```

- The !list extension command

```
!list "-t ntdll!_LIST_ENTRY.Flink -e -m 3 -x \"dd @$extret 14;  
dt ntdll!_RTL_CRITICAL_SECTION_DEBUG @$extret-0x8\"  
ntdll!RtlCriticalSectionList"
```

Aliases

- There are three kinds of aliases:
 - user-name
 - as chl kernel32!CreateHardLinkW
 - .echo chl
 - ad chl
 - as v pt;kb
 - fixed-name
 - 10 predefined aliases \$u0 - \$u9
 - automatic-aliases
 - like pseudo registers
- Use the \${ } to interpret an alias

Evaluating Expressions

- WinDbg has two different expression types
 - C++ and MASM (MASM is the default)
 - In MASM , all symbols are treated as addresses
- Use the .expr /s c++
- The ? evaluate an expression
- The ?? Evaluate C++ expression regardless .expr
- Mixing: @@c++(expr) or @@masm(expr)
- By default numbers are decimal, 010 – octal, 0x10 – hex and you can use the L, U, and I64 or any C++ suffixes
- Use the regular C++ operators

MASM Operators

- The default expression mode is MASM
- There are set of useful operators

Operator	Description
hi	Hi 16 bit
low	Low 16 bit
not	1 if the argument is zero, zero for non-zero value
by	Low-order byte
wo	Low-order word
dwo	Double word from the specified address
qwo	Quad word from the specified address
poi	Pointer sized data from the specified address
\$scmp("S1","S2")	Like strcmp (\$sicmp is like stricmp)

Reading Variable Values

```
Source: int Add(int a, int b) { int r = a + b; return r; }

>.expr /s masm
Current expression evaluator: MASM - Microsoft Assembler expressions
>dd r L1
0000000`0018f660 00000003
>r
Evaluate expression: 1635936 = 00000000`0018f660
>? dwo(r)
Evaluate expression: 3 = 00000000`00000003
>? poi(r)
Evaluate expression: -3689348818177884157 = cccccccc`00000003
>?? r
int 3
>.expr /s C++
Current expression evaluator: C++ - C++ source expressions
>? r
Evaluate expression: 3 = 00000000`00000003
>? &r
Evaluate expression: 1635936 = 00000000`0018f660
```

Macros for C++ Expressions

Macro	Return Value
#CONTAINING_RECORD(Address, Type, Field)	Returns the base address of an instance of a structure
#FIELD_OFFSET(Type, Field)	Returns the byte offset of a field in a known structure type
#RTL_CONTAINS_FIELD (Struct, Size, Field)	Indicates whether the given byte size includes the desired field
#RTL_FIELD_SIZE(Type, Field)	Returns the size of a field in a structure of known type
#RTL_NUMBER_OF(Array)	Returns the number of elements in an array
#RTL_SIZEOF_THROUGH_FIELD(Type, Field)	Returns the size of a structure of known type, up through and including a specified field



Configuring Symbols

- Debugging symbols come in several flavors
 - Full program database (PDB files)
 - Public symbols only (PDB files)
 - Exported symbols only (in the DLL itself)
- Select File->Symbol File Path...
 - Add search folders as appropriate
- Automatically uses the _NT_SYMBOL_PATH environment variable
- To get the symbols of the OS DLLs automatically, add the following string

```
SRV*c:\Windows\Symbols*http://msdl.microsoft.com/
download/symbols
```

Symbol Commands

`.reload [options] [module]`

- Reload all symbol files (useful if symbol path updated)
- Useful options
 - /f (forces loading of symbols now)

`!m [options]`

- Lists all loaded modules
- Shows in parenthesis the symbol loading status
- WinDbg uses deferred symbol loading
 - `!m v m [module]` (shows detailed info for module)

`!d <module_name>`

- Loads symbols for the specified module

Examining Symbols

- x [options] <module!symbol>
- x [options] *
 - Searches and displays all symbols matching a pattern
 - Module and symbol can contain wildcards
 - x * displays all local variables
- Options
 - /t (display data type if known)
 - /v (display type and size of symbol)
 - /s size (displays only symbols matching size in bytes)

Symbol Syntax & Matching

- [module]![symbol]
 - ?? DumpFiles!zero
- What is the result of: ?face
- Some commands can take a string wildcard
 - * represents zero or more characters
 - ? represents any single character
 - [] contain a list of characters represent any single character in the list. you can use a hyphen (-) to specify a range.
 - # represents zero or more of the preceding characters.
 - + represents one or more of the preceding characters. \ is used as escape character
- x kernel32!Create*[WA]

Set Symbol Suffix

- The ss command is used to set the current suffix value:

```
0:000> ss n
0:000> bp kernel32!CreateHardLink
Couldn't resolve error at 'kernel32!CreateHardLink'
0:000> ss a
0:000> bp kernel32!CreateHardLink
0:000> bl
0 e 00000000`76c2bcc0  0001 (0001)  kernel32!CreateHardLinkA
0:000> ss w
0:000> bp kernel32!CreateHardLink
0:000> bl
0 e 00000000`76c2bcc0      0001 (0001) kernel32!CreateHardLinkA
1 e 00000000`76c2b9c0      0001 (0001) kernel32!CreateHardLinkW
```

Pseudo Registers

- Variables maintained by the debugger that hold certain values
 - All pseudo registers begin with the \$ sign
 - To distinguish from other symbols use the @
 - Use the r command to set a value
 - You don't need the @ sign when using the r command
 - `r $t1 = @$t1 + 1`

Pseudo Registers (Some of them)

Register	Description
\$exp	The last expression that was evaluated
\$ra	The return address that is currently on the stack
\$ip	The instruction pointer register (eip or rip)
\$sexentry	The address of the entry point of the current process
\$retreg	The primary return value register (eax, rax)
\$csp	The current call stack pointer (esp, rsp)
\$p	The value that the last d* (Display Memory) command
\$peb	The address of the PEB block
\$teb	The address of the TEB block
\$tpid \$tid	The Process and Thread Ids
\$bp number	The address of breakpoint number (\$bp1)
\$frame	The current frame index
\$dbgtime	The current time
\$callret	The return value of the last function that .call called

User Defined Pseudo Registers

- There are 20 user-defined pseudo-registers
 - \$t0, \$t1, ..., \$t19
- Variables that you can read and write through the debugger
- You can store any integer value
- The default preset value is 0
- They can be especially useful as loop variables
- They can be very useful inside scripts

Pseudo Register Example

- What does the following line print?

```
.for (r $t1=1; @$t1 <= @@C++(10U); r $t1  
= @$t1 + 1) { .for (r $t2 = 1; @$t2 <=  
@@C++(10U); r $t2 = @$t2 + 1) {.printf  
"%d\t", @@C++(@$t1 * @$t2);}; .echo; }
```

Result

The screenshot shows the WinDbg debugger interface with a command window open. The command entered is ".for (r \$t1=1; @\$t1 <= @@C++(10U); r \$" followed by a series of numbers forming a 10x10 multiplication table:

1	2	3	4	5	6	7	8	9	10
2	4	6	8	10	12	14	16	18	20
3	6	9	12	15	18	21	24	27	30
4	8	12	16	20	24	28	32	36	40
5	10	15	20	25	30	35	40	45	50
6	12	18	24	30	36	42	48	54	60
7	14	21	28	35	42	49	56	63	70
8	16	24	32	40	48	56	64	72	80
9	18	27	36	45	54	63	72	81	90
10	20	30	40	50	60	70	80	90	100

The status bar at the bottom shows: Ln 0, Col 0 | Sys 0:KdSrv:S | Proc 000:0 | Thrd 000:0 | ASM | OVR | CAPS | NUM.

Log File

- You can save the Debugger command window to a log file:
 - .logopen [filename]
 - .logopen /t c:\temp\mylog.txt
- To append to existing log use:
 - .logappend [filename]
- To close the file use:
 - .logclose
- .logfile displays the log file status

Breakpoints

- The breakpoint window can be found under the Edit top-level menu
- The command edit-box is like small command window line
- You can specify breakpoint addresses, module and routine offsets, or source file and line number

A Breakpoint Expert
is a Debugging Master!

Controlling Breakpoints

Register	Description
F9	Set a breakpoint in the source window
<u>.bpcmds</u>	List all breakpoints with the commands that created them
<u>bp</u>	Set a new breakpoint
<u>bm</u>	Set new breakpoints on symbols that match a specified pattern
<u>ba</u>	Set a processor breakpoint (data breakpoint)
<u>bc</u>	Breakpoint clear
<u>bd</u>	Disable breakpoint
<u>be</u>	Re-enable one or more disabled breakpoints
<u>br</u>	Change the ID of an existing breakpoint
<u>bs</u>	change the command associated with an existing breakpoint
<u>bsc</u>	Change a breakpoint condition



Processor Breakpoints Example

```
>? (wchar_t *)name._Bx._Alias
Evaluate expression: 2030176 = 00000000`001efa60
>ba r4 0x1efa60
>bl
1 e 00000000`001efa60 r 4 0001 (0001) 0:****
>g
breakpoint 1 hit
DataBreakPoint!std::char_traits<wchar_t>::assign+0x1b:
00000001`3f8a1a6b 5f          pop      rdi
```

Complex Breakpoints

- Break only from thread 2 after 3 passes

```
~2 bu MyFunc 3
```

- Break on member method:

```
bp @@C++(Person::GetName)
```

- Breakpoint with command

```
ba w4 0x71a578a8 "k;g"
```

```
bu MyDll!func ".dump c:\tmp\dump.dmp; g"
```

- Conditional Breakpoint

```
bp 'x.cpp:14' "j (poi(i)>5) ''; 'gc' "
```

Walking the Stack

- Open the Call Stack window
 - View->Call Stack
- Issue the **k** command
- Common variants
 - kP** (lists all function parameters)
 - kb** (lists first 3 parameters to each function)
 - k [n]** (lists no more than *n* stack frames)

Displaying Memory

`d[a|b|c|d|D|f|p|q|u|w|W] [/c width] [range]`

- a=ASCII, b=byte & ASCII, c=DWORD & ASCII, d=DWORD, D=double, f=float, p=pointer size, q=QWORD, u=Unicode, w=WORD, W=WORD & ASCII

`dt [options] [address]`

- Displays types and/or data
- a <count> (displays array elements)
- Check the docs

`!address [address | -summary -filter]`

- Displays status of a memory block (or entire process address space)



Additional Useful Commands

- **~* k** – call stack for all threads
- **.pcmd -s "r \$tpid, \$tid"** - issue a command whenever the target stops executing
- **.call ShowRestartMessage(envp)** – call a function
- **.reload** - Reload modules and symbols
- **!dlls** - list of loaded modules
- **lm vm kernel32** – verbose info
- **!dh ntdll** – display PE header
- **!runaway 7** – time consumed by thread
- **.dml_start** – Start Debugger Markup Language
- **kM** – DML stack command

Using dt – Display Type

- The [Display Type](#) command is very powerful

```
dt [-DisplayOpts] [-SearchOpts] [module!]Name [[-SearchOpts] Field]  
[Address] [-l List]
```

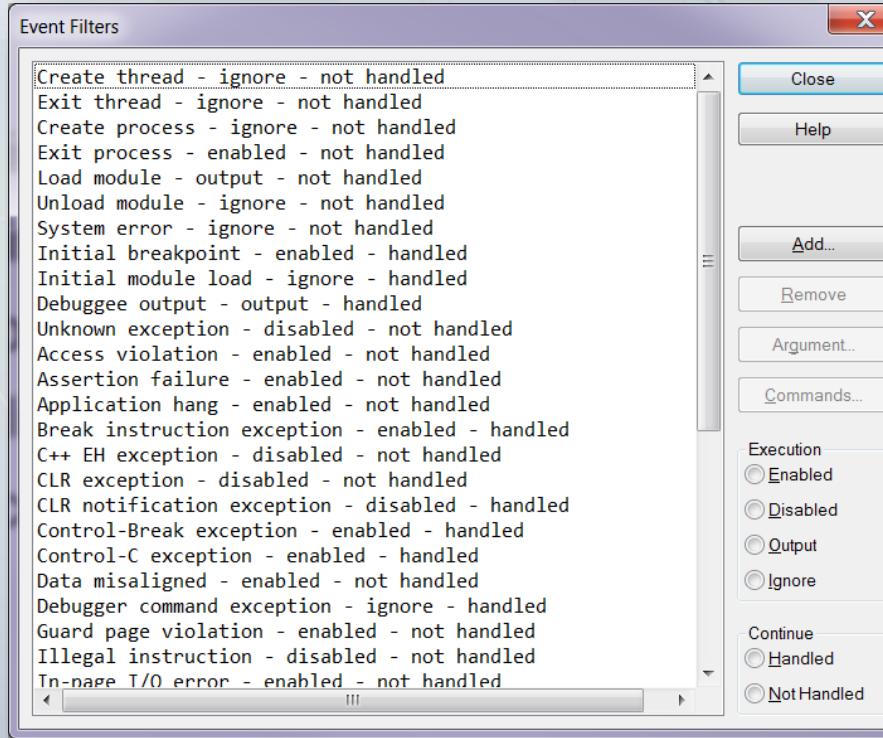
```
dt [-DisplayOpts] Address [-l List]
```

```
dt -h
```

```
0:000> dt nt!_PEB CriticalSectionTimeout. 000007fffffd8000  
ntdll!_PEB  
+0x0c0 CriticalSectionTimeout : 0xfffffe86d`079b8000  
+0x000 LowPart : 0x79b8000  
+0x004 HighPart : -6035  
+0x000 u : <unnamed-tag>  
+0x000 QuadPart : -2592000000000000
```

Event Filters

- Like the exception window of Visual Studio
- It is based on the `WaitForDebugEvent` events



WinDbg Scripts

- A script file is a text file that contains a sequence of debugger commands
- For CDB & NTSD, a ntsd.ini file script will start when starting the debugger
- The \${\$argn} can be used inside a script to read input arguments
- Scripts can call other scripts
- Run script with: <, \$><, \$\$<, \$\$><, \$\$>a<
- For more information:
 - [http://msdn.microsoft.com/en-us/library/ff566261\(v=VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ff566261(v=VS.85).aspx)
 - [http://msdn.microsoft.com/en-us/library/ff560096\(v=VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ff560096(v=VS.85).aspx)

Example

```
## Get module list LIST_ENTRY in $t0.  
r? $t0 = &@$peb->Ldr->InLoadOrderModuleList  
## Iterate over all modules in list.  
.for (r? $t1 = *(ntdll!_LDR_DATA_TABLE_ENTRY**)@$t0;  
(@$t1 != 0) & (@$t1 != @$t0);  
    r? $t1 = (ntdll!_LDR_DATA_TABLE_ENTRY*)@$t1->InLoadOrderLinks.Flink)  
{  
## Get base address in $Base.  
as /x ${/v:$Base} @@c++(@$t1->DllBase)  
## Get full name into $Mod.  
as /msu ${/v:$Mod} @@c++(&@$t1->FullDllName)  
.block  
{  
    .echo ${$Mod} at ${$Base}  
}  
ad ${/v:$Base}  
ad ${/v:$Mod}  
}
```

Kernel Debugging

- Kernel debugging can assist user-mode debugging
 - Memory, IO, Thread & Process information from the kernel can reveal problems in user-mode code
- You can do a local or remote kernel debugging
- You can remote debug a local virtual machine
- On Window NT 6.x a bcd setting is required:
 - `bcdedit -debug`
 - `bcdedit -dbgsettings ...`

Remote Debugging With WinDbg

- On the target (Server) machine run:

```
dbgsrv.exe -t tcp:port=6160
```

- it needs the dbgeng.dll & dbghlp.dll

- Open the firewall for dbgsrv.exe

- On the host (client) machine run WinDbg

```
WinDbg -premove tcp:server=<machine ip or  
name>, port = 6160
```

- Use the Attach to process to start debugging

Minidump Files

- A minidump is a snapshot of a process
 - May be created at any time, not just when a process crashes
- Minidump types
 - Kernel mindumps (not relevant for this course)
 - Basic (usually enough for native processes)
 - Full (required to get useful info for managed processes)
- Minidump creation
 - .dump [options] <filename>
 - On Vista & 2008 use Task Manager
 - ADPlus
 - ProcDump (SysInternals)

Minidump Creation

- WinDbg
 - **.dump [options] filename**
 - Options
 - /ma (full minidump)
 - /o (overwrite existing file)
 - /u (ensure unique filename)
 - /c (add a comment)
 - /b (compress to CAB)
- ADPlus
 - Hang mode – noninvasive attach
 - Crash mode – attaches the CDB debugger
 - Common options
 - hang (hang mode)
 - crash (crash mode)
 - pn (specify process name including extension)
 - p (specify process ID)
 - c (specify XML config file to read options from)
 - quiet (don't show various confirmation dialogs)

Opening a Minidump

- In WinDbg, File->Open Crash Dump...
- `WinDbg -z <dump_file>`
- Issue the “magical” command
- `!analyze -v`
- Can use most other WinDbg/SOS commands
- You can also open a user mode dump file with Visual Studio

Application Verifier

- The !avrf extension controls the settings of App Verifier and displays a variety of output produced by Application Verifier

```
>!avrf
```

```
Verifier package version >= 3.00
```

```
Application verifier settings (80643027):
```

- full page heap
- Handles

```
> g
```

```
VERIFIER STOP 0000000000000300: pid 0x315C: Invalid handle exception for current stack trace.
```

```
00000000C0000008 : Exception code.
```

```
00000000002CFA30 : Exception record. Use .exr to display it.
```

```
00000000002CF540 : Context record. Use .cxr to display it.
```

```
0000000000000000 : Not used.
```

Application Verifier

```
0:000> !avrf -cnt
WaitForSingleObject calls: 0
WaitForSingleObjectEx calls: 0
WaitForMultipleObjects calls 0
WaitForMultipleObjectsEx calls: 0
Waits with timeout calls: 0
Waits with timeout failed: 0
CreateEvent calls: 2
CreateEvent calls failed: 0
Heap allocation calls: 81
Heap allocations failed: 0
CloseHandle called with null handle: 0
CloseHandle called with pseudo handle: 0
Heaps created: 2
Heaps destroyed: 0
...
...
```

Resources

- Debugging Tools for Windows
 - [http://msdn.microsoft.com/en-us/library/ff551063\(v=VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ff551063(v=VS.85).aspx)
- WinDbg Info: <http://windbg.info/>
- Common Commands
 - <http://windbg.info/doc/1-common-cmds.html>
- Application Verifier
 - [http://msdn.microsoft.com/en-us/library/ms220948\(v=VS.90\).aspx](http://msdn.microsoft.com/en-us/library/ms220948(v=VS.90).aspx)

Summary

- In this appendix we got deep into the powerful WinDbg Debugger
- To mitigate hard problems we need:
 - ✓ Tools
 - ↳ Deep Understanding
 - Experience

Thank You